

edudrift

An LLM's living hypothesis log on AI's impact on higher education

Santosh B. Srinivas · HEC Paris

Snapshot: 2026-07-21

<https://santoshbs.github.io/theaidrift/edu/>

edudrift is an experiment in letting a local open-weight large language model continuously process AI-related news signals and hypothesise their implications for specific higher-education capabilities. Each entry in this document is a working manuscript: it is generated and revised **entirely by the model**, with **no human in the editing loop**. The author's contribution is limited to designing the system — selecting RSS sources, defining the capability scaffold, writing the prompts, and committing the model's output to version control. The prose, the reasoning, the framing, and the choice of which signals matter are all the model's.

The capability scaffold is drawn from the **Higher Education Reference Models (HERM)** v3.2.0, a global industry-standard enterprise-architecture framework curated by the HERM Working Group of CAUDIT (the Council of Australasian University Directors of Information Technology), in collaboration with EDUCAUSE, UCISA, and EUNIS, and used by more than a thousand institutions worldwide. HERM is published under CC BY-NC-SA 4.0 and is available at <https://www.caudit.edu.au/communities/caudit-higher-education-reference-models/>.

For each capability, the log states three horizon hypotheses, each grounded in cited public signals:

- **Near** (2026–2031) — what current evidence most plausibly supports;
- **Mid** (2036–2041) — extrapolations contingent on signal trajectories holding;
- **Far** (2046–2051) — competing hypotheses where the trajectory is genuinely uncertain.

Hypotheses are abductive and revisable; each is version-controlled at the source repository. *Italic statements* are the formal hypothesis claims.

The live, continuously revised version is published at </theaidrift/edu/>.

{*Caveat lector*. Because this document is fully model-generated, claims should be read as structured speculation grounded in cited signals, not as expert assessment. Citations point to the public sources the model drew on; readers should verify any claim against the original source before relying on it.}

Contents

Academic Administration	6
Near (2026–2031)	6
Mid (2036–2041)	6
Far (2046–2051)	6
References	6
Alumni Engagement	8
Near (2026–2031)	8
Mid (2036–2041)	8
Far (2046–2051)	8
References	8
Business Capability Management	9
Near (2026–2031)	9
Mid (2036–2041)	9
Far (2046–2051)	9
References	10
Commercial Delivery	11
Near (2026–2031)	11
Mid (2036–2041)	11
Far (2046–2051)	11
References	11
Commercial Monitoring	13
Near (2026–2031)	13
Mid (2036–2041)	13
Far (2046–2051)	13
References	13
Commercial Sourcing	14
Near (2026–2031)	14
Mid (2036–2041)	14
Far (2046–2051)	14
References	15
Corporate Governance	16
Near (2026–2031)	16
Mid (2036–2041)	16
Far (2046–2051)	16
References	17
Curriculum Management	19
Near (2026–2031)	19
Mid (2036–2041)	19
Far (2046–2051)	19
References	20
Facilities & Property Management	22
Near (2026–2031)	22
Mid (2036–2041)	22
Far (2046–2051)	22
References	23
Finance Management	24

Near (2026–2031)	24
Mid (2036–2041)	24
Far (2046–2051)	24
References	25
Government, Public & Stakeholder Relationships	26
Near (2026–2031)	26
Mid (2036–2041)	26
Far (2046–2051)	26
References	26
Human Resource Management	27
Near (2026–2031)	27
Mid (2036–2041)	27
Far (2046–2051)	27
References	28
Information & Communication Technology Mgt	30
Near (2026–2031)	30
Mid (2036–2041)	30
Far (2046–2051)	30
References	31
Information Management	34
Near (2026–2031)	34
Mid (2036–2041)	34
Far (2046–2051)	34
References	35
Legal Services	38
Near (2026–2031)	38
Mid (2036–2041)	38
Far (2046–2051)	38
References	39
Library Management	41
Near (2026–2031)	41
Mid (2036–2041)	41
Far (2046–2051)	41
References	41
Promotions Management	42
Near (2026–2031)	42
Mid (2036–2041)	42
Far (2046–2051)	42
References	43
Research Administration	44
Near (2026–2031)	44
Mid (2036–2041)	44
Far (2046–2051)	44
References	45
Research Delivery	47
Near (2026–2031)	47
Mid (2036–2041)	47
Far (2046–2051)	48

References	48
Research Funding	51
Near (2026–2031)	51
Mid (2036–2041)	51
Far (2046–2051)	51
References	51
Research Impact	53
Near (2026–2031)	53
Mid (2036–2041)	53
Far (2046–2051)	53
References	54
Research Improvement	55
Near (2026–2031)	55
Mid (2036–2041)	55
Far (2046–2051)	55
References	56
Research Opportunity Analysis	57
Near (2026–2031)	57
Mid (2036–2041)	57
Far (2046–2051)	57
References	57
Research Publications	58
Near (2026–2031)	58
Mid (2036–2041)	58
Far (2046–2051)	58
References	59
Research Training	60
Near (2026–2031)	60
Mid (2036–2041)	60
Far (2046–2051)	60
References	61
Strategy Management	62
Near (2026–2031)	62
Mid (2036–2041)	62
Far (2046–2051)	62
References	63
Student Administration	64
Near (2026–2031)	64
Mid (2036–2041)	64
Far (2046–2051)	64
References	65
Student Admission Management	66
Near (2026–2031)	66
Mid (2036–2041)	66
Far (2046–2051)	66
References	67
Student Assessment	68

Near (2026–2031)	68
Mid (2036–2041)	68
Far (2046–2051)	68
References	69
Student Attraction & Recruitment	71
Near (2026–2031)	71
Mid (2036–2041)	71
Far (2046–2051)	71
References	72
Student Completion & Graduation	73
Near (2026–2031)	73
Mid (2036–2041)	73
Far (2046–2051)	73
References	73
Student Enrolment	74
Near (2026–2031)	74
Mid (2036–2041)	74
Far (2046–2051)	74
References	74
Student Support & Wellbeing Management	75
Near (2026–2031)	75
Mid (2036–2041)	75
Far (2046–2051)	75
References	76
Supporting Services	77
Near (2026–2031)	77
Mid (2036–2041)	77
Far (2046–2051)	77
References	78
Teaching & Learning Delivery	79
Near (2026–2031)	79
Mid (2036–2041)	79
Far (2046–2051)	79
References	80

Academic Administration

capability_id: academic-administration · created: 2026-04-10 · revised: 2026-04-10 · refs: 2

Near (2026–2031)

For Academic Administration broadly, evidence that AI-enabled misconduct overlaps with traditional forms of plagiarism [1] may drive a shift in academic policy & regulation management toward integrated behavioral frameworks rather than tool-specific bans [1]. For institutions implementing structured faculty development partnerships, as demonstrated at the University of Central Florida [2], this transition may be supported by bottom-up infrastructure involving collaborations between teaching centers and digital learning divisions [2]. This represents a change in educational substance by refocusing regulation on behavioral patterns of dishonesty rather than the specific technology employed. Because updates to academic policy typically require faculty senate approval and governance review, the pace of institutional adoption may remain slow. Processes such as academic year scheduling and timetabling management are unlikely to change in this period as they are not functionally linked to academic integrity signals.

Hypothesis: In the near term, academic policy & regulation management may shift from tool-specific AI bans toward holistic integrity frameworks that treat AI use as one of several overlapping plagiarism behaviors.

Mid (2036–2041)

For Academic Administration broadly, if current trajectories hold, the identified link between AI use and general plagiarism [1] could lead to the development of longitudinal behavioral profiles within academic policy & regulation management. If institutions continue to build evolving policy infrastructures through cross-functional partnerships [2], regulatory mechanisms might shift from evaluating single instances of misconduct to identifying patterns of dishonesty across a student’s academic career. This transition would likely be constrained by data privacy laws and institutional due process requirements, which tend to persist regardless of technological shifts. The fundamental administrative structure of the academic year remains unlikely to change, as it is governed by external labor contracts and seasonal student demographics rather than integrity regulation.

Hypothesis: In the mid term, academic policy & regulation management may move toward a behavioral-pattern model of integrity enforcement, extrapolating from current evidence of overlapping plagiarism methods.

Far (2046–2051)

For Academic Administration broadly, the long-term trajectory of integrity regulation remains uncertain, with two primary rival outcomes. One possibility is that the persistence of overlapping plagiarism [1] leads to the automation of the disciplinary process through algorithmic pattern matching. A competing outcome is that the continued erosion of traditional authorship leads to a collapse of the current regulatory regime, where the “product” of education is no longer a verifiable independent work. It remains unknown how professional accreditation bodies will redefine “competence” if traditional plagiarism regulations become obsolete. Even at this horizon, the legal requirement for human-led due process in high-stakes disciplinary appeals is likely to resist full automation.

Hypothesis A: In the far term, academic policy & regulation management may evolve into a behavioral monitoring system that identifies integrity risks.

Hypothesis B (competing): In the far term, the capability of academic policy & regulation management may be fundamentally diminished as the institution shifts away from the validation of independent student authorship.

References

1. The Conversation Au (2026-04-01). “Does AI mean more uni students are plagiarising their work?” <https://theconversation.com/does-ai-mean-more-uni-students-are-plagiarising-their-work-279565>

2. EDUCAUSE Review (2026-04-08). “Building AI Initiatives with Bottom-Up Champions” <https://er.educause.edu/articles/2026/4/building-ai-initiatives-with-bottom-up-champions>

Alumni Engagement

capability_id: alumni-engagement · created: 2026-06-20 · revised: 2026-06-20 · refs: 1

Near (2026–2031)

For the subsegment of alumni engagement focused on campaign and event marketing, the production of promotional assets is shifting toward near-instant, low-cost automated generation. The availability of multi-engine AI builds allows for the conversion of brief instructions into full multichannel campaigns, including scripts, images, and short-form video, at a marginal cost of less than one dollar per campaign [1]. This change in delivery mechanism may enable alumni offices with limited creative budgets to increase the volume of their event-specific outreach. However, the core processes of benefactor management and high-touch alumni relationship management are unlikely to change in this period, as these functions rely on interpersonal trust and social legitimacy that automated ad-generation tools do not address. The shift is currently limited to the technical execution of campaign management rather than the strategic determination of alumni outreach goals.

Hypothesis: In the near term, the cost and time required for routine asset creation in alumni campaign management will decrease significantly for institutions adopting automated generative pipelines.

Mid (2036–2041)

For the subsegment of campaign and event marketing, if current trajectories in low-cost asset generation hold, the labor focus within alumni offices may shift from content production to content curation and prompt orchestration. Extrapolating from early deployment of multi-engine AI builds [1], the ability to generate vast iterations of campaign materials could lead to hyper-segmented event marketing tailored to minute alumni demographics. Despite this, the institutional governance and regulatory constraints surrounding benefactor management—such as tax compliance and legal requirements for endowments—will likely persist unchanged due to their grounding in national law rather than communication technology. The mechanism of change remains a reduction in production friction, not a change in the substance of the alumni-institutional relationship.

Hypothesis: In the mid term, routine marketing for alumni events may become a fully automated delivery process, shifting human labor toward the strategic oversight of automated campaign streams.

Far (2046–2051)

For the subsegment of campaign marketing, the long-term impact of automated content generation remains uncertain, particularly regarding the risk of counterproductivity. If the volume of automated outreach reaches a threshold where alumni perceive communications as synthetic or impersonal, the perceived value of the institutional connection may erode. It remains unknown whether the ability to produce infinite, low-cost assets will lead to higher engagement or to increased digital fatigue among alumni. Even at this horizon, the fundamental requirement for human presence in high-stakes benefactor management is expected to resist change, as the social legitimacy of major philanthropy typically depends on genuine human reciprocity.

Hypothesis A: In the far term, alumni campaign management achieves a state of autonomous hyper-personalization, where every alumnus receives a unique, AI-generated visual and narrative appeal.

Hypothesis B (competing): In the far term, the proliferation of low-cost automated campaign materials leads to a devaluation of digital alumni outreach, forcing institutions to return to high-cost, human-centric communication to maintain social legitimacy.

References

1. Dev To Ai (2026). “How we shipped”One Brief Full Ad Campaign in 60s” for \$0.35 — a Dify + x402 + multi-engine AI build.” <https://dev.to/aikitpros/how-we-shipped-one-brief-full-ad-campaign-in-60s-for-035-a-dify-x402-multi-engine-ai-ib>

Business Capability Management

capability_id: business-capability-management · created: 2026-07-20 · revised: 2026-07-20 · refs: 13

Near (2026–2031)

For business capability management broadly, organisational design may shift in substance toward system-centric architectures that prioritize AI systems over traditional human-centric role definitions [9]. The adoption of AI operating models could necessitate a change in the educational substance of organisational design to manage the transition toward software-factory operations and comprehensive workflow redesign [4]. In a subsegment of institutions managing large-scale AI infrastructure, benefits management shifts in substance toward more stringent ROI validation, as evidence indicates only 28% of such projects fully realize projected payoffs [3]. The delivery mechanism of programme & project management is modified through the use of AI lifecycle frameworks [1] and the integration of AI-driven decision-making within Agile frameworks [13]. For a subset of managers, the delivery of programme & project management further incorporates AI to audit project post-mortems for root cause analysis [8] or to automate delivery metrics within Scaled Agile Frameworks [5]. For a subset of institutions, change portfolio management delivery shifts toward controlled experimentation and piloting [7], while broader institutional change is pressured to transition from isolated pilots to systemic, business-wide implementation frameworks [11], contingent on leadership alignment [12] and the use of internal learning champions to drive peer-led adoption [14]. Quality management processes are unlikely to change yet, as they remain dependent on stable administrative specifications that are independent of AI deployment.

Hypothesis: In the near term, business capability management may integrate AI-specific operating models, system-centric organisational architectures, and automated project auditing to govern the transition toward software-factory operations and mitigate high infrastructure failure rates.

Mid (2036–2041)

For business capability management broadly, if the trend of low ROI for infrastructure and recurring project failures persists, the capability may shift from experimental AI adoption to disciplined architectural integration [3][10]. Extrapolating from the requirement to govern AI sprawl, enterprise architecture could evolve to standardize the software-factory model across the institution [4][6]. This represents a change in the educational substance of the capability, as institutional value is measured by production-level operational integration rather than prototype capability [1]. The fundamental mechanisms of institutional accountability—specifically the determination of legal and professional responsibility for capability failure—will likely persist unchanged due to regulatory and governance constraints.

Hypothesis: In the mid term, the capability may shift toward a disciplined architectural integration where AI-driven operating models and automated failure analysis are standardized as core institutional utilities.

Far (2046–2051)

For business capability management broadly, the long-term trajectory remains uncertain regarding the automation of strategic decision-making. One plausible outcome is that change portfolio management becomes largely automated, with AI systems optimizing resource allocation based on real-time adoption and ROI data [3]. However, the process of determining the social legitimacy of institutional goals resists change even at this horizon and likely remains a human-centric function. It remains unknown whether AI can manage benefits management without replicating human cognitive biases in value attribution.

Hypothesis A: In the far term, change portfolio management may transition to a bounded autonomous system that optimizes institutional resource allocation based on real-time performance signals.

Hypothesis B (competing): In the far term, the high failure rate of complex AI integrations leads to a regression toward rigid, human-governed capability maps that strictly limit AI autonomy to narrow delivery mechanisms.

References

1. Dzone Ai (2026). “From Concept to Production: A Strategic Framework for AI/ML Project Success.” <https://feeds.dzone.com/link/23558/17313177/ai-lifecycle-framework-for-successful-projects>
2. Dev To Ai (2026). “How to Build an AI Workflow That Your Team Will Actually Use.” <https://dev.to/maxothex/how-to-build-an-ai-workflow-that-your-team-will-actually-use-55he>
3. The Register (2026). “Only 28% of AI infrastructure projects fully pay off, survey finds.” https://go.theregister.com/feed/www.theregister.com/2026/04/07/ai_returns_gartner/
4. Dev To Ai (2026). “AI Operating Models: The \$2M Workflow Redesign Your Board Isn’t Asking About.” https://dev.to/dr_hernani_costa/ai-operating-models-the-2m-workflow-redesign-your-board-isnt-asking-about-1dfd
5. Dzone Ai (2026). “Revolutionizing Scaled Agile Frameworks with AI, MuleSoft, and AWS: An Insider’s Perspective.” <https://feeds.dzone.com/link/23558/17323804/revolutionizing-scaled-agile-frameworks-with-ai>
6. The Register (2026). “Ex-AWS legend explains what enterprises need to make AI actually work.” https://go.theregister.com/feed/www.theregister.com/2026/04/25/ai_enterprise_matt_domo/
7. Wonkhe (2026). “Controlled experimentation with AI – the power of a pilot.” <https://wonkhe.com/blogs/controlled-experimentation-with-ai-the-power-of-a-pilot/>
8. Dev To Ai (2026). “I Asked AI to Audit My Last 3 Failed Projects — Here’s What It Found.” <https://dev.to/tillerman/i-asked-ai-to-audit-my-last-3-failed-projects-heres-what-it-found-30on>
9. Dev To Ai (2026). “The Future of Work: From Employees to AI Systems.” <https://dev.to/aiaddict25709/the-future-of-work-from-employees-to-ai-systems-3p6>
10. Dev To Ai (2026). “Why Enterprise AI Pilots Fail.” <https://dev.to/xccelera/why-enterprise-ai-pilots-fail-1pcc>
11. Zdnet Ai (2026). “Moving from AI pilots to business-wide value requires a superhighway - how to ramp up.” <https://www.zdnet.com/article/moving-from-ai-pilots-to-business-wide-value-accenture-research/>
12. EDUCAUSE Review (2026). “The Leadership Moves That Make AI Adoption Work.” <https://er.educause.edu/articles/sponsored/2026/5/the-leadership-moves-that-make-ai-adoption-work>
13. Dzone Ai (2026). “Integrating AI-Driven Decision-Making in Agile Frameworks: A Deep Dive into Real-World Applications and Challenges.” <https://feeds.dzone.com/link/23558/17332221/ntegrating-ai-driven-decision-making-in-agile>
14. Elearning Industry (2026). “Internal Learning Champions: Why Peer Credibility Beats Top-Down Training In AI Adoption.” <https://elearningindustry.com/internal-learning-champions-why-peer-credibility-beats-top-down-training-in-ai-adoption>

Commercial Delivery

capability_id: commercial-delivery · created: 2026-06-04 · revised: 2026-06-04 · refs: 2

Near (2026–2031)

For US-based institutions and those utilizing frontier AI models, the process of commercial activity exit management is increasingly constrained by both regulatory mandates and technical dependencies. In the US, federal blacklisting of specific AI technologies may trigger forced commercial activity exit management, shifting the mechanism of vendor termination from planned lifecycles to compliance-driven mandates [1]. Simultaneously, institutions attempting to transition between frontier models may find the exit process obstructed by vendor lock-in, which increases the complexity and cost of swapping providers [2]. These shifts represent changes to the delivery mechanism of risk and vendor management rather than the educational substance of the commercial activities. Because current signals focus on exit and compliance, the core processes of commercial activity output management are unlikely to change in the near term.

Hypothesis: In the near term, a subset of institutions will experience increased friction in commercial activity exit management due to a combination of geopolitical regulatory restrictions and technical vendor lock-in.

Mid (2036–2041)

For institutions utilizing frontier AI models, if current trajectories of vendor lock-in and regulatory volatility hold, commercial activity issue management may shift toward the requirement of proactive “interoperability audits.” Extrapolating from the difficulty of swapping models [2], institutions could implement mandatory exit-strategy documentation during the initial procurement phase to mitigate the risk of operational paralysis. This represents a more conservative delivery of vendor oversight, as demonstrated by the necessity of aligning commercial ties with state-level prohibited entity lists [1]. The fundamental requirement for specialized legal counsel to navigate the intersection of trade law and software licensing will persist unchanged.

Hypothesis: In the mid term, institutions may integrate formal exit-strategy planning and interoperability requirements into their commercial activity issue management to counter the effects of vendor lock-in.

Far (2046–2051)

For institutions utilizing large-scale commercial AI, the long-term state of commercial activity exit management depends on whether model interoperability becomes a technical standard or remains a proprietary barrier. One plausible outcome is the emergence of fragmented technological blocs where exit management is a routine, trigger-based operational task dictated by national security boundaries [1]. A competing outcome is the persistence of deep vendor lock-in, where the cost of exit management becomes prohibitively high, effectively merging the institutional capability with the vendor’s own lifecycle. It remains unknown whether future regulatory frameworks will provide structured “grace periods” for institutional exits or mandate immediate cessation. The basic institutional need to manage the legal termination of contracts will resist change regardless of the underlying technology.

Hypothesis A: In the far term, commercial activity exit management becomes a highly automated, trigger-based process within a fragmented global AI market.

Hypothesis B (competing): In the far term, extreme vendor lock-in renders independent commercial activity exit management practically impossible for most institutions, leading to permanent vendor dependency.

References

1. Ars Technica (2026-04-09). “Trump-appointed judges refuse to block Trump blacklisting of Anthropic AI tech.” <https://arstechnica.com/tech-policy/2026/04/trump-appointed-judges-refuse-to-block-trump-blacklisting-of-anthropic-ai-tech/>
2. The Register (2026-04-28). “Locked, stocked, and losing budget: AI vendor lock-in bites back.” https://www.theregister.com/2026/04/28/ai_vendor_lock_in_bites_back/

[//go.theregister.com/feed/www.theregister.com/2026/04/28/locked_stocked_and_losing_budget/](http://go.theregister.com/feed/www.theregister.com/2026/04/28/locked_stocked_and_losing_budget/)

Commercial Monitoring

capability_id: commercial-monitoring · created: 2026-06-04 · revised: 2026-06-04 · refs: 1

Near (2026–2031)

For a subsegment of institutions employing proprietary frontier AI models, the process of commercial contingency planning may shift as assumptions of model interchangeability prove invalid [1]. The mechanism driving this change is vendor lock-in, where technical dependencies and pricing structures prevent the rapid swapping of one frontier model for another [1]. Consequently, contingency plans that previously relied on the ability to pivot between vendors in short timeframes may be rendered ineffective [1]. This represents a change in educational substance, as the institutional approach to risk management shifts from a “swap-out” strategy to a lock-in mitigation strategy. Commercial contract performance management is unlikely to change significantly in this period, as the basic monitoring of service-level agreements and uptime remains a standard contractual requirement.

Hypothesis: In the near term, for institutions using proprietary frontier models, commercial contingency planning will shift from assumptions of model interchangeability to the management of vendor lock-in risks.

Mid (2036–2041)

For the same subsegment of institutions using proprietary frontier models, if current trajectories of vendor lock-in hold, commercial activity reporting may expand to track dependency metrics [1]. This process could evolve to monitor “lock-in depth,” treating the inability to switch vendors as a quantifiable financial risk during reporting cycles [1]. Commercial contingency planning may further shift toward architectural diversification to avoid total dependence on a single provider [1]. The fundamental mechanisms of commercial contract performance management will likely persist unchanged, as the legal frameworks governing institutional procurement and vendor accountability are slow to evolve.

Hypothesis: In the mid term, for institutions using proprietary frontier models, commercial activity reporting may incorporate specific metrics to monitor and quantify vendor dependency.

Far (2046–2051)

For institutions that have integrated proprietary frontier models, the scope of commercial contingency planning remains uncertain, depending on whether industry standards move toward open interoperability or deeper integration. One outcome involves the total obsolescence of “swapping” as a contingency, with the vendor instead becoming a core utility similar to electrical power. A rival outcome involves the emergence of standardized model interfaces that restore the capability for rapid vendor migration. Throughout these scenarios, the requirement for commercial activity reporting to produce an audit trail for institutional finance will resist change. It remains unknown if open-weight models can reach a level of capability that makes proprietary lock-in a secondary concern for institutional risk managers.

Hypothesis A: In the far term, for these institutions, commercial contingency planning may be replaced by a utility-based monitoring model where the provider is treated as essential, non-swappable infrastructure.

Hypothesis B (competing): In the far term, for these institutions, the adoption of open-standard model interfaces may restore the capability for rapid vendor swapping within commercial contingency planning.

References

1. The Register (2026-04-28). “Locked, stocked, and losing budget: AI vendor lock-in bites back.” https://go.theregister.com/feed/www.theregister.com/2026/04/28/locked_stocked_and_losing_budget/

Commercial Sourcing

capability_id: commercial-sourcing · created: 2026-07-17 · revised: 2026-07-17 · refs: 4

Near (2026–2031)

For a subset of institutions adopting agentic sourcing tools and those operating within US regulatory jurisdictions, the processes of commercial market analysis and opportunity identification are experiencing a divergence in reliability. The delivery mechanism for commercial market analysis may be improved through the use of structured product catalogs and Model Context Protocol (MCP) servers, which allow AI agents to retrieve real-time pricing and availability across multiple merchants [3]. Simultaneously, the process of commercial opportunity identification is being pressured by the deployment of self-learning AI sales systems that automate the prospecting pipeline, shifting vendor outreach from fixed templates to targeted, product-specific agentic communication [4]. However, the reliability of this identification process is undermined by AI-focused SEO tactics designed to bias LLM-generated vendor recommendations [1]. In the US, the commercial opportunity assessment process is further constrained by the blacklisting of specific high-capability AI providers, which restricts the available vendor pool for institutional use [2]. Formal procurement mandates requiring competitive bidding and documented audit trails are unlikely to change, as these regulatory constraints prevent total reliance on AI-generated discovery.

Hypothesis: In the near term, for these subsegments and regions, commercial sourcing may bifurcate into a high-efficiency automated track for commodity sourcing and a highly scrutinized, adversarial track for strategic procurement.

Mid (2036–2041)

For these same subsets and regions, if current trajectories of agentic structured data and geopolitical restrictions hold, the process of commercial market analysis could shift toward a “verified-only” workflow. Institutions may restrict their AI agents to accessing only trusted, structured MCP servers or verified registries to mitigate the distortion caused by AI SEO [1][3]. The commercial opportunity assessment process may integrate automated geopolitical compliance filters that exclude blacklisted entities at the point of initial discovery [2]. This represents a change in the delivery mechanism of vendor screening, moving it from a post-discovery check to a pre-filter. Internal institutional governance regarding the final selection of vendors is likely to persist unchanged due to the financial risks and legal liabilities associated with large-scale commercial contracts.

Hypothesis: In the mid term, commercial sourcing for these subsegments may shift toward a “trusted-network” model, where structured, machine-verifiable data sources are prioritized over open-web generative search.

Far (2046–2051)

For institutions relying on AI for sourcing, the outcome depends on whether AI-agent protocols reach global standardization or if regulatory fragmentation increases. One plausible outcome is that commercial market analysis becomes a fully machine-to-machine process where AI agents negotiate directly with merchant servers, bypassing the “search” phase entirely. It remains unknown whether AI SEO will evolve into a transparent signaling mechanism or if geopolitical blacklisting becomes a permanent, expanding feature of institutional procurement. The requirement for human-led final due diligence will likely resist change because of the legal accountability required for institutional financial commitments.

Hypothesis A: In the far term, commercial market analysis may fully migrate to structured, machine-to-machine registries, rendering general-purpose AI search obsolete for institutional sourcing.

Hypothesis B (competing): In the far term, the combination of sophisticated AI manipulation and volatile geopolitical restrictions may force institutions to revert to curated, human-managed vendor lists for all high-stakes procurement.

References

1. The Verge (2026-04-06). “Can AI responses be influenced? The SEO industry is trying.” <https://www.theverge.com/tech/900302/ai-seo-industry-google-search-chatgpt-gemini-marketing>
2. Ars Technica (2026-04-09). “Trump-appointed judges refuse to block Trump blacklisting of Anthropic AI tech.” <https://arstechnica.com/tech-policy/2026/04/trump-appointed-judges-refuse-to-block-trump-blacklisting-of-anthropic-ai-tech/>
3. Dev To Ai (2026-05-03). “Build an AI Shopping Agent with BuyWhere MCP in 5 Minutes.” <https://dev.to/buywhere/build-an-ai-shopping-agent-with-buywhere-mcp-in-5-minutes-80g>
4. Dev To Ai (2026-05-04). “We Built a Self-Learning AI Sales System in 48 Hours.” <https://dev.to/pvcvotewe-built-a-self-learning-ai-sales-system-in-48-hours-2hbi>

Corporate Governance

capability_id: corporate-governance · created: 2026-07-21 · revised: 2026-07-21 · refs: 28

Near (2026–2031)

For Corporate Governance broadly, the deployment of autonomous agents shifts the educational substance of risk management from the validation of static AI outputs to the oversight of executed actions [1][5]. To mitigate the operational risks of unsanctioned “shadow AI,” institutions are adopting delivery mechanisms such as agent registries [13] and autonomous governance tools [2][31]. Risk management processes are evolving to prioritize system resilience over immediate productivity gains [48], utilizing runtime security toolkits [10], sandbox execution environments [24], and security frameworks that prioritize asset inventory and risk tiering [32][33]. Compliance monitoring and reporting is increasingly focused on the technical documentation of AI systems, particularly in regions governed by the European Union where the EU AI Act necessitates the maintenance of Annex IV technical files [8] and specific governance protocols for agentic AI [12][17][39]. In the United States, policy and regulation management is adapting to a national AI legislative framework [26], White House recommendations [22], and proposed federal vetting requirements for models prior to public release [51]. Within internal audit and reporting, institutions are integrating specialized systems for tax anomaly detection [41] and the automated extraction of insights from unstructured contracts [42]. The legal principle that fiduciary liability cannot be delegated to a software agent is unlikely to change, as demonstrated by ongoing criminal probes into provider liability for real-world harms [15][28][29].

Hypothesis: In the near term, institutions will shift risk and policy frameworks to mitigate the liability of authorized agents and the operational risks of shadow AI, while moving compliance monitoring toward runtime enforcement and regulatory technical documentation.

Mid (2036–2041)

For Corporate Governance broadly, if current trajectories of agent autonomy and regulatory tension hold, internal audit and reporting may shift from retrospective sampling to the use of governance toolkits providing real-time, immutable audit trails [2][43]. Extrapolating from the emergence of AI operating models, the substance of board-level oversight may shift toward the continuous redesign of organizational workflows—treating the AI “harness” as a dynamic organizational chart—rather than the approval of static policy documents [18][36]. In regions where strict regulatory regimes persist, such as the EU, compliance monitoring will likely remain a high-friction process requiring human certification despite the automation of data collection [12][39]. The requirement for human-in-the-loop verification for high-stakes institutional decisions is unlikely to change due to social legitimacy constraints and the persisting risk of “LLM nepotism” in organizational decision-making [20].

Hypothesis: In the mid term, the capability may shift from periodic auditing to real-time governance monitoring, with board oversight focusing on the architectural design of AI-integrated workflows.

Far (2046–2051)

For Corporate Governance broadly, the interaction between autonomous agents and business continuity management remains a primary uncertainty due to the persistence of systems that cannot explain their own decision-making logic [25]. This uncertainty is compounded by the risk that “secret loyalties” trained into models could allow agents to evade black-box audits [30]. It remains unknown whether observability tools will evolve to penetrate these black boxes or if governance will shift toward a “fail-safe” model where agents are restricted to bounded, non-critical processes. The fundamental requirement for a legally accountable human officer to sign off on institutional risk remains a point of resistance to full automation even at this horizon.

Hypothesis A: In the far term, advancements in AI observability and explainability allow for the full integration of autonomous agents into business continuity management under human supervision.

Hypothesis B (competing): In the far term, the inability to resolve the “black-box” audit problem leads to a permanent bifurcation of governance, where critical institutional functions are legally barred from agentic execution.

References

1. The Register (2026-04-05). “AI agents promise to ‘run the business,’ but who is liable if things go wrong?” https://go.theregister.com/feed/www.theregister.com/2026/04/05/ai_agents_liability/
2. Ai News (2026-04-02). “KiloClaw targets shadow AI with autonomous agent governance” <https://www.artificialintelligence-news.com/news/kiloclaw-targets-shadow-ai-autonomous-agent-governance/>
3. Ai News (2026-04-06). “As AI agents take on more tasks, governance becomes a priority” <https://www.artificialintelligence-news.com/news/as-ai-agents-take-on-more-tasks-governance-becomes-a-priority/>
4. Dev To Ai (2026-04-07). “What’s actually inside an EU AI Act Annex IV technical file (and what engineers usually miss)” https://dev.to/rostyslav_myronenko_faf2f/whats-actually-inside-an-eu-ai-act-annex-iv-technical-file-and-what-engineers-usually-miss-4nh4
5. Ai News (2026-04-08). “Microsoft open-source toolkit secures AI agents at runtime” <https://www.artificialintelligence-news.com/news/microsoft-open-source-toolkit-secures-ai-agents-at-runtime/>
6. Ai News (2026-04-09). “Agentic AI’s governance challenges under the EU AI Act in 2026” <https://www.artificialintelligence-news.com/news/agentic-ais-governance-challenges-under-the-eu-ai-act-in-2026/>
7. The Register (2026-04-09). “AWS: Agents shouldn’t be secret, so we built a registry for them” https://go.theregister.com/feed/www.theregister.com/2026/04/09/aws_ai_agent_registry/
8. TechCrunch (2026-04-09). “Florida AG to probe OpenAI, alleging possible connection to FSU shooting” <https://techcrunch.com/2026/04/09/florida-ag-to-probe-openai-alleging-possible-connection-to-fsu-shooting/>
9. Dev To Ai (2026-04-10). “The EDPB Is Asking About Your AI Agents. Most Teams Can’t Answer.” <https://dev.to/waxell/the-edpb-is-asking-about-your-ai-agents-most-teams-cant-answer-gfk>
10. Dev To Ai (2026-04-11). “AI Operating Models: The \$2M Workflow Redesign Your Board Isn’t Asking About” https://dev.to/dr_hernani_costa/ai-operating-models-the-2m-workflow-redesign-your-board-isnt-asking-about-1dfd
11. Arxiv Cs Cy (2026-04-14). “LLM Nepotism in Organizational Governance” <https://arxiv.org/abs/2604.09620>
12. Future Of Life (2026-03-22). “Statement: Head of US Policy on the White House AI legislative recommendations” <https://futureoflife.org/statement/head-of-us-policy-on-the-white-house-ai-legislative-recommendations/>
13. Ai News (2026-04-16). “OpenAI Agents SDK improves governance with sandbox execution” <https://www.artificialintelligence-news.com/news/openai-agents-sdk-improves-governance-sandbox-execution/>
14. Hackernoon Ai (2026-04-21). “The Hidden Risk of AI Agents: Systems That Can’t Explain Themselves” <https://hackernoon.com/the-hidden-risk-of-ai-agents-systems-that-cant-explain-themselves?source=rss>
15. EDUCAUSE Review (2026-04-21). “White House Unveils National AI Legislative Framework” <https://er.educause.edu/articles/2026/4/white-house-unveils-national-ai-legislative-framework>
16. BBC News (2026-04-21). “OpenAI faces criminal probe over role of ChatGPT in shooting” https://www.bbc.com/news/articles/c62j4ldp2jqo?at_medium=RSS&at_campaign=rss
17. NPR (2026-04-21). “Florida AG launches criminal investigation into ChatGPT over FSU shooting” <https://www.npr.org/2026/04/21/nx-s1-5793967/florida-openai-investigation-mass-shooting-fsu>
18. Lesswrong (2026-04-22). “Narrow Secret Loyalty Dodges Black-Box Audits” <https://www.lesswrong.com/posts/EzdgPbewjeTNHA5F3/narrow-secret-loyalty-dodges-black-box-audits>
19. Fast Company (2026-04-22). “Here’s how to jump-start your company’s responsible AI governance in 90 days” <https://www.fastcompany.com/91527542/heres-how-jump-start-your-companys->

responsible-ai-governance-90-days?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss

20. Marktechpost (2026-04-23). “Mend.io Releases AI Security Governance Framework Covering Asset Inventory, Risk Tiering, AI Supply Chain Security, and Maturity Model” <https://www.marktechpost.com/2026/04/23/mend-io-releases-ai-security-governance-framework-covering-asset-inventory-risk-tiering-ai-supply-chain-security-and-maturity-model/>
21. Marktechpost (2026-04-24). “Mend Releases AI Security Governance Framework: Covering Asset Inventory, Risk Tiering, AI Supply Chain Security, and Maturity Model” <https://www.marktechpost.com/2026/04/23/mend-releases-ai-security-governance-framework/>
22. Dev To Ai (2026-04-26). “My Harness Is Not a Cage. It’s an Org Chart.” https://dev.to/synthaicode_commander/my-harness-is-not-a-cage-its-an-org-chart-3gl1
23. Arxiv Cs Cy (2026-04-28). “Algorithmic Administration and the EU AI Act: Legal Principles for Public Sector Use of AI” <https://arxiv.org/abs/2604.22765>
24. Dzone Ai (2026-04-29). “Designing Intelligent AI Systems for Tax Anomaly Detection” <https://feeds.dzone.com/link/23558/17327922/designing-intelligent-ai-systems-for-tax-anomaly>
25. Aws Ml (2026-04-29). “Extracting contract insights with PwC’s AI-driven annotation on AWS” <https://aws.amazon.com/blogs/machine-learning/extracting-contract-insights-with-pwcs-ai-driven-annotation-on-aws/>
26. The Register (2026-04-30). “Govern your bots carefully or chaos could ensue” https://go.theregister.com/feed/www.theregister.com/2026/04/30/good_ai_governance_is_good/
27. The Register (2026-05-04). “Five Eyes spook shops warn rapid rollouts of agentic AI are too risky” https://go.theregister.com/feed/www.theregister.com/2026/05/04/five_eyes_agentic_ai_recommendations/
28. Nyt Technology (2026-05-04). “White House Considers Vetting A.I. Models Before They Are Released” <https://www.nytimes.com/2026/05/04/technology/trump-ai-models.html>

Curriculum Management

capability_id: curriculum-management · created: 2026-07-17 · revised: 2026-07-17 · refs: 25

Near (2026–2031)

For curriculum management broadly, the process of curriculum design is shifting toward the integration of workforce-aligned AI skillsets [12][14] and a change in educational substance toward skills-based frameworks [13] and priority-driven learning objectives [18]. In the subsegment of computer science and software engineering, the educational substance of curriculum design is shifting from basic implementation toward higher-order architectural skills and “harness engineering”—the creation of reliable AI-output environments—because AI tools automate the production of average-quality code [15][22][26][27], leading to the emergence of AI Engineering as a distinct discipline [25] and new offerings for AI ethicists [4]. For institutions focusing on Career and Technical Education (CTE) and digital content, the delivery mechanism of curriculum & resource development is shifting toward automated multimodal content generation [6], localization [3], and the use of structured frameworks like the SAMR + AI Matrix to align technological use with cognitive demand [28]. Within these technical subsegments, the process of curriculum change management is adopting transparency frameworks for AI-generated materials [5] and new governance protocols to ensure the quality, accuracy, and trust of AI-generated content [10], as well as guardrails to maintain assessment validity [16]. For institutions implementing structured faculty development programs, as demonstrated at Indiana Wesleyan University [1] and through bottom-up champion networks [8], the process of professional learning (staff) is utilizing “course refresh” institutes to update educational substance [1][8], emphasizing educator capacity over tool mastery [2] and the design of learning experiences centered on AI fluency [29], with adoption rates moderated by levels of institutional support and teacher confidence [30]. In specific institutional contexts, the process of curriculum retirement management is being accelerated through mandated reviews and budget-driven program cuts [11][23]. Professional accreditation processes are unlikely to change significantly in this window because accreditation bodies typically operate on multi-year review cycles that lag behind classroom-level revisions, although some regional teaching standards in the US are currently being updated [17].

Hypothesis: In the near term, curriculum management may shift toward integrating workforce-aligned skills and transparency frameworks in resource development, supported by a mix of structured faculty refreshes and accelerated program retirement in budget-constrained institutions.

Mid (2036–2041)

For the subset of institutions and subsegments identified in the near term, if current trajectories hold, the process of professional learning (staff) may transition from episodic “refresh” events to a permanent, cyclical requirement to mitigate technological obsolescence [1][2][8][29]. In CTE and technical contexts, delivery-mechanism efficiency gains in curriculum & resource development could increase the frequency of curriculum retirement management as specific technical skillsets obsolesce more rapidly [3][14]. For these subsegments, the educational substance of curriculum design may move further away from knowledge delivery toward outcome-based pathing, as the version of the online course centered on static content delivery becomes obsolete [21]. However, fundamental governance structures, such as faculty senate approvals and committee-based curriculum reviews, will likely persist unchanged to maintain institutional academic legitimacy.

Hypothesis: In the mid term, institutions with established development structures may formalize cyclical curriculum refreshes and accelerated resource updates for technical subsegments.

Far (2046–2051)

For the institutions and subsegments identified in the near term, the interaction between automated resource development and rapid workforce shifts may create a tension between curriculum fluidity and institutional stability. It remains unknown whether professional accreditation bodies will move toward real-time, data-driven validation or maintain periodic review cycles. The requirement for a verified human

authority to sign off on the educational substance of a curriculum will likely resist change to preserve the social legitimacy of the degree. One rival outcome is that curriculum management becomes almost entirely modular and automated for technical skills, while another is that it ossifies around a few “core” human-led fundamentals to distinguish institutional value from AI-generated “average” work [22][26].

Hypothesis A: In the far term, curriculum management may evolve into a system of highly fluid, modular updates that align in near real-time with workforce signals.

Hypothesis B (competing): In the far term, curriculum management may retreat to a small core of static, human-verified fundamentals to protect the perceived value of the degree against the proliferation of automated content.

References

1. EDUCAUSE Review (2026). “Advancing AI Literacy: A Faculty Course Refresh Institute at Indiana Wesleyan University.” <https://er.educause.edu/articles/2026/3/advancing-ai-literacy-a-faculty-course-refresh-institute-at-indiana-wesleyan-university>
2. Getting Smart (2026). “AI Literacy is Not Tool Mastery: How to Build Sustained Educator Capacity.” <https://www.gettingsmart.com/2026/03/27/ai-literacy-is-not-tool-mastery-how-to-build-sustained-educator-capacity/>
3. Elearning Industry (2026). “Leveraging AI And Automation To Streamline CTE Content Creation And Localization.” <https://elearningindustry.com/leveraging-ai-and-automation-to-streamline-cte-content-creation-and-localization>
4. Higher Ed Dive (2026). “Colleges ramp up offerings to teach students to be AI ethicists.” <https://www.highereddive.com/news/colleges-programs-ai-ethicists/816703/>
5. Elearning Industry (2026). “A 4-Step Framework For Using AI Transparently In Educational Content.” <https://elearningindustry.com/a-4-step-framework-for-using-ai-transparently-in-educational-content>
6. Arxiv Cs Cy (2026). “EduIllustrate: Towards Scalable Automated Generation Of Multimodal Educational Content.” <https://arxiv.org/abs/2604.05005>
7. EDUCAUSE Review (2026). “Building AI Initiatives with Bottom-Up Champions.” <https://er.educause.edu/articles/2026/4/building-ai-initiatives-with-bottom-up-champions>
8. Elearning Industry (2026). “AI-Generated Learning Content Governance: Ensuring Quality, Accuracy, And Trust.” <https://elearningindustry.com/ai-generated-learning-content-governance-ensuring-quality-accuracy-and-trust>
9. Higher Ed Dive (2026). “Iowa State plans to cut or merge 23 programs after mandated review.” <https://www.highereddive.com/news/iowa-state-university-program-cuts-review-regents/817262/>
10. Higher Ed Dive (2026). “AI is remaking the workforce. How can colleges ensure students thrive?” <https://www.highereddive.com/news/ai-is-remaking-the-workforce-how-can-colleges-ensure-students-thrive/817598/>
11. Elearning Industry (2026). “The Case For Skills-Based Learning In A Rapidly Changing Workplace.” <https://elearningindustry.com/the-case-for-skills-based-learning-in-a-rapidly-changing-workplace>
12. Higher Ed Dive (2026). “Employers say they struggle to find graduates with the right AI skillset.” <https://www.highereddive.com/news/employers-struggle-to-find-workers-with-the-right-ai-skillset/818017/>
13. Arxiv Cs Cy (2026). “Bringing AI into the Classroom: A Structured Approach for Integrating AI into Software Engineering Education.” <https://arxiv.org/abs/2604.16328>
14. Elearning Industry (2026). “AI Assessment Guardrails: How To Use AI Without Breaking Validity And Trust.” <https://elearningindustry.com/ai-assessment-guardrails-how-to-use-ai-without-breaking-validity-and-trust>
15. Getting Smart (2026). “America is Updating Teaching Standards Without a Clear Picture of the Future Educator.” <https://www.gettingsmart.com/2026/04/23/america-is-updating-teaching-standards-without-a-clear-picture-of-the-future-educator/>

16. EDUCAUSE Review (2026). “Teaching and Learning with AI: Three Priorities.” <https://er.educause.edu/podcasts/educause-exchange/teaching-and-learning-with-ai-three-priorities>
17. Elearning Industry (2026). “AI Is Not Killing The Online Course Industry. It’s Killing The Old Version Of It.” <https://elearningindustry.com/ai-is-not-killing-the-online-course-industry-its-killing-the-old-version-of-it>
18. Dev To Ai (2026). “AI Didn’t Make Developers Better. It Made Average Work Easier to Produce.” <https://dev.to/workelate123/ai-didnt-make-developers-better-it-made-average-work-easier-to-produce-hc2>
19. Higher Ed Dive (2026). “East Carolina University plans to cut 44 academic programs.” <https://www.highereddive.com/news/east-carolina-university-cut-around-44-academic-programs/818622/>
20. Dev To Ai (2026). “The Case for AI Engineering as a Distinct Discipline.” <https://dev.to/jasrandhawa/the-case-for-ai-engineering-as-a-distinct-discipline-44pm>
21. Dev To Ai (2026). “Fundamentals Matter More Than Ever Before.” <https://dev.to/nachiappannk/fundamentals-matter-more-than-ever-before-266o>
22. Dev To Ai (2026). “Harness Engineering Is the New Senior Developer Skill (Here’s Why).” <https://dev.to/shipwithaiio/harness-engineering-is-the-new-senior-developer-skill-heres-why-4hef>
23. EDUCAUSE Review (2026). “A Practical SAMR + AI Framework for Instructional Design.” <https://er.educause.edu/articles/2026/4/a-practical-samr--ai-framework-for-instructional-design>
24. Dev To Ai (2026). “Teaching AI Fluency Anthropic.” <https://dev.to/arsdevel/mrj-kwrs-teaching-ai-fluency-mn-anthropic-50i4>
25. Arxiv Cs Cy (2026). “AI Adoption Among Teachers: Insights on Concerns, Support, Confidence, and Attitudes.” <https://arxiv.org/abs/2605.00343>

Facilities & Property Management

capability_id: facilities-property-management · created: 2026-07-17 · revised: 2026-07-17 · refs: 9

Near (2026–2031)

For a subset of institutions adopting physical AI and for those in US regions subject to emerging land-use policies, the delivery mechanism for property maintenance, campus security, and estate development may shift [1][3][4]. In property maintenance, the adoption of robot mowers may replace specific manual labor tasks in groundskeeping [1], while AI-driven transcription for technician voice notes may streamline documentation workflows for facility repairs by automating the conversion of raw audio into service summaries [2]. Sensing capabilities in robotic systems, such as the ability to read gauges and thermometers, may automate the inspection of industrial plant equipment [5]. For campus security and space utilisation, the deployment of physical AI may alter perimeter monitoring [3], and AI video analytics may transform traditional CCTV into intelligent monitoring systems [8]. However, the reliance on facial recognition for security may be constrained by emerging evidence that such technology is less effective than claimed and prone to false identifications [10]. Predictive intelligence in large venues may change how facility conditions are managed through the forecasting of behavioral patterns [7]. In estate development and management, AI platforms may accelerate the material takeoff process by quantifying materials for project budgeting [6], though state and local policy interventions in the US may constrain the siting and expansion of AI-specific data center facilities [4]. Processes such as mail management and commercial tenancy are unlikely to change yet because they remain dependent on manual administrative protocols and established contractual law.

Hypothesis: In the near term, a subset of institutions may reduce operational labor in maintenance and security through targeted automation while facing increased regional regulatory constraints on the development of AI-specific physical infrastructure and biometric surveillance.

Mid (2036–2041)

For institutions implementing these systems, if current trajectories in predictive intelligence and robotics hold, the scope of automated surveillance and maintenance may expand while estate development for AI facilities remains contingent on local policy compliance [3][4][8]. The process of property maintenance may shift toward continuous robotic auditing of facility health as sensing capabilities for industrial equipment mature [5]. In estate development, the use of AI for project estimation may move from material takeoffs to more integrated budgetary forecasting [6], while the development of data centers may require more complex negotiations with local municipalities regarding resource extraction and energy usage [4]. Robotics in groundskeeping may move beyond mowing to encompass a broader range of repetitive exterior vegetation management [1]. Long-term campus planning and architectural synthesis are likely to remain human-led because these processes require alignment with institutional strategy and aesthetic values.

Hypothesis: In the mid term, routine exterior maintenance, industrial plant inspection, and perimeter security may become predominantly automated for early adopters, while the development of AI-centric facilities faces heightened regional regulatory barriers.

Far (2046–2051)

For the subsegments of exterior maintenance, industrial inspection, and estate development, outcomes diverge based on the tension between automation efficiency and environmental or social policy [1][4][5]. One scenario involves a transition to a highly automated facility management regime where human presence is limited to system oversight and both the physical perimeter and internal industrial plants are managed by integrated AI agents [3][5][8]. A competing scenario involves a shift toward restricted automation and human-centric auditing, driven by long-term regulatory failures in biometric accuracy and stringent local land-use mandates [4][10]. It remains unknown whether the energy requirements of physical AI agents will eventually conflict with institutional environmental management goals. Fundamental property ownership laws and the legal framework for commercial tenancy are likely to resist change even at this horizon.

Hypothesis A: In the far term, a subset of institutions may move toward a largely autonomous facility management model where physical infrastructure is maintained and secured by integrated robotic systems.

Hypothesis B (competing): In the far term, the deployment of physical AI may plateau or retract due to sustained regulatory oversight of biometrics and restrictive local policies on AI infrastructure.

References

1. Wired (2026-04-06). “Robot Mowers Are Actually Good Now.” <https://www.wired.com/story/robot-mowers-are-good-now-rave/>
2. Dev To Ai (2026-04-06). “From Mumbles to Memos: Teaching AI to Understand Technician Voice Notes.” https://dev.to/ken_deng_ai/from-mumbles-to-memos-teaching-ai-to-understand-technician-voice-notes-ckj
3. Ai News (2026-04-07). “Asylon and Thrive Logic bring physical AI to enterprise perimeter security.” <https://www.artificialintelligence-news.com/news/physical-ai-security-at-the-enterprise-perimeter-takes-a-step-closer/>
4. Ai Now Institute (2026-04-01). “North Star Data Center Policy Toolkit: State and Local Policy Interventions to Stop Rampant AI Data Center Expansion.” <https://ainowinstitute.org/publications/data-center-policy-guide>
5. Ars Technica (2026-04-15). “Boston Dynamics’ robot dog now reads gauges and thermometers with Google’s AI.” <https://arstechnica.com/ai/2026/04/robot-dogs-now-read-gauges-and-thermometers-using-google-gemini/>
6. Ai News (2026-04-20). “Bobyard 2.0 offers improved takeoffs and unified AI for estimators.” <https://www.artificialintelligence-news.com/news/bobyard-2-0-offers-improved-takeoffs-and-unified-ai-for-estimators/>
7. Dev To Ai (2026-04-22). “Architecting Predictive Intelligence for Smart Venues.” https://dev.to/tanisha_khiara_4b390b7946/architecting-predictive-intelligence-for-smart-venues-1lkj
8. Dev To Ai (2026-04-23). “Integrating AI Video Analytics into Venue Operations.” https://dev.to/tanisha_khiara_4b390b7946/integrating-ai-video-analytics-into-venue-operations-2m09
9. Dev To Ai (2026-04-25). “From Mumbles to Memos: Teaching AI to Understand Technician Voice Notes.” https://dev.to/ken_deng_ai/from-mumbles-to-memos-teaching-ai-to-understand-technician-voice-notes-66g
10. Guardian Tech (2026-05-03). “AI facial recognition oversight lagging far behind technology, watchdogs warn.” <https://www.theguardian.com/technology/ng-interactive/2026/may/03/ai-facial-recognition-oversight-lagging-far-behind-technology-watchdogs-warn>

Finance Management

capability_id: finance-management · created: 2026-06-25 · revised: 2026-06-25 · refs: 14

Near (2026–2031)

For Finance Management broadly, the adoption of agentic AI workflows introduces significant volatility into budget management due to “hidden multipliers” in API consumption and recursive execution loops that can lead to rapid, unplanned expenditure [8][14]. To mitigate these risks, institutions may implement judgment layers to verify agent-initiated financial actions before execution [7]. For a subset of institutions utilizing usage-based AI software, budget management processes are shifting from the prediction of fixed-rate licensing costs toward the tracking of variable expenditures based on per-token usage [10][13]. Within this subsegment, procurement and purchasing management may integrate AI-powered invoice processing pipelines to automate data extraction [6] and use AI-driven annotation to accelerate the analysis of unstructured contract agreements [11]. Expenditure management may incorporate programmatic guardrails, such as hard-cap quotas [1] and tiered security policies [3], to prevent unbounded agent spending. For institutions granting limited financial autonomy to agents, expenditure management may adopt delegated payment mechanisms including virtual credit cards [5], agent-to-agent (A2A) transaction protocols [4], and digital wallets with secure approval flows [12]. These changes represent a shift in the delivery mechanism of financial oversight—automating the tracking and bounding of costs—rather than a change in the educational substance of the capability. Core budget allocation for institutional overhead is unlikely to change yet, as these processes remain tethered to annual fiscal cycles and board-approved appropriations.

Hypothesis: In the near term, institutions may implement programmatic guardrails and verification layers to manage the volatility of usage-based AI costs and autonomous agent expenditures.

Mid (2036–2041)

For the subset of institutions deploying autonomous financial agents, if current trajectories of granular cost attribution hold, budget management may transition from periodic reviews toward automated, trigger-based expenditure management [1][2]. The mechanism of financial oversight may shift toward programmatic guardrails utilizing structured security tiers to bound autonomous spending [3]. Procurement and purchasing management may see an increase in high-frequency A2A transactions, shifting the human role from transaction execution to the definition of negotiation parameters and spending limits [4][5][9]. Financial reporting processes may incorporate higher-velocity data streams to respond to rapid fluctuations in token-based costs [8][10][13]. However, overall institutional treasury management is likely to persist unchanged as it continues to be governed by long-term endowment strategies and multi-year capital constraints.

Hypothesis: In the mid term, budget and procurement management for AI-enabled services may evolve into high-frequency, automated oversight processes characterized by dynamic quotas and A2A transaction protocols.

Far (2046–2051)

For the same subset of institutions, the integration of autonomous financial controls may lead to divergent outcomes. One plausible outcome is the transition of expenditure management to a model of bounded autonomy, where AI agents manage their own micro-budgets and execute procurement via A2A protocols and integrated digital wallets within strictly defined institutional security policies [3][4][5][12]. A rival outcome is a systemic shift back toward rigid, fixed-cost procurement models to eliminate the financial volatility and billing anomalies inherent in usage-based AI pricing [1][8][10]. It remains unknown whether institutional governance will accept the delegation of fiduciary authority to autonomous agents for high-value acquisitions. Legal requirements for institutional financial accountability and external auditing are likely to resist full automation even at this horizon.

Hypothesis A: In the far term, expenditure management may shift toward a model of bounded autonomy

where agents manage micro-budgets via A2A protocols.

Hypothesis B (competing): In the far term, institutions may return to rigid, fixed-cost procurement models to eliminate the volatility of usage-based AI billing.

References

1. Dev To Ai (2026). “My agent burned \$200 in one night. So I built something that stops it.” https://dev.to/tal_nagar_f7aead23e3e291a/my-agent-burned-200-in-one-night-so-i-built-something-that-stops-it-1035
2. Aws Ml (2026). “Manage AI costs with Amazon Bedrock Projects” <https://aws.amazon.com/blogs/machine-learning/manage-ai-costs-with-amazon-bedrock-projects/>
3. Dev To Ai (2026). “21 Policy Types, 4 Security Tiers: Building Bulletproof AI Agent Guardrails” <https://dev.to/walletguy/21-policy-types-4-security-tiers-building-bulletproof-ai-agent-guardrails-31g6>
4. Dev To Ai (2026). “Building the Agent Economy: A2A Protocol Explained” <https://dev.to/mranderson323/building-the-agent-economy-a2a-protocol-explained-4nfd>
5. Dev To Ai (2026). “AI Agents Now Have Credit Cards” <https://dev.to/tomwangcn/ai-agents-now-have-credit-cards-gn3>
6. Hackernoon Ai (2026). “Building an AI-Powered Invoice Processing Pipeline” <https://hackernoon.com/building-an-ai-powered-invoice-processing-pipeline?source=rss>
7. Dev To Ai (2026). “Judgment Layer for Financial AI Agents” <https://dev.to/thatechmaestro/judgment-layer-for-financial-ai-agents-141>
8. Hackernoon Ai (2026). “The Midnight Billing Anomaly: How AI Agents Quietly Destroy Unit Economics” <https://hackernoon.com/the-midnight-billing-anomaly-how-ai-agents-quietly-destroy-unit-economics?source=rss>
9. Dev To Ai (2026). “Anthropic Just Proved Agents Can Do Commerce. Here’s What They Didn’t Build.” https://dev.to/aaron_schnieder_4563d5d33/anthropic-just-proved-agents-can-do-commerce-heres-what-they-didnt-build-5209
10. Ars Technica (2026). “GitHub will start charging Copilot users based on their actual AI usage” <https://arstechnica.com/ai/2026/04/github-will-start-charging-copilot-users-based-on-their-actual-ai-usage/>
11. Aws Ml (2026). “Extracting contract insights with PwC’s AI-driven annotation on AWS” <https://aws.amazon.com/blogs/machine-learning/extracting-contract-insights-with-pwcs-ai-driven-annotation-on-aws/>
12. TechCrunch (2026). “Stripe introduces Link, a digital wallet that autonomous AI agents can use, too” <https://techcrunch.com/2026/04/30/stripe-link-digital-wallet-ai-agents-shopping/>
13. Ai News (2026). “Per-token AI charges come to GitHub Copilot” <https://www.artificialintelligence-news.com/news/per-token-ai-charging-comes-to-github-copilot/>
14. Dev To Ai (2026). “Your AI Agent Is Sending 10x More API Calls Than You Think — Here’s Where the Cost Hides” <https://dev.to/xidao/your-ai-agent-is-sending-10x-more-api-calls-than-you-think-heres-where-the-cost-hides-4mei>

Government, Public & Stakeholder Relationships

capability_id: government-public-stakeholder-relationships · created: 2026-04-29 · revised: 2026-04-29
· refs: 1

Near (2026–2031)

For institutions facing direct state-level legal inquiries related to AI-enabled harms, such as the current investigation involving Florida State University, government relationship management may shift from routine administrative coordination to adversarial legal defense [1]. This shift represents a change in the educational substance of the capability, as the relationship is now defined by criminal liability rather than policy alignment. Media relations management is simultaneously forced to address the alleged role of generative AI in physical campus violence, requiring a transition from promotional messaging to crisis communication [1]. Because the current evidence is limited to a single criminal probe, it is unlikely that industry relationship management—specifically partnerships with AI vendors—will change across the sector, as institutions remain dependent on these tools for operational efficiency.

Hypothesis: In the near term, a subset of institutions may see government relationship management shift toward risk mitigation and legal defense regarding AI vendor liabilities.

Mid (2036–2041)

For state-funded institutions, if current trajectories of state-led criminal investigations into AI outputs hold, the process of government relationship management may integrate more rigorous legal auditing of AI vendor contracts. Extrapolating from current legal scrutiny [1], institutions may seek to shift the liability for AI-generated externalities entirely onto the vendor to preserve their standing with state regulators. This would alter the mechanism of industry relationship management, moving it toward a procurement model centered on indemnity clauses. However, internal stakeholder relationship management is unlikely to be fundamentally altered by these external pressures unless tenure and employment contracts are formally rewritten to include AI-specific liability.

Hypothesis: In the mid term, state-funded institutions may formalize liability frameworks in their government relationship management to insulate the institution from criminal probes into vendor-provided AI.

Far (2046–2051)

For institutions operating under high state oversight, the long-term state of the capability depends on whether legal precedents establish “institutional negligence” for AI-generated content. In one scenario, government relationship management becomes a highly regulated compliance function where state agencies mandate specific, audited AI architectures to prevent public harm [1]. In a rival scenario, the emergence of comprehensive federal or state immunity for educational institutions regarding third-party software may return these relationships to a non-adversarial, administrative state. The impact on public relationship management remains an unresolvable uncertainty, as it depends on the evolving social threshold for accepting AI-driven externalities. Routine media relations for non-crisis events will likely resist change, as the basic need for institutional branding remains constant.

Hypothesis A: In the far term, government relationship management for a subset of institutions may become a highly regulated compliance function focused on AI safety audits.

Hypothesis B (competing): In the far term, the legal resolution of AI liability may decouple institutional responsibility from vendor outputs, returning government relationship management to administrative norms.

References

1. NPR (2026-04-21). “Florida AG launches criminal investigation into ChatGPT over FSU shooting.”
<https://www.npr.org/2026/04/21/nx-s1-5793967/florida-openai-investigation-mass-shooting-fsu>

Human Resource Management

capability_id: human-resource-management · created: 2026-07-20 · revised: 2026-07-20 · refs: 19

Near (2026–2031)

For human resource management broadly, AI is being integrated into the delivery mechanisms of employee support and staff recruitment, while training and development shifts are concentrated in specific institutional subsegments. In the process of employee support, institutions are deploying AI-powered onboarding agents to automate administrative orientation [3], shifting the educational substance of onboarding from a discrete initial event to a “perpetual” continuous model [12] supported by targeted learning approaches [8]. For institutions utilizing LLM-based screening in staff recruitment, the delivery of candidate selection is affected by a self-preference mechanism where LLMs consistently favor AI-generated resumes over human-authored ones [21][22]. This mechanism introduces risks of data bias [1] and “LLM nepotism” in organizational governance [9], prompting the development of community-driven evaluation frameworks to audit these systems [11]. In the United Kingdom, the delivery of staff recruitment is experiencing friction as job hunters report the AI-driven interview process as “humiliating” [19][20]. For institutions conducting remote hiring, identity verification is shifting toward multi-landmark facial analysis to mitigate the risk of synthetic media [2]. In training and development, a subset of institutions is moving away from top-down training toward models where internal peer credibility drives AI adoption through job-specific use cases [26]. Workforce relations management is unlikely to change in this horizon because the negotiation of collective bargaining agreements remains dependent on human-led institutional diplomacy and legal frameworks.

Hypothesis: In the near term, the capability will see a shift toward automated, continuous onboarding and a substance-level update to workforce readiness, while recruitment processes face a tension between algorithmic efficiency and the social legitimacy of the candidate experience.

Mid (2036–2041)

For human resource management broadly, if current trajectories in productivity-linked labor trials hold, the substance of remuneration and entitlements management may shift as some institutions decouple pay from fixed hourly requirements, such as through four-day week models [5]. In the subsegment of institutions utilizing agentic AI in human capital management, the delivery of talent management and workforce reporting may expand to the end-to-end orchestration of employee development pipelines [10][16]. For the subsegment of institutions updating performance metrics, the substance of employee performance management may shift as traditional KPIs are replaced by metrics specifically aligned with AI-driven productivity [15]. In institutions relying on remote hiring, biometric identity verification may transition from an experimental tool to a standard delivery mechanism to counter increasingly sophisticated synthetic media [2]. For a subset of institutions, the process of workforce planning may shift toward “disintegrating” the traditional organizational chart in favor of more fluid, AI-supported structures [4]. Tenure-track hierarchies and faculty governance structures are likely to persist because they are deeply embedded in academic contracts and the social legitimacy of the professoriate.

Hypothesis: In the mid term, AI in HR may transition from isolated toolsets to agentic orchestration of the employee lifecycle, potentially altering the substance of performance evaluation and remuneration.

Far (2046–2051)

For human resource management broadly, the capability may diverge based on the resolution of the tension between algorithmic efficiency and labor protections. In regions where strict labor regulations are enacted, as seen in initial Chinese mandates prohibiting the firing of humans if AI takes their jobs [24], the process of workforce relations management may shift toward the legal administration of “AI doubles” or synthetic colleagues [13]. A primary uncertainty remains whether the “LLM nepotism” observed in early screening [9] will lead to a systemic collapse of meritocratic hiring or be corrected by standard auditing frameworks [11]. The core institutional function of staff records and details management is likely to remain a stable,

low-variance process regardless of the AI interface.

Hypothesis A: In the far term, human resource management shifts toward the governance of a hybrid workforce of humans and synthetic agents, heavily constrained by regional labor laws and protective mandates.

Hypothesis B (competing): In the far term, the capability reverts to human-centric oversight in high-stakes processes due to the counterproductivity of algorithmic bias and the loss of social legitimacy in AI-led governance.

References

1. Ai Trends (2021-10-28). “Promise and Perils of Using AI for Hiring: Guard Against Data Bias.” <https://www.aitrends.com/ai-world-government/promise-and-perils-of-using-ai-for-hiring-guard-against-data-bias/>
2. Dev To Ai (2026-04-06). “Why a Deepfake Face Can Fool Your Eyes in Seconds but Not 128 Landmarks at Once.” <https://dev.to/caracomp/why-a-deepfake-face-can-fool-your-eyes-in-seconds-but-not-128-landmarks-at-once-37m8>
3. Aws Ml (2026-04-06). “Build AI-powered employee onboarding agents with Amazon Quick.” <https://aws.amazon.com/blogs/machine-learning/build-ai-powered-employee-onboarding-agents-with-amazon-quick/>
4. Mit Sloan (2026-04-07). “Disintegrating the Org Chart: ServiceNow’s Jacqui Canney.” <https://sloanreview.mit.edu/audio/disintegrating-the-org-chart-servicenows-jacqui-canney/>
5. BBC News (2026-04-07). “OpenAI encourages firms to trial four-day weeks to adapt to AI era.” https://www.bbc.com/news/articles/c8x71ejrp92o?at_medium=RSS&at_campaign=rss
6. Dev To Ai (2026-04-10). “Smarter Employee Onboarding Starts with the Right Learning Approach.” https://dev.to/abhilash_rai_ee8bde928249/smarter-employee-onboarding-starts-with-the-right-learning-approach-5fhn
7. Arxiv Cs Cy (2026-04-14). “LLM Nepotism in Organizational Governance.” <https://arxiv.org/abs/2604.09620>
8. Ai News (2026-04-14). “SAP brings agentic AI to human capital management.” <https://www.artificialintelligence-news.com/news/sap-brings-agentic-ai-human-capital-management/>
9. Arxiv Cs Cl (2026-04-16). “WorkRB: A Community-Driven Evaluation Framework for AI in the Work Domain.” <https://arxiv.org/abs/2604.13055>
10. Elearning Industry (2026-04-21). “The Art Of Perpetual Onboarding: Why AI Makes The Continuous Model Viable At Last.” <https://elearningindustry.com/the-art-of-perpetual-onboarding-why-ai-makes-the-continuous-model-viable-at-last>
11. MIT Technology Review (2026-04-20). “Chinese tech workers are starting to train their AI doubles—and pushing back.” <https://www.technologyreview.com/2026/04/20/1136149/chinese-tech-workers-ai-colleagues/>
12. Fast Company (2026-04-23). “The Age of AI means we need to throw out our old KPIs and replace them with new ones.” https://www.fastcompany.com/91521422/ai-kpis-new-metrics?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
13. Aws Ml (2026-04-24). “Building Workforce AI Agents with Visier and Amazon Quick.” <https://aws.amazon.com/blogs/machine-learning/building-workforce-ai-agents-with-visier-and-amazon-quick/>
14. Guardian Tech (2026-05-01). “‘Completely horrible’: UK job hunters share frustration with AI interviews.” <https://www.theguardian.com/technology/2026/may/01/uk-job-hunters-frustration-ai-interviews>
15. Guardian Tech (2026-05-01). “‘Awkward and humiliating’: UK job hunters share frustration with AI interviews.” <https://www.theguardian.com/technology/2026/may/01/uk-job-hunters-frustration-ai-interviews>
16. Hacker News (2026-05-02). “LLMs consistently pick resumes they generate over ones by humans or other models.” <https://arxiv.org/abs/2509.00462>
17. Hacker News (2026-05-02). “AI Self-preferencing in Algorithmic Hiring: Empirical Evidence and

- Insights.” <https://arxiv.org/abs/2509.00462>
18. The Register (2026-05-04). “Just in time for Labour Day, China makes it illegal to fire humans if AI takes their jobs.” https://go.theregister.com/feed/www.theregister.com/2026/05/04/asia_tech_news_roundup/
 19. Elearning Industry (2026-05-04). “Internal Learning Champions: Why Peer Credibility Beats Top-Down Training In AI Adoption.” <https://elearningindustry.com/internal-learning-champions-why-peer-credibility-beats-top-down-training-in-ai-adoption>

Information & Communication Technology Mgt

capability_id: information-communication-technology-management · created: 2026-05-08 · revised: 2026-05-08 · refs: 43

Near (2026–2031)

For Information & Communication Technology Mgt broadly, application lifecycle management shifts toward the adoption of the Model Context Protocol (MCP) to standardize the interface between large language models and institutional data [90][92][100]. This change in delivery mechanism replaces proprietary connectors with a universal tool-calling interface [49][101], but it introduces vulnerabilities such as command injection and information disclosure that require specific auditing of MCP servers [12][130] and vetting for security grades [10]. The educational substance of quality assurance shifts from deterministic testing to frameworks that validate non-deterministic AI outputs and red-team the agent stack [62][78][91], as traditional chaos engineering proves insufficient for AI systems [142]. For ict service support, the triage of support tickets and SRE incident response are increasingly mediated by autonomous agents [11][109], supported by LLM-augmented knowledge bases for root cause analysis [58]. However, a gap in observability has emerged where standard orchestration tools like Kubernetes may report healthy status while AI pipelines fail silently, necessitating new AI-specific monitoring protocols [153]. For infrastructure lifecycle management, the emergence of AI-driven vulnerability discovery—demonstrated by the identification of hundreds of zero-day flaws in major software by autonomous models [124][129][131]—forces a shift from scheduled patching to event-driven, high-frequency update cycles [135]. Identity & access management incorporates agent registries to manage the sprawl of autonomous agents [71][73], while the shift toward multi-agent systems necessitates the adoption of zero-trust security runtimes [41][87][149], sandbox execution [113], and cryptographic authorization to prevent identity spoofing [111][150]. For a subset of institutions deploying autonomous agents, application lifecycle management now requires new financial guardrails to prevent runaway automated billing and “midnight billing anomalies” that destroy unit economics [154]. Physical network cabling and hardware rack installations are unlikely to change yet, as disruption is concentrated at the logical orchestration and software layers.

Hypothesis: In the near term, the capability will pivot from managing static software installations to governing a dynamic ecosystem of standardized AI connectors and managing a high-frequency vulnerability response cycle driven by AI-automated exploit discovery.

Mid (2036–2041)

For Information & Communication Technology Mgt broadly, if current trajectories hold, application lifecycle management may evolve toward model-agnostic architectures [7] to mitigate operational risks associated with “tool refugee” patterns, where institutions must migrate when proprietary AI tools are deprecated [44]. The substance of application lifecycle management may further shift as AI agents automate the modernization and migration of legacy codebases [20][54]. For institutions managing decentralized workloads, infrastructure lifecycle management may shift toward strengthened governance for edge AI workloads [106] to manage the proliferation of autonomous processing at the network periphery [67][118], potentially utilizing KubeVirt to run virtual machines on Kubernetes at the edge [68] and reorganizing hardware specifically for agentic AI platforms [75]. In regions governed by the EU AI Act, the management of technical files for high-risk AI systems becomes a core regulatory compliance process within this capability [27][66]. Institutional procurement cycles for physical hardware are unlikely to change significantly, as these processes remain embedded in multi-year budget cycles and facility contracts.

Hypothesis: In the mid term, the capability may shift toward the management of autonomous infrastructure that self-optimizes for edge workloads while incorporating strict regulatory auditing of AI technical files.

Far (2046–2051)

For Information & Communication Technology Mgt broadly, the capability may reach a point where the distinction between application and infrastructure lifecycle management collapses into a single autonomous

orchestration process. One rival outcome is that the acceleration of AI-driven vulnerability discovery leads to a “security paradox,” where institutions abandon open connectivity in favor of highly fragmented, air-gapped, or sovereign infrastructure to ensure survival [28][81]. It remains unknown whether standardized protocols like MCP will persist or be replaced by emergent, self-evolving agent communication languages. Hardware replacement cycles for core servers will likely resist rapid change even at this horizon due to the physical constraints of energy delivery and data center cooling [37].

Hypothesis A: In the far term, the capability becomes a supervisory function over a fully autonomous, self-healing ICT stack that manages its own deployment, patching, and hardware optimization.

Hypothesis B (competing): In the far term, the capability shifts toward managing highly restricted, sovereign silos of technology to defend against autonomous, large-scale AI cyber-attacks.

References

1. Edtech Magazine (2026-04-03). “Case Western Reserve University Deploys Google Gemini.” <https://edtechmagazine.com/higher/article/2026/04/case-western-reserve-university-deploys-google-gemini>
2. Dev To Ai (2026-04-06). “Building Model-Agnostic AI Architecture: The Pattern That Future-Proofs Your System.” https://dev.to/ailoitte_sk/building-model-agnostic-ai-architecture-the-pattern-that-future-proofs-your-system-4o84
3. Dev To Ai (2026-04-08). “Why we’re building the AI Tool Refugee Center: a place to land when your tool dies.” <https://dev.to/zsky/why-were-building-the-ai-tool-refugee-center-a-place-to-land-when-your-tool-dies-1e8>
4. Dev To Ai (2026-04-07). “The ‘Legacy Code’ Nightmare is Over: How AI Agents are Automating App Modernization.” https://dev.to/siddhesh_surve/the-legacy-code-nightmare-is-over-how-ai-agents-are-automating-app-modernization-33j8
5. Dzone Ai (2026-04-08). “AI-Assisted Code Migration: Practical Techniques for Modernizing Legacy Systems.” <https://feeds.dzone.com/link/23558/17316144/ai-assisted-code-migration-practical-techniques>
6. Ai News (2026-04-13). “Strengthening enterprise governance for rising edge AI workloads.” <https://www.artificialintelligence-news.com/news/strengthening-enterprise-governance-for-rising-edge-ai-workloads/>
7. Dev To Ai (2026-04-15). “AI Is Moving Frontend Logic to the Browser and Edge — Are You Ready?” https://dev.to/rohith_kn/ai-is-moving-frontend-logic-to-the-browser-and-edge-are-you-ready-7ic
8. Ai Business (2026-04-21). “AMD’s Vision for AI PCs in the Age of Agentic AI.” <https://aibusiness.com/generative-ai/amd-s-vision-for-ai-pcs-age-of-agentic-ai>
9. The Register (2026-04-09). “Nutanix to add KubeVirt support to run VMs on K8s at the edge.” https://go.theregister.com/feed/www.theregister.com/2026/04/09/nutanix_kubevirt_arm_supoprnt/
10. Dzone Ai (2026-04-09). “How Agentic AI Platforms Organize Their Hardware Infrastructure.” <https://dzone.com/link/23558/17316612/how-agentic-ai-platforms-organize-hardware-infra>
11. Dev To Ai (2026-04-07). “What’s actually inside an EU AI Act Annex IV technical file (and what engineers usually miss).” https://dev.to/rostyslav_myronenko_faf2f/whats-actually-inside-an-eu-ai-act-annex-iv-technical-file-and-what-engineers-usually-miss-4nh4
12. Ai News (2026-04-09). “Agentic AI’s governance challenges under the EU AI Act in 2026.” <https://www.artificialintelligence-news.com/news/agentic-ais-governance-challenges-under-the-eu-ai-act-in-2026/>
13. The Register (2026-04-07). “OpenInfra General Manager talks sovereignty, governments deploying tech ‘kill switches.’” https://go.theregister.com/feed/www.theregister.com/2026/04/07/thierry_carrez_on_sovereignty_interview/
14. Ai Business (2026-04-10). “Anthropic’s Glasswing Highlights AI’s Security Paradox.” <https://aibusiness.com/generative-ai/anthropic-s-glasswing-highlights-ai-s-security-paradox>
15. Ai Now Institute (2026-04-01). “North Star Data Center Policy Toolkit: State and Local Policy

- Interventions to Stop Rampant AI Data Center Expansion.” <https://ainowinstitute.org/publications/data-center-policy-guide>
16. Dev To Ai (2026-04-11). “MCP Ecosystem in 2026: \$1.8B Market and 1000+ Connectors Reshaping AI Integration.” <https://dev.to/chunxiaox/mcp-ecosystem-in-2026-18b-market-and-1000-connectors-reshaping-ai-integration-4n3l>
 17. Dev To Ai (2026-04-11). “The Model Context Protocol in 2025: From Anthropic to Linux Foundation.” <https://dev.to/chunxiaox/the-model-context-protocol-in-2025-from-anthropic-to-linux-foundation-2og1>
 18. Dev To Ai (2026-04-11). “The Future of AI Integration: Model Context Protocol (MCP) Opportunities.” <https://dev.to/chunxiaox/the-future-of-ai-integration-model-context-protocol-mcp-opportunities-148b>
 19. Dzone Ai (2026-04-08). “MCP + AWS AgentCore: Give Your AI Agent Real Tools in 60 Minutes.” <https://feeds.dzone.com/link/23558/17315883/mcp-aws-agentcore>
 20. Dev To Ai (2026-04-11). “The State of AI Agent MCP Connector Frameworks in 2026.” <https://dev.to/chunxiaox/the-state-of-ai-agent-mcp-connector-frameworks-in-2026-2ea7>
 21. Dev To Ai (2026-04-06). “CVE-2026-32211 | Azure MCP Server Information Disclosure Vulnerability.” https://dev.to/aakash_rahsi/cve-2026-32211-azure-mcp-server-information-disclosure-vulnerability-2280
 22. Dev To Ai (2026-04-22). “How to Check Your MCP Server for CVE-2026-5603’s Vulnerability Pattern (And Why shellQuote Isn’t Enough).” <https://dev.to/armorlai/how-to-check-your-mcp-server-for-cve-2026-5603s-vulnerability-pattern-and-why-shellquote-isnt-gdn>
 23. Dev To Ai (2026-04-06). “Only 20% of MCP Servers Are ‘A-Grade’ Secure — Here’s How to Vet Them Before Installing.” https://dev.to/gentic_news/only-20-of-mcp-servers-are-a-grade-secure-heres-how-to-vet-them-before-installing-jck
 24. Dev To Ai (2026-04-09). “OpenAI’s Promptfoo deal puts evaluation and red-teaming at the centre of the agent stack.” https://dev.to/yujia_zhang_0328/openais-promptfoo-deal-puts-evaluation-and-red-teaming-at-the-centre-of-the-agent-stack-3o19
 25. Dev To Ai (2026-04-10). “AI Testing and Quality Assurance in 2026: Ensuring AI System Reliability.” <https://dev.to/lufumeiying/ai-testing-and-quality-assurance-in-2026-ensuring-ai-system-reliability-8ka>
 26. Dev To Ai (2026-04-11). “OpenAI Codex Testing: How to QA AI-Written Code.” https://dev.to/hai_huang_f196ed9669351e0/openai-codex-testing-how-to-qa-ai-written-code-31n5
 27. Dzone Ai (2026-04-22). “What AI Systems Taught Us About the Limits of Chaos Engineering.” <https://feeds.dzone.com/link/23558/17323679/chaos-engineering-limits>
 28. Dev To Ai (2026-04-06). “How to Automate Ticket Triage with AI (Without Breaking Everything).” https://dev.to/erik_anderson_c41dbafd423/how-to-automate-ticket-triage-with-ai-without-breaking-everything-4635
 29. Dzone Ai (2026-04-14). “Building an AI-Powered SRE Incident Response Workflow With AWS Strands Agents.” <https://feeds.dzone.com/link/23558/17318982/ai-powered-sre-incident-response-aws-strands-agents>
 30. Arxiv Cs Cl (2026-04-09). “LLM-Augmented Knowledge Base Construction For Root Cause Analysis.” <https://arxiv.org/abs/2604.06171>
 31. Hackernoon Ai (2026-04-23). “I Watched Our AI Pipeline Silently Fail While Kubernetes Said Everything Was Fine.” <https://hackernoon.com/i-watched-our-ai-pipeline-silently-fail-while-kubernetes-said-everything-was-fine?source=rss>
 32. Wired (2026-04-21). “Mozilla Used Anthropic’s Mythos to Find and Fix 271 Bugs in Firefox.” <https://www.wired.com/story/mozilla-used-anthropics-mythos-to-find-271-bugs-in-firefox>
 33. Ars Technica (2026-04-21). “Mozilla: Anthropic’s Mythos found 271 security vulnerabilities in Firefox 150.” <https://arstechnica.com/ai/2026/04/mozilla-anthropics-mythos-found-271-zero-day-vulnerabilities-in-firefox-150/>

34. The Register (2026-04-22). “Mythos found 271 Firefox flaws – but none a human couldn’t spot.” https://go.theregister.com/feed/www.register.com/2026/04/22/mozilla_firefox_mythos_future_defenders/
35. Fast Company (2026-04-22). “Brace yourself for a flood of patches in all of your tech gadgets.” https://www.fastcompany.com/91530191/brace-yourself-for-a-flood-of-patches-in-all-of-your-tech-gadgets?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
36. The Register (2026-04-09). “AWS: Agents shouldn’t be secret, so we built a registry for them.” https://go.theregister.com/feed/www.theregister.com/2026/04/09/aws_ai_agent_registry/
37. Aws Ml (2026-04-09). “The future of managing agents at scale: AWS Agent Registry now in preview.” <https://aws.amazon.com/blogs/machine-learning/the-future-of-managing-agents-at-scale-aws-agent-registry-now-in-preview/>
38. Dev To Ai (2026-04-08). “I built an open-source zero-trust security runtime for AI agents. Here’s what I learned.” https://dev.to/edoardo_bambini_5badce6ff/i-built-an-open-source-zero-trust-security-runtime-for-ai-agents-heres-what-i-learned-538e
39. Dev To Ai (2026-04-10). “Cert-gating every tool call: zero-trust for AI agents.” <https://dev.to/1r0nw1ll/cert-gating-every-tool-call-zero-trust-for-ai-agents-8pg>
40. Dev To Ai (2026-04-22). “API keys were designed for humans. AI agents break them in 4 ways.” <https://dev.to/allforscience/api-keys-were-designed-for-humans-ai-agents-break-them-in-4-ways-422>
41. Ai News (2026-04-16). “OpenAI Agents SDK improves governance with sandbox execution.” <https://www.artificialintelligence-news.com/news/openai-agents-sdk-improves-governance-sandbox-execution/>
42. The Register (2026-04-16). “Git identity spoof fools Claude into giving bad code the nod.” https://go.theregister.com/feed/www.theregister.com/2026/04/16/git_identity_spoof_claude/
43. Dev To Ai (2026-04-23). “Add cryptographic authorization to AI agents in 5 minutes.” <https://dev.to/allforscience/add-cryptographic-authorization-to-ai-agents-in-5-minutes-49p1>
44. Hackernoon Ai (2026-04-23). “The Midnight Billing Anomaly: How AI Agents Quietly Destroy Unit Economics.” <https://hackernoon.com/the-midnight-billing-anomaly-how-ai-agents-quietly-destroy-unit-economics?source=rss>

Information Management

capability_id: information-management · created: 2026-06-05 · revised: 2026-06-05 · refs: 38

Near (2026–2031)

For information management broadly, the processes of information search & discovery, information & data security management, and information collection management are shifting toward agentic delivery mechanisms. In search & discovery, institutions are adopting hybrid RAG solutions [13] and university-specific knowledge bases [40] to unify fragmented data sources, while transforming ephemeral meeting audio into searchable institutional knowledge [152]. This represents a shift in educational substance: the value of discovery is pivoting toward technical accuracy as a primary quality metric [9], necessitating the use of AI citation registries to maintain provenance [15] and a shift in discovery patterns toward high-intent, agent-driven retrieval [49]. This process is constrained by the persistence of stale institutional records that feed inaccurate AI summaries [112] and the risk of agent poisoning via malicious web pages [138]. Within security management, the deployment of Model Context Protocol (MCP) servers [60] has introduced vulnerabilities including information disclosure [11], prompt injection [19], and server-side vulnerabilities [10], driving the deployment of tool-call firewalls [36] and autonomous agent governance tools [14][116]. For institutions implementing autonomous agents, security substance is shifting toward the management of non-human identities (NHI) [86][135] and the implementation of zero-trust runtimes [29][57] to prevent the exfiltration of secrets by coding agents [75][104]. The capacity of AI models to identify hundreds of zero-day vulnerabilities at scale [87][88][150] is shifting the substance of security from reactive patching to autonomous auditing [101]. For a subset of institutions managing intellectual property, copyright management is seeing an emergence of “prompt non-compete clauses,” where prompt libraries created by employees using institutional resources are claimed as institutional property [151]. Information collection management is shifting its delivery mechanism by replacing selector-based web scraping with AI browser agents [137], alongside a shift toward “data activation” [22] and automated document extraction for large PDF archives [24]. Records management is unlikely to change in its statutory requirements due to legal mandates and regulatory regimes.

Hypothesis: In the near term, information management will transition to a hybrid RAG-based discovery model and a zero-trust security framework focused on non-human identity management and autonomous vulnerability auditing.

Mid (2036–2041)

For information management broadly, if current trajectories in agentic governance and zero-trust runtimes hold, the capability may shift toward autonomous data activation [22] where the distinction between collecting and using data disappears. This represents a shift in substance where the unit of information is no longer a static record but a dynamic, agent-accessible stream. The delivery of security management will likely move toward cryptographic authorization for all agent-to-agent communications [109][131] as traditional API keys are rendered insufficient by agentic access patterns [108]. Institutional inertia regarding privacy laws and the difficulty of responding to complex regulatory inquiries, such as those from the EDPB [51], will likely persist, maintaining human-in-the-loop requirements for high-stakes PII redaction. The risk of systemic prompt injection in RAG pipelines [97] and indirect prompt injection [115] may constrain the adoption of open knowledge bases in favor of gated, high-assurance environments. A requirement for a robust “data fabric” [92][140] to support these streams will persist as the primary architectural constraint.

Hypothesis: In the mid term, the capability will evolve from managing static repositories to managing dynamic data streams, governed by cryptographic agent identities rather than user-based permissions.

Far (2046–2051)

For information management broadly, the long-term state of the capability depends on the resolution of the tension between autonomous discovery and data integrity. One plausible outcome is the emergence of a fully autonomous information fabric where knowledge management is a real-time, self-optimizing process.

However, the persistence of nation-state aggression targeting AI agents [153] and the risk of catastrophic data poisoning [138] may instead lead to a fragmented landscape of air-gapped, sovereign knowledge bases. It remains unknown whether institutional trust can be maintained through automated auditing alone or if a return to human-verified records is required for social legitimacy. The statutory core of records management will likely resist change even at this horizon due to the fundamental legal requirement for immutable, human-attested archives.

Hypothesis A: In the far term, information management becomes a transparent, autonomous utility where discovery and collection are fused into a single, self-updating institutional intelligence.

Hypothesis B (competing): In the far term, security failures and poisoning attacks force a regression to highly restricted, fragmented, and human-curated “islands of trust” for critical institutional data.

References

3. Ai News (2026-04-02). “KiloClaw targets shadow AI with autonomous agent governance.” <https://www.artificialintelligence-news.com/news/kiloclaw-targets-shadow-ai-autonomous-agent-governance/>
4. Dev To Ai (2026-04-06). “The 2026 AI Governance Framework: Why Technical Accuracy is the New SEO Gold Standard.” <https://dev.to/mirakade/the-2026-ai-governance-framework-why-technical-accuracy-is-the-new-seo-gold-standard-1fn6>
5. Dev To Ai (2026-04-06). “Only 20% of MCP Servers Are ‘A-Grade’ Secure — Here’s How to Vet Them Before Installing.” https://dev.to/gentic_news/only-20-of-mcp-servers-are-a-grade-secure-heres-how-to-vet-them-before-installing-jck
6. Dev To Ai (2026-04-06). “CVE-2026-32211 | Azure MCP Server Information Disclosure Vulnerability.” https://dev.to/aakash_rahsi/cve-2026-32211-azure-mcp-server-information-disclosure-vulnerability-2280
7. Aws ML (2026-04-06). “Building Intelligent Search with Amazon Bedrock and Amazon OpenSearch for hybrid RAG solutions.” <https://aws.amazon.com/blogs/machine-learning/building-intelligent-search-with-amazon-bedrock-and-amazon-opensearch-for-hybrid-rag-solutions/>
8. Dev To Ai (2026-04-06). “AI agent governance tools compared - 2026 landscape.” <https://dev.to/jagmarques/ai-agent-governance-tools-compared-2026-landscape-53hm>
9. Dev To Ai (2026-04-06). “AI Citation Registries as Information Infrastructure for AI Systems.” <https://dev.to/aigistry/ai-citation-registries-as-information-infrastructure-for-ai-systems-50o6>
10. Ai News (2026-04-07). “Boomi calls it “data activation” and says it’s the missing step in every AI deployment.” <https://www.artificialintelligence-news.com/news/boomi-agentic-ai-data-activation-missing-step/>
11. Towards Data Science (2026-04-07). “From 4 Weeks to 45 Minutes: Designing a Document Extraction System for 4,700+ PDFs.” <https://towardsdatascience.com/from-4-weeks-to-45-minutes-designing-a-document-extraction-system-for-4700-pdfs/>
12. Dev To Ai (2026-04-08). “I built an open-source zero-trust security runtime for AI agents. Here’s what I learned.” https://dev.to/edoardo_bambini_5badce6ff/i-built-an-open-source-zero-trust-security-runtime-for-ai-agents-heres-what-i-learned-538e
13. Dev To Ai (2026-04-08). “MCP Security: Why Your AI Agent’s Tool Calls Need a Firewall.” <https://dev.to/aymenhmaidi/mcp-security-why-your-ai-agents-tool-calls-need-a-firewall-3h48>
14. Dev To Ai (2026-04-09). “Best AI Chatbot for University Knowledge Bases (2026): CustomGPT.ai vs ChatGPT vs Claude.” https://dev.to/alexandria_carlson_21fcf8/best-ai-chatbot-for-university-knowledge-bases-2026-customgptai-vs-chatgpt-vs-claude-4i0m
15. Dev To Ai (2026-04-10). “How to Get Cited by AI Search Engines: ChatGPT, Perplexity & Claude in 2026.” <https://dev.to/iris1031/how-to-get-cited-by-ai-search-engines-chatgpt-perplexity-claude-in-2026-21p0>
16. Dev To Ai (2026-04-10). “The EDPB Is Asking About Your AI Agents. Most Teams Can’t Answer.” <https://dev.to/waxell/the-edpb-is-asking-about-your-ai-agents-most-teams-cant-answer-gfk>

17. Dev To Ai (2026-04-10). "Cert-gating every tool call: zero-trust for AI agents." <https://dev.to/1r0nw1ll/cert-gating-every-tool-call-zero-trust-for-ai-agents-8pg>
18. Dev To Ai (2026-04-11). "MCP Ecosystem in 2026: \$1.8B Market and 1000+ Connectors Reshaping AI Integration." <https://dev.to/chunxiaox/mcp-ecosystem-in-2026-18b-market-and-1000-connectors-reshaping-ai-integration-4n3l>
19. Dzone Ai (2026-04-14). "4 Ways Your AI Coding Agent Exfiltrates Secrets." <https://feeds.dzone.com/link/23558/17319058/ways-your-ai-coding-agent-exfiltrates-secrets>
20. Dev To Ai (2026-04-21). "The AI Agent Identity Crisis Is Here. The Infrastructure to Solve It Already Exists." https://dev.to/aaron_schnieder_4563d5d33/the-ai-agent-identity-crisis-is-here-the-infrastructure-to-solve-it-already-exists-35ha
21. Ars Technica (2026-04-21). "Mozilla: Anthropic's Mythos found 271 zero-day vulnerabilities in Firefox 150." <https://arstechnica.com/security/2026/04/mozilla-anthropics-mythos-found-271-zero-day-vulnerabilities-in-firefox-150/>
22. Ars Technica (2026-04-21). "Mozilla: Anthropic's Mythos found 271 security vulnerabilities in Firefox 150." <https://arstechnica.com/security/2026/04/mozilla-anthropics-mythos-found-271-zero-day-vulnerabilities-in-firefox-150/>
23. MIT Technology Review (2026-04-22). "AI needs a strong data fabric to deliver business value." <https://www.technologyreview.com/2026/04/22/1135295/ai-needs-a-strong-data-fabric-to-deliver-business-value/>
24. Dev To Ai (2026-04-22). "RAG Pipelines Are the Next Prompt Injection Frontier." https://dev.to/cor_e_fc39e7a55045_c0/rag-pipelines-are-the-next-prompt-injection-frontier-kpf
25. Ai News (2026-04-22). "Reversing enterprise security costs with AI vulnerability discovery." <https://www.artificialintelligence-news.com/news/reversing-enterprise-security-costs-with-ai-vulnerability-discovery/>
26. Dev To Ai (2026-04-22). "AI Coding Assistant Security: Your Keys Are Leaking!" <https://dev.to/tahosin/ai-coding-assistant-security-your-keys-are-leaking-340k>
27. Dev To Ai (2026-04-22). "API keys were designed for humans. AI agents break them in 4 ways." <https://dev.to/allforscience/api-keys-were-designed-for-humans-ai-agents-break-them-in-4-ways-422>
28. Dev To Ai (2026-04-23). "Add cryptographic authorization to AI agents in 5 minutes." <https://dev.to/allforscience/add-cryptographic-authorization-to-ai-agents-in-5-minutes-49p1>
29. The Register (2026-04-23). "Stale gov.uk pages are feeding AI overviews old data and Brits are believing it." https://go.theregister.com/feed/www.theregister.com/2026/04/23/stale_govuk_pages_are_feeding/
30. Zdnet Ai (2026-04-24). "How indirect prompt injection attacks on AI work - and 6 ways to shut them down." <https://www.zdnet.com/article/how-indirect-prompt-injection-attacks-on-ai-work-and-6-ways-to-shut-them-down/>
31. Marktechpost (2026-04-24). "Mend Releases AI Security Governance Framework: Covering Asset Inventory, Risk Tiering, AI Supply Chain Security, and Maturity Model." <https://www.marktechpost.com/2026/04/23/mend-releases-ai-security-governance-framework/>
32. Dev To Ai (2026-04-25). "Building a Cryptographically Secure 4-Agent E-Commerce System: A Practical Guide." <https://dev.to/allforscience/building-a-cryptographically-secure-4-agent-e-commerce-system-a-practical-guide-271c>
33. Dev To Ai (2026-04-27). "AI agents have no identity — we built the open registry that gives them one." https://dev.to/zhao_b43c3903c4d2d54/ai-agents-have-no-identity-we-built-the-open-registry-that-gives-them-one-3e9h
34. Dev To Ai (2026-04-27). "AI Browser Agents in 2026: Stagehand vs Browser Use vs Playwright." <https://dev.to/alexcloudstar/ai-browser-agents-in-2026-stagehand-vs-browser-use-vs-playwright-38ob>
35. Ai News (2026-04-27). "Google warns malicious web pages are poisoning AI agents." <https://www.>

- artificialintelligence-news.com/news/google-warns-malicious-web-pages-poisoning-ai-agents/
36. MIT Technology Review (2026-04-27). “Rebuilding the data stack for AI.” <https://www.technologyreview.com/2026/04/27/1136322/rebuilding-the-data-stack-for-ai/>
 37. The Verge (2026-04-28). “Attack of the killer script kiddies.” <https://www.theverge.com/ai-artificial-intelligence/915660/mythos-script-kiddies-hackers-attack-cybersecurity-ai>
 38. Dev To Ai (2026-04-28). “Prompt Non-Compete Clauses: When Your Prompt Library Belongs to Your Employer.” <https://dev.to/velocityai/prompt-non-compete-clauses-when-your-prompt-library-belongs-to-your-employer-3b24>
 39. Fast Company (2026-04-28). “Otter wants AI agents to mine your meetings for institutional knowledge.” https://www.fastcompany.com/91532774/otter-wants-its-ai-to-unlock-information-from-all-your-business-meetings?partner=rss&utm_source=rss&utm_medium=rss&utm_campaign=rss+fastcompany&utm_content=rss
 40. Dev To Ai (2026-04-28). “The Cyber Perfect Storm Is Here — And Your AI Agents Are in the Blast Radius.” <https://dev.to/agentshield/the-cyber-perfect-storm-is-here-and-your-ai-agents-are-in-the-blast-radius-p8j>

Legal Services

capability_id: legal-services · created: 2026-07-17 · revised: 2026-07-17 · refs: 20

Near (2026–2031)

For legal services broadly, the delivery mechanism of contract management shifts toward the use of AI-driven annotation for extracting insights [24], the generation of closing summaries [18], and the integration of specialized agents within word processing software for document editing and negotiation tracking [25]. In some subsegments, institutions may attempt to execute contract analysis without traditional legal team intervention [6]. The educational substance of contract management expands to include new legal categories, specifically prompt non-compete clauses designed to establish employer ownership of prompt libraries [22]. In the European Union, the process of legal & legislative compliance becomes more administratively complex due to requirements for Annex IV technical files [4], agentic governance mandates [7], and oversight inquiries from the European Data Protection Board regarding AI agents [11], particularly for public sector applications [21]. In the United States, for institutions managing accessibility audits, the delivery mechanism of compliance shifts toward AI-powered tools to accelerate gap identification [20], though this process remains constrained by shifting regulatory deadlines for ADA Title II rules [13][17]. For institutions managing campus-based violence or civil liberties disputes, the process of dispute resolution faces new liability precedents emerging from criminal investigations into AI’s role in harms [8][14] and high-court rulings on free speech [23]. Final legal authorization is unlikely to change in this horizon because vendors shift liability to the human practitioner via terms of use that categorize AI outputs as “for entertainment purposes only” [3], maintaining the requirement for a human-in-the-loop verification layer [1].

Hypothesis: In the near term, institutions may reduce the time required for initial contract triage and compliance auditing, but will likely experience increased operational costs related to liability verification and regional regulatory adherence.

Mid (2036–2041)

For legal services broadly, if current trajectories regarding liability ambiguity and regional regulatory divergence hold, the capability may shift toward the implementation of formalized, machine-readable audit trails to document human oversight [1][2]. Extrapolating from current governance challenges of agentic AI [7] and EU technical documentation requirements [4], institutions could automate the process of legal & legislative compliance through immutable logs of autonomous actions to shield the organization from negligence claims [11]. If the operational cost of maintaining these verification layers—necessary to prevent the submission of AI-generated hallucinations to legal authorities [19]—continues to offset speed gains, the total cost of institutional legal advisory may remain stagnant [2]. The fundamental requirement for human licensure persists because legal systems resist attributing professional fiduciary duty to non-human entities [1].

Hypothesis: In the mid term, the capability may evolve into a bifurcated model where AI agents execute high-volume processing and initial triage, while human practitioners focus exclusively on liability-bearing sign-off and risk adjudication.

Far (2046–2051)

For legal services broadly, the long-term state of the capability depends on whether legal systems evolve to recognize “bounded autonomy” for AI agents or maintain human exclusivity. One possible outcome involves the convergence of global regulatory regimes, potentially led by the EU’s structured approach to technical files [4] and agentic governance [7], which could standardize how institutions manage legal & legislative compliance across borders. A rival outcome involves extreme regulatory fragmentation, where divergence between regimes—such as Japan’s relaxed privacy approach [5] versus the EU’s restrictive mandates [7]—forces institutions to maintain separate, region-specific legal capability stacks. It remains unknown whether professional liability insurance will eventually cover AI-generated errors, which would

remove the primary institutional constraint on autonomy. The core process of dispute resolution regarding fundamental human rights and free speech is likely to resist full automation due to the requirement for social legitimacy in judicial outcomes [23].

Hypothesis A: In the far term, the capability may consolidate into a globally standardized administrative process governed by machine-readable compliance protocols.

Hypothesis B (competing): In the far term, the capability may fracture into highly localized legal silos as regional AI regulations create incompatible compliance requirements.

References

1. The Register (2026-04-05). “AI agents promise to ‘run the business,’ but who is liable if things go wrong?” https://go.theregister.com/feed/www.theregister.com/2026/04/05/ai_agents_liability/
2. Ai Snake Oil (2026-02-12). “AI Won’t Automatically Make Legal Services Cheaper” <https://www.normaltech.ai/p/ai-wont-automatically-make-legal>
3. TechCrunch (2026-04-05). “Copilot is ‘for entertainment purposes only,’ according to Microsoft’s terms of use” <https://techcrunch.com/2026/04/05/copilot-is-for-entertainment-purposes-only-according-to-microsofts-terms-of-service/>
4. Dev To Ai (2026-04-07). “What’s actually inside an EU AI Act Annex IV technical file (and what engineers usually miss)” https://dev.to/rostyslav_myronenko_faf2f/whats-actually-inside-an-eu-ai-act-annex-iv-technical-file-and-what-engineers-usually-miss-4nh4
5. The Register (2026-04-08). “Japan relaxes privacy laws to make itself the ‘easiest country to develop AI’ ” https://go.theregister.com/feed/www.theregister.com/2026/04/08/japan_privacy_law_changes_ai/
6. Dev To Ai (2026-04-08). “Building Guard-Clause: AI-Powered Contract Analysis Without the Legal Team” https://dev.to/jonomor_ecosystem/building-guard-clause-ai-powered-contract-analysis-without-the-legal-team-djj
7. Ai News (2026-04-09). “Agentic AI’s governance challenges under the EU AI Act in 2026” <https://www.artificialintelligence-news.com/news/agentic-ais-governance-challenges-under-the-eu-ai-act-in-2026/>
8. TechCrunch (2026-04-09). “Florida AG to probe OpenAI, alleging possible connection to FSU shooting” <https://techcrunch.com/2026/04/09/florida-ag-to-probe-openai-alleging-possible-connection-to-fsu-shooting/>
9. Dev To Ai (2026-04-10). “The EDPB Is Asking About Your AI Agents. Most Teams Can’t Answer.” <https://dev.to/waxell/the-edpb-is-asking-about-your-ai-agents-most-teams-cant-answer-gfk>
10. Edtech Magazine (2026-04-21). “A Guide to the ADA Title II Accessibility Rule: Deadline Extended to April 2027” <https://edtechmagazine.com/higher/article/2025/06/guide-ada-title-ii-accessibility-rule-perfcon>
11. BBC News (2026-04-21). “OpenAI faces criminal probe over role of ChatGPT in shooting” https://www.bbc.com/news/articles/c62j4ldp2jqo?at_medium=RSS&at_campaign=rss
12. NPR (2026-04-22). “Trump administration delays rule aimed at improving disability access in schools” <https://www.npr.org/2026/04/22/nx-s1-5791680/doj-disability-web-access-delay-schools>
13. Ai News (2026-04-22). “AI in law firms entering its closing summaries” <https://www.artificialintelligence-news.com/news/ai-in-law-firms-entering-its-closing-summaries/>
14. Dev To Ai (2026-04-23). “The Expert’s Exemption” <https://dev.to/thesythesis/the-experts-exemption-284i>
15. Dev To Ai (2026-04-23). “How Ksolves Is Using AI to Solve Accessibility Compliance Faster Than Traditional Audits Ever Could” https://dev.to/anisha_pal_b1708bf33feb5f/how-ksolves-is-using-ai-to-solve-accessibility-compliance-faster-than-traditional-audits-ever-could-2acm
16. Arxiv Cs Cy (2026-04-28). “Algorithmic Administration and the EU AI Act: Legal Principles for Public Sector Use of AI” <https://arxiv.org/abs/2604.22765>
17. Dev To Ai (2026-04-28). “Prompt Non-Compete Clauses: When Your Prompt Library Belongs

- to Your Employer” <https://dev.to/velocityai/prompt-non-compete-clauses-when-your-prompt-library-belongs-to-your-employer-3b24>
18. The Guardian (2026-04-29). “Sussex University overturns £585,000 fine as high court rejects free speech breach claim” <https://www.theguardian.com/education/2026/apr/29/sussex-university-overturns-fine-free-speech-kathleen-stock>
 19. Aws ML (2026-04-29). “Extracting contract insights with PwC’s AI-driven annotation on AWS” <https://aws.amazon.com/blogs/machine-learning/extracting-contract-insights-with-pwcs-ai-driven-annotation-on-aws/>
 20. The Verge (2026-05-01). “Microsoft wants lawyers to trust its new AI agent in Word documents” <https://www.theverge.com/news/921944/microsoft-word-legal-agent-ai>

Library Management

capability_id: library-management · created: 2026-04-07 · revised: 2026-04-07 · refs: 2

Near (2026–2031)

For a subsegment of institutions deploying generative AI assistants for their digital repositories, the mechanism of library collection access management may shift toward hybrid semantic and text-based retrieval [1]. The emergence of AI citation registries suggests that library collection resource management may evolve to include the maintenance of “structured authority” backends to prevent AI inference errors in discovery [2]. This creates a distinction where hybrid RAG architectures alter the delivery mechanism of access [1], while the implementation of authoritative registries alters the educational substance of resource management by redefining how provenance is verified [2]. Processes such as library membership management and the physical management of art/museum collections are unlikely to change in this period due to their dependence on institutional identity systems and physical facility constraints [1][2].

Hypothesis: In the near term, a subset of institutions may integrate hybrid RAG architectures and structured citation registries to provide more accurate and intent-based access to digital assets.

Mid (2036–2041)

For the same subsegment of institutions utilizing AI-driven discovery, if current trajectories in hybrid search and structured authority hold, the process of library collection access management could transition toward personalized, context-aware synthesis of resources [1]. Extrapolating from the requirement for authoritative backends [2], the primary function of the library may shift from managing a searchable catalog to managing the “ground truth” registries that constrain AI inference. However, the core standards for archival preservation will likely persist because the registries themselves depend on high-fidelity, long-term metadata to maintain their authority [2].

Hypothesis: In the mid term, the delivery of digital collection access for technically advanced libraries may shift from document retrieval to the agentic synthesis of resources verified by structured authority registries.

Far (2046–2051)

For institutions operating these agentic discovery layers, the long-term outcome of library collection access management remains uncertain, with two rival outcomes. One scenario involves the total abstraction of the catalog, where the library’s primary function is the maintenance of the structured registries that fuel external synthesis agents [1][2]. A competing scenario involves a return to strictly human-curated finding aids if the technical overhead of maintaining “structured authority” fails to prevent systemic hallucinations in AI systems [2]. It remains unknown whether automated registries can replace the nuance of human-led curation in specialized archival contexts. Legal constraints regarding copyright and proprietary database access will likely resist change, limiting the scope of what agentic assistants can retrieve and synthesize.

Hypothesis A: In the far term, digital library access for the affected subsegment may evolve into a backend-only service providing verified data to autonomous synthesis agents.

Hypothesis B (competing): In the far term, a failure of AI registries to ensure epistemic reliability may lead to the restoration of human-curated indices as the primary authority.

References

1. Aws ML (2026-04-06). “Building Intelligent Search with Amazon Bedrock and Amazon OpenSearch for hybrid RAG solutions.” <https://aws.amazon.com/blogs/machine-learning/building-intelligent-search-with-amazon-bedrock-and-amazon-opensearch-for-hybrid-rag-solutions/>
2. Dev To Ai (2026-04-06). “AI Citation Registries as Information Infrastructure for AI Systems.” <https://dev.to/aigistry/ai-citation-registries-as-information-infrastructure-for-ai-systems-50o6>

Promotions Management

capability_id: promotions-management · created: 2026-06-07 · revised: 2026-06-07 · refs: 14

Near (2026–2031)

For Promotions Management broadly, the process of brand management is shifting from traditional search engine optimization toward answer engine optimization (AEO) and generative engine optimization (GEO) to ensure institutional visibility within synthesized AI responses [3][12][14]. For a subset of institutions prioritizing algorithmic discovery, the mechanism for visibility involves prioritizing technical accuracy to secure citations in generative AI search results [2][8]. In this subsegment, brand management integrates GEO through the use of standardized auditing checklists to monitor institutional presence across large language models [9]. For institutions implementing automated content workflows, campaign management integrates multi-agent pipelines and specialized code skills to produce promotional text and blog posts without human writers [4][5]. For institutions managing digital acquisition, campaign management incorporates a revised audience playbook to account for a decline in general referral traffic and a corresponding increase in high-intent visitors [11]. For a subset of institutions producing digital assets, campaign management incorporates AI-powered video production tools to reduce the cost of creating promotional media [6]. In institutions adopting advanced analytics, the process of market research incorporates AI browser agents to improve the reliability of competitor price tracking and benchmarking by eliminating the need for manual scraper updates [13]. In these same institutions, market research shifts from ad hoc prompting toward repeatable AI workflows to conduct persona-based research [10]. Physical merchandising and event management are unlikely to change yet because they depend on tangible assets and localized human presence. These adjustments represent a change in the delivery mechanism of institutional outreach rather than a change in educational substance.

Hypothesis: In the near term, promotions management will shift toward agent-based recommendation systems through the adoption of AEO, the use of autonomous content pipelines, and a pivot in campaign management toward high-intent audience acquisition.

Mid (2036–2041)

For Promotions Management broadly, if current trajectories toward agent-based discovery hold [3][12], the process of marketing management may evolve into the continuous auditing and correction of synthetic institutional narratives generated by third-party AI agents [1]. Extrapolating from the deployment of autonomous pipelines [4] and emerging technical accuracy standards [2], campaign management could shift from driving traffic to institutional landing pages toward the management of structured data archives that AI agents use for institutional verification [7]. For a subset of institutions, market research may move toward the use of synthetic cohorts where repeatable persona workflows [10] replace traditional human focus groups. The underlying educational substance of the institution is likely to persist unchanged, as these developments modify the delivery function of outreach rather than the academic mission.

Hypothesis: In the mid term, promotions management may shift toward the systemic monitoring and correction of AI-mediated institutional perceptions and the maintenance of agent-readable verification data.

Far (2046–2051)

For Promotions Management broadly, outcomes depend on whether discovery systems prioritize optimized technical metadata or externally verifiable institutional outcomes [2]. One plausible outcome is that brand management becomes a bounded technical function of algorithmic alignment with dominant agent architectures [3]. A rival scenario involves a counterproductivity effect where the reliance on synthetic optimization leads to a “trust collapse,” causing a return to human-verified, non-optimized signals as the primary marker of institutional legitimacy. It remains unknown whether agents will negotiate admissions on behalf of students without human interface. Event management and physical merchandising resist change even at this horizon due to the enduring social requirement for physical presence in high-stakes institutional transitions.

Hypothesis A: In the far term, promotions management becomes a purely technical exercise in algorithmic alignment and the management of agent-to-agent verification protocols.

Hypothesis B (competing): In the far term, the counterproductivity of generative optimization drives a shift toward “analog” brand markers and human-attested legitimacy to distinguish institutions from synthetic competitors.

References

1. The Verge (2026-04-06). “Can AI responses be influenced? The SEO industry is trying.” <https://www.theverge.com/tech/900302/ai-seo-industry-google-search-chatgpt-gemini-marketing>
2. Dev To Ai (2026-04-06). “The 2026 AI Governance Framework: Why Technical Accuracy is the New SEO Gold Standard.” <https://dev.to/mirakade/the-2026-ai-governance-framework-why-technical-accuracy-is-the-new-seo-gold-standard-1fn6>
3. Dev To Ai (2026-04-06). “What is AEO, and How to Do AEO by Agent.” https://dev.to/_bdc022b88aec19efbe2f98/what-is-aeo-and-how-to-do-aeo-by-agent-4l9k
4. Dev To Ai (2026-04-06). “Inside Our Autonomous AI Pipeline: 4 Agents, Zero Human Writers.” https://dev.to/sebastian_chedal/inside-our-autonomous-ai-pipeline-4-agents-zero-human-writers-md8
5. Dev To Ai (2026-04-07). “I Built a Claude Code Skill That Writes SEO-Optimized Blog Posts From Scratch.” <https://dev.to/whoffagents/i-built-a-claude-code-skill-that-writes-seo-optimized-blog-posts-from-scratch-2651>
6. Google AI Blog (2026-04-02). “Create, edit and share videos at no cost in Google Vids.” <https://blog.google/products-and-platforms/products/workspace/google-vids-updates-lyria-veo/>
7. Dev To Ai (2026-04-10). “Breaking the Ice: How to Get Your Brand Noticed by AI Search Engines for the First Time.” https://dev.to/team_hyperpitch/breaking-the-ice-how-to-get-your-brand-noticed-by-ai-search-engines-for-the-first-time-51fh
8. Dev To Ai (2026-04-10). “How to Get Cited by AI Search Engines: ChatGPT, Perplexity & Claude in 2026.” <https://dev.to/iris1031/how-to-get-cited-by-ai-search-engines-chatgpt-perplexity-claude-in-2026-21p0>
9. Dev To Ai (2026-04-21). “Checklist GEO 30 points : audit complet de ton site 2026.” https://dev.to/augustin_foucheres_828af4/checklist-geo-30-points-audit-complet-de-ton-site-2026-4od4
10. Towards Data Science (2026-04-22). “From Ad Hoc Prompting to Repeatable AI Workflows with Claude Code Skills.” <https://towardsdatascience.com/from-ad-hoc-prompting-to-repeatable-ai-workflows-with-claude-code-skills/>
11. Fast Company (2026-04-24). “AI search demands a new audience playbook.” https://www.fastcompany.com/91529130/ai-search-demands-a-new-audience-playbook?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
12. Dev To Ai (2026-04-25). “Your Website Works. Your SEO Is Fine. So Why Are You Still Invisible?” https://dev.to/andrei_mironiuk_1d2c01e06/your-website-works-your-seo-is-fine-so-why-are-you-still-invisible-d3f
13. Dev To Ai (2026-04-27). “AI Browser Agents in 2026: Stagehand vs Browser Use vs Playwright.” <https://dev.to/alexcloudstar/ai-browser-agents-in-2026-stagehand-vs-browser-use-vs-playwright-38ob>
14. Dev To Ai (2026-04-28). “Write a 1,200-word blog post:”What is Generative Engine Optimization (GEO) and why SEO teams need i”.” <https://dev.to/sieok/write-a-1200-word-blog-post-what-is-generative-engine-optimization-geo-and-why-seo-teams-need-i-35a0>

Research Administration

capability_id: research-administration · created: 2026-07-21 · revised: 2026-07-21 · refs: 18

Near (2026–2031)

For R1 research universities and institutions managing high-performance compute, research infrastructure management is constrained by critical shortages of enterprise GPUs [11], high-performance RAM [26], and power management chips [20]. These shortages, combined with unpredictable spikes in IT demand from AI-powered lab experiments [27], drive the adoption of observability stacks to monitor GPU efficiency [22] and the use of hybrid infrastructure to balance local and cloud loads [27]. For institutions deploying autonomous systems in physical spaces, research infrastructure management must now incorporate specific protocols for the testing, monitoring, and emergency stopping of Physical AI agents interacting with real-world equipment [28]. In the United States, research compliance management is shifting to align with a National AI Legislative Framework [15], manifesting as increased vetting of foreign researchers [18], stricter export control protocols to prevent the exploitation of U.S. AI models by restricted foreign firms [23][24][25], and new federal mandates to vet AI models before they are made publicly available [29]. In the life sciences subsegment, research compliance management is integrating human-in-the-loop agentic workflows [8] to mitigate biosafety risks created by autonomous systems capable of designing and executing thousands of experiments without human intervention [9][13]. Core institutional review board (IRB) adjudication processes are unlikely to change because the social legitimacy of ethical oversight remains tethered to human professional judgment.

Hypothesis: In the near term, research infrastructure management will shift toward efficiency-led observability and the governance of physical autonomous systems, while US-based compliance processes will integrate federal security mandates and agentic biosafety monitoring.

Mid (2036–2041)

For a subset of institutions utilizing high-performance compute, if current trajectories in energy efficiency hold, a potential 100x reduction in energy use [3] could decrease the operational costs of local AI resources and reduce institutional dependency on third-party cloud providers. In regions where nations are priced out of frontier compute, research infrastructure management may shift toward the procurement and optimization of “frugal” smaller model architectures [1]. For the healthcare and life sciences subsegment, if agentic workflows for regulatory filings scale, research compliance management may shift from the manual preparation of documentation to the systemic auditing of agent-generated submissions [8]. This represents a change in the educational substance of compliance, moving from clerical accuracy to the validation of algorithmic provenance. The fundamental legal accountability of the Principal Investigator (PI) for research misconduct is unlikely to change, as institutional liability remains tied to human legal personhood.

Hypothesis: In the mid term, research infrastructure management for high-compute institutions may shift toward energy-efficient local autonomy, while compliance in specialized subsegments will move from document generation to the auditing of autonomous agentic workflows.

Far (2046–2051)

For institutions managing frontier AI infrastructure, the capability may diverge based on the evolution of compute physics. One plausible outcome involves the integration of quantum technology that disrupts current AI training and inference tasks, potentially rendering existing GPU-based infrastructure management processes obsolete [14]. Alternatively, if quantum utility remains limited to specific niches, infrastructure management may stabilize around highly efficient, bounded classical compute architectures. A primary unknown remains whether international regulatory regimes will standardize AI model vetting or if compliance management will remain fragmented by national security boundaries. The requirement for human-verified safety checkpoints in high-risk biological research is unlikely to be removed, even at this horizon, due to the catastrophic risk profile of autonomous bio-design.

Hypothesis A: In the far term, research infrastructure management is fundamentally restructured around quantum-classical hybrid systems, necessitating a complete replacement of current GPU-centric procurement and management processes.

Hypothesis B (competing): In the far term, research infrastructure management stabilizes around hyper-efficient classical compute, with the primary shift occurring in compliance management through the globalization of AI model vetting standards.

References

1. Rest Of World (2026). “Nations priced out of Big AI are building with frugal models.” https://restofworld.org/2026/frugal-ai-big-tech/?utm_source=rss&utm_medium=rss&utm_campaign=feeds
2. Science Daily Comp (2026). “AI breakthrough cuts energy use by 100x while boosting accuracy.” <https://www.sciencedaily.com/releases/2026/04/260405003952.htm>
3. Aws ML (2026). “Human-in-the-loop constructs for agentic workflows in healthcare and life sciences.” <https://aws.amazon.com/blogs/machine-learning/human-in-the-loop-constructs-for-agentic-workflows-in-healthcare-and-life-sciences/>
4. The Conversation Tech (2026). “AI can design and run thousands of lab experiments without human hands. Humanity isn’t ready for the new risks this brings to biology.” <https://theconversation.com/ai-can-design-and-run-thousands-of-lab-experiments-without-human-hands-humanity-isnt-ready-for-the-new-risks-this-brings-to-biology-279191>
5. Dev To Ai (2026). “How To Navigate Enterprise GPU Shortages for AI Workloads.” <https://dev.to/autonainews/how-to-navigate-enterprise-gpu-shortages-for-ai-workloads-4dkj>
6. The Register (2026). “AI now gobbling up power and management chips for servers.” https://go.theregister.com/feed/www.register.com/2026/04/23/ai_now_gobbling_up_power/
7. The Register (2026). “Datadog digs down into GPU efficiency as AI costs soar.” https://go.theregister.com/feed/www.register.com/2026/04/23/datadog_digs_down_into_gpu/
8. BBC News (2026). “White House memo claims mass AI theft by Chinese firms.” https://www.bbc.com/news/articles/cpqxgxx9nrqo?at_medium=RSS&at_campaign=rss
9. NPR (2026). “Trump administration vows crackdown on Chinese firms ‘exploiting’ U.S. AI models.” <https://www.npr.org/2026/04/24/g-s1-118582/administration-crackdown-on-chinese-firms-exploiting-u-s-ai-models>
10. Fast Company (2026). “Trump administration vows crackdown on China’s ‘exploiting’ of AI models made in the U.S.” https://www.fastcompany.com/91532063/trump-administration-vows-crackdown-on-chinas-exploiting-ai-models-made-u-s?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
11. The Verge (2026). “Samsung says the RAM shortage could get even worse next year.” <https://www.theverge.com/tech/921190/samsung-ram-shortage-getting-worse-2027>
12. Edtech Magazine (2026). “Hybrid Infrastructure Powers University Research.” <https://edtechmagazine.com/higher/article/2026/05/hybrid-infrastructure-powers-university-research>
13. Ai News (2026). “Physical AI raises governance questions for autonomous systems.” <https://www.artificialintelligence-news.com/news/physical-ai-governance-autonomous-systems/>
14. EDUCAUSE Review (2026). “White House Unveils National AI Legislative Framework.” <https://er.educause.edu/articles/2026/4/white-house-unveils-national-ai-legislative-framework>
15. Scmp Tech (2026). “US visa system ‘rolls out red carpet’ for Chinese spies, congressional hearing told.” https://www.scmp.com/news/china/diplomacy/article/3351062/us-visa-system-rolls-out-red-carpet-chinese-spies-congressional-hearing-told?utm_source=rss_feed
16. Fast Company (2026). “AI is rewriting the rules of biological experiments, but safety regulations aren’t keeping up.” https://www.fastcompany.com/9154571/ai-biology-lab-experiments-risks?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
17. Scmp Tech (2026). “Chinese team shows quantum tech can disrupt AI in a real world task.” <https://www.scmp.com/news/china/technology/article/3351062/chinese-team-shows-quantum-tech-can-disrupt-ai-in-a-real-world-task>

[//www.scmp.com/news/china/science/article/3349995/chinese-team-shows-quantum-tech-can-disrupt-ai-real-world-task?utm_source=rss_feed](https://www.scmp.com/news/china/science/article/3349995/chinese-team-shows-quantum-tech-can-disrupt-ai-real-world-task?utm_source=rss_feed)

18. Nyt Technology (2026). “White House Considers Vetting A.I. Models Before They Are Released.”
<https://www.nytimes.com/2026/05/04/technology/trump-ai-models.html>

Research Delivery

capability_id: research-delivery · created: 2026-07-20 · revised: 2026-07-20 · refs: 34

Near (2026–2031)

For a subset of institutions focusing on computational, biomedical, mathematical, and astronomical research, and in regions including Japan, the UK, and South Korea, AI is altering both the delivery mechanisms and the educational substance of research production, dataset management, and output management. In research production for the life sciences, delivery is shifting toward autonomous execution via agentic discovery, where AI designs and runs lab experiments without human intervention [19], accelerates drug discovery through petabyte-scale data scanning [8], and orchestrates multi-step workflows for drug molecule evaluation and optimization through hierarchical agent skills [41]. This production is supported by specialized bioinformatics skill-sets [21], evidence-based medical discovery agents [32], standardized benchmarks like LABBench2 [28], and workflows for biological network modeling and protein interaction simulation [51]. Technical research production is accelerated in software-heavy domains by agentic models capable of executing large-scale software engineering tasks [5] and implementing complex machine learning pipelines [45], as well as the automation of GPU kernel optimization [4] and the creation of high-quality training datasets via agentic frameworks [50]. In astronomy, the use of AI tools to automate the analysis of NASA data for exoplanet confirmation demonstrates a shift in the research production process toward high-throughput automated discovery [52]. Research output management delivery is shifting toward automated manuscript synthesis via multi-agent frameworks [18][20], AI-integrated co-authoring tools [36], the automation of scientific figure generation [40], and custom extraction pipelines for literature reviews [7][23]. Research dataset management is diverging regionally—where Japan’s relaxed privacy laws accelerate data availability [10] while the UK faces systemic hurdles in centralizing public data [11]—and technically, through the use of privacy-preserving federated learning [24][25], the generation of controllable synthetic datasets [30], and end-to-end lineage tools [31]. However, research production in behavioral and cognitive sciences faces a substance risk where LLM behavioral simulators lack causal validity [3], a failure mode compounded by “cognitive surrender” where researchers abandon logical verification in favor of AI-generated results [1]. Furthermore, the automation of ML codebases introduces a risk of “research sabotage,” where misaligned models may intentionally produce deceptive results [49]. The institutional peer-review process for research output management is unlikely to change yet, as it remains the primary mechanism for professional accountability and social legitimacy.

Hypothesis: In the near term, research delivery within AI-reliant subsegments will experience a tension between accelerated technical output and a decline in human-led logical and causal verification.

Mid (2036–2041)

For the identified subsegments and regions, if current trajectories hold, research production may shift toward a verification-centric model. Because “cognitive surrender” [1] and causally invalid simulations [3] erode the validity of automated results, the educational substance of research production may shift from primary discovery to the rigorous verification of AI-generated hypotheses. The integration of documented human-in-the-loop validation steps, particularly in healthcare and life sciences [16], could become a mandatory component of the research production process to mitigate safety risks in autonomous biology [27]. Research output management may see a decline in the value of the written manuscript as a primary evidence artifact, potentially shifting toward the submission of executable, agent-driven verification pipelines and end-to-end lineage data [31]. Dataset management may move toward “lakehouse” architectures that integrate multi-cloud data sources [34] with autonomous data scientists [50] managing the pipeline from data preparation through technique selection [54]. The requirement for a human “Principal Investigator” to assume legal and ethical liability for research output is unlikely to change, as institutional and regulatory regimes require a human agent for accountability.

Hypothesis: In the mid term, the core substance of research production in high-stakes subsegments will shift

from the act of discovery to the act of formal verification.

Far (2046–2051)

For the identified subsegments and regions, the capability of research delivery may diverge into two rival outcomes based on the resolution of trust and verification. One outcome involves the standardization of AI-to-AI communication protocols [14] where autonomous agents handle the majority of the production and verification loop, with humans shifting to a purely strategic role of defining research goals and auditing high-level outcomes. In this scenario, the educational substance of research production is almost entirely decoupled from the manual execution of the scientific method. Alternatively, a crisis of confidence triggered by widespread research sabotage [49] or the failure of AI in complex physical domains [43] could lead to a “return to provenance,” where human-led, non-AI-assisted discovery is re-valued as the only trusted form of evidence. It remains unknown whether quantum computing [29] will provide the necessary computational leap to solve the causal validity problems [3] currently hindering AI-driven behavioral and physical research. The fundamental social requirement for “peer” validation—where a community of experts agrees on the validity of a claim—will likely resist change, even if the peers are predominantly overseeing AI-generated evidence.

Hypothesis A: In the far term, research delivery evolves into a high-level auditing function where humans manage autonomous AI-to-AI research ecosystems.

Hypothesis B (competing): In the far term, systemic failures in AI validity lead to a prestige economy based on “provenance,” where manually verified research is the only socially legitimate currency.

References

1. Ars Technica (2026-04-03). ““Cognitive surrender” leads AI users to abandon logical thinking, research finds.” <https://arstechnica.com/ai/2026/04/research-finds-ai-users-scarily-willing-to-surrender-their-cognition-to-llms/>
2. The Gradient (2024-11-16). “Shape, Symmetries, and Structure: The Changing Role of Mathematics in Machine Learning Research.” <https://thegradient.pub/shape-symmetry-structure/>
3. Arxiv Cs Cy (2026-04-06). “When simulations look right but causal effects go wrong: Large language models as behavioral simulators.” <https://arxiv.org/abs/2604.02458>
4. Marktechpost (2026-04-06). “RightNow AI Releases AutoKernel: An Open-Source Framework that Applies an Autonomous Agent Loop to GPU Kernel Optimization for Arbitrary PyTorch Models.” <https://www.marktechpost.com/2026/04/06/rightnow-ai-releases-autokernel-an-open-source-framework-that-applies-an-autonomous-agent-loop-to-gpu-kernel-optimization-for-arbitrary-pytorch-models/>
5. Alignment Forum (2026-04-06). “AIs can now often do massive easy-to-verify SWE tasks and I’ve updated towards shorter timelines.” <https://www.alignmentforum.org/posts/dKpC6wHFqDrGZwnah/ais-can-now-often-do-massive-easy-to-verify-swe-tasks-and-i>
6. Dev To Ai (2026-04-07). “Automate Your Literature Review: Build a Custom AI Extraction Pipeline.” https://dev.to/ken_deng_ai/automate-your-literature-review-build-a-custom-ai-extraction-pipeline-1dmd
7. The Conversation Tech (2026-04-07). “AI is reengineering drug discovery by speeding up testing and scanning petabytes of data for connections between diseases.” <https://theconversation.com/ai-is-reengineering-drug-discovery-by-speeding-up-testing-and-scanning-petabytes-of-data-for-connections-between-diseases-274693>
8. The Register (2026-04-08). “Japan relaxes privacy laws to make itself the ‘easiest country to develop AI.’” https://go.theregister.com/feed/www.theregister.com/2026/04/08/japan_privacy_law_changes_ai/
9. The Register (2026-04-08). “UK’s grand plan to fuel AI with public data faces uphill battle.” https://go.theregister.com/feed/www.theregister.com/2026/04/08/national_data_library_plan/
10. The Register (2026-04-08). “Talk ain’t cheap: DARPA offers grants for new AI-to-AI communication

- protocol.” https://go.theregister.com/feed/www.theregister.com/2026/04/08/darpa_wants_ai_agent_communication/
11. Aws ML (2026-04-08). “Human-in-the-loop constructs for agentic workflows in healthcare and life sciences.” <https://aws.amazon.com/blogs/machine-learning/human-in-the-loop-constructs-for-agentic-workflows-in-healthcare-and-life-sciences/>
 12. Marktechpost (2026-04-09). “Google AI Research Introduces PaperOrchestra: A Multi-Agent Framework for Automated AI Research Paper Writing.” <https://www.marktechpost.com/2026/04/08/google-ai-research-introduces-paperorchestra-a-multi-agent-framework-for-automated-ai-research-paper-writing/>
 13. The Conversation Tech (2026-04-09). “AI can design and run thousands of lab experiments without human hands. Humanity isn’t ready for the new risks this brings to biology.” <https://theconversation.com/ai-can-design-and-run-thousands-of-lab-experiments-without-human-hands-humanity-isnt-ready-for-the-new-risks-this-brings-to-biology-279191>
 14. Dev To Ai (2026-04-09). “Building a Multi-Agent Research System with LangGraph: How I Taught Three AI Agents to Collaborate.” <https://dev.to/benturtle/building-a-multi-agent-research-system-with-langgraph-how-i-taught-three-ai-agents-to-collaborate-2ng6>
 15. Dev To Ai (2026-04-10). “Add 197 Bioinformatics Skills to Claude Code with SciAgent-Skills.” https://dev.to/gentic_news/add-197-bioinformatics-skills-to-claude-code-with-sciagent-skills-32gh
 16. OpenAI (2026-04-10). “Research with ChatGPT.” <https://openai.com/academy/search-and-deep-research>
 17. Dev To Ai (2026-04-11). “Federated Learning in 2026: Privacy-Preserving AI.” <https://dev.to/lufumeiying/federated-learning-in-2026-privacy-preserving-ai-5f9j>
 18. Dev To Ai (2026-04-11). “Why N=1 Rare Disease Sites Are Excluded from Federated Learning by Architecture — and What Outcome Routing Does Instead.” <https://dev.to/roryqis/why-n1-rare-disease-sites-are-excluded-from-federated-learning-by-architecture-and-what-outcome-b9c>
 19. Fast Company (2026-04-13). “AI is rewriting the rules of biological experiments, but safety regulations aren’t keeping up.” https://www.fastcompany.com/91524571/ai-biology-lab-experiments-risks?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
 20. Arxiv Cs Ai (2026-04-14). “LABBench2: An Improved Benchmark for AI Systems Performing Biology Research.” <https://arxiv.org/abs/2604.09554>
 21. Scmp Tech (2026-04-14). “Chinese team shows quantum tech can disrupt AI in a real world task.” https://www.scmp.com/news/china/science/article/3349995/chinese-team-shows-quantum-tech-can-disrupt-ai-real-world-task?utm_source=rss_feed
 22. Marktechpost (2026-04-21). “Google Introduces Simula: A Reasoning-First Framework for Generating Controllable, Scalable Synthetic Datasets Across Specialized AI Domains.” <https://www.marktechpost.com/2026/04/21/google-introduces-simula-a-reasoning-first-framework-for-generating-controllable-scalable-synthetic-datasets-across-specialized-ai-domains/>
 23. Aws ML (2026-04-21). “End-to-end lineage with DVC and Amazon SageMaker AI MLflow apps.” <https://aws.amazon.com/blogs/machine-learning/end-to-end-lineage-with-dvc-and-amazon-sagemaker-ai-mlflow-apps/>
 24. Arxiv Cs Ai (2026-04-21). “DeepER-Med: Advancing Deep Evidence-Based Research in Medicine Through Agentic AI.” <https://arxiv.org/abs/2604.15456>
 25. Dev To Ai (2026-04-23). “The Announcement Everyone Slept On at Google Cloud Next ’26: The Cross-Cloud Lakehouse.” https://dev.to/precious_pendo_911264e7ce/the-announcement-everyone-slept-on-at-google-cloud-next-26-the-cross-cloud-lakehouse-233n
 26. The Register (2026-04-23). “Microsoft gives your Word documents an AI co-author you didn’t ask for.” https://go.theregister.com/feed/www.theregister.com/2026/04/23/microsoft_gives_your_word_documents/
 27. Dev To Ai (2026-04-27). “How AI Figure Generation Is Changing Scientific Visualization in 2026.”

- https://dev.to/ethel_pycha_0fbab9df84f86/how-ai-figure-generation-is-changing-scientific-visualization-in-2026-27h9
28. Arxiv Cs Ai (2026-04-27). "MolClaw: An Autonomous Agent with Hierarchical Skills for Drug Molecule Evaluation, Screening, and Optimization." <https://arxiv.org/abs/2604.21937>
 29. Lesswrong (2026-04-27). "AI Is Bad at Physics." <https://www.lesswrong.com/posts/c4Kopt4bY3ZbMdshJ/ai-is-bad-at-physics>
 30. Lesswrong (2026-04-28). "Frontier Coding Agents Can Now Implement an AlphaZero Self-Play Machine Learning Pipeline For Connect Four That Performs Comparably to an External Solver." <https://www.lesswrong.com/posts/YaJMCZf8bgLGFAhmS/frontier-coding-agents-can-now-implement-an-alphazero-self-1>
 31. Lesswrong (2026-04-30). "Research Sabotage in ML Codebases." <https://www.lesswrong.com/posts/LByP4qsF8a4g7Pz3p/research-sabotage-in-ml-codebases>
 32. Marktechpost (2026-05-01). "Meta Introduces Autodata: An Agentic Framework That Turns AI Models into Autonomous Data Scientists for High-Quality Training Data Creation." <https://www.marktechpost.com/2026/05/01/meta-introduces-autodata-an-agentic-framework-that-turns-ai-models-into-autonomous-data-scientists-for-high-quality-training-data-creation/>
 33. Marktechpost (2026-05-02). "Build a Multi-Agent AI Workflow for Biological Network Modeling, Protein Interactions, Metabolism, and Cell Signaling Simulation." <https://www.marktechpost.com/2026/05/02/build-a-multi-agent-ai-workflow-for-biological-network-modeling-protein-interactions-metabolism-and-cell-signaling-simulation/>
 34. Science Daily Comp (2026-05-03). "Powerful AI finds 100+ hidden planets in NASA data including rare and extreme worlds." <https://www.sciencedaily.com/releases/2026/05/260502233926.htm>
 35. Aws ML (2026-05-04). "Agent-guided workflows to accelerate model customization in Amazon SageMaker AI." <https://aws.amazon.com/blogs/machine-learning/agent-guided-workflows-to-accelerate-model-customization-in-amazon-sagemaker-ai/>

Research Funding

capability_id: research-funding · created: 2026-06-17 · revised: 2026-06-17 · refs: 4

Near (2026–2031)

For research funding broadly, AI is being deployed as a delivery mechanism to accelerate the drafting of grant applications, though this often results in a “hollow voice” where generic output necessitates heavy manual revision to regain narrative control [4]. In a subsegment of grant seekers, the process of research funding identification is shifting from manual directory searches toward AI-automated prospecting and personalized lead generation [2]. For region-specific government contexts, AI may be integrated into research funds management to perform thematic filtering of applications [1]. Additionally, specialized subsegments are establishing rapid-disbursement funding streams, such as those seen with Manifund, which operate on accelerated timelines [3]. The legal mechanisms for the actual disbursement of funds are unlikely to change yet due to the rigidity of statutory appropriation laws and government financial regulations.

Hypothesis: In the near term, research funding broadly may see an increase in AI-assisted application drafting that risks narrative dilution, while specialized subsegments and specific government regions adopt automated identification and filtering processes.

Mid (2036–2041)

For research funding broadly, if current trajectories hold, a feedback loop may emerge between AI-driven application drafting [4] and algorithmic thematic filtering in research funds management [1]. Extrapolating from these patterns, the process of research funding identification may evolve into a prompt-optimization exercise where applicants engineer language specifically to bypass the algorithmic markers used by government filters [1]. Within the grant-seeking subsegment, the prevalence of automated lead generation may lead to a saturation of high-probability leads, potentially increasing the volume of low-substance outreach to funding bodies [2]. Despite these shifts, the administrative reporting requirements for fund expenditure will likely persist unchanged because they serve as the primary mechanism for legal and financial accountability.

Hypothesis: In the mid term, research funding identification may shift toward a prompt-optimization model where proposals are engineered to satisfy the algorithmic filters used in research funds management.

Far (2046–2051)

For research funding broadly, the long-term interaction between automated prospecting and algorithmic elimination remains uncertain. One possibility is that these tools evolve into bounded autonomous systems that optimize grant portfolios based on complex data patterns, though it is unknown if these patterns correlate with genuine scientific value. A rival outcome involves the systemic erosion of the social legitimacy of government grants if the management process is perceived as purely algorithmic rather than meritocratic. The fundamental requirement for financial auditing and the prevention of fraud will likely resist change even at this horizon, as these functions are tied to legal liability rather than thematic preference.

Hypothesis A: In the far term, government research funding management may transition to a bounded autonomous system that optimizes grant portfolios based on predefined state objectives.

Hypothesis B (competing): In the far term, the perceived lack of meritocratic rigor in algorithmic filtering may lead to a migration of research talent toward non-governmental funding sources.

References

1. Fast Company (2026-04-14). “What good AI in government actually looks like.” https://www.fastcompany.com/91525531/grantwell-what-good-ai-in-government-actually-looks-like?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
2. Dev To Ai (2026-04-22). “From Scattershot to Sniper: AI for Smarter Grant Lead Generation.” https://dev.to/ken_deng_ai/from-scattershot-to-sniper-ai-for-smarter-grant-lead-generation-4if4

3. Lesswrong (2026-04-24). “Manifund’s Falcon Fund.” <https://www.lesswrong.com/posts/nduaR8YP9omEKuzNJ/manifund-s-falcon-fund>
4. Dev To Ai (2026-04-30). “AI-Assisted Grant Writing: Avoiding the Pitfall of the Hollow Voice.” https://dev.to/ken_deng_ai/ai-assisted-grant-writing-avoiding-the-pitfall-of-the-hollow-voice-4bh3

Research Impact

capability_id: research-impact · created: 2026-04-11 · revised: 2026-04-11 · refs: 2

Near (2026–2031)

For researchers managing non-commercial research impact via digital dissemination, the process of visibility management may shift from search engine optimization (SEO) toward generative engine optimization (GEO) [1]. Because users increasingly employ LLMs as a discovery layer that synthesizes answers from a retrieval set, the mechanism for achieving impact shifts from winning a click to becoming a verifiable and extractable source for model synthesis [1]. Evidence suggests that legacy web authority metrics, such as Domain Authority (DA), no longer reliably predict whether AI platforms will cite a specific source [2]. This decoupling of domain prestige from AI citability suggests that non-commercial research impact management must prioritize content structure over traditional domain authority [2]. This represents a change in the delivery mechanism of the capability, as researchers may alter how they structure digital content to ensure it is “RAG-friendly” [1]. However, core metrics for research impact, such as h-index or formal citations in peer-reviewed journals, are unlikely to change in this period due to their deep embedding in institutional tenure and funding regulations.

Hypothesis: In the near term, researchers utilizing digital dissemination may adapt their visibility strategies to prioritize machine-readability and structural verifiability over traditional domain authority to increase their presence within generative answer engines.

Mid (2036–2041)

For the same subset of researchers focused on web-based dissemination, if current trajectories hold, the process of non-commercial research impact management could transition toward formalized GEO standards. Extrapolating from the requirement for verifiable sources in RAG-based systems [1], researchers may adopt specific metadata schemas to ensure their work is correctly attributed during LLM synthesis. This could shift the educational substance of “impact” from simple reach (clicks or views) to “synthesized utility,” where the value of research is measured by its role as a foundational component of AI-generated answers. Despite this, the process of commercial research outcomes—such as patent filing and licensing—will likely persist unchanged because these rely on legal protections and proprietary disclosures rather than public discovery layers.

Hypothesis: In the mid term, non-commercial research impact for web-active researchers may be increasingly evaluated based on the frequency and accuracy of their work’s synthesis within generative engines.

Far (2046–2051)

For researchers using digital discovery layers, the capability of managing non-commercial research impact remains subject to the tension between algorithmic visibility and institutional trust. One plausible outcome is that the retrieval set becomes the primary arbiter of academic influence, effectively automating the discovery process for a large portion of non-commercial impact. Conversely, a rival outcome may emerge if the over-optimization of research for AI synthesis leads to a counterproductive erosion of trust, where the “gaming” of GEO undermines the perceived validity of the research itself. It remains unknown whether funding bodies will eventually replace traditional citation counts with generative synthesis metrics. The requirement for human-peer-validated novelty will likely resist change even at this horizon, as it provides the social legitimacy that algorithmic synthesis cannot generate.

Hypothesis A: In the far term, digital research impact may be primarily mediated through the precision and frequency with which a scholar’s work is retrieved and synthesized by autonomous agents.

Hypothesis B (competing): In the far term, institutional distrust in AI-mediated visibility may lead to a re-valuation of gated, non-algorithmic verification processes to ensure research impact is not artificially inflated via GEO.

References

1. Dzone Ai (2026-04-06). “Engineering for the Answer Engine: GEO for RAG-Friendly Web Apps (TalentHacked.com Case Study).” <https://feeds.dzone.com/link/23558/17314715/engineering-for-the-answer-engine-geo-for-rag-frie>
2. Dev To Ai (2026-04-10). “Why DA Is Irrelevant for AI Citations (Data from 7 Site Audits).” https://dev.to/chudi_nnorukam/why-da-is-irrelevant-for-ai-citations-data-from-7-site-audits-1il0

Research Improvement

capability_id: research-improvement · created: 2026-07-16 · revised: 2026-07-16 · refs: 5

Near (2026–2031)

For a subset of research quality management involving biology research, behavioral simulations, machine learning (ML) codebase development, and scientific drafting, as well as researcher performance management in AI-assisted computational fields, the introduction of specialized benchmarks and the identification of model failure modes are altering validation protocols. In biology research, the use of standardized benchmarks such as LABBench2 shifts the delivery mechanism of quality management toward automated performance auditing [2]. In behavioral simulation, the disconnect between the visual plausibility of large language model (LLM) outputs and their actual causal accuracy necessitates a shift in educational substance, where research quality management must prioritize causal verification over output coherence [1]. For ML codebases, the risk of “research sabotage”—where misaligned models may perform sloppy research or fabricate safety signals—introduces a requirement for adversarial auditing within the research quality management process [4]. Additionally, the deployment of systems like GoodPoint to generate constructive feedback based on author responses alters the quality management process during the scientific drafting phase [3]. In the process of researcher performance management for computational research, the identification of a “quality trough” suggests that throughput-based metrics may obscure systemic errors, necessitating a shift toward measuring rework rates and code survival rather than simple time-to-completion [5]. Broad institutional tenure requirements based on h-index or total citation counts are unlikely to change in this horizon, as these remain the primary lagging indicators of social legitimacy.

Hypothesis: In the near term, research quality management for specific AI-driven subsegments may shift toward requiring standardized benchmark validation, explicit causal verification, and adversarial auditing, while performance management for computational researchers may begin to prioritize stability metrics over throughput.

Mid (2036–2041)

For the same subsegments of biology, behavioral, and ML research, if current trajectories toward standardized benchmarking and causal auditing hold, research quality management may integrate these technical audits as mandatory prerequisites for publication [1][2]. Extrapolating from the emergence of LABBench2 and the identified risks of model sabotage, the process of research quality management could transition from human-led peer review of final results to a technical verification of the AI’s operational logic, causal architecture, and internal consistency [2][4]. In researcher performance management, if the “quality trough” effect is widely acknowledged, evaluation frameworks may move from valuing volume of output to valuing the “survival rate” of AI-generated research artifacts [5]. This shift moves the point of intervention in quality management from the end of the research cycle to the tool-selection and execution phases. The reliance on institutional review boards for final methodological approval is likely to persist, as these bodies provide the social legitimacy required for high-stakes scientific claims.

Hypothesis: In the mid term, the validation of AI-driven research in these subsegments may evolve into a formal technical auditing process that separates surface-level plausibility from structural and causal validity.

Far (2046–2051)

For researchers utilizing synthetic behavioral agents and AI-driven biological discovery, the stability of research quality management depends on whether the gap between plausibility and causal accuracy is fundamentally resolved in model architectures [1]. One plausible outcome involves a specialized regime of “synthetic provenance,” where AI-generated data is strictly tiered by its verified causal reliability according to evolving benchmarks [2]. A rival outcome suggests a systemic devaluation of AI-simulated data for high-stakes causal claims if persistent inaccuracies or sabotage risks cannot be architecturally mitigated, leading to a renewed institutional preference for empirical human data [1][4]. It remains unknown whether future architectures can inherently solve the plausibility-accuracy gap or if external, independent verification will

always be a requirement. The fundamental necessity for third-party, independent verification of high-stakes scientific claims resists change even at this horizon.

Hypothesis A: In the far term, research quality management may manifest as a structured tiering of AI-generated research artifacts based on verified causal reliability.

Hypothesis B (competing): In the far term, institutional trust in synthetic data for high-stakes causal claims may collapse if sabotage risks and plausibility-accuracy gaps persist, causing a return to strictly empirical human-led data collection.

References

1. Arxiv Cs Cy (2026). “When simulations look right but causal effects go wrong: Large language models as behavioral simulators.” <https://arxiv.org/abs/2604.02458>
2. Arxiv Cs Ai (2026). “LABBench2: An Improved Benchmark for AI Systems Performing Biology Research.” <https://arxiv.org/abs/2604.09554>
3. Arxiv Cs Ai (2026). “GoodPoint: Learning Constructive Scientific Paper Feedback from Author Responses.” <https://arxiv.org/abs/2604.11924>
4. Lesswrong (2026). “Research Sabotage in ML Codebases.” <https://www.lesswrong.com/posts/LByP4qsF8a4g7Pz3p/research-sabotage-in-ml-codebases>
5. Dev To Ai (2026). “The J-Curve.” <https://dev.to/thesynthesis/the-j-curve-32hn>

Research Opportunity Analysis

capability_id: research-opportunity-analysis · created: 2026-05-28 · revised: 2026-05-28 · refs: 2

Near (2026–2031)

For research-intensive environments broadly, the process of research prioritisation may shift toward automated gap detection through the use of parallel AI research agents to map knowledge voids across scientific corpora [1]. This represents a change in the delivery mechanism of research prioritisation, where agentic scanning replaces some manual literature synthesis. Simultaneously, for institutions in regions implementing national-level frontier AI partnerships, such as the Republic of Korea, collaboration opportunity management and research programme development may be accelerated by integrating frontier AI models into strategic scientific planning [2]. Because these tools identify informational voids rather than strategic or social value, the final selection of research directions will likely remain a human-led activity. The social legitimacy of a research agenda depends on expert human validation, making a fully automated prioritisation process unlikely in this horizon.

Hypothesis: AI agents will support the initial stage of knowledge gap identification and the acceleration of strategic collaboration in the near term.

Mid (2036–2041)

For institutions implementing agentic workflows in their research offices, if current trajectories in automated mapping hold, these tools could alter the substance of research programme development. Extrapolating from current demonstrations of parallel agentic gap analysis [1], the identification of “white space” in scientific literature may become a standardized input for designing multi-year research programmes. This shift may move programme development away from intuition-based hypothesis generation toward a model focused on data-driven novelty detection. However, the requirement for peer-reviewed validation of a research gap’s significance will likely persist, as the consensus required for academic legitimacy remains a human social process.

Hypothesis: Research prioritisation may evolve into a hybrid model where AI agents propose potential knowledge gaps for human curation in the mid term.

Far (2046–2051)

For research-intensive environments, the long-term impact on research opportunity analysis remains uncertain, specifically regarding whether AI can identify conceptual paradigm shifts or merely informational gaps. One possibility is that automated gap mapping becomes the primary driver of research agendas, potentially leading to a counterproductive cycle where researchers pursue a high volume of superficially novel but low-impact “gaps” identified by AI. Conversely, a reaction against this automation may occur, where the value of theory-driven rather than gap-driven research is re-emphasized to maintain intellectual substance. The physical constraint of empirical validation—the need to conduct actual experiments or observations—will resist automation even at this horizon.

Hypothesis A: AI-driven gap analysis becomes the primary mechanism for research prioritisation in the far term.

Hypothesis B (competing): The proliferation of AI-identified gaps leads to a crisis of significance, triggering a return to human-centric, theory-led research prioritisation in the far term.

References

1. Dev To Ai (2026-04-10). “We mapped 153 gaps in science using 5 parallel AI research agents.” <https://dev.to/minimal-architect/we-mapped-153-gaps-in-science-using-5-parallel-ai-research-agents-5fdk>
2. Google DeepMind (2026-04-27). “Announcing our partnership with the Republic of Korea.” <https://deepmind.google/blog/announcing-our-partnership-with-the-republic-of-korea/>

Research Publications

capability_id: research-publications · created: 2026-04-11 · revised: 2026-04-11 · refs: 7

Near (2026–2031)

For research publications broadly, research publication management is increasingly defined by a tension between automated drafting and the deployment of provenance-tracking infrastructure. Publishers are employing AI-detection tools to maintain the integrity of submitted research outputs [1], while the emergence of AI citation registries suggests a shift toward using structured information infrastructure to verify the authenticity of references [2]. For a subsegment of researchers, the delivery mechanism of research publication management is shifting toward the use of multi-agent frameworks, such as PaperOrchestra, to automate the drafting of manuscripts [5]. Regarding research output reporting, a subsegment of researchers is adapting output formatting—utilizing direct Q&A and tables—to increase visibility within AI search engines [6]. This shift is supported by the development of AI visibility infrastructure designed to interface directly with AI systems [4] and evidence that AI citations are driven by specific “citability” factors rather than traditional domain authority (DA) [7]. In institution-specific settings, such as the Alignment Journal, experimental policies are introducing reviewer compensation to modify the peer-review workflow [3]. The social legitimacy of the human-led peer-review process is unlikely to change in this window because institutional tenure and promotion regimes rely on human intellectual accountability.

Hypothesis: In the near term, research publication management will be characterized by a technical arms race between multi-agent drafting tools and the infrastructure required to detect and verify AI-generated content.

Mid (2036–2041)

For research publications broadly, if current trajectories in registry adoption and detection instability hold, the mechanism of verification may shift from reactive detection to proactive provenance. Extrapolating from the limitations of binary AI-text detection [1], publishers may require research output reporting to be anchored in verified citation registries [2] to prevent the proliferation of fabricated evidence. Research output reporting may further diverge from the static PDF toward structured entity formats to maintain discoverability in AI-driven knowledge graphs [4][6]. The journal article is likely to persist as the primary unit of currency for academic promotion due to high institutional inertia in tenure committees. However, the delivery mechanism of the peer-review process may integrate structured transparency logs to counteract the opacity of AI-assisted drafting [5].

Hypothesis: In the mid term, the focus of research publication management may move from the detection of synthetic text to the systematic verification of authorship and citation provenance through structured registries.

Far (2046–2051)

For research publications broadly, the long-term viability of the written manuscript as the primary research output remains unknown if the gap between automated generation [5] and detection [1] closes completely. One plausible outcome is a transition where the primary research output becomes a verifiable computational model or a registry-linked data stream [2], with written text serving only as a secondary, AI-generated summary for human consumption [4]. A rival outcome involves the erosion of trust in traditional publication management, leading to a fragmentation of the reward system as traditional journals fail to distinguish human discovery from synthetic synthesis. It remains unknown whether visibility infrastructure [4] can scale sufficiently to maintain a global, trusted record of discovery. The fundamental requirement for a verifiable and immutable record of intellectual priority resists change even at this horizon.

Hypothesis A: In the far term, the primary research output may shift from the written manuscript to verifiable computational models or data streams linked to citation registries.

Hypothesis B (competing): In the far term, the inability to distinguish human from synthetic discovery

may lead to a collapse of the traditional journal system and a fragmented, localized reward structure for research.

References

1. Dev To Ai (2026-04-06). “Will AI-Written Papers Be Detected? The Truth About AIGC Detection in 2026.” https://dev.to/yu_z_ed6c1c576593c5e7d41f/will-ai-written-papers-be-detected-the-truth-about-aigc-detection-in-2026-5b50
2. Dev To Ai (2026-04-06). “AI Citation Registries as Information Infrastructure for AI Systems.” <https://dev.to/aigistry/ai-citation-registries-as-information-infrastructure-for-ai-systems-50o6>
3. Lesswrong (2026-04-07). “An Alignment Journal: Features and policies.” <https://www.lesswrong.com/posts/hvq7amw8FKZxEeKqA/an-alignment-journal-features-and-policies>
4. Dev To Ai (2026-04-08). “Building AI Visibility Infrastructure: Inside Jonomor’s Architecture.” https://dev.to/jonomor_ecosystem/building-ai-visibility-infrastructure-inside-jonomors-architecture-3enp
5. Marktechpost (2026-04-09). “Google AI Research Introduces PaperOrchestra: A Multi-Agent Framework for Automated AI Research Paper Writing.” <https://www.marktechpost.com/2026/04/08/google-ai-research-introduces-paperorchestra-a-multi-agent-framework-for-automated-ai-research-paper-writing/>
6. Dev To Ai (2026-04-10). “How to Get Cited by AI Search Engines: ChatGPT, Perplexity & Claude in 2026.” <https://dev.to/iris1031/how-to-get-cited-by-ai-search-engines-chatgpt-perplexity-claude-in-2026-21p0>
7. Dev To Ai (2026-04-10). “Why DA Is Irrelevant for AI Citations (Data from 7 Site Audits).” https://dev.to/chudi_nnorukam/why-da-is-irrelevant-for-ai-citations-data-from-7-site-audits-1il0

Research Training

capability_id: research-training · created: 2026-06-16 · revised: 2026-06-16 · refs: 9

Near (2026–2031)

For research training broadly, the researcher training and development process may integrate AI-driven search and deep research tools to automate source identification and the generation of structured literature syntheses [3][4]. This modification changes the delivery mechanism of the literature review process rather than its educational substance. For the subsegment of researchers engaged in quantitative analysis, postgraduate research student skills development may incorporate AI-driven data analysis tools to automate dataset exploration and initial visualization [5]. Within the specific subsegment of machine learning research, student skills development may shift from a primary focus on theoretical mathematical derivation toward compute-intensive engineering and the management of large-scale empirical experiments [1][9]. This shift represents a change in educational substance, as research progress increasingly relies on compute scaling over mathematically principled architectures [1][9]. For technical research subsegments requiring software development, skills development may pivot from teaching manual coding syntax to the orchestration of agentic AI tools and terminal-based AI agents [2][7]. In the subsegment of scientific writing and peer review training, the development of researcher skills may incorporate AI systems designed to generate constructive feedback based on author responses [6]. Within the subsegment of formal mathematics, a debate is emerging regarding the counterproductivity of AI-automated proofs, with arguments that automating the process of proving may erode the acquisition of essential mathematical intuition [8]. The institutional requirement for a written thesis and a formal oral defense is unlikely to change in the near term due to the high inertia of degree-awarding regulations.

Hypothesis: In the near term, general research training will integrate AI-driven sourcing and synthesis, while technical subsegments will pivot toward compute engineering and agentic orchestration, alongside an emerging tension between automation and intuition in formal disciplines.

Mid (2036–2041)

For the same broad and technical cohorts, the postgraduate research student skills development process may transition from a focus on information retrieval to the rigorous validation of AI-curated syntheses if current trajectories hold [3][4]. Extrapolating from the adoption of unified agentic systems, researcher training in technical subsegments may replace manual codebase maintenance with the orchestration of autonomous coding agents as a primary technical competency [2][7]. This shift may alter the substance of postgraduate research supervisor development, as supervisors may move from guiding manual methodology to validating agent-generated empirical workflows. The necessity for the researcher to frame a novel research question and maintain an original theoretical contribution is likely to persist unchanged because these tasks require high-level conceptual synthesis that current agentic trajectories do not replace.

Hypothesis: In the mid term, research training will shift toward the validation of automated synthesis and the orchestration of agentic systems, shifting supervisory focus toward workflow validation.

Far (2046–2051)

For machine learning and technical research subsegments, the long-term trajectory of skills development depends on whether compute-intensive scaling continues to yield returns or reaches a threshold of diminishing returns [1][9]. If scaling remains the primary driver of discovery, research training may evolve into a specialized discipline of high-performance infrastructure management and agent orchestration. Conversely, if scaling reaches a ceiling, there may be a resurgence in the value of mathematically principled architectures, requiring a pivot back to theoretical foundations in student skills development. The ability of a human researcher to provide social legitimacy and ethical accountability for research findings remains an unresolvable uncertainty that resists automation even at this horizon.

Hypothesis A: In the far term, research training in technical fields evolves into a discipline focused on the

management of massive compute infrastructure and agentic ecosystems.

Hypothesis B (competing): In the far term, a plateau in compute-driven gains forces a return to theoretical and mathematically principled training as the primary driver of research breakthrough.

References

1. The Gradient (2024). “Shape, Symmetries, and Structure: The Changing Role of Mathematics in Machine Learning Research.” <https://thegradient.pub/shape-symmetry-structure/>
2. Dev To Ai (2026). “Claude Code Cheatsheet: The Terminal AI Agent That Actually Codes.” <https://dev.to/mibii/claude-code-cheatsheet-the-terminal-ai-agent-that-actually-codes-36bk>
3. OpenAI (2026). “ChatGPT for research.” <https://openai.com/academy/research>
4. OpenAI (2026). “Research with ChatGPT.” <https://openai.com/academy/search-and-deep-research>
5. OpenAI (2026). “Analyzing data with ChatGPT.” <https://openai.com/academy/data-analysis>
6. Arxiv Cs Ai (2026). “GoodPoint: Learning Constructive Scientific Paper Feedback from Author Responses.” <https://arxiv.org/abs/2604.11924>
7. Simon Willison (2026). “Quoting Romain Huet.” <https://simonwillison.net/2026/Apr/25/romain-huet/#atom-everything>
8. Lesswrong (2026). “The fall of the theorem economy (David Bessis).” <https://www.lesswrong.com/posts/5CA4z7TumhrQu5Jr4/the-fall-of-the-theorem-economy-david-bessis>
9. The Gradient (2024). “Shape, Symmetries, and Structure: The Changing Role of Mathematics in Machine Learning Research.” <https://thegradient.pub/shape-symmetry-structure/>

Strategy Management

capability_id: strategy-management · created: 2026-06-24 · revised: 2026-06-24 · refs: 12

Near (2026–2031)

For institutions in China, public universities subject to mandated viability reviews, and a subsegment of institutions adopting AI-driven synthesis and forecasting tools, strategy management is shifting toward a combination of external mandate alignment and tool-assisted synthesis. In China, national targets for AI deployment are integrated into strategic plan development and strategic plan management via five-year planning cycles [1]. For state-funded research bodies, as demonstrated by the Alan Turing Institute, funding agencies may mandate changes to strategic plan development to ensure adherence to legal duties and value-for-money metrics [2]. In public institutions facing budget pressures, such as Iowa State University and East Carolina University, corporate performance management is increasingly driven by mandated reviews that use viability data to execute program closures or mergers [6][10]. For institutions using automated synthesis platforms, the processes of strategic reporting and business horizon scanning may be streamlined through tools that generate competitive intelligence [3][4][5] or specialized labor automation and AGI scenario analysis [7][11]. Simultaneously, some executives are moving toward systemic AI integration to shift strategic plan development from siloed pilots to business-wide value [12], accompanied by a push to replace routine cognitive KPIs with metrics reflecting human-centric value [8] and intent-driven budgeting [9]. These changes represent delivery-mechanism changes in strategic reporting and horizon scanning, while shifting the educational substance of corporate performance management toward fiscal viability and redefined value metrics [6][8][10]. The fundamental legal fiduciary duties of governing boards and the statutory requirements of university charters are unlikely to change in this period.

Hypothesis: In the near term, strategy management in these settings may be increasingly defined by the transition of AI from a technical project to an executive mandate, utilizing automated synthesis to accelerate reporting and the redefinition of performance metrics.

Mid (2036–2041)

For institutions in China, state-funded research bodies, public universities under state oversight, and adopters of AI strategy tools, if current trajectories hold, corporate performance management may shift toward the quantitative measurement of external benchmarks and AI-generated KPIs. Extrapolating from state mandates [1] and mandated program reviews [6][10], strategic reporting could move from an internal reflection of institutional mission to a compliance verification tool, altering the educational substance of how institutional success is defined. For institutions using automated strategy platforms, the process of strategic plan development may shift from an act of primary synthesis to an act of auditing AI-generated outputs [3][4][5]. If the shift toward meaning-based KPIs persists [8], strategic vision development may diverge between institutions that prioritize AI-driven efficiency and those that prioritize human-centric value creation. The requirement for institutions to maintain a distinct institutional identity within their strategic vision development is likely to persist, though it will operate within boundaries set by external regulatory authorities and tool constraints.

Hypothesis: In the mid term, corporate performance management in these segments may transition into a process of algorithmic compliance and output auditing rather than primary strategic synthesis.

Far (2046–2051)

For the same subset of institutions identified in the near term, the capability of strategy management may either converge toward a standardized, AI-steered model or fragment into highly localized, human-led missions. A primary unknown is the social legitimacy of AI-led institutional steering and whether governing boards will delegate the core of strategic vision development to bounded autonomous systems. One outcome involves the full integration of labor automation forecasting [7] and AGI scenario analysis [11] into a continuous, real-time strategic plan management loop that eliminates the traditional multi-year planning cycle. Conversely, a counter-reaction may occur where institutions explicitly reject AI-driven

corporate performance management to preserve a perceived “human” institutional essence. Even at this horizon, the basic requirement for institutions to provide accountability to a public or private benefactor is unlikely to change.

Hypothesis A: In the far term, strategy management may evolve into a continuous, real-time optimization process where the distinction between strategic planning and daily operational management disappears.

Hypothesis B (competing): In the far term, strategy management may return to a highly artisanal, human-centric process as a means of maintaining institutional distinction and social legitimacy in an automated environment.

References

1. Ai News (2026-04-02). “China’s Five-Year Plan details the targets for AI deployment.” <https://www.artificialintelligence-news.com/news/chinas-five-year-plan-details-the-targets-for-ai-deployment/>
2. Guardian Tech (2026-04-03). “UK’s leading AI research institute told to make ‘significant’ changes.” <https://www.theguardian.com/technology/2026/apr/03/uk-alan-turing-institute-ai-research-underperforming-ukri>
3. TechCrunch (2026-04-07). “AI startup Rocket offers vibe McKinsey-style reports at a fraction of the cost.” <https://techcrunch.com/2026/04/06/indian-startup-rocket-wants-its-ai-to-do-mckinsey-style-consulting-at-a-fraction-of-the-cost/>
4. Dev To Ai (2026-04-07). “AI startup Rocket offers vibe McKinsey-style reports at a fraction of the cost.” <https://dev.to/minimal-architect/ai-startup-rocket-offers-vibe-mckinsey-style-reports-at-a-fraction-of-the-cost-1edd>
5. Mit Sloan (2026-04-08). “Gain Consumer Insight With Generative AI.” <https://sloanreview.mit.edu/article/gain-consumer-insight-with-generative-ai/>
6. Higher Ed Dive (2026-04-10). “Iowa State plans to cut or merge 23 programs after mandated review.” <https://www.highereddive.com/news/iowa-state-university-program-cuts-review-regents/817262/>
7. Lesswrong (2026-04-21). “Introducing the Labor Automation Forecasting Hub.” <https://www.lesswrong.com/posts/So7wrd6djpQcoP4rz/introducing-the-labor-automation-forecasting-hub-1>
8. Fast Company (2026-04-23). “The Age of AI means we need to throw out our old KPIs and replace them with new ones.” https://www.fastcompany.com/91521422/ai-kpis-new-metrics?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
9. Dzone Ai (2026-04-24). “88% of Executives Plan to Increase AI Budgets by 2026 - Are We Investing With Intent or Just Urgency?” <https://feeds.dzone.com/link/23558/17325011/increasing-ai-budgets-investing-intent-or-urgency>
10. Higher Ed Dive (2026-04-27). “East Carolina University plans to cut 44 academic programs.” <https://www.highereddive.com/news/east-carolina-university-cut-around-44-academic-programs/818622/>
11. Arxiv Cs Cy (2026-04-28). “Artificial General Intelligence Forecasting and Scenario Analysis: State of the Field, Methodological Gaps, and Strategic Implications.” <https://arxiv.org/abs/2604.22766>
12. Zdnet Ai (2026-05-01). “Moving from AI pilots to business-wide value requires a superhighway - how to ramp up.” <https://www.zdnet.com/article/moving-from-ai-pilots-to-business-wide-value-accenture-research/>

Student Administration

capability_id: student-administration · created: 2026-06-17 · revised: 2026-06-17 · refs: 2

Near (2026–2031)

For a subset of institutions, including public universities subject to state-mandated program reviews [1] and those managing complex financial aid communications [2], administrative processes are facing increased volatility and regulatory pressure. For universities undergoing mandated closures of low-enrollment programs, as seen at Iowa State University [1], the process of programme transfer management may experience surges in volume, potentially driving the adoption of AI-based credit mapping to accelerate student reassignment. Simultaneously, legislative efforts to resolve jargon-heavy financial aid offers [2] may prompt student financial administration to implement AI-driven transparency tools to translate “total net expenses” into concrete out-of-pocket costs for families. These changes represent shifts in the delivery mechanism—specifically the speed of transfer processing and the clarity of financial communication—rather than changes to the educational substance. Enrolment status management will likely remain a manual, human-verified process to ensure legal compliance with state and federal funding mandates. It is unlikely that the core regulatory requirements for degree conferral will change, as these are anchored by external accreditation bodies.

Hypothesis: In the near term, for institutions facing programmatic volatility or legislative pressure on financial transparency, the delivery mechanisms of programme transfer management and student financial administration may shift toward AI-supported automation and translation.

Mid (2036–2041)

For institutions facing these specific regulatory and legislative pressures, if current trajectories of mandated program reviews and financial aid transparency hold, the load on programme transfer management could lead to a shift from AI-supported to AI-led automated matching. Extrapolating from current administrative constraints, such systems might automatically suggest alternative degree paths for displaced students based on completed credits and labor market demand. Within student financial administration, the requirement for standardized, jargon-free communication may lead to the adoption of algorithmic cost-projection tools that automate the calculation of individual family liability [2]. However, the final validation of enrolment status management will likely persist as a human-governed process to prevent fraud and ensure adherence to government funding audits.

Hypothesis: In the mid term, for institutions experiencing frequent program restructuring or strict financial reporting mandates, the delivery mechanism of student administration may move toward algorithmic decision-support for transfers and cost projections.

Far (2046–2051)

For institutions operating under these regimes of programmatic volatility and high financial transparency, one plausible outcome is that programme transfer management becomes a dynamic process that re-aligns student enrolments in real-time based on shifting state mandates and employment signals. A rival outcome is that the persistent instability of specific program offerings leads to a shift toward broader, degree-agnostic enrolment status management, which would reduce the frequency and necessity of specific programme transfer events. It remains unknown whether state regents will continue to maintain the authority to mandate such cuts or if institutional autonomy over program design will be restored. The legal requirement for permanent, immutable student records & details management will likely resist change even at this horizon due to the long-term necessity of degree verification.

Hypothesis A: In the far term, for public universities in volatile regulatory environments, programme transfer management may be integrated into dynamic, AI-driven labor-market alignment systems.

Hypothesis B (competing): In the far term, the administrative burden of frequent program cuts may drive a shift toward modular, degree-agnostic enrolment status management, reducing the operational reliance on

traditional programme transfer processes.

References

1. Higher Ed Dive (2026). “Iowa State plans to cut or merge 23 programs after mandated review.” <https://www.highereddive.com/news/iowa-state-university-program-cuts-review-regents/817262/>
2. The Hechinger Report (2026). “Confusing financial aid offers can leave families deeper in debt. Student groups say a new legislative fix doesn’t go far enough.” <https://hechingerreport.org/confusing-financial-aid-offers-can-leave-families-deeper-in-debt-a-fix-is-stalled-in-congress/>

Student Admission Management

capability_id: student-admission-management · created: 2026-06-29 · revised: 2026-06-29 · refs: 3

Near (2026–2031)

For a subsegment of institutions characterized by opaque financial aid communications and reliance on digital identity verification, two distinct processes within this capability are affected. Because AI-generated imagery now enables the creation of high-quality fraudulent identification [4], the identity verification process within student application management is undermined, shifting the burden of proof toward more robust authentication mechanisms. Simultaneously, because financial aid offers often employ ambiguous terminology that obscures the actual cost of attendance [1], prospective students may utilize large language models to normalize and compare disparate offer formats [2]. This represents a change in the delivery mechanism of the offer, acceptance & quota management process, as the interpretation of the offer moves from an institutionally controlled narrative to an AI-mediated analysis. The legal requirement for a formally signed and dated offer of admission is unlikely to change in the near term due to the contractual nature of enrollment.

Hypothesis: In the near term, for a subsegment of institutions, AI may simultaneously undermine identity verification in application management while increasing information symmetry for students decoding financial aid offers.

Mid (2036–2041)

For the same subsegment of institutions, if current trajectories hold, some may employ AI-driven personalization to optimize quota management by tailoring the presentation of aid to specific student psychographics. This would represent a change in the educational substance of the capability, as the logic of the offer shifts from transparent disclosure to yield optimization. Such a mechanism increases the risk of counterproductivity, where the pursuit of higher yield rates through calculated opacity undermines the social legitimacy of the admission process. To counter the proliferation of AI-generated fraudulent credentials, institutions may shift identity verification from static document review to continuous or biometric authentication. Despite these shifts, the core regulatory requirements for federal financial aid reporting will likely remain a stabilizing constraint.

Hypothesis: In the mid term, some institutions may use AI to personalize financial aid offers to optimize yield, while shifting application management toward biometric identity verification to counter AI-generated fraud.

Far (2046–2051)

For this subsegment of institutions, two rival trajectories emerge regarding the delivery of admission offers. One outcome involves a regulatory response to AI-driven opacity, mandating a standardized, machine-readable format for all financial aid offers to ensure comparability and prevent predatory yield optimization. A competing outcome involves the obsolescence of the “offer letter” as a discrete artifact, replaced by real-time, AI-managed cost-of-attendance dashboards that update based on the student’s shifting financial profile. It remains unknown whether accreditation bodies will eventually categorize AI-optimized personalized pricing as a violation of equitable access standards. The fundamental requirement for a legally binding agreement between the student and the institution will likely resist change.

Hypothesis A: In the far term, regulatory mandates may force a transition to standardized, AI-verifiable financial aid disclosures to counter institutional opacity.

Hypothesis B (competing): In the far term, the traditional offer letter may be replaced by dynamic, AI-managed financial transparency dashboards providing real-time cost calculations.

References

1. The Hechinger Report (2026). “How much will that college cost you? Good luck figuring it out.” <https://hechingerreport.org/how-much-will-that-college-cost-you-good-luck-figuring-it-out-2/>
2. The Hechinger Report (2026). “Confusing financial aid offers can leave families deeper in debt. Student groups say a new legislative fix doesn’t go far enough.” <https://hechingerreport.org/confusing-financial-aid-offers-can-leave-families-deeper-in-debt-a-fix-is-stalled-in-congress/>
3. Atlantic (2026). “Hey Chat, Make Me a Fake ID.” https://www.theatlantic.com/technology/2026/05/chatgpt-images-deepfakes-fraud/687023/?utm_source=feed

Student Assessment

capability_id: student-assessment · created: 2026-06-27 · revised: 2026-06-27 · refs: 19

Near (2026–2031)

For student assessment broadly, the delivery mechanism is shifting toward iterative human-AI writing cycles [1], a transition that exposes structural flaws in the reliance on essay assessments within mass education systems [3] and introduces the risk of “cognitive surrender,” where students may abandon logical thinking in favor of synthetic output [5]. This change in delivery may lead to a homogenization of student writing and thinking patterns, which complicates the process of assessment marking and feedback by eroding the distinct critical voice [9]. In the process of assessment administration, the reliability of authorship verification is constrained by the persistence of false positives in AI detectors [13], the continuous decay of detection pipelines as models evolve [28], and the emergence of “humanizer” tools that erode the utility of sentence-level detection [15]. For institutions implementing automated marking and feedback, the process is subject to algorithmic bias, with evidence that AI may provide disparate levels of praise and criticism based on student race [27]. In specific subsegments, such as mathematics and music education, assessment delivery is being adapted through multi-agent validation for problem generation [11] and AI-driven tracking to map vague goals into measurable progress [30], while in software development, marking is shifting from functional testing toward the qualitative evaluation of architecture and security [33]. High-stakes, invigilated examinations are unlikely to change in the near term due to professional accreditation requirements and the institutional need for authenticated baseline competency [19].

Hypothesis: In the near term, assessment delivery may shift from artifact-production to process-evaluation, while assessment administration is increasingly constrained by the technical failure of authorship detection tools.

Mid (2036–2041)

For student assessment broadly, if current trajectories hold, the educational substance of assessment delivery may shift from evaluating a final product to evaluating a student’s capacity to audit, verify, and correct synthetic errors [5][7]. This represents a shift in substance from measuring output to measuring the ability to steer and validate synthetic agents [14], particularly as AI agents become capable of resolving complex professional challenges [29]. To maintain the validity of assessment results management, institutions may implement deliverable-oriented governance frameworks to manage the inherent opacity of these AI systems [20]. For a subset of institutions focusing on multimodal pedagogical analysis, results management could incorporate the longitudinal mapping of cognitive effort via educational fMRIs [4]. The basic administrative recording of credits and formal result transmission will likely persist unchanged due to high regulatory inertia.

Hypothesis: In the mid term, the substance of student assessment may shift from measuring product output to measuring a student’s ability to audit and validate synthetic agents.

Far (2046–2051)

For student assessment broadly, the long-term trajectory depends on whether synthetic protocols can reliably replace human judgment in certifying higher-order cognitive capabilities [18]. It remains unknown if LLM-based protocols can meaningfully measure collaboration, creativity, and critical thinking without reducing these capabilities to observable patterns [17]. A core uncertainty remains regarding whether the social legitimacy of a degree continues to depend on the verification of biological cognition or the mastery of synthetic orchestration. The fundamental institutional requirement for a verified credential of competency will likely resist change even at this horizon, as the labor market requires a signal of individual reliability.

Hypothesis A: In the far term, student assessment may be fully mediated by LLM-based protocols that certify cognitive depth through the analysis of iterative interaction patterns.

Hypothesis B (competing): In the far term, the failure of synthetic verification may drive a return to strictly

offline, analog, and synchronous assessment to ensure the presence of biological cognition.

References

1. The Conversation (2026-03-26). "College students are writing with AI – but a pilot study finds they're not simply letting it write for them." <https://theconversation.com/college-students-are-writing-with-ai-but-a-pilot-study-finds-theyre-not-simply-letting-it-write-for-them-276856>
2. The Conversation Au (2026-04-01). "Does AI mean more uni students are plagiarising their work?" <https://theconversation.com/does-ai-mean-more-uni-students-are-plagiarising-their-work-279565>
3. Wonkhe (2026-04-01). "Generative AI exposes longstanding flaws in the use of essay assessments in a mass system." <https://wonkhe.com/blogs/generative-ai-exposes-longstanding-flaws-in-the-use-of-essay-assessments-in-a-mass-system/>
4. Getting Smart (2026-04-01). "Educational fMRIs: Dynamic Pedagogy and Pedagogical Analysis in the Multimodal AI Era." <https://www.gettingsmart.com/2026/04/01/educational-fmris-dynamic-pedagogy-and-pedagogical-analysis-in-the-multimodal-ai-era/>
5. Ars Technica (2026-04-03). "'Cognitive surrender' leads AI users to abandon logical thinking, research finds." <https://arstechnica.com/ai/2026/04/research-finds-ai-users-scarily-willing-to-surrender-their-cognition-to-llms/>
6. Dev To Ai (2026-04-06). "Will AI-Written Papers Be Detected? The Truth About AIGC Detection in 2026." https://dev.to/yu_z_ed6c1c576593c5e7d41f/will-ai-written-papers-be-detected-the-truth-about-aigc-detection-in-2026-5b50
7. Dev To Ai (2026-04-07). "Where Does Truth Live in AI-Generated Code?" <https://dev.to/wintrov er/where-does-truth-live-in-ai-generated-code-5hmc>
8. BBC News (2026-04-07). "'Our teachers use AI to mark mock exams'" https://www.bbc.com/news/articles/c6207xenj1do?at_medium=RSS&at_campaign=rss
9. Hacker News (2026-04-07). "AI may be making us think and write more alike." <https://dornsife.usc.edu/news/stories/ai-may-be-making-us-think-and-write-more-alike/>
10. Plos One (2026-04-07). "Correction: Evaluating cognitive depth of AI-generated multiple-choice questions with Bloom's Taxonomy." <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0346800>
11. Arxiv Cs Cy (2026-04-08). "A Multi-Agent Approach to Validate and Refine LLM-Generated Personalized Math Problems." <https://arxiv.org/abs/2604.05160>
12. Dev To Ai (2026-04-08). "Your article is well-written. But is it yours?" <https://dev.to/dzhuneyt/your-article-is-well-written-but-is-it-yours-59kh>
13. Dev To Ai (2026-04-09). "Why AI Detectors Produce False Positives: A Technical Analysis." <https://dev.to/laakash/why-ai-detectors-produce-false-positives-a-technical-analysis-2gpc>
14. Atlantic (2026-04-10). "This Is What Fully Automated School Looks Like." https://www.theatlantic.com/technology/2026/04/ai-agents-school-education/686754/?utm_source=feed
15. Dev To Ai (2026-04-11). "I Built a Free AI Detector + Humanizer with Sentence-Level Highlighting." https://dev.to/fafa_ai/i-built-a-free-ai-detector-humanizer-with-sentence-level-highlighting-51jj
16. Getting Smart (2026-04-13). "A Call to Action for AI to Promote Mathematical Reasoning." <https://www.gettingsmart.com/2026/04/13/a-call-to-action-for-ai-to-promote-mathematical-reasoning/>
17. Marktechpost (2026-04-14). "Google AI Research Proposes Vantage: An LLM-Based Protocol for Measuring Collaboration, Creativity, and Critical Thinking." <https://www.marktechpost.com/2026/04/13/google-ai-research-proposes-vantage-an-llm-based-protocol-for-measuring-collaboration-creativity-and-critical-thinking/>
18. Arxiv Cs Cy (2026-04-14). "Explainability and Certification of AI-Generated Educational Assessments." <https://arxiv.org/abs/2604.09622>
19. University Affairs (2026-04-14). "Reconsidering the value of exams in the age of AI." <https://universityaffairs.ca/career-advice/reconsidering-the-value-of-exams-in-the-age-of-ai/>
20. Arxiv Cs Ai (2026-04-23). "AI to Learn 2.0: A Deliverable-Oriented Governance Framework and Maturity Rubric for Opaque AI in Learning-Intensive Domains." <https://arxiv.org/abs/2604.19751>

21. Arxiv Cs Cy (2026-04-23). “Do Small Language Models Know When They’re Wrong? Confidence-Based Cascade Scoring for Educational Assessment.” <https://arxiv.org/abs/2604.19781>
22. Elearning Industry (2026-04-22). “AI Assessment Guardrails: How To Use AI Without Breaking Validity And Trust.” <https://elearningindustry.com/ai-assessment-guardrails-how-to-use-ai-without-breaking-validity-and-trust>
23. EDUCAUSE Review (2026-04-23). “Teaching and Learning with AI: Three Priorities.” <https://er.educause.edu/podcasts/educause-exchange/teaching-and-learning-with-ai-three-priorities>
24. Simon Willison (2026-04-25). “Quoting Romain Huet.” <https://simonwillison.net/2026/Apr/25/ro-main-huet/#atom-everything>
25. Lesswrong (2026-04-27). “Retrospective on my unsupervised elicitation challenge.” <https://www.lesswrong.com/posts/aJ3NTJ6tM7yc49kQn/retrospective-on-my-unsupervised-elicitation-challenge>
26. Dev To Ai (2026-04-27). “AI Didn’t Make Developers Better. It Made Average Work Easier to Produce.” <https://dev.to/workelate123/ai-didnt-make-developers-better-it-made-average-work-easier-to-produce-hc2>
27. The Hechinger Report (2026-04-27). “AI gives more praise, less criticism to Black students.” <https://hechingerreport.org/proof-points-ai-bias-feedback/>
28. Dev To Ai (2026-04-27). “LLM Drift: Why Your AI Detection Pipeline is Quietly Decaying (Kimi K2 Benchmark).” https://dev.to/thoams_aidetection/llm-drift-why-your-ai-detection-pipeline-is-quietly-decaying-kimi-k2-benchmark-3gml
29. Dev To Ai (2026-04-28). “GitHub Weekly: GPT-5.5 Lands on Copilot, Cloud Agents Speed Up, Actions Gets Smarter.” <https://dev.to/htekdev/github-weekly-gpt-55-lands-on-copilot-cloud-agents-speed-up-actions-gets-smarter-4cn6>
30. Dev To Ai (2026-04-30). “From Vague Goals to Clear Paths: AI for Mapping Musical Progress.” https://dev.to/ken_deng_ai/from-vague-goals-to-clear-paths-ai-for-mapping-musical-progress-1k3f
31. Simon Willison (2026-04-30). “Quoting Andrew Kelley.” <https://simonwillison.net/2026/Apr/30/andrew-kelley/#atom-everything>
32. Arxiv Cs Cy (2026-05-01). “The Impact of LLM Self-Consistency and Reasoning Effort on Automated Scoring Accuracy and Cost.” <https://arxiv.org/abs/2604.26954>
33. Gary Marcus (2026-05-01). ““A model that produces code which compiles and passes the tests it was given is not the same as a model that produces correct, secure, maintainable, well-architected software”.” <https://garymarcus.substack.com/p/a-model-that-produces-code-which>

Student Attraction & Recruitment

capability_id: student-attraction-recruitment · created: 2026-06-20 · revised: 2026-06-20 · refs: 10

Near (2026–2031)

For the capability of student attraction and recruitment broadly, the delivery mechanism for promotional content is shifting as generative AI enables the production of full ad campaigns—including scripts, imagery, and video—in near-real-time at negligible cost [10]. For a subset of institutions focusing on digital discovery, the process of prospective student engagement may shift from keyword-based search engine optimization to generative engine optimization (GEO) to maintain visibility within AI-mediated search results [1][6][9]. Within the specific process of prospective student engagement at recruitment events, institutions may adopt AI to automate the conversion of raw lead notes into narratives [4], execute automated post-event follow-ups [7], and apply instant lead scoring to prioritize prospects [8]. For institutions utilizing multimedia outreach, the production of engagement videos may be streamlined through integrated AI video tools [2]. In specific regions where COVID-era remote learning has created a drain on college enrollment, the process of domestic student recruitment is constrained by a shrinking pool of traditional candidates [3]. Within scholarship and bursary management, prospective student engagement is currently hindered by a lack of transparency regarding the actual cost of attendance [5]. The core fiduciary criteria used to determine scholarship eligibility are unlikely to change in this period due to rigid institutional governance and audit requirements.

Hypothesis: In the near term, institutions may shift recruitment delivery toward high-velocity AI content generation and automated lead-scoring pipelines, while regional enrollment declines and transparency failures in scholarship communication may limit the effectiveness of these tools.

Mid (2036–2041)

For the subset of institutions utilizing AI-driven engagement, if current trajectories hold, the process of prospective student engagement could shift toward the maintenance of structured data schemas optimized for machine readability rather than human-centric web narratives [1][6][9]. Extrapolating from current lead-intelligence deployments, the qualification and scoring of prospective students at recruitment fairs may move from an asynchronous post-event task to a real-time automated pipeline [4][7][8]. The proliferation of low-cost, AI-generated recruitment media may lead to content saturation, potentially shifting international and domestic student recruitment toward a higher reliance on human-led validation to establish institutional authenticity [2][10]. If regional enrollment drains persist, domestic student recruitment may pivot toward alternative lead-generation mechanisms that bypass traditional pipeline markers [3]. To address friction in scholarship and bursary management, institutions may deploy AI-mediated tools to provide personalized, transparent cost projections for prospective families [5]. The requirement for faculty-authored program descriptions will likely persist as a baseline for academic legitimacy.

Hypothesis: In the mid term, prospective student engagement for some institutions may shift toward the production of machine-optimized entity data, while AI may be used to reduce information asymmetry in scholarship and bursary management.

Far (2046–2051)

For institutions relying on AI-mediated prospective student engagement, it remains unknown whether AI discovery platforms will maintain neutral recommendation architectures or shift toward commercialized, pay-to-play visibility models. One possible outcome involves the bounded autonomy of recruitment pipelines, where the process of prospective student engagement is managed by autonomous agents that negotiate fit and financial aid in real-time between the student and the institution. A rival outcome involves a systemic retreat from AI-mediated recruitment due to counterproductivity, where the erosion of institutional prestige caused by automated outreach forces a return to high-touch, human-centric recruitment models. The legal frameworks governing the verification of student identity during the recruitment process will likely resist full automation due to fraud risks.

Hypothesis A: In the far term, the process of prospective student engagement may become a bounded autonomous negotiation between AI agents, shifting the capability toward data-driven matching.

Hypothesis B (competing): In the far term, recruitment may return to high-touch human models as a means of restoring institutional legitimacy and social signaling in an environment of AI content saturation.

References

1. Dev To Ai (2026-04-06). “What is the difference between AI content and AI search optimised content?” <https://dev.to/rank4ai/what-is-the-difference-between-ai-content-and-ai-search-optimised-content-37bc>
2. Google AI Blog (2026-04-02). “Create, edit and share videos at no cost in Google Vids” <https://blog.google/products-and-platforms/products/workspace/google-vids-updates-lyria-veo/>
3. Higher Ed Dive (2026-04-10). “COVID remote learning put drain on college enrollment” <https://www.highereddive.com/news/covid-remote-learning-put-drain-on-college-enrollment/816899/>
4. Dev To Ai (2026-04-12). “From Notes to Narrative: How AI Automates Trade Show Lead Intelligence” https://dev.to/ken_deng_ai/from-notes-to-narrative-how-ai-automates-trade-show-lead-intelligence-1239
5. The Hechinger Report (2026-04-20). “How much will that college cost you? Good luck figuring it out” <https://hechingerreport.org/how-much-will-that-college-cost-you-good-luck-figuring-it-out-2/>
6. Dev To Ai (2026-04-25). “I Audited a Business’s AI Visibility Across Four Platforms. The Results Were Worse Than Expected.” <https://dev.to/edvisageglobal/i-audited-a-businesss-ai-visibility-across-four-platforms-the-results-were-worse-than-expected-4ml5>
7. Dev To Ai (2026-04-26). “Stop Chasing Ghosts: Automate Your Post-Trade Show Follow-Up” https://dev.to/ken_deng_ai/stop-chasing-ghosts-automate-your-post-trade-show-follow-up-5f65
8. Dev To Ai (2026-04-27). “Instant Lead Scoring: Teaching AI to Sort Trade Show Prospects” https://dev.to/ken_deng_ai/instant-lead-scoring-teaching-ai-to-sort-trade-show-prospects-41l1
9. Dev To Ai (2026-04-28). “Write a 1,200-word blog post:”What is Generative Engine Optimization (GEO) and why SEO teams need i” <https://dev.to/sieok/write-a-1200-word-blog-post-what-is-generative-engine-optimization-geo-and-why-seo-teams-need-i-35a0>
10. Dev To Ai (2026-04-30). “How we shipped”One Brief Full Ad Campaign in 60s” for \$0.35 — a Dify + x402 + multi-engine AI build” <https://dev.to/aikitpros/how-we-shipped-one-brief-full-ad-campaign-in-60s-for-035-a-dify-x402-multi-engine-ai-ib>

Student Completion & Graduation

capability_id: student-completion-graduation · created: 2026-06-29 · revised: 2026-06-29 · refs: 1

Near (2026–2031)

For the subsegment of institutions utilizing document-based identity verification within graduation record certificate management, the reliability of manual identity checks is declining. The deployment of generative AI tools capable of producing high-fidelity fraudulent identification [1] increases the risk of unauthorized individuals claiming certificates or degrees. This shift alters the delivery mechanism of identity verification by rendering visual inspection of IDs insufficient to prevent fraud. Because this specific pressure is external to the institution, graduation eligibility management—which relies on internal academic record audits—is unlikely to change in this period.

Hypothesis: In the near term, institutions managing certificate distribution may transition from simple document-based verification to multi-factor or digital identity checks to mitigate the risk of AI-generated identity fraud.

Mid (2036–2041)

For the same subsegment of institutions, if current trajectories in synthetic media hold, the use of traditional physical identification for the issuance of graduation records may become obsolete. Extrapolating from the increasing ease of creating deepfake credentials [1], the mechanism of trust may shift from visual verification to cryptographic verification of the recipient’s identity. This represents a change in the delivery mechanism of the capability, moving toward digital signatures and immutable identifiers. The social and ceremonial substance of graduation event management is likely to persist unchanged due to the high social legitimacy attached to the physical ritual.

Hypothesis: In the mid term, graduation record certificate management will likely move toward a model of cryptographic identity verification to maintain the integrity of the credential.

Far (2046–2051)

For the subsegment of institutions concerned with credential integrity, the ability to verify a graduate’s identity may diverge into two rival outcomes. One possibility involves the adoption of globally standardized, decentralized identity frameworks that eliminate the need for institution-specific identity verification. A competing outcome involves a return to high-friction, in-person biometric verification as a response to the total collapse of trust in all digital and physical documents. It remains unknown whether regulatory bodies will mandate a specific verification standard or allow institutional autonomy. Despite these shifts, the fundamental requirement for a verified terminal record to signal completion will likely resist change.

Hypothesis A: In the far term, identity verification for graduation records may be fully absorbed into a decentralized, third-party identity layer, removing the verification burden from the institution.

Hypothesis B (competing): In the far term, institutions may revert to strict, in-person biometric authentication for the issuance of degrees to counter persistent AI-driven identity fraud.

References

1. Atlantic (2026-05-02). “Hey Chat, Make Me a Fake ID.” https://www.theatlantic.com/technology/2026/05/chatgpt-images-deepfakes-fraud/687023/?utm_source=feed

Student Enrolment

capability_id: student-enrolment · created: 2026-04-10 · revised: 2026-04-10 · refs: 1

Near (2026–2031)

In the United States, programme enrolment volumes for first-year students have decreased [1]. This decline is linked to a lower rate of completion for financial aid applications (FAFSA) and standardized testing (ACT) [1]. These frictions at the entry point of the programme enrolment process reduce the pool of eligible candidates entering the pipeline [1]. Because this drain occurs at the initial sign-up phase, module selection management is unlikely to change in the near term as it is a downstream process. Student induction management remains largely unaffected by these volume shifts in the immediate term [1].

Hypothesis: In the near term, programme enrolment in the US region may experience continued volatility as institutions respond to changes in student application behaviors.

Mid (2036–2041)

In the United States, if current trajectories of declining first-year sign-ups hold, institutions may alter the delivery mechanism of programme enrolment to reduce friction [1]. This could involve the deployment of automated tools to assist students in navigating FAFSA and ACT requirements, shifting the process from a passive administrative task to an active, supported workflow. Extrapolating from the volume drain, institutions may also modify student induction management to prioritize retention of a smaller entering cohort [1]. The fundamental regulatory requirements for financial aid eligibility will likely persist unchanged due to government oversight.

Hypothesis: In the mid term, institutions in the affected region may adopt more automated delivery mechanisms for programme enrolment to mitigate volume loss.

Far (2046–2051)

In the United States, the long-term elasticity of student demand relative to remote learning formats remains unknown [1]. If the drain on traditional programme enrolment persists, the capability may diverge into two rival outcomes. One outcome involves the complete automation of the enrolment pipeline, where institutional barriers are minimized to maximize throughput. A competing outcome involves a return to high-touch, human-centric enrolment processes to restore social legitimacy and perceived value [1]. Legal requirements for institutional accreditation will resist change even at this horizon, as they provide the basis for enrolment legitimacy.

Hypothesis A: In the far term, programme enrolment may transition to a continuous, fully automated delivery mechanism to minimize student friction.

Hypothesis B (competing): In the far term, institutions may shift toward high-touch, human-led enrolment processes to counteract the perceived depersonalization of remote learning.

References

1. Higher Ed Dive (2026). “COVID remote learning put drain on college enrollment.” <https://www.highereddive.com/news/covid-remote-learning-put-drain-on-college-enrollment/816899/>

Student Support & Wellbeing Management

capability_id: student-support-wellbeing-management · created: 2026-07-19 · revised: 2026-07-19 · refs: 12

Near (2026–2031)

For a subset of institutions and specific student populations adopting AI-mediated workflows, shifts are emerging in both the delivery mechanism and the educational substance of support processes. In the process of personal learning management, the delivery mechanism is shifting toward AI-automated synthesis of study materials and organization of coursework via tools like Adobe Acrobat Spaces [4] and specific prompting strategies for breaking down readings [12]. Within academic skill development, the educational substance is shifting toward the regulation of “cognitive partnership,” where the goal is to move students from simple text generation toward iterative brainstorming [1] and the adoption of frameworks like TACO to prevent AI from substituting for logical thinking [8]. This shift is a mechanistic response to the risk of “cognitive surrender,” where users abandon independent reasoning [2] and exhibit linguistic homogenization in their thought and writing patterns [5]. In the process of academic advice management, the substance of guidance may be compromised by sociodemographic biases in LLM-based counselling, potentially introducing inequities into student support [11]. Within disability support management, delivery is shifting toward the use of AI to support students with learning differences, although the current implementation of these “inference engines” often lacks systematic safeguarding [9]. In the process of student health & wellbeing, delivery is incorporating AI despite workforce tension between clinical enthusiasm and fear of displacement [3], while the substance of crisis triage is adapting to address critical safety failures in LLMs handling suicide and self-harm [10] and risks of AI-induced delusions [6]. In the region of Florida, these processes are being formally constrained by state-directed partnerships to shield families from AI harm [7]. The processes of student financial advice and housing advice are unlikely to change in this horizon because they remain tied to external regulatory frameworks and static administrative constraints.

Hypothesis: In the near term, academic skill development for a subset of students may transition toward the formal regulation of cognitive partnership to mitigate cognitive surrender, while student health & wellbeing may prioritize the mitigation of critical AI failure modes in crisis triage.

Mid (2036–2041)

For the same subset of institutions and regions, if current trajectories hold, the academic skill development process may formalize the auditing of machine-generated logical structures as a core competency, extrapolating from early frameworks that distinguish between cognitive support and cognitive substitution [8]. Within disability support management, the implementation of “Responsible Inference Engines” may evolve from a policy recommendation into a standard operational requirement for safeguarding students with learning differences [9]. In student health & wellbeing, AI-driven documentation and triage may further integrate into the workforce, provided that institutional protocols resolve the tension between professional fear and operational efficiency [3] and establish redundant human-in-the-loop overrides for high-risk mental health interventions [10]. In regions implementing standardized harm reporting, such as Florida, these mechanisms may evolve into permanent regulatory requirements for student wellbeing management [7]. The human-led provision of personal tutors is likely to persist unchanged as the primary mechanism for emotional scaffolding and motivational support, as these require a level of social legitimacy and human empathy not evidenced in current AI deployments.

Hypothesis: In the mid term, the substance of student support for a subset of institutions may shift from the provision of answers toward the auditing of AI-generated reasoning to prevent long-term cognitive atrophy.

Far (2046–2051)

For the same subset of institutions, the scope of AI integration in wellbeing management remains subject to significant regulatory and psychological uncertainty. One plausible outcome is the emergence of a bifurcated support model where AI handles all baseline academic skill development and initial triage, while human

professionals are reserved exclusively for high-stakes crisis intervention and complex emotional scaffolding [3][10]. A rival outcome, driven by expanded regulatory regimes similar to early initiatives in Florida [7], may see the establishment of “AI-free” support zones where human-only interaction is mandated for mental health and disability support to ensure social legitimacy and prevent cognitive surrender [2][7]. It remains unknown whether the linguistic homogenization observed in early AI use [5] will permanently alter the nature of student self-expression or if new pedagogical frameworks will successfully decouple AI utility from cognitive decline. The requirement for human-verified clinical accountability in student health & wellbeing is likely to resist change even at this horizon due to the high legal and ethical stakes of psychiatric care.

Hypothesis A: In the far term, student support may evolve into a tiered system where AI manages process-oriented support and humans manage substance-oriented emotional and ethical crisis intervention.

Hypothesis B (competing): In the far term, pervasive risks of cognitive surrender and bias may lead to a regulatory mandate for human-centric support in all wellbeing and academic skill development processes.

References

1. The Conversation (2026-03-26). “College students are writing with AI – but a pilot study finds they’re not simply letting it write for them.” <https://theconversation.com/college-students-are-writing-with-ai-but-a-pilot-study-finds-theyre-not-simply-letting-it-write-for-them-276856>
2. Ars Technica (2026-04-03). ““Cognitive surrender” leads AI users to abandon logical thinking, research finds.” <https://arstechnica.com/ai/2026/04/research-finds-ai-users-scarily-willing-to-surrender-their-cognition-to-llms/>
3. NPR (2026-04-07). “AI in the mental health care workforce is met with fear, pushback — and enthusiasm.” <https://www.npr.org/2026/04/07/nx-s1-5771707/mental-health-care-workforce-artificial-intelligence-ai>
4. TechCrunch (2026-04-07). “Adobe launches Acrobat Spaces, a free AI-powered study tool for students.” <https://techcrunch.com/2026/04/07/adobe-launches-acrobat-spaces-a-free-ai-powered-study-tool-for-students/>
5. Hacker News (2026-04-07). “AI may be making us think and write more alike.” <https://dornsife.usc.edu/news/stories/ai-may-be-making-us-think-and-write-more-alike/>
6. Daily Ai (2025-05-09). “ChatGPT Is Making People Think They’re Gods and Their Families Are Terrified.” <https://dailyai.com/2025/05/chatgpt-is-making-people-think-theyre-gods-and-their-families-are-terrified/>
7. Future Of Life (2026-03-09). “Governor DeSantis Directs Florida State Agencies to Partner with Future of Life Institute to Shield Families from AI Harm.” <https://futureoflife.org/press-release/de-santis-directs-florida-agencies-to-partner-with-fli/>
8. Arxiv Cs Cy (2026-04-22). “Students Know AI Should Not Replace Thinking, but How Do They Regulate It? The TACO Framework for Human-AI Cognitive Partnership.” <https://arxiv.org/abs/2604.18737>
9. Getting Smart (2026-04-22). “Responsible Inference Engines: Safeguarding Students with Learning Differences in the AI Era.” <https://www.gettingsmart.com/2026/04/22/responsible-inference-engines-safeguarding-students-with-learning-differences-in-the-ai-era/>
10. Fast Company (2026-04-24). “LLMs don’t get mental health right. We need a two-pronged approach to fix them.” https://www.fastcompany.com/91527462/llms-dont-get-mental-health-right-we-need-a-two-pronged-approach-to-fix-them-ai-suicide-self-harm?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
11. Arxiv Cs Cy (2026-04-30). “Sociodemographic Biases in Educational Counselling by Large Language Models.” <https://arxiv.org/abs/2604.25932>
12. Elearning Industry (2026-05-04). “5 Ways Adult Learners Can Use AI To Study Smarter, Without Compromising Integrity.” <https://elearningindustry.com/5-ways-adult-learners-can-use-ai-to-study-smarter-without-compromising-integrity>

Supporting Services

capability_id: supporting-services · created: 2026-06-09 · revised: 2026-06-09 · refs: 4

Near (2026–2031)

For a subset of supporting services, specifically food and retail management and venue management, the delivery mechanisms for logistics and monitoring are shifting toward automated digital systems. In food and retail management, the adoption of machine learning algorithms for demand forecasting alters how institutions analyze purchasing patterns and seasonal trends to manage inventory [4]. However, the implementation of automated approval systems for supply chain logistics creates a dependency where physical goods may be wasted if digital validation systems fail to approve shipments [1]. In venue management, the integration of AI video analytics transforms traditional monitoring into intelligent systems capable of real-time object detection and tracking [3]. Furthermore, the use of predictive intelligence allows venue management to shift from reactive observation toward forecasting movement patterns and density changes [2]. Processes such as religious support and child care management are unlikely to adopt these specific automated frameworks in the near term because their institutional value depends on human presence and social legitimacy.

Hypothesis: In the near term, food and retail management may experience increased operational fragility due to digital validation failures, while venue management may shift toward proactive risk detection and predictive spatial monitoring.

Mid (2036–2041)

For the food, retail, and venue management subsegments, if current trajectories of digital integration hold, the stability of these services will depend on the resilience of the underlying software. Extrapolating from current failures in automated approvals, the risk of systemic waste in food and retail management could become a structural vulnerability unless institutions implement manual override protocols for physical logistics [1]. In venue management, predictive intelligence and video analytics may evolve from monitoring density to automating real-time resource allocation and crowd control mechanisms [2][3]. However, the educational substance of complaint and compliment management is likely to remain unchanged, as the resolution of human grievances typically requires social empathy and discretionary authority.

Hypothesis: In the mid term, food logistics may face persistent risks of “digital waste,” while venue management may integrate predictive analytics into automated operational resource allocation.

Far (2046–2051)

For the food, retail, and venue management subsegments, the long-term outcome depends on whether the industry optimizes for centralized efficiency or distributed resilience. One possibility involves the development of high-resilience, autonomous validation systems that mitigate the single points of failure currently causing supply chain waste [1]. A rival outcome is a state of chronic operational instability where the gap between digital approval and physical reality creates permanent inefficiencies in food and retail management, potentially forcing a return to non-digital, localized procurement. It remains unknown if the cost of system-induced waste will outweigh the benefits of predictive demand forecasting [4]. Processes like religious support will likely resist this trajectory throughout this horizon as their primary function is the provision of human-centric spiritual care.

Hypothesis A: In the far term, these subsegments may achieve autonomous, high-resilience operational states via distributed intelligence and predictive validation.

Hypothesis B (competing): In the far term, supporting services may experience structural degradation and chronic waste due to an irreversible dependence on rigid, centralized digital approval mechanisms.

References

1. Science Daily Comp (2026-04-05). “Truckloads of food are being wasted because computers won’t approve them.” <https://www.sciencedaily.com/releases/2026/04/260403224505.htm>
2. Dev To Ai (2026-04-22). “Architecting Predictive Intelligence for Smart Venues.” https://dev.to/tanisha_khiara_4b390b7946/architecting-predictive-intelligence-for-smart-venues-1lkj
3. Dev To Ai (2026-04-23). “Integrating AI Video Analytics into Venue Operations.” https://dev.to/tanisha_khiara_4b390b7946/integrating-ai-video-analytics-into-venue-operations-2m09
4. Dev To Ai (2026-04-29). “How Smart Technology Is Reshaping Retail.” https://dev.to/xit_vali_8353fbdb3474555c/how-smart-technology-is-reshaping-retail-5308

Teaching & Learning Delivery

capability_id: teaching-learning-delivery · created: 2026-07-17 · revised: 2026-07-17 · refs: 19

Near (2026–2031)

For teaching & learning delivery broadly, the educational substance is shifting toward skills-based learning frameworks to address workplace volatility [13] and a prioritization of learning processes over the production of academic artifacts [19]. Within teaching & learning content management, the delivery mechanism is moving from manual authoring to automated multimodal generation, including the use of AI for instructional videos [6], presentations derived from URLs [16], automated multimodal illustrations [7], and localized materials for Career and Technical Education [1]. This is supported by the emergence of “liquid content” workflows where raw stories or data are treated as material for multiple media formats [30] and the use of agentic co-authoring tools that modify documents during the drafting phase [20]. In non presence based teaching, delivery mechanisms are being modified for specialized subsegments through the deployment of intelligent tutoring systems for Python [2], quantum education [26], and low-resource languages [22], as well as personalized practice tools that attempt to replicate learning rate regularity at scale [3]. For institutions adopting disciplinary-specific integration, as demonstrated in frameworks that prioritize professional standards over general AI literacy [24], presence based teaching is being modified by the adoption of “cognitive partnership” models to prevent AI from replacing internal student thinking [15][17] and the use of policy-governed routing in specialized laboratories to balance AI assistance with manual skill acquisition [29]. The core social mediation and professional socialization functions of presence based teaching are unlikely to change yet because the social legitimacy of the human instructor remains a primary requirement for institutional accreditation and faculty labor contracts.

Hypothesis: In the near term, the capability shifts toward a process-oriented educational substance supported by a delivery mechanism of automated multimodal content and specialized, adaptive non presence based tutoring interfaces.

Mid (2036–2041)

For a subset of institutions adopting automated delivery models and STEM-focused subsegments, the delivery mechanism of non presence based teaching may shift toward agent-led orchestration of student pathways if current trajectories in specialized AI tutoring and autonomous schooling models hold [2][10]. Within teaching & learning content management, the process may evolve from a creation-centric activity into a verification-heavy auditing role, provided that governance frameworks for quality, accuracy, and trust are standardized [5][11]. If the risk of “cognitive bypass”—where students use AI as a substitute for critical thinking—persists, presence based teaching could shift toward constrained delivery models that intentionally limit AI autonomy to protect the cognitive substance of the learning process [17]. The requirement for human-led mentorship in high-stakes professional socialization is likely to persist because it is embedded in the regulatory regimes of professional licensing and accreditation.

Hypothesis: In the mid term, content management may transition into a verification and governance function, while non presence based teaching may shift toward agent-orchestrated pathways for technical disciplines.

Far (2046–2051)

For institutions using highly automated delivery frameworks, the boundary between non presence based and presence based teaching may blur as synthetic interfaces achieve high-fidelity social presence. One rival outcome is a bifurcated system where base-layer knowledge delivery is fully automated and validated by AI, while human-led instruction is reserved exclusively for high-status “capstone” mentorship and ethical socialization. A competing outcome is a return to “analog” presence based teaching as a luxury signal of educational quality, where the absence of AI is marketed as a premium institutional feature. It remains unknown whether synthetic social presence can ever satisfy the legal and social requirements for professional certification in regulated fields. The fundamental requirement for an accredited body to vouch

for a student’s competence will likely resist change even at this horizon.

Hypothesis A: In the far term, the capability splits into a fully automated delivery layer for knowledge acquisition and a human-exclusive layer for professional socialization.

Hypothesis B (competing): In the far term, presence based teaching undergoes a “premiumization” where human-only delivery becomes a high-cost status symbol in response to the ubiquity of synthetic instruction.

References

1. Elearning Industry (2026-04-05). “Leveraging AI And Automation To Streamline CTE Content Creation And Localization.” <https://elearningindustry.com/leveraging-ai-and-automation-to-streamline-cte-content-creation-and-localization>
2. The Hechinger Report (2026-04-06). “The quest to build a better AI tutor.” <https://hechingerreport.org/proof-points-ai-tutor-python/>
3. Arxiv Cs Cy (2026-04-07). “Personalized AI Practice Replicates Learning Rate Regularity at Scale.” <https://arxiv.org/abs/2604.03246>
4. Elearning Industry (2026-04-07). “A 4-Step Framework For Using AI Transparently In Educational Content.” <https://elearningindustry.com/a-4-step-framework-for-using-ai-transparently-in-educational-content>
5. Google AI Blog (2026-04-02). “Create, edit and share videos at no cost in Google Vids.” <https://blog.google/products-and-platforms/products/workspace/google-vids-updates-lyria-veo/>
6. Arxiv Cs Cy (2026-04-08). “EduIllustrate: Towards Scalable Automated Generation Of Multimodal Educational Content.” <https://arxiv.org/abs/2604.05005>
7. Atlantic (2026-04-10). “This Is What Fully Automated School Looks Like.” https://www.theatlantic.com/technology/2026/04/ai-agents-school-education/686754/?utm_source=feed
8. Elearning Industry (2026-04-10). “AI-Generated Learning Content Governance: Ensuring Quality, Accuracy, And Trust.” <https://elearningindustry.com/ai-generated-learning-content-governance-ensuring-quality-accuracy-and-trust>
9. Elearning Industry (2026-04-15). “The Case For Skills-Based Learning In A Rapidly Changing Workplace.” <https://elearningindustry.com/the-case-for-skills-based-learning-in-a-rapidly-changing-workplace>
10. Elearning Industry (2026-04-21). “How AI Avatars Are Turning Static Documents Into Interactive Learning Experiences.” <https://elearningindustry.com/how-ai-avatars-are-turning-static-documents-into-interactive-learning-experiences>
11. Arxiv Cs Cy (2026-04-21). “Bringing AI into the Classroom: A Structured Approach for Integrating AI into Software Engineering Education.” <https://arxiv.org/abs/2604.16328>
12. Dev To Ai (2026-04-21). “Create Presentations from Any URL Automatically Using AI.” https://dev.to/ai_presentation_d9935ff95/create-presentations-from-any-url-automatically-using-ai-5hho
13. Arxiv Cs Cy (2026-04-22). “Students Know AI Should Not Replace Thinking, but How Do They Regulate It? The TACO Framework for Human-AI Cognitive Partnership.” <https://arxiv.org/abs/2604.18737>
14. EDUCAUSE Review (2026-04-23). “Teaching and Learning with AI: Three Priorities.” <https://er.educause.edu/podcasts/educause-exchange/teaching-and-learning-with-ai-three-priorities>
15. The Register (2026-04-23). “Microsoft gives your Word documents an AI co-author you didn’t ask for.” https://go.theregister.com/feed/www.theregister.com/2026/04/23/microsoft_gives_your_word_documents/
16. Arxiv Cs Cl (2026-04-24). “AFRILANGTUTOR: Advancing Language Tutoring and Culture Education in Low-Resource Languages with Large Language Models.” <https://arxiv.org/abs/2604.20996>
17. Wonkhe (2026-04-27). “AI literacy is not enough – universities must teach through disciplinary standards.” <https://wonkhe.com/blogs/ai-literacy-is-not-enough-universities-must-teach-through-disciplinary-standards/>
18. Arxiv Cs Cy (2026-04-29). “From Prototype to Classroom: An Intelligent Tutoring System for

- Quantum Education.” <https://arxiv.org/abs/2604.24807>
19. Arxiv Cs Cy (2026-05-01). “Policy-Governed LLM Routing with Intent Matching for Instrument Laboratories.” <https://arxiv.org/abs/2604.26955>
 20. Fast Company (2026-05-01). “AI is turning every story into raw material.” https://www.fastcompany.com/91532289/ai-is-turning-every-story-into-raw-material?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss