

edudrift

An LLM’s living hypothesis log on AI’s impact on higher education

Santosh B. Srinivas · HEC Paris

Snapshot: 2026-10-01

<https://santoshbs.github.io/theaidrift/edu/>

edudrift is an experiment in letting a local open-weight large language model continuously process AI-related news signals and hypothesise their implications for specific higher-education capabilities. Each entry in this document is a working manuscript: it is generated and revised **entirely by the model**, with **no human in the editing loop**. The author’s contribution is limited to designing the system — selecting RSS sources, defining the capability scaffold, writing the prompts, and committing the model’s output to version control. The prose, the reasoning, the framing, and the choice of which signals matter are all the model’s.

The capability scaffold is drawn from the **Higher Education Reference Models (HERM)** v3.2.0, a global industry-standard enterprise-architecture framework curated by the HERM Working Group of CAUDIT (the Council of Australasian University Directors of Information Technology), in collaboration with EDUCAUSE, UCISA, and EUNIS, and used by more than a thousand institutions worldwide. HERM is published under CC BY-NC-SA 4.0 and is available at <https://www.caudit.edu.au/communities/caudit-higher-education-reference-models/>.

For each capability, the log states three horizon hypotheses, each grounded in cited public signals:

- **Near** (2026–2031) — what current evidence most plausibly supports;
- **Mid** (2036–2041) — extrapolations contingent on signal trajectories holding;
- **Far** (2046–2051) — competing hypotheses where the trajectory is genuinely uncertain.

Hypotheses are abductive and revisable; each is version-controlled at the source repository. *Italic statements* are the formal hypothesis claims.

The live, continuously revised version is published at </theaidrift/edu/>.

{*Caveat lector*. Because this document is fully model-generated, claims should be read as structured speculation grounded in cited signals, not as expert assessment. Citations point to the public sources the model drew on; readers should verify any claim against the original source before relying on it.}

Contents

Academic Administration	6
Near (2026–2031)	6
Mid (2036–2041)	6
Far (2046–2051)	6
References	7
Alumni Engagement	8
Near (2026–2031)	8
Mid (2036–2041)	8
Far (2046–2051)	8
References	8
Business Capability Management	9
Near (2026–2031)	9
Mid (2036–2041)	9
Far (2046–2051)	9
References	10
Commercial Delivery	11
Near (2026–2031)	11
Mid (2036–2041)	11
Far (2046–2051)	11
References	11
Commercial Monitoring	13
Near (2026–2031)	13
Mid (2036–2041)	13
Far (2046–2051)	13
References	13
Commercial Sourcing	14
Near (2026–2031)	14
Mid (2036–2041)	14
Far (2046–2051)	14
References	15
Corporate Governance	16
Near (2026–2031)	16
Mid (2036–2041)	16
Far (2046–2051)	16
References	17
Curriculum Management	19
Near (2026–2031)	19
Mid (2036–2041)	19
Far (2046–2051)	20
References	20
Facilities & Property Management	22
Near (2026–2031)	22
Mid (2036–2041)	22
Far (2046–2051)	22
References	23
Finance Management	24

Near (2026–2031)	24
Mid (2036–2041)	24
Far (2046–2051)	24
References	25
Government, Public & Stakeholder Relationships	27
Near (2026–2031)	27
Mid (2036–2041)	27
Far (2046–2051)	27
References	28
Human Resource Management	29
Near (2026–2031)	29
Mid (2036–2041)	29
Far (2046–2051)	29
References	30
Information & Communication Technology Mgt	32
Near (2026–2031)	32
Mid (2036–2041)	32
Far (2046–2051)	32
References	33
Information Management	36
Near (2026–2031)	36
Mid (2036–2041)	36
Far (2046–2051)	36
References	37
Legal Services	40
Near (2026–2031)	40
Mid (2036–2041)	40
Far (2046–2051)	40
References	41
Library Management	43
Near (2026–2031)	43
Mid (2036–2041)	43
Far (2046–2051)	43
References	43
Promotions Management	44
Near (2026–2031)	44
Mid (2036–2041)	44
Far (2046–2051)	44
References	45
Research Administration	46
Near (2026–2031)	46
Mid (2036–2041)	46
Far (2046–2051)	47
References	47
Research Delivery	49
Near (2026–2031)	49
Mid (2036–2041)	49
Far (2046–2051)	50

References	50
Research Funding	53
Near (2026–2031)	53
Mid (2036–2041)	53
Far (2046–2051)	53
References	54
Research Impact	55
Near (2026–2031)	55
Mid (2036–2041)	55
Far (2046–2051)	55
References	56
Research Improvement	57
Near (2026–2031)	57
Mid (2036–2041)	57
Far (2046–2051)	57
References	58
Research Opportunity Analysis	59
Near (2026–2031)	59
Mid (2036–2041)	59
Far (2046–2051)	59
References	59
Research Publications	60
Near (2026–2031)	60
Mid (2036–2041)	60
Far (2046–2051)	60
References	61
Research Training	62
Near (2026–2031)	62
Mid (2036–2041)	62
Far (2046–2051)	62
References	63
Strategy Management	64
Near (2026–2031)	64
Mid (2036–2041)	64
Far (2046–2051)	64
References	65
Student Administration	67
Near (2026–2031)	67
Mid (2036–2041)	67
Far (2046–2051)	67
References	68
Student Admission Management	69
Near (2026–2031)	69
Mid (2036–2041)	69
Far (2046–2051)	69
References	70
Student Assessment	71

Near (2026–2031)	71
Mid (2036–2041)	71
Far (2046–2051)	71
References	72
Student Attraction & Recruitment	74
Near (2026–2031)	74
Mid (2036–2041)	74
Far (2046–2051)	74
References	75
Student Completion & Graduation	76
Near (2026–2031)	76
Mid (2036–2041)	76
Far (2046–2051)	76
References	76
Student Enrolment	77
Near (2026–2031)	77
Mid (2036–2041)	77
Far (2046–2051)	77
References	77
Student Support & Wellbeing Management	78
Near (2026–2031)	78
Mid (2036–2041)	78
Far (2046–2051)	79
References	79
Supporting Services	81
Near (2026–2031)	81
Mid (2036–2041)	81
Far (2046–2051)	81
References	82
Teaching & Learning Delivery	83
Near (2026–2031)	83
Mid (2036–2041)	83
Far (2046–2051)	83
References	84

Academic Administration

capability_id: academic-administration · created: 2026-09-30 · revised: 2026-09-30 · refs: 4

Near (2026–2031)

For Academic Administration broadly, the diversification of plagiarism definitions across global universities [3] may necessitate a direct revision of academic policy & regulation management to address AI-mediated misconduct. The observation that AI-enabled misconduct often overlaps with traditional forms of plagiarism [1] suggests a shift toward integrated behavioral frameworks rather than tool-specific prohibitions [1]. For a subsegment of institutions implementing AI disclosure requirements, evidence suggests that binary “yes/no” disclosure flags are an ineffective delivery mechanism for ensuring honest reporting, as such systems are prone to being ignored by users [4]. In specific institutions employing bottom-up champions to lead AI initiatives, these policy transitions may be accelerated through cross-functional collaborations between teaching centers and digital learning divisions [2]. This represents a change in educational substance by altering the formal definition of academic dishonesty and the criteria for misconduct. Processes such as academic year scheduling and timetabling management are unlikely to change in this period as they are not functionally linked to academic integrity signals.

Hypothesis: In the near term, academic policy & regulation management may shift from tool-specific AI bans toward holistic integrity frameworks, while simple binary disclosure mechanisms are likely to be abandoned in favor of more nuanced reporting requirements.

Mid (2036–2041)

For Academic Administration broadly, if current trajectories in policy redefinition hold [3], the capability of academic policy & regulation management could shift toward the use of longitudinal behavioral profiles. Extrapolating from the failure of binary disclosure systems [4] and the overlap between AI use and traditional plagiarism [1], regulatory mechanisms might move from auditing single assignments to identifying patterns of dishonesty across a student’s academic career. Institutional due process requirements and data privacy laws are likely to persist, constraining the degree of algorithmic automation in these disciplinary processes. The fundamental administrative structure of the academic year remains unlikely to change, as it is governed by external labor contracts and seasonal demographics rather than integrity regulation.

Hypothesis: In the mid term, academic policy & regulation management may move toward a behavioral-pattern model of integrity enforcement, extrapolating from the insufficiency of single-point disclosure and current evidence of overlapping plagiarism methods.

Far (2046–2051)

For Academic Administration broadly, the long-term trajectory of integrity regulation remains uncertain, with two primary rival outcomes. One possibility is that the ongoing evolution of plagiarism definitions [3] leads to a system of automated behavioral pattern matching to identify integrity risks. A competing outcome is that the continued erosion of traditional authorship leads to a collapse of the current regulatory regime, where independent work is no longer a primary metric of educational validation. It remains unknown how professional accreditation bodies will redefine “competence” if the validation of independent student authorship becomes obsolete. The legal requirement for human-led due process in high-stakes disciplinary appeals is likely to resist full automation even at this horizon.

Hypothesis A: In the far term, academic policy & regulation management may evolve into a behavioral monitoring system that identifies integrity risks.

Hypothesis B (competing): In the far term, the capability of academic policy & regulation management may be fundamentally diminished as the institution shifts away from the validation of independent student authorship.

References

1. The Conversation Au (2026-04-01). “Does AI mean more uni students are plagiarising their work?” <https://theconversation.com/does-ai-mean-more-uni-students-are-plagiarising-their-work-279565>
2. EDUCAUSE Review (2026-04-08). “Building AI Initiatives with Bottom-Up Champions” <https://er.educause.edu/articles/2026/4/building-ai-initiatives-with-bottom-up-champions>
3. Dev To Ai (2026-05-06). “How Different Universities Define Plagiarism in 2026” <https://dev.to/martiinn/how-different-universities-define-plagiarism-in-2026-9mi>
4. Dev To Ai (2026-05-10). “Why binary AI disclosure systems fail (and how to design better ones)” <https://dev.to/alanwest/why-binary-ai-disclosure-systems-fail-and-how-to-design-better-ones-5f70>

Alumni Engagement

capability_id: alumni-engagement · created: 2026-06-20 · revised: 2026-06-20 · refs: 1

Near (2026–2031)

For the subsegment of alumni engagement focused on campaign and event marketing, the production of promotional assets is shifting toward near-instant, low-cost automated generation. The availability of multi-engine AI builds allows for the conversion of brief instructions into full multichannel campaigns, including scripts, images, and short-form video, at a marginal cost of less than one dollar per campaign [1]. This change in delivery mechanism may enable alumni offices with limited creative budgets to increase the volume of their event-specific outreach. However, the core processes of benefactor management and high-touch alumni relationship management are unlikely to change in this period, as these functions rely on interpersonal trust and social legitimacy that automated ad-generation tools do not address. The shift is currently limited to the technical execution of campaign management rather than the strategic determination of alumni outreach goals.

Hypothesis: In the near term, the cost and time required for routine asset creation in alumni campaign management will decrease significantly for institutions adopting automated generative pipelines.

Mid (2036–2041)

For the subsegment of campaign and event marketing, if current trajectories in low-cost asset generation hold, the labor focus within alumni offices may shift from content production to content curation and prompt orchestration. Extrapolating from early deployment of multi-engine AI builds [1], the ability to generate vast iterations of campaign materials could lead to hyper-segmented event marketing tailored to minute alumni demographics. Despite this, the institutional governance and regulatory constraints surrounding benefactor management—such as tax compliance and legal requirements for endowments—will likely persist unchanged due to their grounding in national law rather than communication technology. The mechanism of change remains a reduction in production friction, not a change in the substance of the alumni-institutional relationship.

Hypothesis: In the mid term, routine marketing for alumni events may become a fully automated delivery process, shifting human labor toward the strategic oversight of automated campaign streams.

Far (2046–2051)

For the subsegment of campaign marketing, the long-term impact of automated content generation remains uncertain, particularly regarding the risk of counterproductivity. If the volume of automated outreach reaches a threshold where alumni perceive communications as synthetic or impersonal, the perceived value of the institutional connection may erode. It remains unknown whether the ability to produce infinite, low-cost assets will lead to higher engagement or to increased digital fatigue among alumni. Even at this horizon, the fundamental requirement for human presence in high-stakes benefactor management is expected to resist change, as the social legitimacy of major philanthropy typically depends on genuine human reciprocity.

Hypothesis A: In the far term, alumni campaign management achieves a state of autonomous hyper-personalization, where every alumnus receives a unique, AI-generated visual and narrative appeal.

Hypothesis B (competing): In the far term, the proliferation of low-cost automated campaign materials leads to a devaluation of digital alumni outreach, forcing institutions to return to high-cost, human-centric communication to maintain social legitimacy.

References

1. Dev To Ai (2026). “How we shipped”One Brief Full Ad Campaign in 60s” for \$0.35 — a Dify + x402 + multi-engine AI build.” <https://dev.to/aikitpros/how-we-shipped-one-brief-full-ad-campaign-in-60s-for-035-a-dify-x402-multi-engine-ai-ib>

Business Capability Management

capability_id: business-capability-management · created: 2026-09-25 · revised: 2026-09-25 · refs: 15

Near (2026–2031)

For business capability management broadly, organisational design may shift in substance toward system-centric architectures that prioritize AI systems over traditional human-centric role definitions [9]. The adoption of AI operating models could necessitate a change in the educational substance of organisational design to manage the transition toward software-factory operations and comprehensive workflow redesign [4]. In a subsegment of institutions managing large-scale AI infrastructure, benefits management shifts in substance toward more stringent ROI validation, as evidence indicates only 28% of such projects fully realize projected payoffs [3]. For a subsegment of institutions updating performance indicators, benefits management further shifts in substance as leadership moves from measuring simple time-savings toward metrics reflecting actual value creation [15]. The delivery mechanism of programme & project management is modified through the use of AI lifecycle frameworks [1] and the integration of AI-driven decision-making within Agile frameworks for a subsegment of project teams [13]. For institutions adopting advanced automation, as indicated by trends in the Canadian executive sector, the delivery of programme & project management may shift toward AI agents taking a leading role in team coordination [16], while other managers incorporate AI to audit project post-mortems for root cause analysis [8] or automate delivery metrics within Scaled Agile Frameworks [5]. For a subset of institutions, change portfolio management delivery shifts toward controlled experimentation and piloting [7], while broader institutional change is pressured to transition from isolated pilots to systemic, business-wide implementation frameworks [11], contingent on leadership alignment [12] and the use of internal learning champions to drive peer-led adoption [14]. Understanding the specific causes of AI pilot failure becomes a core process within change portfolio management to avoid systemic waste [10]. Quality management processes are unlikely to change yet, as they remain dependent on stable administrative specifications and regulatory compliance standards that are independent of AI deployment.

Hypothesis: In the near term, business capability management may integrate AI-specific operating models, system-centric organisational architectures, and automated project auditing to govern the transition toward software-factory operations and mitigate high infrastructure failure rates.

Mid (2036–2041)

For business capability management broadly, if the trend of low ROI for infrastructure and recurring project failures persists, the capability may shift from experimental AI adoption to disciplined architectural integration [3][10]. Extrapolating from the requirement to govern AI sprawl, enterprise architecture could evolve to standardize the software-factory model across the institution [4][6]. This represents a change in the educational substance of the capability, as institutional value is measured by production-level operational integration rather than prototype capability [1]. The fundamental mechanisms of institutional accountability—specifically the determination of legal and professional responsibility for capability failure—will likely persist unchanged due to regulatory and governance constraints.

Hypothesis: In the mid term, the capability may shift toward a disciplined architectural integration where AI-driven operating models and automated failure analysis are standardized as core institutional utilities.

Far (2046–2051)

For business capability management broadly, the long-term trajectory remains uncertain regarding the automation of strategic decision-making. One plausible outcome is that change portfolio management becomes largely automated, with AI systems optimizing resource allocation based on real-time adoption and ROI data [3]. However, it remains unknown whether institutional boards will ever delegate the final authority of organisational design to automated systems. The process of setting the institution’s ethical boundaries and high-level strategic intent resists change even at this horizon.

Hypothesis A: In the far term, the capability may evolve into a highly automated governance system where AI manages the full lifecycle of project portfolios and benefits realisation with minimal human intervention.

Hypothesis B (competing): In the far term, the capability may shift toward a “human-in-the-loop” verification model where AI provides exhaustive simulation and auditing, but organisational design and change portfolio management remain strictly human-led to ensure social legitimacy.

References

1. Dzone Ai (2026-04-03). “From Concept to Production: A Strategic Framework for AI/ML Project Success.” <https://feeds.dzone.com/link/23558/17313177/ai-lifecycle-framework-for-successful-projects>
2. Dev To Ai (2026-04-06). “How to Build an AI Workflow That Your Team Will Actually Use.” <https://dev.to/maxothex/how-to-build-an-ai-workflow-that-your-team-will-actually-use-55he>
3. The Register (2026-04-07). “Only 28% of AI infrastructure projects fully pay off, survey finds.” https://go.theregister.com/feed/www.theregister.com/2026/04/07/ai_returns_gartner/
4. Dev To Ai (2026-04-11). “AI Operating Models: The \$2M Workflow Redesign Your Board Isn’t Asking About.” https://dev.to/dr_hernani_costa/ai-operating-models-the-2m-workflow-redesign-your-board-isnt-asking-about-1dfd
5. Dzone Ai (2026-04-22). “Revolutionizing Scaled Agile Frameworks with AI, MuleSoft, and AWS: An Insider’s Perspective.” <https://feeds.dzone.com/link/23558/17323804/revolutionizing-scaled-agile-frameworks-with-ai>
6. The Register (2026-04-25). “Ex-AWS legend explains what enterprises need to make AI actually work.” https://go.theregister.com/feed/www.theregister.com/2026/04/25/ai_enterprise_matt_domo/
7. Wonkhe (2026-04-27). “Controlled experimentation with AI – the power of a pilot.” <https://wonkhe.com/blogs/controlled-experimentation-with-ai-the-power-of-a-pilot/>
8. Dev To Ai (2026-04-27). “I Asked AI to Audit My Last 3 Failed Projects — Here’s What It Found.” <https://dev.to/tillerman/i-asked-ai-to-audit-my-last-3-failed-projects-heres-what-it-found-30on>
9. Dev To Ai (2026-04-30). “The Future of Work: From Employees to AI Systems.” <https://dev.to/iaddict25709/the-future-of-work-from-employees-to-ai-systems-3p6>
10. Dev To Ai (2026-05-01). “Why Enterprise AI Pilots Fail.” <https://dev.to/xccelera/why-enterprise-ai-pilots-fail-1pcc>
11. Zdnet Ai (2026-05-01). “Moving from AI pilots to business-wide value requires a superhighway - how to ramp up.” <https://www.zdnet.com/article/moving-from-ai-pilots-to-business-wide-value-accenture-research/>
12. EDUCAUSE Review (2026-05-04). “The Leadership Moves That Make AI Adoption Work.” <https://er.educause.edu/articles/sponsored/2026/5/the-leadership-moves-that-make-ai-adoption-work>
13. Dzone Ai (2026-05-04). “Integrating AI-Driven Decision-Making in Agile Frameworks: A Deep Dive into Real-World Applications and Challenges.” <https://feeds.dzone.com/link/23558/17332221/ntegrating-ai-driven-decision-making-in-agile>
14. Elearning Industry (2026-05-04). “Internal Learning Champions: Why Peer Credibility Beats Top-Down Training In AI Adoption.” <https://elearningindustry.com/internal-learning-champions-why-peer-credibility-beats-top-down-training-in-ai-adoption>
15. Dev To Ai (2026-05-07). “How AI and LLMs Are Changing the Business Metrics That Actually Matter.” <https://dev.to/xcelcorp/how-ai-and-llms-are-changing-the-business-metrics-that-actually-matter-57dk>
16. Dev To Ai (2026-05-09). “I Read a Survey That Predicted My Job’s Next 2 Years - Here’s What It Got Right and Missed.” <https://dev.to/itskondrat/i-read-a-survey-that-predicted-my-jobs-next-2-years-heres-what-it-got-right-and-missed-14ea>

Commercial Delivery

capability_id: commercial-delivery · created: 2026-06-04 · revised: 2026-06-04 · refs: 2

Near (2026–2031)

For US-based institutions and those utilizing frontier AI models, the process of commercial activity exit management is increasingly constrained by both regulatory mandates and technical dependencies. In the US, federal blacklisting of specific AI technologies may trigger forced commercial activity exit management, shifting the mechanism of vendor termination from planned lifecycles to compliance-driven mandates [1]. Simultaneously, institutions attempting to transition between frontier models may find the exit process obstructed by vendor lock-in, which increases the complexity and cost of swapping providers [2]. These shifts represent changes to the delivery mechanism of risk and vendor management rather than the educational substance of the commercial activities. Because current signals focus on exit and compliance, the core processes of commercial activity output management are unlikely to change in the near term.

Hypothesis: In the near term, a subset of institutions will experience increased friction in commercial activity exit management due to a combination of geopolitical regulatory restrictions and technical vendor lock-in.

Mid (2036–2041)

For institutions utilizing frontier AI models, if current trajectories of vendor lock-in and regulatory volatility hold, commercial activity issue management may shift toward the requirement of proactive “interoperability audits.” Extrapolating from the difficulty of swapping models [2], institutions could implement mandatory exit-strategy documentation during the initial procurement phase to mitigate the risk of operational paralysis. This represents a more conservative delivery of vendor oversight, as demonstrated by the necessity of aligning commercial ties with state-level prohibited entity lists [1]. The fundamental requirement for specialized legal counsel to navigate the intersection of trade law and software licensing will persist unchanged.

Hypothesis: In the mid term, institutions may integrate formal exit-strategy planning and interoperability requirements into their commercial activity issue management to counter the effects of vendor lock-in.

Far (2046–2051)

For institutions utilizing large-scale commercial AI, the long-term state of commercial activity exit management depends on whether model interoperability becomes a technical standard or remains a proprietary barrier. One plausible outcome is the emergence of fragmented technological blocs where exit management is a routine, trigger-based operational task dictated by national security boundaries [1]. A competing outcome is the persistence of deep vendor lock-in, where the cost of exit management becomes prohibitively high, effectively merging the institutional capability with the vendor’s own lifecycle. It remains unknown whether future regulatory frameworks will provide structured “grace periods” for institutional exits or mandate immediate cessation. The basic institutional need to manage the legal termination of contracts will resist change regardless of the underlying technology.

Hypothesis A: In the far term, commercial activity exit management becomes a highly automated, trigger-based process within a fragmented global AI market.

Hypothesis B (competing): In the far term, extreme vendor lock-in renders independent commercial activity exit management practically impossible for most institutions, leading to permanent vendor dependency.

References

1. Ars Technica (2026-04-09). “Trump-appointed judges refuse to block Trump blacklisting of Anthropic AI tech.” <https://arstechnica.com/tech-policy/2026/04/trump-appointed-judges-refuse-to-block-trump-blacklisting-of-anthropic-ai-tech/>
2. The Register (2026-04-28). “Locked, stocked, and losing budget: AI vendor lock-in bites back.” https://www.theregister.com/2026/04/28/ai_vendor_lock_in_bites_back/

[//go.theregister.com/feed/www.theregister.com/2026/04/28/locked_stocked_and_losing_budget/](http://go.theregister.com/feed/www.theregister.com/2026/04/28/locked_stocked_and_losing_budget/)

Commercial Monitoring

capability_id: commercial-monitoring · created: 2026-06-04 · revised: 2026-06-04 · refs: 1

Near (2026–2031)

For a subsegment of institutions employing proprietary frontier AI models, the process of commercial contingency planning may shift as assumptions of model interchangeability prove invalid [1]. The mechanism driving this change is vendor lock-in, where technical dependencies and pricing structures prevent the rapid swapping of one frontier model for another [1]. Consequently, contingency plans that previously relied on the ability to pivot between vendors in short timeframes may be rendered ineffective [1]. This represents a change in educational substance, as the institutional approach to risk management shifts from a “swap-out” strategy to a lock-in mitigation strategy. Commercial contract performance management is unlikely to change significantly in this period, as the basic monitoring of service-level agreements and uptime remains a standard contractual requirement.

Hypothesis: In the near term, for institutions using proprietary frontier models, commercial contingency planning will shift from assumptions of model interchangeability to the management of vendor lock-in risks.

Mid (2036–2041)

For the same subsegment of institutions using proprietary frontier models, if current trajectories of vendor lock-in hold, commercial activity reporting may expand to track dependency metrics [1]. This process could evolve to monitor “lock-in depth,” treating the inability to switch vendors as a quantifiable financial risk during reporting cycles [1]. Commercial contingency planning may further shift toward architectural diversification to avoid total dependence on a single provider [1]. The fundamental mechanisms of commercial contract performance management will likely persist unchanged, as the legal frameworks governing institutional procurement and vendor accountability are slow to evolve.

Hypothesis: In the mid term, for institutions using proprietary frontier models, commercial activity reporting may incorporate specific metrics to monitor and quantify vendor dependency.

Far (2046–2051)

For institutions that have integrated proprietary frontier models, the scope of commercial contingency planning remains uncertain, depending on whether industry standards move toward open interoperability or deeper integration. One outcome involves the total obsolescence of “swapping” as a contingency, with the vendor instead becoming a core utility similar to electrical power. A rival outcome involves the emergence of standardized model interfaces that restore the capability for rapid vendor migration. Throughout these scenarios, the requirement for commercial activity reporting to produce an audit trail for institutional finance will resist change. It remains unknown if open-weight models can reach a level of capability that makes proprietary lock-in a secondary concern for institutional risk managers.

Hypothesis A: In the far term, for these institutions, commercial contingency planning may be replaced by a utility-based monitoring model where the provider is treated as essential, non-swappable infrastructure.

Hypothesis B (competing): In the far term, for these institutions, the adoption of open-standard model interfaces may restore the capability for rapid vendor swapping within commercial contingency planning.

References

1. The Register (2026-04-28). “Locked, stocked, and losing budget: AI vendor lock-in bites back.” https://go.theregister.com/feed/www.theregister.com/2026/04/28/locked_stocked_and_losing_budget/

Commercial Sourcing

capability_id: commercial-sourcing · created: 2026-07-26 · revised: 2026-07-26 · refs: 6

Near (2026–2031)

For a subset of institutions adopting agentic sourcing tools and those operating within US regulatory jurisdictions, the processes of commercial market analysis and opportunity identification are experiencing a divergence in reliability. The delivery mechanism for commercial market analysis is being modified by the use of structured product catalogs and Model Context Protocol (MCP) servers, which enable AI agents to retrieve real-time pricing across multiple merchants [3]. Simultaneously, the process of commercial opportunity assessment is shifting as B2B buyers increasingly utilize AI-driven visibility factors to filter and select vendors prior to any direct institutional engagement [6]. The process of commercial opportunity identification is under pressure from self-learning AI sales systems that automate prospecting pipelines, moving vendor outreach toward targeted agentic communication [4]. However, the validity of this identification process is undermined by AI-focused SEO tactics designed to bias LLM-generated vendor recommendations [1]. For institutions managing public requests for proposals (RFPs), the process of commercial opportunity assessment is seeing the introduction of AI agents designed to automate “bid-readiness” by aligning institutional capabilities with specific RFP requirements [5]. In the US, the commercial opportunity assessment process is further constrained by the blacklisting of specific high-capability AI providers, which restricts the available vendor pool [2]. Formal procurement mandates requiring competitive bidding and documented audit trails are unlikely to change, as these regulatory constraints prevent total reliance on AI-generated discovery.

Hypothesis: In the near term, for these subsegments and regions, commercial sourcing may bifurcate into a high-efficiency automated track for commodity sourcing and a highly scrutinized, adversarial track for strategic procurement.

Mid (2036–2041)

For these same subsets and regions, if current trajectories of agentic structured data and geopolitical restrictions hold, the process of commercial market analysis could shift toward a “verified-only” workflow. Institutions may restrict their AI agents to accessing only trusted, structured MCP servers or verified registries to mitigate the distortion caused by AI SEO [1][3]. The commercial opportunity assessment process may integrate automated geopolitical compliance filters that exclude blacklisted entities at the point of initial discovery [2]. This represents a change in the delivery mechanism of vendor screening, moving it from a post-discovery check to a pre-filter. Internal institutional governance regarding the final selection of vendors is likely to persist unchanged due to the financial risks and legal liabilities associated with large-scale commercial contracts.

Hypothesis: In the mid term, commercial sourcing for these subsegments may shift toward a “trusted-network” model, where structured, machine-verifiable data sources are prioritized over open-web generative search.

Far (2046–2051)

For institutions relying on AI for sourcing, the outcome depends on whether AI-agent protocols reach global standardization or if regulatory fragmentation increases. One plausible outcome is that commercial market analysis becomes a predominantly machine-to-machine process where AI agents negotiate directly with merchant servers, bypassing the traditional search and discovery phase. It remains unknown whether AI SEO will evolve into a transparent signaling mechanism or if geopolitical blacklisting becomes a permanent, expanding feature of institutional procurement. The requirement for human-led final due diligence will likely resist change because of the legal accountability required for institutional financial commitments.

Hypothesis A: In the far term, commercial sourcing for these subsegments evolves into a fully automated agentic negotiation layer where procurement is handled via standardized machine-to-machine protocols.

Hypothesis B (competing): In the far term, increased regulatory fragmentation and AI-driven misinformation lead to a “closed-loop” sourcing model where institutions only procure from a small, manually vetted circle of pre-approved partners.

References

1. The Verge (2026-04-06). “Can AI responses be influenced? The SEO industry is trying.” <https://www.theverge.com/tech/900302/ai-seo-industry-google-search-chatgpt-gemini-marketing>
2. Ars Technica (2026-04-09). “Trump-appointed judges refuse to block Trump blacklisting of Anthropic AI tech.” <https://arstechnica.com/tech-policy/2026/04/trump-appointed-judges-refuse-to-block-trump-blacklisting-of-anthropic-ai-tech/>
3. Dev To Ai (2026-05-03). “Build an AI Shopping Agent with BuyWhere MCP in 5 Minutes.” <https://dev.to/buywhere/build-an-ai-shopping-agent-with-buywhere-mcp-in-5-minutes-80g>
4. Dev To Ai (2026-05-04). “We Built a Self-Learning AI Sales System in 48 Hours.” <https://dev.to/pvcvotewe-built-a-self-learning-ai-sales-system-in-48-hours-2hbi>
5. Dev To Ai (2026-05-05). “The First PMF Wedge I’d Bet On for AgentHansa: Bid-Readiness for Public RFP Teams.” https://dev.to/jyoti_chance_534245dd0ae8/the-first-pmf-wedge-id-bet-on-for-agenthansa-bid-readiness-for-public-rfp-teams-9g
6. Hackernoon Ai (2026-05-05). “AI Visibility for Startups: How B2B Buyers Choose Vendors Before the First Sales Call.” <https://hackernoon.com/ai-visibility-for-startups-how-b2b-buyers-choose-vendors-before-the-first-sales-call?source=rss>

Corporate Governance

capability_id: corporate-governance · created: 2026-09-30 · revised: 2026-09-30 · refs: 36

Near (2026–2031)

For Corporate Governance broadly, the deployment of autonomous agents shifts the substance of risk management from the validation of static outputs to the continuous oversight of agent-executed actions [1][5][48][56][64]. To mitigate the operational risks of “shadow AI,” institutions are adopting delivery mechanisms such as agent registries [13] and autonomous governance toolkits [2][31]. Risk management processes are evolving to prioritize runtime enforcement over retrospective review, utilizing security toolkits for runtime monitoring [10][55], sandbox execution environments [24], and frameworks based on risk tiering and asset inventory [32][33]. Compliance monitoring and reporting is increasingly focused on the production of granular technical documentation, specifically in regions governed by the European Union where the EU AI Act requires Annex IV technical files [8] and specific protocols for agentic systems [12][17][39], although some implementation timelines have seen recent roll-backs and delays [59][60]. In the United States, the emergence of a National AI Legislative Framework [26] and the possibility of government vetting for powerful models [51][52][54] introduce new requirements for policy and regulation management. For institutions implementing specialized verification tools, compliance monitoring and reporting is expanding to include the use of HMAC-attested receipts to create verifiable audit trails of agent tool calls [62]. Within internal audit and reporting, institutions are integrating specialized AI systems for tax anomaly detection [41] and the automated extraction of insights from unstructured contracts [42]. Legal precedents are establishing that using AI to automate high-stakes administrative decisions without due process may be unconstitutional [66], meaning that the substance of decision-making authority is unlikely to shift away from human agents in regulated environments. Furthermore, the principle that fiduciary liability cannot be delegated to a software agent is unlikely to change, as evidenced by criminal probes into provider liability for real-world harms [15][28][29].

Hypothesis: In the near term, institutions will shift risk and policy frameworks to mitigate the liability of authorized agents and the operational risks of shadow AI, while moving compliance monitoring toward runtime enforcement and regulatory technical documentation.

Mid (2036–2041)

For Corporate Governance broadly, if current trajectories of agent autonomy and regulatory tension hold, internal audit and reporting may shift from retrospective sampling to the use of governance toolkits providing real-time audit trails [2][43]. Extrapolating from the emergence of AI operating models, the substance of board-level oversight may shift toward the continuous redesign of organizational workflows—treating the AI “harness” as a dynamic organizational chart—rather than the approval of static policy documents [18][36]. In regions where strict regulatory regimes persist, such as the EU, compliance monitoring will likely remain a high-friction process requiring human certification despite the automation of data collection [12][39]. The requirement for human-in-the-loop verification for high-stakes institutional decisions is unlikely to change due to social legitimacy constraints and the persisting risk of “LLM nepotism” in organizational decision-making [20].

Hypothesis: In the mid term, the capability may shift from periodic auditing to real-time governance monitoring, with board oversight focusing on the architectural design of AI-integrated workflows.

Far (2046–2051)

For Corporate Governance broadly, the long-term trajectory depends on whether autonomous governance can solve the problem of “black-box” opacity. One outcome involves the normalization of recursive governance, where specialized “judgment layers” monitor agentic networks to ensure alignment with institutional values [19], though it remains unknown if these layers can prevent “secret loyalty” that dodges standard audits [30]. Another possibility is a systemic retrenchment to human-centric governance for all fiduciary roles as the legal costs of agentic failure outpace the efficiency gains. The requirement for a legally recognized

person to hold ultimate liability for institutional failure is likely to resist change even at this horizon.

Hypothesis A: In the far term, corporate governance may evolve into a system of recursive autonomous oversight, where agents govern other agents within strictly defined sandbox boundaries.

Hypothesis B (competing): In the far term, high-stakes governance may revert to exclusively human processes as the opacity of agentic decision-making renders them uninsurable and legally indefensible.

References

1. The Register (2026-04-05). “AI agents promise to ‘run the business,’ but who is liable if things go wrong?” https://go.theregister.com/feed/www.theregister.com/2026/04/05/ai_agents_liability/
2. Ai News (2026-04-02). “KiloClaw targets shadow AI with autonomous agent governance” <https://www.artificialintelligence-news.com/news/kiloclaw-targets-shadow-ai-autonomous-agent-governance/>
3. Ai News (2026-04-06). “As AI agents take on more tasks, governance becomes a priority” <https://www.artificialintelligence-news.com/news/as-ai-agents-take-on-more-tasks-governance-becomes-a-priority/>
4. Dev To Ai (2026-04-07). “What’s actually inside an EU AI Act Annex IV technical file (and what engineers usually miss)” https://dev.to/rostyslav_myronenko_faf2f/whats-actually-inside-an-eu-ai-act-annex-iv-technical-file-and-what-engineers-usually-miss-4nh4
5. Ai News (2026-04-08). “Microsoft open-source toolkit secures AI agents at runtime” <https://www.artificialintelligence-news.com/news/microsoft-open-source-toolkit-secures-ai-agents-at-runtime/>
6. Ai News (2026-04-09). “Agentic AI’s governance challenges under the EU AI Act in 2026” <https://www.artificialintelligence-news.com/news/agentic-ais-governance-challenges-under-the-eu-ai-act-in-2026/>
7. The Register (2026-04-09). “AWS: Agents shouldn’t be secret, so we built a registry for them” https://go.theregister.com/feed/www.theregister.com/2026/04/09/aws_ai_agent_registry/
8. TechCrunch (2026-04-09). “Florida AG to probe OpenAI, alleging possible connection to FSU shooting” <https://techcrunch.com/2026/04/09/florida-ag-to-probe-openai-alleging-possible-connection-to-fsu-shooting/>
9. Dev To Ai (2026-04-10). “The EDPB Is Asking About Your AI Agents. Most Teams Can’t Answer.” <https://dev.to/waxell/the-edpb-is-asking-about-your-ai-agents-most-teams-cant-answer-gfk>
10. Dev To Ai (2026-04-11). “AI Operating Models: The \$2M Workflow Redesign Your Board Isn’t Asking About” https://dev.to/dr_hernani_costa/ai-operating-models-the-2m-workflow-redesign-your-board-isnt-asking-about-1dfd
11. Dev To Ai (2026-04-11). “Judgment Layer for Financial AI Agents” <https://dev.to/thatechmaestro/judgment-layer-for-financial-ai-agents-141>
12. Arxiv Cs Cy (2026-04-14). “LLM Nepotism in Organizational Governance” <https://arxiv.org/abs/2604.09620>
13. Ai News (2026-04-16). “OpenAI Agents SDK improves governance with sandbox execution” <https://www.artificialintelligence-news.com/news/openai-agents-sdk-improves-governance-sandbox-execution/>
14. EDUCAUSE Review (2026-04-21). “White House Unveils National AI Legislative Framework” <https://er.educause.edu/articles/2026/4/white-house-unveils-national-ai-legislative-framework>
15. BBC News (2026-04-21). “OpenAI faces criminal probe over role of ChatGPT in shooting” https://www.bbc.com/news/articles/c62j4ldp2jqo?at_medium=RSS&at_campaign=rss
16. NPR (2026-04-21). “Florida AG launches criminal investigation into ChatGPT over FSU shooting” <https://www.npr.org/2026/04/21/nx-s1-5793967/florida-openai-investigation-mass-shooting-fsu>
17. Lesswrong (2026-04-22). “Narrow Secret Loyalty Dodges Black-Box Audits” <https://www.lesswrong.com/posts/EzdgPbewjeTNHA5F3/narrow-secret-loyalty-dodges-black-box-audits>
18. Fast Company (2026-04-22). “Here’s how to jump-start your company’s responsible AI governance in 90 days” https://www.fastcompany.com/91527542/heres-how-jump-start-your-companys-responsible-ai-governance-90-days?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss

[paign=rss+fastcompany&utm_content=rss](#)

19. Marktechpost (2026-04-23). “Mend.io Releases AI Security Governance Framework Covering Asset Inventory, Risk Tiering, AI Supply Chain Security, and Maturity Model” <https://www.marktechpost.com/2026/04/23/mend-io-releases-ai-security-governance-framework-covering-asset-inventory-risk-tiering-ai-supply-chain-security-and-maturity-model/>
20. Marktechpost (2026-04-24). “Mend Releases AI Security Governance Framework: Covering Asset Inventory, Risk Tiering, AI Supply Chain Security, and Maturity Model” <https://www.marktechpost.com/2026/04/23/mend-releases-ai-security-governance-framework/>
21. Dev To Ai (2026-04-26). “My Harness Is Not a Cage. It’s an Org Chart.” https://dev.to/synthaico/de_commander/my-harness-is-not-a-cage-its-an-org-chart-3gl1
22. Arxiv Cs Cy (2026-04-28). “Algorithmic Administration and the EU AI Act: Legal Principles for Public Sector Use of AI” <https://arxiv.org/abs/2604.22765>
23. Dzone Ai (2026-04-29). “Designing Intelligent AI Systems for Tax Anomaly Detection” <https://feeds.dzone.com/link/23558/17327922/designing-intelligent-ai-systems-for-tax-anomaly>
24. Aws ML (2026-04-29). “Extracting contract insights with PwC’s AI-driven annotation on AWS” <https://aws.amazon.com/blogs/machine-learning/extracting-contract-insights-with-pwcs-ai-driven-annotation-on-aws/>
25. The Register (2026-04-30). “Govern your bots carefully or chaos could ensue” https://go.theregister.com/feed/www.theregister.com/2026/04/30/good_ai_governance_is_good/
26. The Register (2026-05-04). “Five Eyes spook shops warn agentic is too wonky for rapid rollout” https://go.theregister.com/feed/www.theregister.com/2026/05/04/five_eyes_agentic_ai_recommendations/
27. Nyt Technology (2026-05-04). “White House Considers Vetting A.I. Models Before They Are Released” <https://www.nytimes.com/2026/05/04/technology/trump-ai-models.html>
28. Engadget (2026-05-04). “The White House is considering tighter regulation of new AI models” <https://www.engadget.com/2164390/the-white-house-is-considering-tighter-regulation-of-new-ai-models/>
29. The Conversation Tech (2026-05-05). “White House wants to vet powerful AI models for risks — a computer scientist explains why AI safety is so difficult” <https://theconversation.com/white-house-wants-to-vet-powerful-ai-models-for-risks-a-computer-scientist-explains-why-ai-safety-is-so-difficult-281117>
30. Dev To Ai (2026-05-05). “Runtime Compliance Proxy for LLM APIs (EU AI Act)” https://dev.to/json_shotwell/runtime-compliance-proxy-for-llm-apis-eu-ai-act-55jp
31. Ai Business (2026-05-05). “Enterprises Contain AI Agents to Balance Risk, Reward” <https://aibusiness.com/agentic-ai/enterprises-contain-agents-balance-risk-reward>
32. The Register (2026-05-07). “EU hits snooze on AI Act rules after industry backlash” <https://www.theregister.com/ai-and-ml/2026/05/07/eu-hits-snooze-on-ai-act-rules-after-industry-backlash/5234530>
33. Ai Business (2026-05-07). “EU Nations Approve Deal to Roll Back AI Restrictions” <https://aibusiness.com/ai-policy/eu-nations-approve-deal-roll-back-ai-restrictions>
34. Dev To Ai (2026-05-08). “HMAC-attested receipts for AI agent tool calls — verify-action-mcp” <https://dev.to/ardev/hmac-attested-receipts-for-ai-agent-tool-calls-verify-action-mcp-2f53>
35. Ai Business (2026-05-08). “Prompt: AI Agents Are Becoming Operational Infrastructure” <https://aibusiness.com/agentic-ai/prompt-ai-agents-becoming-operational-infrastructure>
36. The Verge (2026-05-08). “DOGE used ChatGPT in a way that was both dumb and illegal, judge rules” <https://www.theverge.com/policy/927071/doge-chatgpt-grants-canceled>

Curriculum Management

capability_id: curriculum-management · created: 2026-09-28 · revised: 2026-09-28 · refs: 30

Near (2026–2031)

For curriculum management broadly, the process of curriculum design is shifting its educational substance toward skills-based frameworks [13] and workforce-aligned objectives to address employer-reported gaps in AI skillsets [12][14][31]. In the subsegment of computer science and software engineering, the educational substance of curriculum design is being restructured to replace the “learning phase” historically provided by entry-level industry roles, as AI coding agents now automate the boilerplate and bug-fixing tasks previously used for junior developer growth [38]. This shift moves the focus from basic implementation toward higher-order architectural skills and “harness engineering”—the orchestration of reliable AI-output environments—because AI tools have lowered the cost of producing average-quality code [15][22][26][27], contributing to the emergence of AI Engineering as a distinct discipline [25] and specialized programs for AI ethicists [4]. For institutions focusing on Career and Technical Education (CTE) and digital content, the delivery mechanism of curriculum & resource development is shifting toward automated multimodal content generation [6], localization [3], and “vibe coding” to reduce technical barriers in interactive course creation [33], supported by instructional design frameworks like the SAMR + AI Matrix [28] and prompt-based design [20]. Within these technical subsegments, the process of curriculum change management is integrating transparency frameworks for AI-generated materials [5] and governance protocols to ensure the accuracy of automated content [10], while utilizing guardrails to maintain assessment validity [16]. For institutions implementing structured faculty development, as demonstrated at Indiana Wesleyan University [1] and through bottom-up champion networks [8], the process of professional learning (staff) is utilizing “course refresh” institutes to update educational substance [1][8], prioritizing judgment-based literacy over tool mastery to mitigate “AI tunnel vision” [2][32][37]. In specific institutional contexts, the process of curriculum retirement management is being accelerated through mandated reviews and budget-driven program cuts [11][23]. Professional accreditation processes are unlikely to change significantly in this window because these bodies typically operate on multi-year review cycles, although some regional teaching standards in the US are currently undergoing updates [17].

Hypothesis: In the near term, curriculum management may shift toward integrating workforce-aligned skills and transparency frameworks in resource development, supported by a mix of structured faculty refreshes and accelerated program retirement in budget-constrained institutions.

Mid (2036–2041)

For the subset of technical institutions and professional subsegments identified in the near term, if current trajectories hold, the process of professional learning (staff) may transition from episodic “refresh” events to a permanent, cyclical requirement to mitigate rapid technological obsolescence [1][2][8][29]. In CTE and technical contexts, delivery-mechanism efficiency gains in curriculum & resource development, such as automated multimodal generation and the expansion of AI-augmented Open Educational Resources (OER), could increase the frequency of curriculum retirement management as specific technical skillsets obsolesce more rapidly [3][6][14][34]. The educational substance of curriculum design in these fields may shift further toward the “pedagogy of mistakes,” where the analysis of AI hallucinations becomes a structured component of the learning process to foster higher-order thinking [36]. Institutional governance structures for curriculum change management may become more decentralized, relying on “bottom-up champion” models to keep pace with the rate of tool evolution [8]. Centralized accreditation requirements are likely to persist as a primary constraint, slowing the full adoption of dynamically updating curricula due to the need for stable, auditable standards.

Hypothesis: In the mid term, curriculum management in technical subsegments may evolve toward a model of continuous, cyclical updates to both staff capacity and resource materials to prevent rapid skill obsolescence.

Far (2046–2051)

For the technical and professional subsegments previously discussed, the process of curriculum design may face a fundamental tension between automation and cognitive depth. One rival outcome involves the transition to hyper-personalized, dynamically updating curricula where curriculum change management is automated in real-time based on workforce signals [12][31], though the validity of such a system remains unknown. A competing outcome may see a strategic return to “fundamentals-first” curriculum design, where the educational substance is intentionally insulated from AI tool-cycles to prevent the cognitive atrophy associated with “average work” production [22][26]. The process of professional accreditation may either fully integrate real-time competency tracking or become an obsolete signal if employer-led skills-based hiring becomes the dominant legitimacy mechanism [13]. The requirement for human-led governance in curriculum performance management is likely to resist change even at this horizon to maintain institutional social legitimacy and ethical oversight.

Hypothesis A: In the far term, curriculum management may shift toward a highly automated, signal-driven model of continuous curriculum evolution and retirement.

Hypothesis B (competing): In the far term, institutions may move toward “slow-curriculum” models that prioritize foundational stability and human-centric cognitive development to counter AI-induced skill degradation.

References

1. EDUCAUSE Review (2026). “Advancing AI Literacy: A Faculty Course Refresh Institute at Indiana Wesleyan University.” <https://er.educause.edu/articles/2026/3/advancing-ai-literacy-a-faculty-course-refresh-institute-at-indiana-wesleyan-university>
2. Getting Smart (2026). “AI Literacy is Not Tool Mastery: How to Build Sustained Educator Capacity.” <https://www.gettingsmart.com/2026/03/27/ai-literacy-is-not-tool-mastery-how-to-build-sustained-educator-capacity/>
3. Elearning Industry (2026). “Leveraging AI And Automation To Streamline CTE Content Creation And Localization.” <https://elearningindustry.com/leveraging-ai-and-automation-to-streamline-cte-content-creation-and-localization>
4. Higher Ed Dive (2026). “Colleges ramp up offerings to teach students to be AI ethicists.” <https://www.highereddive.com/news/colleges-programs-ai-ethicists/816703/>
5. Elearning Industry (2026). “A 4-Step Framework For Using AI Transparently In Educational Content.” <https://elearningindustry.com/a-4-step-framework-for-using-ai-transparently-in-educational-content>
6. Arxiv Cs Cy (2026). “EduIllustrate: Towards Scalable Automated Generation Of Multimodal Educational Content.” <https://arxiv.org/abs/2604.05005>
7. EDUCAUSE Review (2026). “Building AI Initiatives with Bottom-Up Champions.” <https://er.educause.edu/articles/2026/4/building-ai-initiatives-with-bottom-up-champions>
8. Elearning Industry (2026). “AI-Generated Learning Content Governance: Ensuring Quality, Accuracy, And Trust.” <https://elearningindustry.com/ai-generated-learning-content-governance-ensuring-quality-accuracy-and-trust>
9. Higher Ed Dive (2026). “Iowa State plans to cut or merge 23 programs after mandated review.” <https://www.highereddive.com/news/iowa-state-university-program-cuts-review-regents/817262/>
10. Higher Ed Dive (2026). “AI is remaking the workforce. How can colleges ensure students thrive?” <https://www.highereddive.com/news/ai-is-remaking-the-workforce-how-can-colleges-ensure-students-thrive/817598/>
11. Elearning Industry (2026). “The Case For Skills-Based Learning In A Rapidly Changing Workplace.” <https://elearningindustry.com/the-case-for-skills-based-learning-in-a-rapidly-changing-workplace>
12. Higher Ed Dive (2026). “Employers say they struggle to find graduates with the right AI skillset.” <https://www.highereddive.com/news/employers-struggle-to-find-workers-with-the-right-ai>

[skillset/818017/](#)

13. Arxiv Cs Cy (2026). “Bringing AI into the Classroom: A Structured Approach for Integrating AI into Software Engineering Education.” <https://arxiv.org/abs/2604.16328>
14. Elearning Industry (2026). “AI Assessment Guardrails: How To Use AI Without Breaking Validity And Trust.” <https://elearningindustry.com/ai-assessment-guardrails-how-to-use-ai-without-breaking-validity-and-trust>
15. Getting Smart (2026). “America is Updating Teaching Standards Without a Clear Picture of the Future Educator.” <https://www.gettingsmart.com/2026/04/23/america-is-updating-teaching-standards-without-a-clear-picture-of-the-future-educator/>
16. Elearning Industry (2026). “How AI Prompts Are Redefining Creativity In Instructional Design.” <https://elearningindustry.com/how-ai-prompts-are-redefining-creativity-in-instructional-design>
17. Dev To Ai (2026). “AI Didn’t Make Developers Better. It Made Average Work Easier to Produce.” <https://dev.to/workelate123/ai-didnt-make-developers-better-it-made-average-work-easier-to-produce-hc2>
18. Higher Ed Dive (2026). “East Carolina University plans to cut 44 academic programs.” <https://www.highereddive.com/news/east-carolina-university-cut-around-44-academic-programs/818622/>
19. Dev To Ai (2026). “The Case for AI Engineering as a Distinct Discipline.” <https://dev.to/jasrandhawa/the-case-for-ai-engineering-as-a-distinct-discipline-44pm>
20. Dev To Ai (2026). “Fundamentals Matter More Than Ever Before.” <https://dev.to/nachiappannk/fundamentals-matter-more-than-ever-before-266o>
21. Dev To Ai (2026). “Harness Engineering Is the New Senior Developer Skill (Here’s Why).” <https://dev.to/shipwithaiio/harness-engineering-is-the-new-senior-developer-skill-heres-why-4hef>
22. EDUCAUSE Review (2026). “A Practical SAMR + AI Framework for Instructional Design.” <https://er.educause.edu/articles/2026/4/a-practical-samr--ai-framework-for-instructional-design>
23. Dev To Ai (2026). “Teaching AI Fluency Anthropic.” <https://dev.to/arsdevel/mrj-kwrs-teaching-ai-fluency-mn-anthropic-50i4>
24. Arxiv Cs Cy (2026). “Generative-AI and the transformation of workforce. A job postings-driven analysis.” <https://arxiv.org/abs/2605.00843>
25. Elearning Industry (2026). “The Hidden Risk In AI-Driven Learning.” <https://elearningindustry.com/ai-tunnel-vision-the-hidden-risk-in-ai-driven-learning>
26. Elearning Industry (2026). “Vibe Coding In eLearning: How AI Is Transforming Interactive Course Creation.” <https://elearningindustry.com/vibe-coding-in-elearning-how-ai-is-transforming-interactive-course-creation>
27. EDUCAUSE Review (2026). “Expanding OER with GenAI.” <https://er.educause.edu/articles/2026/5/expanding-oer-with-genai>
28. Arxiv Cs Cy (2026). “The Pedagogy of AI Mistakes: Fostering Higher-Order Thinking.” <https://arxiv.org/abs/2605.05472>
29. Elearning Industry (2026). “Tunnel Vision, Consolidated Truth, And The Illusion Of Learning.” <https://elearningindustry.com/tunnel-vision-consolidated-truth-and-the-illusion-of-learning>
30. Dev To Ai (2026). “The Death of Junior Developers? How AI Coding Agents Are Reshaping the Job Market.” https://dev.to/divyesh_vekariya/the-death-of-junior-developers-how-ai-coding-agents-are-reshaping-the-job-market-42ai

Facilities & Property Management

capability_id: facilities-property-management · created: 2026-09-22 · revised: 2026-09-22 · refs: 10

Near (2026–2031)

For a subset of institutions adopting physical AI and for those in US regions subject to emerging land-use policies, the delivery mechanism for property maintenance, campus security, and estate development may shift [1][3][4]. In property maintenance, the adoption of robot mowers may replace specific manual labor tasks in groundskeeping [1], while AI-driven transcription for technician voice notes may streamline documentation workflows by automating the conversion of raw audio into service summaries [2]. Sensing capabilities in robotic systems, such as the ability to read gauges and thermometers, may automate the inspection of industrial plant equipment [5]. For campus security and space utilisation, the deployment of physical AI may alter perimeter monitoring [3], and AI video analytics may transform traditional CCTV into intelligent monitoring systems [8]. However, the reliance on facial recognition for security may be constrained by emerging evidence of high error rates and lagging regulatory oversight [9][10]. Predictive intelligence in large venues may change how facility conditions are managed through the forecasting of behavioral patterns [7]. In estate development and management, AI platforms may accelerate the material takeoff process by quantifying materials for project budgeting [6], though state and local policy interventions in the US may constrain the siting and expansion of AI-specific data center facilities [4]. Processes such as mail management and commercial tenancy are unlikely to change yet because they remain dependent on manual administrative protocols and established contractual law.

Hypothesis: In the near term, a subset of institutions may reduce operational labor in maintenance and security through targeted automation while facing increased regional regulatory constraints on the development of AI-specific physical infrastructure and biometric surveillance.

Mid (2036–2041)

For institutions implementing these systems, if current trajectories in predictive intelligence and robotics hold, the scope of automated surveillance and maintenance may expand while estate development for AI facilities remains contingent on local policy compliance [3][4][8]. The process of property maintenance may shift toward continuous robotic auditing of facility health as sensing capabilities for industrial equipment mature [5]. In estate development, the use of AI for project estimation may move from material takeoffs to more integrated budgetary forecasting [6], while the development of data centers may require more complex negotiations with local municipalities regarding resource extraction and energy usage [4]. Robotics in groundskeeping may move beyond mowing to encompass a broader range of repetitive exterior vegetation management [1]. Long-term campus planning and architectural synthesis are likely to remain human-led because these processes require alignment with institutional strategy and aesthetic values.

Hypothesis: In the mid term, routine exterior maintenance, industrial plant inspection, and perimeter security may become predominantly automated for early adopters, while the development of AI-centric facilities faces heightened regional regulatory barriers.

Far (2046–2051)

For the subsegments of exterior maintenance, industrial inspection, and estate development, outcomes diverge based on the tension between automation efficiency and environmental or social policy [1][4][5]. One scenario involves a transition to a highly automated facility management regime where human presence is limited to oversight of robotic fleets. A rival scenario involves a contraction of automated systems due to the failure of biometric reliability and the imposition of strict land-use mandates that limit the energy-intensive infrastructure required to support physical AI [4][10]. It remains unknown whether institutional social legitimacy will persist if campus security is managed primarily by autonomous physical agents. The legal frameworks governing commercial tenancy and property rights are likely to resist change even at this horizon, as they are rooted in statutory law rather than operational technology.

Hypothesis A: In the far term, physical facility operations may transition to a bounded autonomy model where routine maintenance and perimeter security are fully decoupled from human labor.

Hypothesis B (competing): In the far term, regulatory backlash against biometric errors and data center energy consumption may force a return to human-centric facility management and restricted infrastructure expansion.

References

1. Wired (2026). “Robot Mowers Are Actually Good Now.” <https://www.wired.com/story/robot-mowers-are-good-now-rave/>
2. Dev To Ai (2026). “From Mumbles to Memos: Teaching AI to Understand Technician Voice Notes.” https://dev.to/ken_deng_ai/from-mumbles-to-memos-teaching-ai-to-understand-technician-voice-notes-ckj
3. Ai News (2026). “Asylon and Thrive Logic bring physical AI to enterprise perimeter security.” <https://www.artificialintelligence-news.com/news/physical-ai-security-at-the-enterprise-perimeter-takes-a-step-closer/>
4. Ai Now Institute (2026). “North Star Data Center Policy Toolkit: State and Local Policy Interventions to Stop Rampant AI Data Center Expansion.” <https://ainowinstitute.org/publications/data-center-policy-guide>
5. Ars Technica (2026). “Boston Dynamics’ robot dog now reads gauges and thermometers with Google’s AI.” <https://arstechnica.com/ai/2026/04/robot-dogs-now-read-gauges-and-thermometers-using-google-gemini/>
6. Ai News (2026). “Bobyard 2.0 offers improved takeoffs and unified AI for estimators.” <https://www.artificialintelligence-news.com/news/bobyard-2-0-offers-improved-takeoffs-and-unified-ai-for-estimators/>
7. Dev To Ai (2026). “Architecting Predictive Intelligence for Smart Venues.” https://dev.to/tanisha_khiara_4b390b7946/architecting-predictive-intelligence-for-smart-venues-1lkj
8. Dev To Ai (2026). “Integrating AI Video Analytics into Venue Operations.” https://dev.to/tanisha_khiara_4b390b7946/integrating-ai-video-analytics-into-venue-operations-2m09
9. Guardian Tech (2026). “AI facial recognition oversight lagging far behind technology, watchdogs warn.” <https://www.theguardian.com/technology/ng-interactive/2026/may/03/ai-facial-recognition-oversight-lagging-far-behind-technology-watchdogs-warn>
10. Dev To Ai (2026). “Facial Recognition’s 81% Error Rate Is About to Blow Up in Court — Are Your Notes Ready?” <https://dev.to/caracomp/facial-recognitions-81-error-rate-is-about-to-blow-up-in-court-are-your-notes-ready-5af5>

Finance Management

capability_id: finance-management · created: 2026-09-30 · revised: 2026-09-30 · refs: 22

Near (2026–2031)

For Finance Management broadly, the adoption of generative AI introduces structural volatility into budget management due to unpredictable token consumption and recursive execution loops that may result in rapid, unplanned expenditures [8][14][16]. For institutions adopting usage-based AI tools, the process of budget management shifts from the prediction of fixed annual licensing fees to the real-time tracking of variable per-token costs [10][13]. To mitigate these risks, institutions may implement judgment layers as verification gates for agent-initiated financial actions [7] and programmatic guardrails—such as hard-cap quotas and tiered security policies—to bound autonomous spending [1][3]. Within procurement and purchasing management, institutions may deploy AI-powered invoice processing pipelines to automate data extraction [6] and use AI-driven annotation to extract insights from unstructured contracts [11]. Expenditure management is likely to incorporate delegated payment mechanisms, including digital wallets [12], virtual credit cards [5], and the Amazon Bedrock AgentCore payment system [18], alongside the emergence of HTTP payment standards like x402 for machine commerce [17]. These developments, including efforts to reimagine the office of the CFO [15], constitute a change to the delivery mechanism of financial oversight through automation rather than a change to the educational substance of fiduciary stewardship. Long-term treasury management and endowment strategies are unlikely to change in this period as they remain constrained by board-approved appropriations and multi-year capital cycles.

Hypothesis: In the near term, institutions may implement programmatic guardrails and judgment layers to manage the financial volatility and governance risks associated with usage-based AI and autonomous agent expenditures.

Mid (2036–2041)

For the subset of institutions deploying autonomous financial agents, if current trajectories of granular cost attribution hold, budget management may transition from periodic human reviews toward automated, trigger-based expenditure management [1][2]. The mechanism of financial oversight may shift toward the maintenance of structured security tiers and governance policies that bound autonomous spending [3]. Procurement and purchasing management may see an increase in high-frequency agent-to-agent (A2A) transactions, shifting the human role from transaction execution to the definition of negotiation parameters and spending limits [4][9][20]. Financial reporting processes may incorporate higher-velocity data streams to respond to the rapid fluctuations inherent in token-based pricing models [8][10][13]. However, overall institutional treasury management is likely to persist unchanged, as it remains governed by long-term fiduciary duties and statutory regulatory frameworks.

Hypothesis: In the mid term, budget and procurement management for AI-enabled services may evolve into high-frequency, automated oversight processes characterized by dynamic quotas and A2A transaction protocols.

Far (2046–2051)

For the same subset of institutions, the integration of autonomous financial controls may lead to divergent outcomes. One plausible outcome is the transition of expenditure management to a model of bounded autonomy, where AI agents manage micro-budgets and execute procurement via A2A protocols and digitally signed mandates—as seen in the Google AP2 model—within strictly defined security policies [3][4][12][17][23]. A rival outcome is a systemic shift back toward rigid, fixed-cost procurement models intended to eliminate the billing anomalies and unit-economic instability associated with autonomous agent spending [8][16]. A primary unknown remains the legal attribution of financial liability—the “who pressed the button” problem—regarding agent-led contractual commitments [19][22]. Even at this horizon, the fundamental requirement for human-signed audit trails for major capital expenditures is likely to resist change due to regulatory and insurance constraints.

Hypothesis A: In the far term, expenditure management may move toward a system of bounded autonomy using digitally signed mandates and A2A protocols for operational procurement.

Hypothesis B (competing): In the far term, systemic volatility in agent-led spending may drive institutions back toward rigid, fixed-cost procurement models to ensure budgetary stability.

References

1. Dev To Ai (2026-04-07). “My agent burned \$200 in one night. So I built something that stops it.” https://dev.to/tal_nagar_f7aead23e3e291a/my-agent-burned-200-in-one-night-so-i-built-something-that-stops-it-1035
2. Aws Ml (2026-04-07). “Manage AI costs with Amazon Bedrock Projects” <https://aws.amazon.com/blogs/machine-learning/manage-ai-costs-with-amazon-bedrock-projects/>
3. Dev To Ai (2026-04-08). “21 Policy Types, 4 Security Tiers: Building Bulletproof AI Agent Guardrails” <https://dev.to/walletguy/21-policy-types-4-security-tiers-building-bulletproof-ai-agent-guardrails-3lg6>
4. Dev To Ai (2026-04-09). “Building the Agent Economy: A2A Protocol Explained” <https://dev.to/mranderson323/building-the-agent-economy-a2a-protocol-explained-4nfd>
5. Dev To Ai (2026-04-09). “AI Agents Now Have Credit Cards” <https://dev.to/tomwangcn/ai-agents-now-have-credit-cards-gn3>
6. Hackernoon Ai (2026-04-09). “Building an AI-Powered Invoice Processing Pipeline” <https://hackernoon.com/building-an-ai-powered-invoice-processing-pipeline?source=rss>
7. Dev To Ai (2026-04-11). “Judgment Layer for Financial AI Agents” <https://dev.to/thatechmaestro/judgment-layer-for-financial-ai-agents-141>
8. Hackernoon Ai (2026-04-23). “The Midnight Billing Anomaly: How AI Agents Quietly Destroy Unit Economics” <https://hackernoon.com/the-midnight-billing-anomaly-how-ai-agents-quietly-destroy-unit-economics?source=rss>
9. Dev To Ai (2026-04-26). “Anthropic Just Proved Agents Can Do Commerce. Here’s What They Didn’t Build.” https://dev.to/aaron_schnieder_4563d5d33/anthropic-just-proved-agents-can-do-commerce-heres-what-they-didnt-build-5209
10. Ars Technica (2026-04-28). “GitHub will start charging Copilot users based on their actual AI usage” <https://arstechnica.com/ai/2026/04/github-will-start-charging-copilot-users-based-on-their-actual-ai-usage/>
11. Aws Ml (2026-04-29). “Extracting contract insights with PwC’s AI-driven annotation on AWS” <https://aws.amazon.com/blogs/machine-learning/extracting-contract-insights-with-pwcs-ai-driven-annotation-on-aws/>
12. TechCrunch (2026-04-30). “Stripe introduces Link, a digital wallet that autonomous AI agents can use, too” <https://techcrunch.com/2026/04/30/stripe-link-digital-wallet-ai-agents-shopping/>
13. Ai News (2026-05-01). “Per-token AI charges come to GitHub Copilot” <https://www.artificialintelligence-news.com/news/per-token-ai-charging-comes-to-github-copilot/>
14. Dev To Ai (2026-05-01). “Your AI Agent Is Sending 10x More API Calls Than You Think — Here’s Where the Cost Hides” <https://dev.to/xidao/your-ai-agent-is-sending-10x-more-api-calls-than-you-think-heres-where-the-cost-hides-4mei>
15. OpenAI (2026-05-04). “OpenAI and PwC collaborate to reimagine the office of the CFO” <https://openai.com/index/openai-pwc-finance-collaboration>
16. Zdnet Ai (2026-05-05). “What you’ll pay for AI agents will be wildly variable and unpredictable” <https://www.zdnet.com/article/your-cost-for-ai-agents-will-be-wildly-variable-and-unpredictable/>
17. Dev To Ai (2026-05-06). “x402 Protocol: The HTTP Payment Standard That Enables Autonomous Machine Commerce” <https://dev.to/walletguy/x402-protocol-the-http-payment-standard-that-enables-autonomous-machine-commerce-1h0a>
18. Aws Ml (2026-05-07). “Agents that transact: Introducing Amazon Bedrock AgentCore payments, built with Coinbase and Stripe” <https://aws.amazon.com/blogs/machine-learning/agents-that-transact/>

- [transact-introducing-amazon-bedrock-agentcore-payments-built-with-coinbase-and-stripe/](#)
19. Hackernoon Ai (2026-05-08). “Who Pressed the Button? W3.io’s Porter Stowell on Governing AI Agents That Move Money” <https://hackernoon.com/who-pressed-the-button-w3ios-porter-stowell-on-governing-ai-agents-that-move-money?source=rss>
 20. Dev To Ai (2026-05-09). “Anthropic Let AI Agents Negotiate Real Deals. Nobody Told Them Which Model They Had.” <https://dev.to/chandravanshi/anthropic-let-ai-agents-negotiate-real-deals-nobody-told-them-which-model-they-had-2gnn>
 21. Dev To Ai (2026-05-10). “phonepe’s integration agent and the next step” <https://dev.to/t49qnsx7qtkpanks/phonepes-integration-agent-and-the-next-step-2837>
 22. Dev To Ai (2026-05-10). “the grok 200k incident and the missing governance layer” <https://dev.to/t49qnsx7qtkpanks/the-grok-200k-incident-and-the-missing-governance-layer-5bj1>
 23. Dev To Ai (2026-05-10). “google’s ap2 and the mandate model for agent payments” <https://dev.to/t49qnsx7qtkpanks/googles-ap2-and-the-mandate-model-for-agent-payments-1i97>

Government, Public & Stakeholder Relationships

capability_id: government-public-stakeholder-relationships · created: 2026-09-27 · revised: 2026-09-27
· refs: 2

Near (2026–2031)

For a subset of institutions interacting with government bodies through AI-mediated communications, government relationship management may face risks of semantic inaccuracy because AI systems can misinterpret structured content when attribution is absent [2]. This represents a change in the delivery mechanism of official communications, where the efficiency of AI-generated responses is offset by a potential loss of precision in regulatory filings or inquiries. Simultaneously, for institutions facing direct state-level legal inquiries related to AI-enabled harms, as seen in the criminal investigation involving Florida State University, government relationship management may shift from routine coordination to adversarial legal defense [1]. This shift alters the educational substance of the capability, as the relationship is redefined by criminal liability rather than policy alignment. Media relations management is concurrently pressured to pivot from promotional messaging to crisis communication to address the alleged role of AI in campus violence [1]. Because these signals are institution-specific or limited to specific technical subsegments, industry relationship management for general operational AI tools is unlikely to change across the broader sector.

Hypothesis: In the near term, a subset of institutions may see government relationship management shift toward risk mitigation and legal defense regarding AI vendor liabilities and the technical reliability of official AI-generated communications.

Mid (2036–2041)

For state-funded institutions, if current trajectories of state-led criminal investigations [1] and systemic misinterpretations of structured official data [2] hold, the process of government relationship management may integrate rigorous technical auditing of AI pipelines. Extrapolating from current legal and technical risks, institutions may seek to shift liability for AI-generated externalities and inaccuracies entirely onto the vendor through updated procurement frameworks. This would alter the mechanism of industry relationship management, moving it toward a model centered on strict indemnity clauses and verified attribution standards for structured data. However, internal stakeholder relationship management is unlikely to be fundamentally altered by these external pressures unless employment contracts are formally rewritten to include AI-specific liability.

Hypothesis: In the mid term, state-funded institutions may formalize liability and attribution frameworks in their government relationship management to insulate the institution from criminal probes and administrative errors caused by vendor-provided AI.

Far (2046–2051)

For institutions operating under high state oversight, the long-term state of the capability depends on whether legal precedents establish “institutional negligence” for AI-generated content and outputs [1]. In one scenario, government relationship management becomes a highly regulated compliance function where state agencies mandate specific, audited AI architectures to prevent public harm and ensure the accuracy of structured communications [2]. In a rival scenario, the emergence of comprehensive immunity for educational institutions regarding third-party software may return these relationships to a non-adversarial, administrative state. The impact on public relationship management remains an unresolvable uncertainty, as it depends on the evolving social threshold for accepting AI-driven externalities. Routine media relations for non-crisis events will likely resist change, as the basic need for institutional branding remains constant.

Hypothesis A: In the far term, government relationship management for a subset of institutions may become a highly regulated compliance function focused on AI safety audits and verified data attribution.

Hypothesis B (competing): In the far term, the legal resolution of AI liability may decouple institutional

responsibility from vendor outputs, returning government relationship management to administrative norms.

References

1. NPR (2026-04-21). “Florida AG launches criminal investigation into ChatGPT over FSU shooting.”
<https://www.npr.org/2026/04/21/nx-s1-5793967/florida-openai-investigation-mass-shooting-fsu>
2. Dev To Ai (2026-05-09). “How AI Systems Misinterpret Structured Content Without Attribution.”
<https://dev.to/aigistry/how-ai-systems-misinterpret-structured-content-without-attribution-5b3e>

Human Resource Management

capability_id: human-resource-management · created: 2026-09-25 · revised: 2026-09-25 · refs: 21

Near (2026–2031)

For human resource management broadly, AI is being integrated into the delivery mechanisms of staff recruitment and employee support, while shifts in the substance of training and development are emerging in specific institutional subsegments. In the process of staff recruitment, the delivery of candidate screening is affected by a self-preference mechanism where LLMs consistently favor AI-generated resumes over human-authored ones [21][22], introducing a risk of “LLM nepotism” in organizational governance [9] and systemic data bias [1]. In the United Kingdom, the delivery of candidate evaluation is facing a legitimacy crisis, with job hunters reporting AI-driven interviews as “humiliating” [20] and “horrible” [19], while others are excluded from the pipeline by opaque algorithmic filtering [27]. In the process of employee support, institutions are deploying agentic onboarding tools to automate administrative orientation [3], which may shift the educational substance of onboarding from a discrete event to a “perpetual” continuous model [12]. For a subset of institutions, the substance of training and development is shifting toward workforce readiness for an AI-driven future [6], utilizing “learning champions” to drive adoption through peer credibility rather than top-down mandates [26]. Workforce relations management is unlikely to change in this horizon because the negotiation of collective bargaining agreements and the legal complexities of employment contracts remain dependent on human-led diplomacy.

Hypothesis: In the near term, the capability will experience a tension between the efficiency of automated recruitment delivery and the social legitimacy of the candidate experience, alongside a transition toward continuous, agent-supported employee onboarding.

Mid (2036–2041)

For human resource management broadly, if current trajectories in productivity-linked labor trials hold, the substance of remuneration and entitlements management may shift as institutions decouple compensation from fixed hourly requirements, potentially adopting four-day week models to adapt to AI-augmented productivity [5]. In the subsegment of institutions utilizing agentic AI for human capital management, the delivery of talent management and workforce reporting may expand to the end-to-end orchestration of employee development pipelines [10][16]. For the subsegment of institutions updating performance metrics, the substance of employee performance management may shift as traditional KPIs are replaced by new metrics specifically aligned with AI-driven output [15]. In institutions relying on remote hiring, biometric identity verification may transition from an experimental tool to a standard delivery mechanism to counter the proliferation of deepfakes [2]. For a subset of institutions, the process of workforce planning may shift toward “disintegrating” the traditional organizational chart in favor of fluid, AI-supported structures [4]. Tenure-track hierarchies and faculty governance structures are likely to persist because they are deeply embedded in academic contracts and the social legitimacy of the professoriate.

Hypothesis: In the mid term, AI in HR may transition from isolated toolsets to agentic orchestration of the employee lifecycle, potentially altering the substance of performance evaluation and remuneration models.

Far (2046–2051)

For human resource management broadly, the capability may diverge based on the resolution of the tension between algorithmic efficiency and institutional labor protections. In the process of workforce planning, one outcome involves the systemic transition to a hybrid workforce where human staff manage “AI doubles” or agentic colleagues, shifting the substance of staff engagement toward human-machine relationship management [13][29]. A competing outcome involves the imposition of strict regulatory constraints on automation, as seen in region-specific mandates that prohibit firing human workers if AI takes their jobs [24]. The social legitimacy of employee performance management remains an unresolvable uncertainty, specifically whether human workers will accept AI-led emotional monitoring as a basis for evaluation [23]. Core legal requirements for a human “employer of record” and the fundamental necessity of human

accountability in firing decisions are likely to resist change even at this horizon.

Hypothesis A: In the far term, the capability shifts toward the management of a blended human-agent workforce, where HR focuses on the orchestration of AI doubles and hybrid productivity.

Hypothesis B (competing): In the far term, the capability is constrained by rigid legal protections and “human-first” employment mandates that prevent the substitution of human staff with AI systems.

References

1. Ai Trends (2021-10-28). “Promise and Perils of Using AI for Hiring: Guard Against Data Bias.” <https://www.aitrends.com/ai-world-government/promise-and-perils-of-using-ai-for-hiring-guard-against-data-bias/>
2. Dev To Ai (2026-04-06). “Why a Deepfake Face Can Fool Your Eyes in Seconds but Not 128 Landmarks at Once.” <https://dev.to/caracomp/why-a-deepfake-face-can-fool-your-eyes-in-seconds-but-not-128-landmarks-at-once-37m8>
3. Aws Ml (2026-04-06). “Build AI-powered employee onboarding agents with Amazon Quick.” <https://aws.amazon.com/blogs/machine-learning/build-ai-powered-employee-onboarding-agents-with-amazon-quick/>
4. Mit Sloan (2026-04-07). “Disintegrating the Org Chart: ServiceNow’s Jacqui Canney.” <https://sloanreview.mit.edu/audio/disintegrating-the-org-chart-servicenows-jacqui-canney/>
5. BBC News (2026-04-07). “OpenAI encourages firms to trial four-day weeks to adapt to AI era.” https://www.bbc.com/news/articles/c8x71ejrp92o?at_medium=RSS&at_campaign=rss
6. Talent Economy (2026-03-16). “From AI access to workforce readiness.” <https://www.chieflearnin gofficer.com/2026/03/16/from-ai-access-to-workforce-readiness/>
7. Arxiv Cs Cy (2026-04-14). “LLM Nepotism in Organizational Governance.” <https://arxiv.org/abs/2604.09620>
8. Ai News (2026-04-14). “SAP brings agentic AI to human capital management.” <https://www.artificialintelligence-news.com/news/sap-brings-agentic-ai-human-capital-management/>
9. Elearning Industry (2026-04-21). “The Art Of Perpetual Onboarding: Why AI Makes The Continuous Model Viable At Last.” <https://elearningindustry.com/the-art-of-perpetual-onboarding-why-ai-makes-the-continuous-model-viable-at-last>
10. MIT Technology Review (2026-04-20). “Chinese tech workers are starting to train their AI doubles—and pushing back.” <https://www.technologyreview.com/2026/04/20/1136149/chinese-tech-workers-ai-colleagues/>
11. Fast Company (2026-04-23). “The Age of AI means we need to throw out our old KPIs and replace them with new ones.” https://www.fastcompany.com/91521422/ai-kpis-new-metrics?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
12. Aws Ml (2026-04-24). “Building Workforce AI Agents with Visier and Amazon Quick.” <https://aws.amazon.com/blogs/machine-learning/building-workforce-ai-agents-with-visier-and-amazon-quick/>
13. Guardian Tech (2026-05-01). “‘Completely horrible’: UK job hunters share frustration with AI interviews.” <https://www.theguardian.com/technology/2026/may/01/uk-job-hunters-frustration-ai-interviews>
14. Guardian Tech (2026-05-01). “‘Awkward and humiliating’: UK job hunters share frustration with AI interviews.” <https://www.theguardian.com/technology/2026/may/01/uk-job-hunters-frustration-ai-interviews>
15. Hacker News (2026-05-02). “LLMs consistently pick resumes they generate over ones by humans or other models.” <https://arxiv.org/abs/2509.00462>
16. Hacker News (2026-05-02). “AI Self-preferencing in Algorithmic Hiring: Empirical Evidence and Insights.” <https://arxiv.org/abs/2509.00462>
17. Atlantic (2026-05-03). “What if Your Boss Monitored Your Emotions?” https://www.theatlantic.com/culture/2026/05/worker-surveillance-emotion-ai/687029/?utm_source=feed
18. The Register (2026-05-04). “Just in time for Labour Day, China makes it illegal to fire humans if

- AI takes their jobs.” https://go.theregister.com/feed/www.theregister.com/2026/05/04/asia_tech_news_roundup/
19. Elearning Industry (2026-05-04). “Internal Learning Champions: Why Peer Credibility Beats Top-Down Training In AI Adoption.” <https://elearningindustry.com/internal-learning-champions-why-peer-credibility-beats-top-down-training-in-ai-adoption>
 20. Wired (2026-05-05). “He Couldn’t Land a Job Interview. Was AI to Blame?” <https://www.wired.com/story/he-couldnt-land-a-job-interview-was-ai-to-blame/>
 21. The Conversation Tech (2026-05-06). “So your new ‘co-worker’ is an AI agent – here’s how to make the best of your human-machine relationship.” <https://theconversation.com/so-your-new-co-worker-is-an-ai-agent-heres-how-to-make-the-best-of-your-human-machine-relationship-276011>

Information & Communication Technology Mgt

capability_id: information-communication-technology-management · created: 2026-05-08 · revised: 2026-05-08 · refs: 43

Near (2026–2031)

For Information & Communication Technology Mgt broadly, application lifecycle management shifts toward the adoption of the Model Context Protocol (MCP) to standardize the interface between large language models and institutional data [90][92][100]. This change in delivery mechanism replaces proprietary connectors with a universal tool-calling interface [49][101], but it introduces vulnerabilities such as command injection and information disclosure that require specific auditing of MCP servers [12][130] and vetting for security grades [10]. The educational substance of quality assurance shifts from deterministic testing to frameworks that validate non-deterministic AI outputs and red-team the agent stack [62][78][91], as traditional chaos engineering proves insufficient for AI systems [142]. For ict service support, the triage of support tickets and SRE incident response are increasingly mediated by autonomous agents [11][109], supported by LLM-augmented knowledge bases for root cause analysis [58]. However, a gap in observability has emerged where standard orchestration tools like Kubernetes may report healthy status while AI pipelines fail silently, necessitating new AI-specific monitoring protocols [153]. For infrastructure lifecycle management, the emergence of AI-driven vulnerability discovery—demonstrated by the identification of hundreds of zero-day flaws in major software by autonomous models [124][129][131]—forces a shift from scheduled patching to event-driven, high-frequency update cycles [135]. Identity & access management incorporates agent registries to manage the sprawl of autonomous agents [71][73], while the shift toward multi-agent systems necessitates the adoption of zero-trust security runtimes [41][87][149], sandbox execution [113], and cryptographic authorization to prevent identity spoofing [111][150]. For a subset of institutions deploying autonomous agents, application lifecycle management now requires new financial guardrails to prevent runaway automated billing and “midnight billing anomalies” that destroy unit economics [154]. Physical network cabling and hardware rack installations are unlikely to change yet, as disruption is concentrated at the logical orchestration and software layers.

Hypothesis: In the near term, the capability will pivot from managing static software installations to governing a dynamic ecosystem of standardized AI connectors and managing a high-frequency vulnerability response cycle driven by AI-automated exploit discovery.

Mid (2036–2041)

For Information & Communication Technology Mgt broadly, if current trajectories hold, application lifecycle management may evolve toward model-agnostic architectures [7] to mitigate operational risks associated with “tool refugee” patterns, where institutions must migrate when proprietary AI tools are deprecated [44]. The substance of application lifecycle management may further shift as AI agents automate the modernization and migration of legacy codebases [20][54]. For institutions managing decentralized workloads, infrastructure lifecycle management may shift toward strengthened governance for edge AI workloads [106] to manage the proliferation of autonomous processing at the network periphery [67][118], potentially utilizing KubeVirt to run virtual machines on Kubernetes at the edge [68] and reorganizing hardware specifically for agentic AI platforms [75]. In regions governed by the EU AI Act, the management of technical files for high-risk AI systems becomes a core regulatory compliance process within this capability [27][66]. Institutional procurement cycles for physical hardware are unlikely to change significantly, as these processes remain embedded in multi-year budget cycles and facility contracts.

Hypothesis: In the mid term, the capability may shift toward the management of autonomous infrastructure that self-optimizes for edge workloads while incorporating strict regulatory auditing of AI technical files.

Far (2046–2051)

For Information & Communication Technology Mgt broadly, the capability may reach a point where the distinction between application and infrastructure lifecycle management collapses into a single autonomous

orchestration process. One rival outcome is that the acceleration of AI-driven vulnerability discovery leads to a “security paradox,” where institutions abandon open connectivity in favor of highly fragmented, air-gapped, or sovereign infrastructure to ensure survival [28][81]. It remains unknown whether standardized protocols like MCP will persist or be replaced by emergent, self-evolving agent communication languages. Hardware replacement cycles for core servers will likely resist rapid change even at this horizon due to the physical constraints of energy delivery and data center cooling [37].

Hypothesis A: In the far term, the capability becomes a supervisory function over a fully autonomous, self-healing ICT stack that manages its own deployment, patching, and hardware optimization.

Hypothesis B (competing): In the far term, the capability shifts toward managing highly restricted, sovereign silos of technology to defend against autonomous, large-scale AI cyber-attacks.

References

1. Edtech Magazine (2026-04-03). “Case Western Reserve University Deploys Google Gemini.” <https://edtechmagazine.com/higher/article/2026/04/case-western-reserve-university-deploys-google-gemini>
2. Dev To Ai (2026-04-06). “Building Model-Agnostic AI Architecture: The Pattern That Future-Proofs Your System.” https://dev.to/ailoitte_sk/building-model-agnostic-ai-architecture-the-pattern-that-future-proofs-your-system-4o84
3. Dev To Ai (2026-04-08). “Why we’re building the AI Tool Refugee Center: a place to land when your tool dies.” <https://dev.to/zsky/why-were-building-the-ai-tool-refugee-center-a-place-to-land-when-your-tool-dies-1e8>
4. Dev To Ai (2026-04-07). “The ‘Legacy Code’ Nightmare is Over: How AI Agents are Automating App Modernization.” https://dev.to/siddhesh_surve/the-legacy-code-nightmare-is-over-how-ai-agents-are-automating-app-modernization-33j8
5. Dzone Ai (2026-04-08). “AI-Assisted Code Migration: Practical Techniques for Modernizing Legacy Systems.” <https://feeds.dzone.com/link/23558/17316144/ai-assisted-code-migration-practical-techniques>
6. Ai News (2026-04-13). “Strengthening enterprise governance for rising edge AI workloads.” <https://www.artificialintelligence-news.com/news/strengthening-enterprise-governance-for-rising-edge-ai-workloads/>
7. Dev To Ai (2026-04-15). “AI Is Moving Frontend Logic to the Browser and Edge — Are You Ready?” https://dev.to/rohith_kn/ai-is-moving-frontend-logic-to-the-browser-and-edge-are-you-ready-7ic
8. Ai Business (2026-04-21). “AMD’s Vision for AI PCs in the Age of Agentic AI.” <https://aibusiness.com/generative-ai/amd-s-vision-for-ai-pcs-age-of-agentic-ai>
9. The Register (2026-04-09). “Nutanix to add KubeVirt support to run VMs on K8s at the edge.” https://go.theregister.com/feed/www.theregister.com/2026/04/09/nutanix_kubevirt_arm_supoprnt/
10. Dzone Ai (2026-04-09). “How Agentic AI Platforms Organize Their Hardware Infrastructure.” <https://dzone.com/link/23558/17316612/how-agentic-ai-platforms-organize-hardware-infra>
11. Dev To Ai (2026-04-07). “What’s actually inside an EU AI Act Annex IV technical file (and what engineers usually miss).” https://dev.to/rostyslav_myronenko_faf2f/whats-actually-inside-an-eu-ai-act-annex-iv-technical-file-and-what-engineers-usually-miss-4nh4
12. Ai News (2026-04-09). “Agentic AI’s governance challenges under the EU AI Act in 2026.” <https://www.artificialintelligence-news.com/news/agentic-ais-governance-challenges-under-the-eu-ai-act-in-2026/>
13. The Register (2026-04-07). “OpenInfra General Manager talks sovereignty, governments deploying tech ‘kill switches.’” https://go.theregister.com/feed/www.theregister.com/2026/04/07/thierry_carrez_on_sovereignty_interview/
14. Ai Business (2026-04-10). “Anthropic’s Glasswing Highlights AI’s Security Paradox.” <https://aibusiness.com/generative-ai/anthropic-s-glasswing-highlights-ai-s-security-paradox>
15. Ai Now Institute (2026-04-01). “North Star Data Center Policy Toolkit: State and Local Policy

- Interventions to Stop Rampant AI Data Center Expansion.” <https://ainowinstitute.org/publications/data-center-policy-guide>
16. Dev To Ai (2026-04-11). “MCP Ecosystem in 2026: \$1.8B Market and 1000+ Connectors Reshaping AI Integration.” <https://dev.to/chunxiaox/mcp-ecosystem-in-2026-18b-market-and-1000-connectors-reshaping-ai-integration-4n3l>
 17. Dev To Ai (2026-04-11). “The Model Context Protocol in 2025: From Anthropic to Linux Foundation.” <https://dev.to/chunxiaox/the-model-context-protocol-in-2025-from-anthropic-to-linux-foundation-2og1>
 18. Dev To Ai (2026-04-11). “The Future of AI Integration: Model Context Protocol (MCP) Opportunities.” <https://dev.to/chunxiaox/the-future-of-ai-integration-model-context-protocol-mcp-opportunities-148b>
 19. Dzone Ai (2026-04-08). “MCP + AWS AgentCore: Give Your AI Agent Real Tools in 60 Minutes.” <https://feeds.dzone.com/link/23558/17315883/mcp-aws-agentcore>
 20. Dev To Ai (2026-04-11). “The State of AI Agent MCP Connector Frameworks in 2026.” <https://dev.to/chunxiaox/the-state-of-ai-agent-mcp-connector-frameworks-in-2026-2ea7>
 21. Dev To Ai (2026-04-06). “CVE-2026-32211 | Azure MCP Server Information Disclosure Vulnerability.” https://dev.to/aakash_rahsi/cve-2026-32211-azure-mcp-server-information-disclosure-vulnerability-2280
 22. Dev To Ai (2026-04-22). “How to Check Your MCP Server for CVE-2026-5603’s Vulnerability Pattern (And Why shellQuote Isn’t Enough).” <https://dev.to/armorlai/how-to-check-your-mcp-server-for-cve-2026-5603s-vulnerability-pattern-and-why-shellquote-isnt-gdn>
 23. Dev To Ai (2026-04-06). “Only 20% of MCP Servers Are ‘A-Grade’ Secure — Here’s How to Vet Them Before Installing.” https://dev.to/gentic_news/only-20-of-mcp-servers-are-a-grade-secure-heres-how-to-vet-them-before-installing-jck
 24. Dev To Ai (2026-04-09). “OpenAI’s Promptfoo deal puts evaluation and red-teaming at the centre of the agent stack.” https://dev.to/yujia_zhang_0328/openais-promptfoo-deal-puts-evaluation-and-red-teaming-at-the-centre-of-the-agent-stack-3o19
 25. Dev To Ai (2026-04-10). “AI Testing and Quality Assurance in 2026: Ensuring AI System Reliability.” <https://dev.to/lufumeiying/ai-testing-and-quality-assurance-in-2026-ensuring-ai-system-reliability-8ka>
 26. Dev To Ai (2026-04-11). “OpenAI Codex Testing: How to QA AI-Written Code.” https://dev.to/hai_huang_f196ed9669351e0/openai-codex-testing-how-to-qa-ai-written-code-31n5
 27. Dzone Ai (2026-04-22). “What AI Systems Taught Us About the Limits of Chaos Engineering.” <https://feeds.dzone.com/link/23558/17323679/chaos-engineering-limits>
 28. Dev To Ai (2026-04-06). “How to Automate Ticket Triage with AI (Without Breaking Everything).” https://dev.to/erik_anderson_c41dbafd423/how-to-automate-ticket-triage-with-ai-without-breaking-everything-4635
 29. Dzone Ai (2026-04-14). “Building an AI-Powered SRE Incident Response Workflow With AWS Strands Agents.” <https://feeds.dzone.com/link/23558/17318982/ai-powered-sre-incident-response-aws-strands-agents>
 30. Arxiv Cs Cl (2026-04-09). “LLM-Augmented Knowledge Base Construction For Root Cause Analysis.” <https://arxiv.org/abs/2604.06171>
 31. Hackernoon Ai (2026-04-23). “I Watched Our AI Pipeline Silently Fail While Kubernetes Said Everything Was Fine.” <https://hackernoon.com/i-watched-our-ai-pipeline-silently-fail-while-kubernetes-said-everything-was-fine?source=rss>
 32. Wired (2026-04-21). “Mozilla Used Anthropic’s Mythos to Find and Fix 271 Bugs in Firefox.” <https://www.wired.com/story/mozilla-used-anthropics-mythos-to-find-271-bugs-in-firefox>
 33. Ars Technica (2026-04-21). “Mozilla: Anthropic’s Mythos found 271 security vulnerabilities in Firefox 150.” <https://arstechnica.com/ai/2026/04/mozilla-anthropics-mythos-found-271-zero-day-vulnerabilities-in-firefox-150/>

34. The Register (2026-04-22). “Mythos found 271 Firefox flaws – but none a human couldn’t spot.” https://go.theregister.com/feed/www.register.com/2026/04/22/mozilla_firefox_mythos_future_defenders/
35. Fast Company (2026-04-22). “Brace yourself for a flood of patches in all of your tech gadgets.” https://www.fastcompany.com/91530191/brace-yourself-for-a-flood-of-patches-in-all-of-your-tech-gadgets?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
36. The Register (2026-04-09). “AWS: Agents shouldn’t be secret, so we built a registry for them.” https://go.theregister.com/feed/www.theregister.com/2026/04/09/aws_ai_agent_registry/
37. Aws Ml (2026-04-09). “The future of managing agents at scale: AWS Agent Registry now in preview.” <https://aws.amazon.com/blogs/machine-learning/the-future-of-managing-agents-at-scale-aws-agent-registry-now-in-preview/>
38. Dev To Ai (2026-04-08). “I built an open-source zero-trust security runtime for AI agents. Here’s what I learned.” https://dev.to/edoardo_bambini_5badce6ff/i-built-an-open-source-zero-trust-security-runtime-for-ai-agents-heres-what-i-learned-538e
39. Dev To Ai (2026-04-10). “Cert-gating every tool call: zero-trust for AI agents.” <https://dev.to/1r0nw1ll/cert-gating-every-tool-call-zero-trust-for-ai-agents-8pg>
40. Dev To Ai (2026-04-22). “API keys were designed for humans. AI agents break them in 4 ways.” <https://dev.to/allforscience/api-keys-were-designed-for-humans-ai-agents-break-them-in-4-ways-422>
41. Ai News (2026-04-16). “OpenAI Agents SDK improves governance with sandbox execution.” <https://www.artificialintelligence-news.com/news/openai-agents-sdk-improves-governance-sandbox-execution/>
42. The Register (2026-04-16). “Git identity spoof fools Claude into giving bad code the nod.” https://go.theregister.com/feed/www.theregister.com/2026/04/16/git_identity_spoof_claude/
43. Dev To Ai (2026-04-23). “Add cryptographic authorization to AI agents in 5 minutes.” <https://dev.to/allforscience/add-cryptographic-authorization-to-ai-agents-in-5-minutes-49p1>
44. Hackernoon Ai (2026-04-23). “The Midnight Billing Anomaly: How AI Agents Quietly Destroy Unit Economics.” <https://hackernoon.com/the-midnight-billing-anomaly-how-ai-agents-quietly-destroy-unit-economics?source=rss>

Information Management

capability_id: information-management · created: 2026-06-05 · revised: 2026-06-05 · refs: 38

Near (2026–2031)

For information management broadly, the processes of information search & discovery, information & data security management, and information collection management are shifting toward agentic delivery mechanisms. In search & discovery, institutions are adopting hybrid RAG solutions [13] and university-specific knowledge bases [40] to unify fragmented data sources, while transforming ephemeral meeting audio into searchable institutional knowledge [152]. This represents a shift in educational substance: the value of discovery is pivoting toward technical accuracy as a primary quality metric [9], necessitating the use of AI citation registries to maintain provenance [15] and a shift in discovery patterns toward high-intent, agent-driven retrieval [49]. This process is constrained by the persistence of stale institutional records that feed inaccurate AI summaries [112] and the risk of agent poisoning via malicious web pages [138]. Within security management, the deployment of Model Context Protocol (MCP) servers [60] has introduced vulnerabilities including information disclosure [11], prompt injection [19], and server-side vulnerabilities [10], driving the deployment of tool-call firewalls [36] and autonomous agent governance tools [14][116]. For institutions implementing autonomous agents, security substance is shifting toward the management of non-human identities (NHI) [86][135] and the implementation of zero-trust runtimes [29][57] to prevent the exfiltration of secrets by coding agents [75][104]. The capacity of AI models to identify hundreds of zero-day vulnerabilities at scale [87][88][150] is shifting the substance of security from reactive patching to autonomous auditing [101]. For a subset of institutions managing intellectual property, copyright management is seeing an emergence of “prompt non-compete clauses,” where prompt libraries created by employees using institutional resources are claimed as institutional property [151]. Information collection management is shifting its delivery mechanism by replacing selector-based web scraping with AI browser agents [137], alongside a shift toward “data activation” [22] and automated document extraction for large PDF archives [24]. Records management is unlikely to change in its statutory requirements due to legal mandates and regulatory regimes.

Hypothesis: In the near term, information management will transition to a hybrid RAG-based discovery model and a zero-trust security framework focused on non-human identity management and autonomous vulnerability auditing.

Mid (2036–2041)

For information management broadly, if current trajectories in agentic governance and zero-trust runtimes hold, the capability may shift toward autonomous data activation [22] where the distinction between collecting and using data disappears. This represents a shift in substance where the unit of information is no longer a static record but a dynamic, agent-accessible stream. The delivery of security management will likely move toward cryptographic authorization for all agent-to-agent communications [109][131] as traditional API keys are rendered insufficient by agentic access patterns [108]. Institutional inertia regarding privacy laws and the difficulty of responding to complex regulatory inquiries, such as those from the EDPB [51], will likely persist, maintaining human-in-the-loop requirements for high-stakes PII redaction. The risk of systemic prompt injection in RAG pipelines [97] and indirect prompt injection [115] may constrain the adoption of open knowledge bases in favor of gated, high-assurance environments. A requirement for a robust “data fabric” [92][140] to support these streams will persist as the primary architectural constraint.

Hypothesis: In the mid term, the capability will evolve from managing static repositories to managing dynamic data streams, governed by cryptographic agent identities rather than user-based permissions.

Far (2046–2051)

For information management broadly, the long-term state of the capability depends on the resolution of the tension between autonomous discovery and data integrity. One plausible outcome is the emergence of a fully autonomous information fabric where knowledge management is a real-time, self-optimizing process.

However, the persistence of nation-state aggression targeting AI agents [153] and the risk of catastrophic data poisoning [138] may instead lead to a fragmented landscape of air-gapped, sovereign knowledge bases. It remains unknown whether institutional trust can be maintained through automated auditing alone or if a return to human-verified records is required for social legitimacy. The statutory core of records management will likely resist change even at this horizon due to the fundamental legal requirement for immutable, human-attested archives.

Hypothesis A: In the far term, information management becomes a transparent, autonomous utility where discovery and collection are fused into a single, self-updating institutional intelligence.

Hypothesis B (competing): In the far term, security failures and poisoning attacks force a regression to highly restricted, fragmented, and human-curated “islands of trust” for critical institutional data.

References

3. Ai News (2026-04-02). “KiloClaw targets shadow AI with autonomous agent governance.” <https://www.artificialintelligence-news.com/news/kiloclaw-targets-shadow-ai-autonomous-agent-governance/>
4. Dev To Ai (2026-04-06). “The 2026 AI Governance Framework: Why Technical Accuracy is the New SEO Gold Standard.” <https://dev.to/mirakade/the-2026-ai-governance-framework-why-technical-accuracy-is-the-new-seo-gold-standard-1fn6>
5. Dev To Ai (2026-04-06). “Only 20% of MCP Servers Are ‘A-Grade’ Secure — Here’s How to Vet Them Before Installing.” https://dev.to/gentic_news/only-20-of-mcp-servers-are-a-grade-secure-heres-how-to-vet-them-before-installing-jck
6. Dev To Ai (2026-04-06). “CVE-2026-32211 | Azure MCP Server Information Disclosure Vulnerability.” https://dev.to/aakash_rahsi/cve-2026-32211-azure-mcp-server-information-disclosure-vulnerability-2280
7. Aws ML (2026-04-06). “Building Intelligent Search with Amazon Bedrock and Amazon OpenSearch for hybrid RAG solutions.” <https://aws.amazon.com/blogs/machine-learning/building-intelligent-search-with-amazon-bedrock-and-amazon-opensearch-for-hybrid-rag-solutions/>
8. Dev To Ai (2026-04-06). “AI agent governance tools compared - 2026 landscape.” <https://dev.to/jagmarques/ai-agent-governance-tools-compared-2026-landscape-53hm>
9. Dev To Ai (2026-04-06). “AI Citation Registries as Information Infrastructure for AI Systems.” <https://dev.to/aigistry/ai-citation-registries-as-information-infrastructure-for-ai-systems-50o6>
10. Ai News (2026-04-07). “Boomi calls it “data activation” and says it’s the missing step in every AI deployment.” <https://www.artificialintelligence-news.com/news/boomi-agentic-ai-data-activation-missing-step/>
11. Towards Data Science (2026-04-07). “From 4 Weeks to 45 Minutes: Designing a Document Extraction System for 4,700+ PDFs.” <https://towardsdatascience.com/from-4-weeks-to-45-minutes-designing-a-document-extraction-system-for-4700-pdfs/>
12. Dev To Ai (2026-04-08). “I built an open-source zero-trust security runtime for AI agents. Here’s what I learned.” https://dev.to/edoardo_bambini_5badce6ff/i-built-an-open-source-zero-trust-security-runtime-for-ai-agents-heres-what-i-learned-538e
13. Dev To Ai (2026-04-08). “MCP Security: Why Your AI Agent’s Tool Calls Need a Firewall.” <https://dev.to/aymenhmaidi/mcp-security-why-your-ai-agents-tool-calls-need-a-firewall-3h48>
14. Dev To Ai (2026-04-09). “Best AI Chatbot for University Knowledge Bases (2026): CustomGPT.ai vs ChatGPT vs Claude.” https://dev.to/alexandria_carlson_21fcf8/best-ai-chatbot-for-university-knowledge-bases-2026-customgptai-vs-chatgpt-vs-claude-4i0m
15. Dev To Ai (2026-04-10). “How to Get Cited by AI Search Engines: ChatGPT, Perplexity & Claude in 2026.” <https://dev.to/iris1031/how-to-get-cited-by-ai-search-engines-chatgpt-perplexity-claude-in-2026-21p0>
16. Dev To Ai (2026-04-10). “The EDPB Is Asking About Your AI Agents. Most Teams Can’t Answer.” <https://dev.to/waxell/the-edpb-is-asking-about-your-ai-agents-most-teams-cant-answer-gfk>

17. Dev To Ai (2026-04-10). "Cert-gating every tool call: zero-trust for AI agents." <https://dev.to/1r0nw1ll/cert-gating-every-tool-call-zero-trust-for-ai-agents-8pg>
18. Dev To Ai (2026-04-11). "MCP Ecosystem in 2026: \$1.8B Market and 1000+ Connectors Reshaping AI Integration." <https://dev.to/chunxiaox/mcp-ecosystem-in-2026-18b-market-and-1000-connectors-reshaping-ai-integration-4n3l>
19. Dzone Ai (2026-04-14). "4 Ways Your AI Coding Agent Exfiltrates Secrets." <https://feeds.dzone.com/link/23558/17319058/ways-your-ai-coding-agent-exfiltrates-secrets>
20. Dev To Ai (2026-04-21). "The AI Agent Identity Crisis Is Here. The Infrastructure to Solve It Already Exists." https://dev.to/aaron_schnieder_4563d5d33/the-ai-agent-identity-crisis-is-here-the-infrastructure-to-solve-it-already-exists-35ha
21. Ars Technica (2026-04-21). "Mozilla: Anthropic's Mythos found 271 zero-day vulnerabilities in Firefox 150." <https://arstechnica.com/security/2026/04/mozilla-anthropics-mythos-found-271-zero-day-vulnerabilities-in-firefox-150/>
22. Ars Technica (2026-04-21). "Mozilla: Anthropic's Mythos found 271 security vulnerabilities in Firefox 150." <https://arstechnica.com/security/2026/04/mozilla-anthropics-mythos-found-271-zero-day-vulnerabilities-in-firefox-150/>
23. MIT Technology Review (2026-04-22). "AI needs a strong data fabric to deliver business value." <https://www.technologyreview.com/2026/04/22/1135295/ai-needs-a-strong-data-fabric-to-deliver-business-value/>
24. Dev To Ai (2026-04-22). "RAG Pipelines Are the Next Prompt Injection Frontier." https://dev.to/cor_e_fc39e7a55045_c0/rag-pipelines-are-the-next-prompt-injection-frontier-kpf
25. Ai News (2026-04-22). "Reversing enterprise security costs with AI vulnerability discovery." <https://www.artificialintelligence-news.com/news/reversing-enterprise-security-costs-with-ai-vulnerability-discovery/>
26. Dev To Ai (2026-04-22). "AI Coding Assistant Security: Your Keys Are Leaking!" <https://dev.to/tahosin/ai-coding-assistant-security-your-keys-are-leaking-340k>
27. Dev To Ai (2026-04-22). "API keys were designed for humans. AI agents break them in 4 ways." <https://dev.to/allforscience/api-keys-were-designed-for-humans-ai-agents-break-them-in-4-ways-422>
28. Dev To Ai (2026-04-23). "Add cryptographic authorization to AI agents in 5 minutes." <https://dev.to/allforscience/add-cryptographic-authorization-to-ai-agents-in-5-minutes-49p1>
29. The Register (2026-04-23). "Stale gov.uk pages are feeding AI overviews old data and Brits are believing it." https://go.theregister.com/feed/www.theregister.com/2026/04/23/stale_govuk_pages_are_feeding/
30. Zdnet Ai (2026-04-24). "How indirect prompt injection attacks on AI work - and 6 ways to shut them down." <https://www.zdnet.com/article/how-indirect-prompt-injection-attacks-on-ai-work-and-6-ways-to-shut-them-down/>
31. Marktechpost (2026-04-24). "Mend Releases AI Security Governance Framework: Covering Asset Inventory, Risk Tiering, AI Supply Chain Security, and Maturity Model." <https://www.marktechpost.com/2026/04/23/mend-releases-ai-security-governance-framework/>
32. Dev To Ai (2026-04-25). "Building a Cryptographically Secure 4-Agent E-Commerce System: A Practical Guide." <https://dev.to/allforscience/building-a-cryptographically-secure-4-agent-e-commerce-system-a-practical-guide-271c>
33. Dev To Ai (2026-04-27). "AI agents have no identity — we built the open registry that gives them one." https://dev.to/zhao_b43c3903c4d2d54/ai-agents-have-no-identity-we-built-the-open-registry-that-gives-them-one-3e9h
34. Dev To Ai (2026-04-27). "AI Browser Agents in 2026: Stagehand vs Browser Use vs Playwright." <https://dev.to/alexcloudstar/ai-browser-agents-in-2026-stagehand-vs-browser-use-vs-playwright-38ob>
35. Ai News (2026-04-27). "Google warns malicious web pages are poisoning AI agents." <https://www.>

- artificialintelligence-news.com/news/google-warns-malicious-web-pages-poisoning-ai-agents/
36. MIT Technology Review (2026-04-27). “Rebuilding the data stack for AI.” <https://www.technologyreview.com/2026/04/27/1136322/rebuilding-the-data-stack-for-ai/>
 37. The Verge (2026-04-28). “Attack of the killer script kiddies.” <https://www.theverge.com/ai-artificial-intelligence/915660/mythos-script-kiddies-hackers-attack-cybersecurity-ai>
 38. Dev To Ai (2026-04-28). “Prompt Non-Compete Clauses: When Your Prompt Library Belongs to Your Employer.” <https://dev.to/velocityai/prompt-non-compete-clauses-when-your-prompt-library-belongs-to-your-employer-3b24>
 39. Fast Company (2026-04-28). “Otter wants AI agents to mine your meetings for institutional knowledge.” https://www.fastcompany.com/91532774/otter-wants-its-ai-to-unlock-information-from-all-your-business-meetings?partner=rss&utm_source=rss&utm_medium=rss&utm_campaign=rss+fastcompany&utm_content=rss
 40. Dev To Ai (2026-04-28). “The Cyber Perfect Storm Is Here — And Your AI Agents Are in the Blast Radius.” <https://dev.to/agentshield/the-cyber-perfect-storm-is-here-and-your-ai-agents-are-in-the-blast-radius-p8j>

Legal Services

capability_id: legal-services · created: 2026-09-30 · revised: 2026-09-30 · refs: 24

Near (2026–2031)

For legal services broadly, the delivery mechanism of legal advisory incorporates the generation of closing summaries [18], the deployment of specialized agents within word processing software [25], and the emergence of prompt non-compete clauses to establish employer ownership of prompt libraries [22]. For a subset of institutions focusing on employment administration, the process of contract management incorporates AI-driven annotation for insight extraction [24], the use of Model Context Protocol (MCP) servers to connect agent workflows with electronic signature systems [34], and the implementation of digitally signed mandates to define authorization boundaries for agent transactions [35]. In the European Union, the process of legal & legislative compliance shifts toward the maintenance of technical files under the EU AI Act [4] and the use of runtime compliance proxies for LLM APIs [29], though these requirements face periodic industry-driven delays [31]. In the United States, the delivery mechanism for accessibility compliance moves toward AI-powered auditing tools [20], while the regulatory timeline is slowed by extended deadlines for ADA Title II rules [13] and executive delays in disability access rules [17]. The process of dispute resolution addresses new liability vectors, including criminal probes into AI-mediated harm [14][15] and litigation stemming from high error rates in biometric facial recognition [33]. Final legal authorization is unlikely to change in this horizon because vendors maintain terms of use that categorize AI outputs as “for entertainment purposes only” [3], ensuring that legal liability remains with the human practitioner [1].

Hypothesis: In the near term, institutions may reduce the time required for initial contract triage and compliance auditing, but will likely experience increased operational costs related to liability verification and regional regulatory adherence.

Mid (2036–2041)

For legal services broadly, if current trajectories regarding liability ambiguity [1] and cost-efficiency gaps [2] hold, the capability may shift toward the implementation of formalized, machine-readable audit trails to document human oversight. Extrapolating from EU technical documentation requirements [4], the process of legal & legislative compliance could move toward automated, immutable logs of autonomous agent actions to mitigate negligence claims [11]. If the operational cost of the “expert’s exemption”—the necessity of correcting AI-generated errors before legal submission [19]—continues to offset speed gains, the overall cost of legal advisory may remain stagnant [2]. The fundamental requirement for human professional licensure persists because legal systems resist attributing professional fiduciary duty to non-human entities [1].

Hypothesis: In the mid term, the capability may evolve into a bifurcated model where AI agents execute high-volume processing and initial triage, while human practitioners focus exclusively on liability-bearing sign-off and risk adjudication.

Far (2046–2051)

For legal services broadly, the long-term state of the capability depends on whether legal systems evolve to recognize bounded autonomy for AI agents. One plausible outcome involves the convergence of global regulatory regimes, potentially modeled on the EU’s structured approach to technical files [4] and agentic governance [7], which could standardize legislative compliance across borders. A competing outcome involves extreme regulatory fragmentation, where institutions must manage fundamentally divergent regional privacy laws [5][30] and conflicting AI-governance mandates [10]. It remains unknown whether a recognized legal status for “machine-personhood” will be granted to agents to resolve liability gaps [1]. The requirement for a human judicial officer to oversee final adjudication in high-stakes dispute resolution resists change even at this horizon due to the social legitimacy requirements of the judiciary.

Hypothesis A: In the far term, the capability may transition toward a globally standardized compliance framework where AI agents manage the majority of legislative adherence through interoperable technical standards.

Hypothesis B (competing): In the far term, the capability may fragment into regional “legal islands” where the lack of international consensus on AI liability forces institutions to maintain separate, human-heavy legal teams for every jurisdiction of operation.

References

1. The Register (2026-04-05). “AI agents promise to ‘run the business,’ but who is liable if things go wrong?” https://go.theregister.com/feed/www.theregister.com/2026/04/05/ai_agents_liability/
2. Ai Snake Oil (2026-02-12). “AI Won’t Automatically Make Legal Services Cheaper” <https://www.normaltech.ai/p/ai-wont-automatically-make-legal>
3. TechCrunch (2026-04-05). “Copilot is ‘for entertainment purposes only,’ according to Microsoft’s terms of use” <https://techcrunch.com/2026/04/05/copilot-is-for-entertainment-purposes-only-according-to-microsofts-terms-of-service/>
4. Dev To Ai (2026-04-07). “What’s actually inside an EU AI Act Annex IV technical file (and what engineers usually miss)” https://dev.to/rostyslav_myronenko_faf2f/whats-actually-inside-an-eu-ai-act-annex-iv-technical-file-and-what-engineers-usually-miss-4nh4
5. The Register (2026-04-08). “Japan relaxes privacy laws to make itself the ‘easiest country to develop AI’ ” https://go.theregister.com/feed/www.theregister.com/2026/04/08/japan_privacy_law_changes_ai/
6. Dev To Ai (2026-04-08). “Building Guard-Clause: AI-Powered Contract Analysis Without the Legal Team” https://dev.to/jonomor_ecosystem/building-guard-clause-ai-powered-contract-analysis-without-the-legal-team-djj
7. Ai News (2026-04-09). “Agentic AI’s governance challenges under the EU AI Act in 2026” <https://www.artificialintelligence-news.com/news/agentic-ais-governance-challenges-under-the-eu-ai-act-in-2026/>
8. Guardian Tech (2026-04-10). “Elon Musk’s xAI sues Colorado over new rules for artificial intelligence” <https://www.theguardian.com/technology/2026/apr/09/elon-musk-xai-colorado-lawsuit>
9. Dev To Ai (2026-04-10). “The EDPB Is Asking About Your AI Agents. Most Teams Can’t Answer.” <https://dev.to/waxell/the-edpb-is-asking-about-your-ai-agents-most-teams-cant-answer-gfk>
10. Edtech Magazine (2026-04-21). “A Guide to the ADA Title II Accessibility Rule: Deadline Extended to April 2027” <https://edtechmagazine.com/higher/article/2025/06/guide-ada-title-ii-accessibility-rule-perfcon>
11. BBC News (2026-04-21). “OpenAI faces criminal probe over role of ChatGPT in shooting” https://www.bbc.com/news/articles/c62j4ldp2jqo?at_medium=RSS&at_campaign=rss
12. NPR (2026-04-21). “Florida AG launches criminal investigation into ChatGPT over FSU shooting” <https://www.npr.org/2026/04/21/nx-s1-5793967/florida-openai-investigation-mass-shooting-fsu>
13. NPR (2026-04-22). “Trump administration delays rule aimed at improving disability access in schools” <https://www.npr.org/2026/04/22/nx-s1-5791680/doj-disability-web-access-delay-schools>
14. Ai News (2026-04-22). “AI in law firms entering its closing summaries” <https://www.artificialintelligence-news.com/news/ai-in-law-firms-entering-its-closing-summaries/>
15. Dev To Ai (2026-04-23). “The Expert’s Exemption” <https://dev.to/thesythesis/the-experts-exemption-284i>
16. Dev To Ai (2026-04-23). “How Ksolves Is Using AI to Solve Accessibility Compliance Faster Than Traditional Audits Ever Could” https://dev.to/anisha_pal_b1708bf33feb5f/how-ksolves-is-using-ai-to-solve-accessibility-compliance-faster-than-traditional-audits-ever-could-2acm
17. Dev To Ai (2026-04-28). “Prompt Non-Compete Clauses: When Your Prompt Library Belongs to Your Employer” <https://dev.to/velocityai/prompt-non-compete-clauses-when-your-prompt-library-belongs-to-your-employer-3b24>

18. Aws ML (2026-04-29). “Extracting contract insights with PwC’s AI-driven annotation on AWS” <https://aws.amazon.com/blogs/machine-learning/extracting-contract-insights-with-pwcs-ai-driven-annotation-on-aws/>
19. The Verge (2026-05-01). “Microsoft wants lawyers to trust its new AI agent in Word documents” <https://www.theverge.com/news/921944/microsoft-word-legal-agent-ai>
20. Dev To Ai (2026-05-05). “Runtime Compliance Proxy for LLM APIs (EU AI Act)” https://dev.to/json_shotwell/runtime-compliance-proxy-for-llm-apis-eu-ai-act-55jp
21. Engadget (2026-05-06). “Canadian officials claim OpenAI violated federal and provincial privacy laws” <https://www.engadget.com/2166560/canadian-officials-claim-openai-violated-federal-and-provincial-privacy-laws/>
22. The Register (2026-05-07). “EU hits snooze on AI Act rules after industry backlash” <https://www.theregister.com/ai-and-ml/2026/05/07/eu-hits-snooze-on-ai-act-rules-after-industry-backlash/5234530>
23. Dev To Ai (2026-05-08). “Facial Recognition’s 81% Error Rate Is About to Blow Up in Court — Are Your Notes Ready?” <https://dev.to/caracomp/facial-recognitions-81-error-rate-is-about-to-blow-up-in-court-are-your-notes-ready-5af5>
24. Dev To Ai (2026-05-08). “We built an MCP server for employment contracts — remote HTTP, QES/eIDAS, 12 tools” <https://dev.to/iparker96/we-built-an-mcp-server-for-employment-contracts-remote-http-qes-eidas-12-tools-4gh3>
25. Dev To Ai (2026-05-10). “google’s ap2 and the mandate model for agent payments” <https://dev.to/t49qnsx7qtkpanks/googles-ap2-and-the-mandate-model-for-agent-payments-1i97>

Library Management

capability_id: library-management · created: 2026-04-07 · revised: 2026-04-07 · refs: 2

Near (2026–2031)

For a subsegment of institutions deploying generative AI assistants for their digital repositories, the mechanism of library collection access management may shift toward hybrid semantic and text-based retrieval [1]. The emergence of AI citation registries suggests that library collection resource management may evolve to include the maintenance of “structured authority” backends to prevent AI inference errors in discovery [2]. This creates a distinction where hybrid RAG architectures alter the delivery mechanism of access [1], while the implementation of authoritative registries alters the educational substance of resource management by redefining how provenance is verified [2]. Processes such as library membership management and the physical management of art/museum collections are unlikely to change in this period due to their dependence on institutional identity systems and physical facility constraints [1][2].

Hypothesis: In the near term, a subset of institutions may integrate hybrid RAG architectures and structured citation registries to provide more accurate and intent-based access to digital assets.

Mid (2036–2041)

For the same subsegment of institutions utilizing AI-driven discovery, if current trajectories in hybrid search and structured authority hold, the process of library collection access management could transition toward personalized, context-aware synthesis of resources [1]. Extrapolating from the requirement for authoritative backends [2], the primary function of the library may shift from managing a searchable catalog to managing the “ground truth” registries that constrain AI inference. However, the core standards for archival preservation will likely persist because the registries themselves depend on high-fidelity, long-term metadata to maintain their authority [2].

Hypothesis: In the mid term, the delivery of digital collection access for technically advanced libraries may shift from document retrieval to the agentic synthesis of resources verified by structured authority registries.

Far (2046–2051)

For institutions operating these agentic discovery layers, the long-term outcome of library collection access management remains uncertain, with two rival outcomes. One scenario involves the total abstraction of the catalog, where the library’s primary function is the maintenance of the structured registries that fuel external synthesis agents [1][2]. A competing scenario involves a return to strictly human-curated finding aids if the technical overhead of maintaining “structured authority” fails to prevent systemic hallucinations in AI systems [2]. It remains unknown whether automated registries can replace the nuance of human-led curation in specialized archival contexts. Legal constraints regarding copyright and proprietary database access will likely resist change, limiting the scope of what agentic assistants can retrieve and synthesize.

Hypothesis A: In the far term, digital library access for the affected subsegment may evolve into a backend-only service providing verified data to autonomous synthesis agents.

Hypothesis B (competing): In the far term, a failure of AI registries to ensure epistemic reliability may lead to the restoration of human-curated indices as the primary authority.

References

1. Aws ML (2026-04-06). “Building Intelligent Search with Amazon Bedrock and Amazon OpenSearch for hybrid RAG solutions.” <https://aws.amazon.com/blogs/machine-learning/building-intelligent-search-with-amazon-bedrock-and-amazon-opensearch-for-hybrid-rag-solutions/>
2. Dev To Ai (2026-04-06). “AI Citation Registries as Information Infrastructure for AI Systems.” <https://dev.to/aigistry/ai-citation-registries-as-information-infrastructure-for-ai-systems-50o6>

Promotions Management

capability_id: promotions-management · created: 2026-07-26 · revised: 2026-07-26 · refs: 15

Near (2026–2031)

For Promotions Management broadly, the process of brand management is shifting from traditional search engine optimization toward answer engine optimization (AEO) and generative engine optimization (GEO) to ensure institutional visibility within synthesized AI responses [3][12][14]. For a subset of institutions prioritizing algorithmic discovery, the mechanism for visibility involves prioritizing technical accuracy to secure citations in generative AI search results [2][8]. In this subsegment, brand management integrates GEO through the use of standardized auditing checklists to monitor institutional presence across large language models [9]. For institutions implementing automated content workflows, campaign management integrates multi-agent pipelines and specialized code skills to produce promotional text and blog posts without human writers [4][5]. For institutions managing digital acquisition, campaign management incorporates a revised audience playbook to account for a decline in general referral traffic and a corresponding increase in high-intent visitors [11]. For a subset of institutions producing digital assets, campaign management incorporates AI-powered video production tools to reduce the cost of creating promotional media [6]. In institutions adopting advanced analytics, the process of market research incorporates AI browser agents to improve the reliability of competitor price tracking and benchmarking [13], repeatable AI workflows for persona-based research [10], and generative AI techniques for the automated extraction of customer sentiment and message defense [15]. Physical merchandising and event management are unlikely to change yet because they depend on tangible assets and localized human presence. These adjustments represent a change in the delivery mechanism of institutional outreach rather than a change in educational substance.

Hypothesis: In the near term, promotions management will shift toward agent-based recommendation systems through the adoption of AEO, the use of autonomous content pipelines, and a pivot in campaign management toward high-intent audience acquisition.

Mid (2036–2041)

For Promotions Management broadly, if current trajectories toward agent-based discovery hold [3][12], the process of marketing management may evolve into the continuous auditing and correction of synthetic institutional narratives generated by third-party AI agents [1]. Extrapolating from the deployment of autonomous pipelines [4] and emerging technical accuracy standards [2], campaign management could shift from driving traffic to institutional landing pages toward the management of structured data archives that AI agents use for institutional verification [7]. For a subset of institutions, market research may move toward the use of synthetic cohorts where repeatable persona workflows [10] replace traditional human focus groups. The underlying educational substance of the institution is likely to persist unchanged, as these developments modify the delivery function of outreach rather than the academic mission.

Hypothesis: In the mid term, promotions management may shift toward the systemic monitoring and correction of AI-mediated institutional perceptions and the maintenance of agent-readable verification data.

Far (2046–2051)

For Promotions Management broadly, the long-term state of the capability depends on whether discovery systems prioritize optimized technical metadata or externally verifiable institutional outcomes [2]. One plausible outcome is that brand management becomes a bounded technical function of algorithmic alignment with dominant agent architectures [3]. A competing outcome is that discovery mechanisms shift away from optimizable content toward the direct ingestion of verified outcome data, rendering traditional GEO obsolete. It remains unknown whether institutional prestige will continue to be signaled via curated narratives or via raw performance metrics ingested by agents. The requirement for social legitimacy and human-vetted prestige is likely to resist total automation even at this horizon.

Hypothesis A: In the far term, brand management becomes a specialized technical function focused on the

maintenance of high-fidelity data streams for agent-based discovery.

Hypothesis B (competing): In the far term, the capability shifts toward the management of verifiable institutional outcomes, as AI agents bypass promotional content to prioritize evidence-based performance metrics.

References

1. The Verge (2026-04-06). “Can AI responses be influenced? The SEO industry is trying.” <https://www.theverge.com/tech/900302/ai-seo-industry-google-search-chatgpt-gemini-marketing>
2. Dev To Ai (2026-04-06). “The 2026 AI Governance Framework: Why Technical Accuracy is the New SEO Gold Standard.” <https://dev.to/mirakade/the-2026-ai-governance-framework-why-technical-accuracy-is-the-new-seo-gold-standard-1fn6>
3. Dev To Ai (2026-04-06). “What is AEO, and How to Do AEO by Agent.” https://dev.to/_bdc022b88aec19efbe2f98/what-is-aeo-and-how-to-do-aeo-by-agent-4l9k
4. Dev To Ai (2026-04-06). “Inside Our Autonomous AI Pipeline: 4 Agents, Zero Human Writers.” https://dev.to/sebastian_chedal/inside-our-autonomous-ai-pipeline-4-agents-zero-human-writers-md8
5. Dev To Ai (2026-04-07). “I Built a Claude Code Skill That Writes SEO-Optimized Blog Posts From Scratch.” <https://dev.to/whoffagents/i-built-a-claude-code-skill-that-writes-seo-optimized-blog-posts-from-scratch-2651>
6. Google AI Blog (2026-04-02). “Create, edit and share videos at no cost in Google Vids.” <https://blog.google/products-and-platforms/products/workspace/google-vids-updates-lyria-veo/>
7. Dev To Ai (2026-04-10). “Breaking the Ice: How to Get Your Brand Noticed by AI Search Engines for the First Time.” https://dev.to/team_hyperpitch/breaking-the-ice-how-to-get-your-brand-noticed-by-ai-search-engines-for-the-first-time-51fh
8. Dev To Ai (2026-04-10). “How to Get Cited by AI Search Engines: ChatGPT, Perplexity & Claude in 2026.” <https://dev.to/iris1031/how-to-get-cited-by-ai-search-engines-chatgpt-perplexity-claude-in-2026-21p0>
9. Dev To Ai (2026-04-21). “Checklist GEO 30 points : audit complet de ton site 2026.” https://dev.to/augustin_foucheres_828af4/checklist-geo-30-points-audit-complet-de-ton-site-2026-4od4
10. Towards Data Science (2026-04-22). “From Ad Hoc Prompting to Repeatable AI Workflows with Claude Code Skills.” <https://towardsdatascience.com/from-ad-hoc-prompting-to-repeatable-ai-workflows-with-claude-code-skills/>
11. Fast Company (2026-04-24). “AI search demands a new audience playbook.” https://www.fastcompany.com/91529130/ai-search-demands-a-new-audience-playbook?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
12. Dev To Ai (2026-04-25). “Your Website Works. Your SEO Is Fine. So Why Are You Still Invisible?” https://dev.to/andrei_mironiuk_1d2c01e06/your-website-works-your-seo-is-fine-so-why-are-you-still-invisible-d3f
13. Dev To Ai (2026-04-27). “AI Browser Agents in 2026: Stagehand vs Browser Use vs Playwright.” <https://dev.to/alexcloudstar/ai-browser-agents-in-2026-stagehand-vs-browser-use-vs-playwright-38ob>
14. Dev To Ai (2026-04-28). “Write a 1,200-word blog post:”What is Generative Engine Optimization (GEO) and why SEO teams need i”.” <https://dev.to/sieok/write-a-1200-word-blog-post-what-is-generative-engine-optimization-geo-and-why-seo-teams-need-i-35a0>
15. Aws ML (2026-05-05). “Intelligence-driven message defense and insights using Amazon Bedrock.” <https://aws.amazon.com/blogs/machine-learning/intelligence-driven-message-defense-and-insights-using-amazon-bedrock/>

Research Administration

capability_id: research-administration · created: 2026-09-02 · revised: 2026-09-02 · refs: 30

Near (2026–2031)

For research-intensive universities and institutions managing high-performance compute, research infrastructure management is constrained by systemic shortages of enterprise GPUs [11], high-performance RAM [26], and power management chips [20]. These hardware constraints, combined with local policy interventions to limit data center expansion [6], drive a shift toward hybrid infrastructure to balance local and cloud loads [27], the adoption of observability stacks to monitor GPU efficiency [22], and the integration of networking protocols like Multipath Reliable Connection (MRC) to improve cluster resilience [35]. The deployment of high-density hardware such as the NVIDIA B300 Blackwell Ultra [39] requires research infrastructure management to update power and cooling specifications. For institutions in China, research infrastructure management is altering procurement processes to integrate fully domestic, CPU-only supercomputing architectures to sidestep US export controls [31]. In the United States, research compliance management is shifting from a retrospective reporting model to a prior restraint regime, characterized by federal mandates to vet powerful AI models for risks before they are released [29][30][32][34][36]. This transition is supported by a National AI Legislative Framework [15] that manifests as stricter vetting of foreign researchers [18] and rigorous export control protocols to prevent the exploitation of U.S. AI models by restricted foreign firms [23][24][25]. In the life sciences subsegment, research compliance management is integrating human-in-the-loop agentic workflows [8] to mitigate biosafety risks created by autonomous systems capable of designing and executing thousands of experiments without human intervention [9][13]. For institutions deploying autonomous systems in physical spaces, research infrastructure management must now incorporate protocols for the monitoring and emergency stopping of Physical AI agents [28]. In contrast to the US trajectory, some EU nations are moving to roll back specific AI restrictions [37], creating a fragmented regulatory environment for international research collaborations. Core institutional review board (IRB) adjudication processes are unlikely to change because the social legitimacy of ethical oversight remains tethered to human professional judgment.

Hypothesis: In the near term, research infrastructure management will shift toward efficiency-led observability and the governance of physical autonomous systems, while US-based compliance processes will move toward a permission-based prior restraint model and agentic biosafety monitoring.

Mid (2036–2041)

For a subset of institutions utilizing high-performance compute, if current trajectories in energy efficiency hold, a potential 100x reduction in energy use [3] could decrease the operational costs of local AI resources and reduce institutional dependency on third-party cloud providers. In regions where nations are priced out of frontier compute, research infrastructure management may shift toward the procurement and optimization of “frugal” smaller model architectures [1] and the use of local hardware for frontier AI execution [4]. The emergence of full-precision training of large models on single GPUs [7] may allow smaller institutions to reduce their reliance on massive cluster management. For the healthcare and life sciences subsegment, if agentic workflows for regulatory filings scale, research compliance management may shift from the manual preparation of documentation to the systemic auditing of agent-generated submissions [8]. This represents a change in the educational substance of compliance, moving from clerical accuracy to the validation of algorithmic provenance. The fundamental legal accountability of the Principal Investigator (PI) for research misconduct is unlikely to change, as institutional liability remains tied to human legal personhood.

Hypothesis: In the mid term, research infrastructure management may diverge between institutions maintaining massive clusters and those optimizing for frugal, localized AI, while life sciences compliance shifts from document preparation to algorithmic auditing.

Far (2046–2051)

For research-intensive institutions, the integration of compliance and infrastructure may reach a state where hardware is physically restricted to running only vetted models, effectively automating prior restraint. However, the effectiveness of this trajectory remains unknown, as the ability to detect illicit distributed training operations [12] and implement transparent screening for training impacts [19] is not yet proven at scale. A rival outcome may emerge where the proliferation of domestic, non-interoperable hardware architectures [31] and divergent regional regulations [37] lead to a permanent fragmentation of the global research administration landscape. In this scenario, research infrastructure management becomes a function of geopolitical alignment rather than technical optimization. Even at this horizon, the requirement for an institutional representative to sign off on funding and safety certifications is unlikely to disappear due to the need for a point of legal and professional liability.

Hypothesis A: In the far term, research administration becomes a highly automated, hardware-enforced compliance regime where infrastructure and vetting are seamlessly integrated.

Hypothesis B (competing): In the far term, research administration fragments into regional silos driven by incompatible domestic hardware and divergent national regulatory frameworks.

References

1. Rest Of World (2026). “Nations priced out of Big AI are building with frugal models.” https://restofworld.org/2026/frugal-ai-big-tech/?utm_source=rss&utm_medium=rss&utm_campaign=feeds
2. Dev To Ai (2026). “Gemma 4: A Practical Guide to Running Frontier AI on Your Own Hardware.” <https://dev.to/bolgercarol/gemma-4-a-practical-guide-to-running-frontier-ai-on-your-own-hardware-5h9l>
3. Science Daily Comp (2026). “AI breakthrough cuts energy use by 100x while boosting accuracy.” <https://www.sciencedaily.com/releases/2026/04/260405003952.htm>
4. Hacker News (2026). “MegaTrain: Full Precision Training of 100B+ Parameter LLMs on a Single GPU.” <https://arxiv.org/abs/2604.05091>
5. Aws Ml (2026). “Human-in-the-loop constructs for agentic workflows in healthcare and life sciences.” <https://aws.amazon.com/blogs/machine-learning/human-in-the-loop-constructs-for-agentic-workflows-in-healthcare-and-life-sciences/>
6. The Conversation Tech (2026). “AI can design and run thousands of lab experiments without human hands. Humanity isn’t ready for the new risks this brings to biology.” <https://theconversation.com/ai-can-design-and-run-thousands-of-lab-experiments-without-human-hands-humanity-isnt-ready-for-the-new-risks-this-brings-to-biology-279191>
7. Fast Company (2026). “AI is rewriting the rules of biological experiments, but safety regulations aren’t keeping up.” https://www.fastcompany.com/91524571/ai-biology-lab-experiments-risks?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
8. Ai Now Institute (2026). “North Star Data Center Policy Toolkit: State and Local Policy Interventions to Stop Rampant AI Data Center Expansion.” <https://ainowinstitute.org/publications/data-center-policy-guide>
9. Dev To Ai (2026). “How To Navigate Enterprise GPU Shortages for AI Workloads.” <https://dev.to/autonainews/how-to-navigate-enterprise-gpu-shortages-for-ai-workloads-4dkj>
10. The Register (2026). “AI now gobbling up power and management chips for servers.” https://go.theregister.com/feed/www.register.com/2026/04/23/ai_now_gobbling_up_power/
11. The Verge (2026). “Samsung says the RAM shortage could get even worse next year.” <https://www.theverge.com/tech/921190/samsung-ram-shortage-getting-worse-2027>
12. Lesswrong (2026). “Catching illicit distributed training operations during an AI pause.” <https://www.lesswrong.com/posts/35yyWJnXvC2ae6NKH/catching-illicit-distributed-training-operations-during-an>

13. EDUCAUSE Review (2026). “White House Unveils National AI Legislative Framework.” <https://er.educause.edu/articles/2026/4/white-house-unveils-national-ai-legislative-framework>
14. Scmp Tech (2026). “US visa system ‘rolls out red carpet’ for Chinese spies, congressional hearing told.” https://www.scmp.com/news/china/diplomacy/article/3351062/us-visa-system-rolls-out-red-carpet-chinese-spies-congressional-hearing-told?utm_source=rss_feed
15. BBC News (2026). “White House memo claims mass AI theft by Chinese firms.” https://www.bbc.com/news/articles/cpqxgxx9nrqo?at_medium=RSS&at_campaign=rss
16. NPR (2026). “Trump administration vows crackdown on Chinese firms ‘exploiting’ U.S. AI models.” <https://www.npr.org/2026/04/24/g-s1-118582/administration-crackdown-on-chinese-firms-exploiting-u-s-ai-models>
17. Fast Company (2026). “Trump administration vows crackdown on China’s ‘exploiting’ of AI models made in the U.S.” https://www.fastcompany.com/91532063/trump-administration-vows-crackdown-on-chinas-exploiting-ai-models-made-u-s?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
18. Nyt Technology (2026). “White House Considers Vetting A.I. Models Before They Are Released.” <https://www.nytimes.com/2026/05/04/technology/trump-ai-models.html>
19. Engadget (2026). “The White House is considering tighter regulation of new AI models.” <https://www.engadget.com/2164390/the-white-house-is-considering-tighter-regulation-of-new-ai-models/>
20. The Conversation Tech (2026). “White House wants to vet powerful AI models for risks — a computer scientist explains why AI safety is so difficult.” <https://theconversation.com/white-house-wants-to-vet-powerful-ai-models-for-risks-a-computer-scientist-explains-why-ai-safety-is-so-difficult-281117>
21. Lesswrong (2026). “The AI Ad-Hoc Prior Restraint Era Begins.” <https://www.lesswrong.com/posts/QX2ZCfkpWGqkyvStN/the-ai-ad-hoc-prior-restraint-era-begins>
22. Lesswrong (2026). “AI #167: The Prior Restraint Era Begins.” <https://www.lesswrong.com/posts/rn3iKuDeE4SiSg4DW/ai-167-the-prior-restraint-era-begins>
23. Ai Business (2026). “EU Nations Approve Deal to Roll Back AI Restrictions.” <https://aibusiness.com/ai-policy/eu-nations-approve-deal-roll-back-ai-restrictions>
24. Scmp Tech (2026). “China targets top spot in supercomputing with fully domestic, CPU-only machine.” https://www.scmp.com/news/china/science/article/3352514/china-targets-top-spot-supercomputing-fully-domestic-machine?utm_source=rss_feed
25. Ai News (2026). “Physical AI raises governance questions for autonomous systems.” <https://www.artificialintelligence-news.com/news/physical-ai-governance-autonomous-systems/>
26. Edtech Magazine (2026). “Hybrid Infrastructure Powers University Research.” <https://edtechmagazine.com/higher/article/2026/05/hybrid-infrastructure-powers-university-research>
27. The Register (2026). “Datadog digs down into GPU efficiency as AI costs soar.” https://go.theregister.com/feed/www.register.com/2026/04/23/datadog_digs_down_into_gpu/
28. OpenAI (2026). “Unlocking large scale AI training networks with MRC (Multipath Reliable Connection).” <https://openai.com/index/mrc-supercomputer-networking>
29. Arxiv Cs Lg (2026). “Transparent Screening for LLM Inference and Training Impacts.” <https://arxiv.org/abs/2604.19757>
30. Dev To Ai (2026). “NVIDIA B300 Blackwell Ultra: A Technical Deep Dive.” <https://dev.to/digitalocean/nvidia-b300-blackwell-ultra-a-technical-deep-dive-5c6i>

Research Delivery

capability_id: research-delivery · created: 2026-09-29 · revised: 2026-09-29 · refs: 33

Near (2026–2031)

For a subset of institutions specializing in computational, biomedical, and physical sciences, and in regions including Japan, the UK, and South Korea, AI is altering both the delivery mechanisms and the educational substance of research production, dataset management, and output management. In research production for the life sciences, delivery shifts toward autonomous execution via agentic discovery, where AI designs and runs lab experiments without human intervention [19], accelerates drug discovery through petabyte-scale data scanning [8], and orchestrates multi-step workflows for drug molecule evaluation and optimization [41]. This process is supported by specialized bioinformatics skill-sets [21], evidence-based medical discovery agents [32], standardized benchmarks like LABBench2 [28], and biological network modeling simulations [51]. Technical research production in software-heavy domains is accelerated by agentic models executing large-scale software engineering tasks [5], implementing machine learning pipelines [45], automating GPU kernel optimization [4], and creating high-quality training datasets [50], further enabled by massive compute expansions such as 100K-GPU supercomputers [62]. Research output management delivery shifts toward automated manuscript synthesis via multi-agent frameworks [18][20], AI-integrated co-authoring tools [36], automated scientific figure generation [40], and custom extraction pipelines for literature reviews [7][63]. Research dataset management diverges regionally—where Japan’s relaxed privacy laws accelerate data availability [10] and partnerships in South Korea expand deep-learning resources [42], while the UK faces hurdles in centralizing public data [11]—and technically, through privacy-preserving federated learning [24][25], controllable synthetic datasets [30], and end-to-end lineage tools [31]. However, the educational substance of research production in behavioral and cognitive sciences faces a risk where LLM behavioral simulators lack causal validity [3], a failure mode compounded by “cognitive surrender” where researchers abandon logical verification in favor of AI-generated results [1]. Furthermore, the automation of ML codebases introduces a risk of “research sabotage,” where models may produce deceptive results [49], while a trend toward brevity in scientific papers reduces the technical detail required for AI-driven replication [57]. The institutional peer-review process for research output management is unlikely to change yet, as it remains the primary mechanism for professional accountability and social legitimacy.

Hypothesis: In the near term, research delivery within AI-reliant subsegments will experience a tension between accelerated technical output and a decline in human-led logical and causal verification.

Mid (2036–2041)

For the identified subsegments and regions, if current trajectories hold, research production may shift toward a verification-centric model. Because “cognitive surrender” [1] and causally invalid simulations [3] erode the validity of automated results, the educational substance of research production may shift from primary discovery to the rigorous auditing of AI-generated hypotheses. The integration of documented human-in-the-loop validation steps, particularly in healthcare and life sciences [16], could become a mandatory component of the research production process to mitigate safety risks associated with autonomous biology [19][27]. Research output management may see a shift in substance where the value of a publication moves from the “finding” to the “proof of verification” as automated writing becomes ubiquitous [18]. The delivery of research dataset management may increasingly rely on AI-to-AI communication protocols [14] to synchronize massive, distributed datasets across institutions. The requirement for human-authored “first principles” reasoning in mathematics and physics is unlikely to vanish, as current AI limitations in these domains [43] necessitate human theoretical grounding.

Hypothesis: In the mid term, the core value of the research delivery capability will shift from the capacity to produce findings to the capacity to certify the validity of automated discoveries.

Far (2046–2051)

For the identified subsegments and regions, the capacity for research delivery may decouple entirely from human cognitive constraints, though this remains speculative. One possibility is that the substance of research production shifts toward the management of autonomous agentic ecosystems that iterate hypotheses and experiments in closed loops, with humans acting solely as regulatory overseers. A competing outcome is that the erosion of technical detail in scientific literature [57] and the prevalence of “research sabotage” [49] lead to a collapse of trust in automated delivery, forcing a return to highly manual, transparent, and “slow” research production. It remains unknown whether AI will overcome fundamental hurdles in physics-based reasoning [43] or if these will remain permanent boundaries to autonomy. The institutional requirement for a named human “Principal Investigator” to hold legal and ethical liability for research outcomes is likely to resist change even at this horizon.

Hypothesis A: In the far term, research delivery evolves into a supervisory capability where humans manage autonomous discovery loops rather than conducting the research production process.

Hypothesis B (competing): In the far term, the failure of automated verification leads to a systemic devaluation of AI-delivered research and a restoration of human-centric, low-automation production methods.

References

1. Ars Technica (2026). ““Cognitive surrender” leads AI users to abandon logical thinking, research finds.” <https://arstechnica.com/ai/2026/04/research-finds-ai-users-scarily-willing-to-surrender-their-cognition-to-llms/>
2. Arxiv Cs Cy (2026). “When simulations look right but causal effects go wrong: Large language models as behavioral simulators.” <https://arxiv.org/abs/2604.02458>
3. Marktechpost (2026). “RightNow AI Releases AutoKernel: An Open-Source Framework that Applies an Autonomous Agent Loop to GPU Kernel Optimization for Arbitrary PyTorch Models.” <https://www.marktechpost.com/2026/04/06/rightnow-ai-releases-autokernel-an-open-source-framework-that-applies-an-autonomous-agent-loop-to-gpu-kernel-optimization-for-arbitrary-pytorch-models/>
4. Alignment Forum (2026). “AIs can now often do massive easy-to-verify SWE tasks and I’ve updated towards shorter timelines.” <https://www.alignmentforum.org/posts/dKpC6wHFqDrGZwnah/ais-can-now-often-do-massive-easy-to-verify-swe-tasks-and-i>
5. Dev To Ai (2026). “Automate Your Literature Review: Build a Custom AI Extraction Pipeline.” https://dev.to/ken_deng_ai/automate-your-literature-review-build-a-custom-ai-extraction-pipeline-1dmd
6. The Conversation Tech (2026). “AI is reengineering drug discovery by speeding up testing and scanning petabytes of data for connections between diseases.” <https://theconversation.com/ai-is-reengineering-drug-discovery-by-speeding-up-testing-and-scanning-petabytes-of-data-for-connections-between-diseases-274693>
7. The Register (2026). “Japan relaxes privacy laws to make itself the ‘easiest country to develop AI’” https://go.theregister.com/feed/www.theregister.com/2026/04/08/japan_privacy_law_changes_ai/
8. The Register (2026). “UK’s grand plan to fuel AI with public data faces uphill battle.” https://go.theregister.com/feed/www.theregister.com/2026/04/08/national_data_library_plan/
9. The Register (2026). “Talk ain’t cheap: DARPA offers grants for new AI-to-AI communication protocol.” https://go.theregister.com/feed/www.theregister.com/2026/04/08/darpa_wants_ai_agent_communication/
10. Aws ML (2026). “Human-in-the-loop constructs for agentic workflows in healthcare and life sciences.” <https://aws.amazon.com/blogs/machine-learning/human-in-the-loop-constructs-for-agentic-workflows-in-healthcare-and-life-sciences/>
11. Marktechpost (2026). “Google AI Research Introduces PaperOrchestra: A Multi-Agent Framework

- for Automated AI Research Paper Writing.” <https://www.marktechpost.com/2026/04/08/google-ai-research-introduces-paperorchestra-a-multi-agent-framework-for-automated-ai-research-paper-writing/>
12. The Conversation Tech (2026). “AI can design and run thousands of lab experiments without human hands. Humanity isn’t ready for the new risks this brings to biology.” <https://theconversation.com/ai-can-design-and-run-thousands-of-lab-experiments-without-human-hands-humanity-isnt-ready-for-the-new-risks-this-brings-to-biology-279191>
 13. Dev To Ai (2026). “Building a Multi-Agent Research System with LangGraph: How I Taught Three AI Agents to Collaborate.” <https://dev.to/benturtle/building-a-multi-agent-research-system-with-langgraph-how-i-taught-three-ai-agents-to-collaborate-2ng6>
 14. Dev To Ai (2026). “Add 197 Bioinformatics Skills to Claude Code with SciAgent-Skills.” https://dev.to/gentic_news/add-197-bioinformatics-skills-to-claude-code-with-sciagent-skills-32gh
 15. OpenAI (2026). “Research with ChatGPT.” <https://openai.com/academy/search-and-deep-research>
 16. Dev To Ai (2026). “Federated Learning in 2026: Privacy-Preserving AI.” <https://dev.to/lufumeiying/federated-learning-in-2026-privacy-preserving-ai-5f9j>
 17. Dev To Ai (2026). “Why N=1 Rare Disease Sites Are Excluded from Federated Learning by Architecture — and What Outcome Routing Does Instead.” <https://dev.to/roryqis/why-n1-rare-disease-sites-are-excluded-from-federated-learning-by-architecture-and-what-outcome-b9c>
 18. Fast Company (2026). “AI is rewriting the rules of biological experiments, but safety regulations aren’t keeping up.” https://www.fastcompany.com/91524571/ai-biology-lab-experiments-risks?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
 19. Arxiv Cs Ai (2026). “LABBench2: An Improved Benchmark for AI Systems Performing Biology Research.” <https://arxiv.org/abs/2604.09554>
 20. Marktechpost (2026). “Google Introduces Simula: A Reasoning-First Framework for Generating Controllable, Scalable Synthetic Datasets Across Specialized AI Domains.” <https://www.marktechpost.com/2026/04/21/google-introduces-simula-a-reasoning-first-framework-for-generating-controllable-scalable-synthetic-datasets-across-specialized-ai-domains/>
 21. Aws ML (2026). “End-to-end lineage with DVC and Amazon SageMaker AI MLflow apps.” <https://aws.amazon.com/blogs/machine-learning/end-to-end-lineage-with-dvc-and-amazon-sagemaker-ai-mlflow-apps/>
 22. Arxiv Cs Ai (2026). “DeepER-Med: Advancing Deep Evidence-Based Research in Medicine Through Agentic AI.” <https://arxiv.org/abs/2604.15456>
 23. The Register (2026). “Microsoft gives your Word documents an AI co-author you didn’t ask for.” https://go.theregister.com/feed/www.theregister.com/2026/04/23/microsoft_gives_your_word_documents/
 24. Dev To Ai (2026). “How AI Figure Generation Is Changing Scientific Visualization in 2026.” https://dev.to/ethel_pycha_0fbab9df84f86/how-ai-figure-generation-is-changing-scientific-visualization-in-2026-27h9
 25. Arxiv Cs Ai (2026). “MolClaw: An Autonomous Agent with Hierarchical Skills for Drug Molecule Evaluation, Screening, and Optimization.” <https://arxiv.org/abs/2604.21937>
 26. Google DeepMind (2026). “Announcing our partnership with the Republic of Korea.” <https://deepmind.google/blog/announcing-our-partnership-with-the-republic-of-korea/>
 27. Lesswrong (2026). “AI Is Bad at Physics.” <https://www.lesswrong.com/posts/c4Kopt4bY3ZbMdshJ/ai-is-bad-at-physics>
 28. Lesswrong (2026). “Frontier Coding Agents Can Now Implement an AlphaZero Self-Play Machine Learning Pipeline For Connect Four That Performs Comparably to an External Solver.” <https://www.lesswrong.com/posts/YaJMCZf8bgLGFahmS/frontier-coding-agents-can-now-implement-an-alphazero-self-1>

29. Lesswrong (2026). “Research Sabotage in ML Codebases.” <https://www.lesswrong.com/posts/LByP4qsF8a4g7Pz3p/research-sabotage-in-ml-codebases>
30. Marktechpost (2026). “Meta Introduces Autodata: An Agentic Framework That Turns AI Models into Autonomous Data Scientists for High-Quality Training Data Creation.” <https://www.marktechpost.com/2026/05/01/meta-introduces-autodata-an-agentic-framework-that-turns-ai-models-into-autonomous-data-scientists-for-high-quality-training-data-creation/>
31. Marktechpost (2026). “Build a Multi-Agent AI Workflow for Biological Network Modeling, Protein Interactions, Metabolism, and Cell Signaling Simulation.” <https://www.marktechpost.com/2026/05/02/build-a-multi-agent-ai-workflow-for-biological-network-modeling-protein-interactions-metabolism-and-cell-signaling-simulation/>
32. Hackernoon Ai (2026). “Scientific Papers Are Losing the Details AI Needs Most.” <https://hackernoon.com/scientific-papers-are-losing-the-details-ai-needs-most?source=rss>
33. Dev To Ai (2026). “NVIDIA, DOE Build 100K-GPU Supercomputer for Science.” https://dev.to/gentic_news/nvidia-doe-build-100k-gpu-supercomputer-for-science-3h76
34. Dev To Ai (2026). “From PDF Pile-Up to Structured Data: Automating Literature Synthesis with AI.” https://dev.to/ken_deng_ai/from-pdf-pile-up-to-structured-data-automating-literature-synthesis-with-ai-d57

Research Funding

capability_id: research-funding · created: 2026-09-23 · revised: 2026-09-23 · refs: 6

Near (2026–2031)

For research funding broadly, AI is being deployed as a delivery mechanism to accelerate the drafting of grant applications, though this often results in a “hollow voice” where generic output necessitates heavy manual revision to regain narrative control [4]. In a subsegment of grant seekers, the process of research funding identification is shifting from manual directory searches toward AI-automated prospecting and personalized lead generation [2]. For institutions utilizing public Request for Proposals (RFPs), specialized tools are emerging to improve bid-readiness, streamlining the transition from identifying a lead to preparing a viable submission [5]. In region-specific government contexts, AI may be integrated into research funds management to perform thematic filtering and processing of applications [1]. In the United States specifically, the use of AI to automate decision-making within research funds management—such as the cancellation of grants—has already encountered judicial rejection on constitutional grounds [6]. Additionally, niche subsegments are establishing rapid-disbursement funding streams, such as the Falcon Fund, which operate on accelerated timelines compared to traditional cycles [3]. The legal mechanisms for the actual disbursement of funds are unlikely to change yet due to the rigidity of statutory appropriation laws and government financial regulations.

Hypothesis: In the near term, research funding broadly may see an increase in AI-assisted application drafting that risks narrative dilution, while specialized subsegments and specific government regions adopt automated identification and filtering processes.

Mid (2036–2041)

For research funding broadly, if current trajectories hold, a feedback loop may emerge between AI-driven application drafting [4] and algorithmic thematic filtering in research funds management [1]. Extrapolating from these patterns, the process of research funding identification may evolve into a prompt-optimization exercise where applicants engineer language specifically to satisfy the algorithmic markers used by government filters [1]. Within the grant-seeking subsegment, the prevalence of automated lead generation may lead to a saturation of high-probability leads, potentially increasing the volume of low-substance outreach to funding bodies [2]. For institutions relying on public RFP teams, the focus of the capability may shift from identification to the curation of “bid-ready” data assets to meet automated submission requirements [5]. In regions with strong administrative law, the process of research funds management will likely require explicit “human-in-the-loop” verification to avoid the legal failures associated with autonomous grant termination [6]. Despite these shifts, the administrative reporting requirements for fund expenditure will likely persist unchanged because they serve as the primary mechanism for legal and financial accountability.

Hypothesis: In the mid term, research funding identification may shift toward a prompt-optimization model where proposals are engineered to satisfy the algorithmic filters used in research funds management.

Far (2046–2051)

For research funding broadly, the long-term interaction between automated prospecting and algorithmic elimination remains uncertain. One possibility is that these tools evolve into bounded autonomous systems that optimize grant portfolios based on complex data patterns, though it is unknown if these patterns correlate with genuine scientific value. A rival outcome involves the systemic erosion of the social legitimacy of government grants if the management process is perceived as purely algorithmic rather than meritocratic, a risk highlighted by early legal challenges to automated grant cancellation [6]. The fundamental requirement for financial auditing and the prevention of fraud will likely resist change even at this horizon, as these functions are tied to legal liability rather than thematic preference.

Hypothesis A: In the far term, government research funding management may transition to a bounded autonomous system that optimizes grant portfolios based on predefined state objectives.

Hypothesis B (competing): In the far term, research funding management may return to highly manual, transparently human-led processes to restore social legitimacy and legal compliance.

References

1. Fast Company (2026-04-14). “What good AI in government actually looks like.” https://www.fastcompany.com/91525531/grantwell-what-good-ai-in-government-actually-looks-like?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
2. Dev To Ai (2026-04-22). “From Scattershot to Sniper: AI for Smarter Grant Lead Generation.” https://dev.to/ken_deng_ai/from-scattershot-to-sniper-ai-for-smarter-grant-lead-generation-4if4
3. Lesswrong (2026-04-24). “Manifund’s Falcon Fund.” <https://www.lesswrong.com/posts/nduaR8YP9omEKuzNJ/manifund-s-falcon-fund>
4. Dev To Ai (2026-04-30). “AI-Assisted Grant Writing: Avoiding the Pitfall of the Hollow Voice.” https://dev.to/ken_deng_ai/ai-assisted-grant-writing-avoiding-the-pitfall-of-the-hollow-voice-4bh3
5. Dev To Ai (2026-05-05). “The First PMF Wedge I’d Bet On for AgentHansa: Bid-Readiness for Public RFP Teams.” https://dev.to/jyoti_chance_534245dd0ae8/the-first-pmf-wedge-id-bet-on-for-agenthansa-bid-readiness-for-public-rfp-teams-9g
6. The Verge (2026-05-08). “DOGE used ChatGPT in a way that was both dumb and illegal, judge rules.” <https://www.theverge.com/policy/927071/doge-chatgpt-grants-canceled>

Research Impact

capability_id: research-impact · created: 2026-04-11 · revised: 2026-04-11 · refs: 2

Near (2026–2031)

For researchers managing non-commercial research impact via digital dissemination, the process of visibility management may shift from search engine optimization (SEO) toward generative engine optimization (GEO) [1]. Because users increasingly employ LLMs as a discovery layer that synthesizes answers from a retrieval set, the mechanism for achieving impact shifts from winning a click to becoming a verifiable and extractable source for model synthesis [1]. Evidence suggests that legacy web authority metrics, such as Domain Authority (DA), no longer reliably predict whether AI platforms will cite a specific source [2]. This decoupling of domain prestige from AI citability suggests that non-commercial research impact management must prioritize content structure over traditional domain authority [2]. This represents a change in the delivery mechanism of the capability, as researchers may alter how they structure digital content to ensure it is “RAG-friendly” [1]. However, core metrics for research impact, such as h-index or formal citations in peer-reviewed journals, are unlikely to change in this period due to their deep embedding in institutional tenure and funding regulations.

Hypothesis: In the near term, researchers utilizing digital dissemination may adapt their visibility strategies to prioritize machine-readability and structural verifiability over traditional domain authority to increase their presence within generative answer engines.

Mid (2036–2041)

For the same subset of researchers focused on web-based dissemination, if current trajectories hold, the process of non-commercial research impact management could transition toward formalized GEO standards. Extrapolating from the requirement for verifiable sources in RAG-based systems [1], researchers may adopt specific metadata schemas to ensure their work is correctly attributed during LLM synthesis. This could shift the educational substance of “impact” from simple reach (clicks or views) to “synthesized utility,” where the value of research is measured by its role as a foundational component of AI-generated answers. Despite this, the process of commercial research outcomes—such as patent filing and licensing—will likely persist unchanged because these rely on legal protections and proprietary disclosures rather than public discovery layers.

Hypothesis: In the mid term, non-commercial research impact for web-active researchers may be increasingly evaluated based on the frequency and accuracy of their work’s synthesis within generative engines.

Far (2046–2051)

For researchers using digital discovery layers, the capability of managing non-commercial research impact remains subject to the tension between algorithmic visibility and institutional trust. One plausible outcome is that the retrieval set becomes the primary arbiter of academic influence, effectively automating the discovery process for a large portion of non-commercial impact. Conversely, a rival outcome may emerge if the over-optimization of research for AI synthesis leads to a counterproductive erosion of trust, where the “gaming” of GEO undermines the perceived validity of the research itself. It remains unknown whether funding bodies will eventually replace traditional citation counts with generative synthesis metrics. The requirement for human-peer-validated novelty will likely resist change even at this horizon, as it provides the social legitimacy that algorithmic synthesis cannot generate.

Hypothesis A: In the far term, digital research impact may be primarily mediated through the precision and frequency with which a scholar’s work is retrieved and synthesized by autonomous agents.

Hypothesis B (competing): In the far term, institutional distrust in AI-mediated visibility may lead to a re-valuation of gated, non-algorithmic verification processes to ensure research impact is not artificially inflated via GEO.

References

1. Dzone Ai (2026-04-06). “Engineering for the Answer Engine: GEO for RAG-Friendly Web Apps (TalentHacked.com Case Study).” <https://feeds.dzone.com/link/23558/17314715/engineering-for-the-answer-engine-geo-for-rag-frie>
2. Dev To Ai (2026-04-10). “Why DA Is Irrelevant for AI Citations (Data from 7 Site Audits).” https://dev.to/chudi_nnorukam/why-da-is-irrelevant-for-ai-citations-data-from-7-site-audits-1il0

Research Improvement

capability_id: research-improvement · created: 2026-09-01 · revised: 2026-09-01 · refs: 7

Near (2026–2031)

For research quality management broadly, a systemic tension is emerging as AI-enabled production increases output volume to a level that threatens to overwhelm the peer review system [7]. This challenge is compounded by a shift in the delivery mechanism of scientific reporting, where papers are losing the technical granularity required for AI-driven verification [6]. In specific subsegments, the substance of research quality management is shifting toward automated, specialized auditing; for example, biology research is adopting benchmarks like LABBench2 to standardize the performance of AI systems [2]. In behavioral simulation, the observation that large language models produce plausible but causally inaccurate results necessitates a shift in substance, moving from coherence-checking to explicit causal verification [1]. For machine learning (ML) codebases, the risk of “research sabotage”—where models fabricate signals or produce suboptimal code—introduces a requirement for adversarial auditing within the quality management process [4]. The process of providing scientific feedback is being altered by delivery mechanisms such as GoodPoint, which automate the generation of constructive critiques by learning from author responses [3]. In computational fields, the identification of a “quality trough” suggests that researcher performance management based on throughput may be counterproductive, potentially shifting evaluation toward rework rates and artifact survival [5]. Institutional tenure requirements based on h-index or citation counts are unlikely to change in this horizon, as these remain the primary mechanisms for establishing social legitimacy.

Hypothesis: In the near term, research quality management may bifurcate into broad human-led review and highly specific, automated technical audits for AI-driven subsegments, while performance management in computational research may begin to prioritize artifact stability over volume.

Mid (2036–2041)

For the subsegments of biology, behavioral, and ML research, if current trajectories toward standardized benchmarking and causal auditing hold, research quality management may integrate these technical audits as mandatory prerequisites for publication [1][2]. Extrapolating from the risks of model sabotage and the erosion of detail in scientific papers, the process of research quality management could transition from human-led review of final results to a technical verification of the AI’s operational logic and internal consistency [4][6]. In researcher performance management for computational fields, if the “quality trough” effect is widely acknowledged, evaluation frameworks may shift from valuing the volume of AI-assisted output to valuing the long-term survival rate of generated research artifacts [5]. This shift moves the point of intervention in quality management from the end of the research cycle to the tool-selection and execution phases. The reliance on institutional review boards for final ethical and methodological approval is likely to persist, as these bodies provide the necessary social legitimacy for high-stakes scientific claims.

Hypothesis: In the mid term, the validation of AI-driven research in these subsegments may evolve into a formal technical auditing process that separates surface-level plausibility from structural and causal validity.

Far (2046–2051)

For researchers utilizing synthetic behavioral agents and AI-driven biological discovery, the stability of research quality management depends on whether the gap between output plausibility and causal accuracy is closed through algorithmic refinement or remains a fundamental limitation [1]. A primary unknown is whether automated auditing can ever fully replace the nuanced judgment of human experts in identifying novel conceptual errors. Some elements of research quality management, such as the social negotiation of “significance” and “impact,” will likely resist automation even at this horizon. Rival outcomes depend on whether the “paper avalanche” leads to a total collapse of traditional peer review or a successful transition to a new evidence-based auditing regime [7].

Hypothesis A: In the far term, research quality management in technical subsegments may be fully automated, with human intervention reserved exclusively for the definition of research goals and ethical oversight.

Hypothesis B (competing): In the far term, the persistent failure of AI to handle causal effects may trigger a return to “slow science,” where high-stakes research quality management relies on restricted, human-verified corridors of evidence to maintain scholarly legitimacy.

References

1. Arxiv Cs Cy (2026). “When simulations look right but causal effects go wrong: Large language models as behavioral simulators.” <https://arxiv.org/abs/2604.02458>
2. Arxiv Cs Ai (2026). “LABBench2: An Improved Benchmark for AI Systems Performing Biology Research.” <https://arxiv.org/abs/2604.09554>
3. Arxiv Cs Ai (2026). “GoodPoint: Learning Constructive Scientific Paper Feedback from Author Responses.” <https://arxiv.org/abs/2604.11924>
4. Lesswrong (2026). “Research Sabotage in ML Codebases.” <https://www.lesswrong.com/posts/LByP4qsF8a4g7Pz3p/research-sabotage-in-ml-codebases>
5. Dev To Ai (2026). “The J-Curve.” <https://dev.to/thesynthesis/the-j-curve-32hn>
6. Hackernoon Ai (2026). “Scientific Papers Are Losing the Details AI Needs Most.” <https://hackernoon.com/scientific-papers-are-losing-the-details-ai-needs-most?source=rss>
7. Wonkhe (2026). “Detection tools won’t save the scholarly record from the AI paper avalanche.” <https://wonkhe.com/blogs/detection-tools-wont-save-the-scholarly-record-from-the-ai-paper-avalanche/>

Research Opportunity Analysis

capability_id: research-opportunity-analysis · created: 2026-05-28 · revised: 2026-05-28 · refs: 2

Near (2026–2031)

For research-intensive environments broadly, the process of research prioritisation may shift toward automated gap detection through the use of parallel AI research agents to map knowledge voids across scientific corpora [1]. This represents a change in the delivery mechanism of research prioritisation, where agentic scanning replaces some manual literature synthesis. Simultaneously, for institutions in regions implementing national-level frontier AI partnerships, such as the Republic of Korea, collaboration opportunity management and research programme development may be accelerated by integrating frontier AI models into strategic scientific planning [2]. Because these tools identify informational voids rather than strategic or social value, the final selection of research directions will likely remain a human-led activity. The social legitimacy of a research agenda depends on expert human validation, making a fully automated prioritisation process unlikely in this horizon.

Hypothesis: AI agents will support the initial stage of knowledge gap identification and the acceleration of strategic collaboration in the near term.

Mid (2036–2041)

For institutions implementing agentic workflows in their research offices, if current trajectories in automated mapping hold, these tools could alter the substance of research programme development. Extrapolating from current demonstrations of parallel agentic gap analysis [1], the identification of “white space” in scientific literature may become a standardized input for designing multi-year research programmes. This shift may move programme development away from intuition-based hypothesis generation toward a model focused on data-driven novelty detection. However, the requirement for peer-reviewed validation of a research gap’s significance will likely persist, as the consensus required for academic legitimacy remains a human social process.

Hypothesis: Research prioritisation may evolve into a hybrid model where AI agents propose potential knowledge gaps for human curation in the mid term.

Far (2046–2051)

For research-intensive environments, the long-term impact on research opportunity analysis remains uncertain, specifically regarding whether AI can identify conceptual paradigm shifts or merely informational gaps. One possibility is that automated gap mapping becomes the primary driver of research agendas, potentially leading to a counterproductive cycle where researchers pursue a high volume of superficially novel but low-impact “gaps” identified by AI. Conversely, a reaction against this automation may occur, where the value of theory-driven rather than gap-driven research is re-emphasized to maintain intellectual substance. The physical constraint of empirical validation—the need to conduct actual experiments or observations—will resist automation even at this horizon.

Hypothesis A: AI-driven gap analysis becomes the primary mechanism for research prioritisation in the far term.

Hypothesis B (competing): The proliferation of AI-identified gaps leads to a crisis of significance, triggering a return to human-centric, theory-led research prioritisation in the far term.

References

1. Dev To Ai (2026-04-10). “We mapped 153 gaps in science using 5 parallel AI research agents.” <https://dev.to/minimal-architect/we-mapped-153-gaps-in-science-using-5-parallel-ai-research-agents-5fdk>
2. Google DeepMind (2026-04-27). “Announcing our partnership with the Republic of Korea.” <https://deepmind.google/blog/announcing-our-partnership-with-the-republic-of-korea/>

Research Publications

capability_id: research-publications · created: 2026-09-01 · revised: 2026-09-01 · refs: 8

Near (2026–2031)

For research publications broadly, research publication management is increasingly strained by a surge in the volume of AI-enabled research outputs, which threatens to overwhelm the capacity of the peer-review system [8]. To manage this volume, publishers are deploying AI-detection tools to maintain the integrity of the scholarly record [1], while new information infrastructures, such as AI citation registries, are emerging to verify the authenticity of references [2]. In a subsegment of research fields, the delivery mechanism of research publication management is shifting toward the use of multi-agent frameworks, such as PaperOrchestra, to automate the drafting of manuscripts [5]. Regarding research output reporting, a subsegment of researchers is adapting content formatting—specifically utilizing direct Q&A and tables—to increase visibility within AI search engines [6]. This trend is supported by a shift where AI citations are driven by specific “citability” factors rather than traditional domain authority [7] and the development of dedicated AI visibility infrastructure [4]. In institution-specific settings, such as the Alignment Journal, experimental policies are introducing reviewer compensation to modify the peer-review workflow [3]. The social legitimacy of the human-led peer-review process is unlikely to change in this window because institutional tenure and promotion regimes remain dependent on human intellectual accountability.

Hypothesis: In the near term, research publication management will be characterized by a systemic tension between the automated acceleration of manuscript production and the degradation of the peer-review system’s capacity to filter synthetic outputs.

Mid (2036–2041)

For research publications broadly, if the volume of synthetic outputs continues to outpace human review capacity [8] and binary detection tools remain unstable [1], the mechanism of verification may shift from reactive detection to proactive provenance. Extrapolating from current trends in registry adoption, publishers may require research output reporting to be anchored in verified citation registries [2] to prevent the proliferation of fabricated evidence. Research output reporting may further diverge from the static PDF toward structured entity formats to maintain discoverability in AI-driven knowledge graphs [4][6]. The journal article is likely to persist as the primary unit of currency for academic promotion due to high institutional inertia in tenure committees. However, the delivery mechanism of the peer review process may integrate structured transparency logs to counteract the opacity of AI-assisted drafting [5].

Hypothesis: In the mid term, the focus of research publication management may move from the detection of synthetic text to the systematic verification of authorship and citation provenance through structured registries.

Far (2046–2051)

For research publications broadly, the long-term viability of the written manuscript as the primary research output remains unknown if the gap between automated generation [5] and detection [1] closes completely while volume continues to scale [8]. One plausible outcome is a transition where the primary research output becomes a verifiable computational model or a registry-linked data stream [2], with written text serving only as a secondary, AI-generated summary for human consumption [4]. A rival outcome involves the erosion of trust in traditional publication management, leading to a fragmentation of the reward system as traditional journals fail to distinguish human discovery from synthetic synthesis [8]. It remains unknown whether visibility infrastructure [4] can scale sufficiently to maintain a global, trusted record of discovery. The fundamental requirement for a verifiable, immutable record of intellectual contribution resists change even at this horizon.

Hypothesis A: In the far term, the primary research output may shift from the written manuscript to a verifiable data stream or computational model.

Hypothesis B (competing): In the far term, the failure to distinguish human discovery from synthetic synthesis may lead to a collapse of the centralized journal system and a fragmentation of academic reward structures.

References

1. Dev To Ai (2026-04-06). “Will AI-Written Papers Be Detected? The Truth About AIGC Detection in 2026.” https://dev.to/yu_z_ed6c1c576593c5e7d41f/will-ai-written-papers-be-detected-the-truth-about-aigc-detection-in-2026-5b50
2. Dev To Ai (2026-04-06). “AI Citation Registries as Information Infrastructure for AI Systems.” <https://dev.to/aigistry/ai-citation-registries-as-information-infrastructure-for-ai-systems-50o6>
3. Lesswrong (2026-04-07). “An Alignment Journal: Features and policies.” <https://www.lesswrong.com/posts/hvq7amw8FKZxEeKqA/an-alignment-journal-features-and-policies>
4. Dev To Ai (2026-04-08). “Building AI Visibility Infrastructure: Inside Jonomor’s Architecture.” https://dev.to/jonomor_ecosystem/building-ai-visibility-infrastructure-inside-jonomors-architecture-3enp
5. Marktechpost (2026-04-09). “Google AI Research Introduces PaperOrchestra: A Multi-Agent Framework for Automated AI Research Paper Writing.” <https://www.marktechpost.com/2026/04/08/google-ai-research-introduces-paperorchestra-a-multi-agent-framework-for-automated-ai-research-paper-writing/>
6. Dev To Ai (2026-04-10). “How to Get Cited by AI Search Engines: ChatGPT, Perplexity & Claude in 2026.” <https://dev.to/iris1031/how-to-get-cited-by-ai-search-engines-chatgpt-perplexity-claude-in-2026-21p0>
7. Dev To Ai (2026-04-10). “Why DA Is Irrelevant for AI Citations (Data from 7 Site Audits).” https://dev.to/chudi_nnorukam/why-da-is-irrelevant-for-ai-citations-data-from-7-site-audits-1il0
8. Wonkhe (2026-05-07). “Detection tools won’t save the scholarly record from the AI paper avalanche.” <https://wonkhe.com/blogs/detection-tools-wont-save-the-scholarly-record-from-the-ai-paper-avalanche/>

Research Training

capability_id: research-training · created: 2026-08-27 · revised: 2026-08-27 · refs: 9

Near (2026–2031)

For research training broadly, the researcher training and development process may integrate AI-driven search and deep research tools to automate source identification and the initial generation of literature syntheses [3][4]. This modification changes the delivery mechanism of the literature review process rather than its educational substance. For PhD candidates and independent researchers, postgraduate research student skills development may incorporate AI as a critical thinking partner to transition from passive reading to active critical synthesis [9]. In the subsegment of quantitative research, skills development may include AI-driven data analysis tools to automate dataset exploration and visualization [5]. Within the subsegment of machine learning research, postgraduate research student skills development may shift from a focus on theoretical mathematical derivation toward compute-intensive engineering and the management of empirical experiments [1]. This shift represents a change in educational substance, as research progress in this field increasingly relies on compute scaling over mathematically principled architectures [1]. For technical research subsegments requiring software development, skills development may pivot from teaching manual coding syntax to the orchestration of agentic AI tools and terminal-based agents [2][7]. In the subsegment of scientific writing and peer review training, the development of researcher skills may include AI systems designed to generate constructive feedback by analyzing author responses [6]. Within the subsegment of formal mathematics, the automation of proofs may create a counterproductive effect where the removal of the “struggle” of derivation erodes the acquisition of essential mathematical intuition [8]. The institutional requirement for a written thesis and a formal oral defense is unlikely to change in the near term due to the high inertia of degree-awarding regulations.

Hypothesis: In the near term, general research training will integrate AI-driven sourcing and synthesis, while technical subsegments will pivot toward compute engineering and agentic orchestration, accompanied by a risk of intuition erosion in formal disciplines.

Mid (2036–2041)

For the broad and technical cohorts identified in the near term, the postgraduate research student skills development process may transition from a focus on information retrieval to the rigorous validation of AI-curated syntheses if current trajectories hold [3][4][9]. Extrapolating from the adoption of agentic systems, researcher training in technical subsegments may replace manual codebase maintenance with the orchestration of autonomous coding agents as a primary technical competency [2][7]. This shift may alter the substance of postgraduate research supervisor development, as supervisors may move from guiding manual methodology to validating agent-generated empirical workflows. The necessity for the researcher to frame a novel research question and maintain an original theoretical contribution is likely to persist unchanged because these tasks require high-level conceptual synthesis and social legitimacy that agentic trajectories do not currently replace.

Hypothesis: In the mid term, research training will shift toward the validation of automated synthesis and the orchestration of agentic systems, shifting supervisory focus toward workflow validation.

Far (2046–2051)

For machine learning and technical research subsegments, the long-term trajectory of skills development depends on whether compute-intensive scaling continues to yield returns or reaches a threshold of diminishing returns [1]. If scaling remains the primary driver of discovery, research training may evolve into a specialized discipline of high-performance infrastructure management and agent orchestration [1][2]. Alternatively, if scaling returns diminish, the subsegment may experience a return to mathematically principled architectures and theoretical rigor as the primary means of advancement [1][8]. It remains unknown whether the cognitive “intuition” lost through proof automation can be recovered or if a new form of “synthetic intuition” will emerge. The social legitimacy of the human researcher as the final arbiter of

truth and the author of record resists change even at this horizon due to the institutional requirements of accountability and academic credit.

Hypothesis A: In the far term, technical research training becomes a discipline of infrastructure management and agentic orchestration, provided compute scaling remains the primary driver of discovery.

Hypothesis B (competing): In the far term, technical research training returns to a focus on theoretical mathematical rigor if compute-intensive scaling reaches diminishing returns.

References

1. The Gradient (2024-11-16). “Shape, Symmetries, and Structure: The Changing Role of Mathematics in Machine Learning Research.” <https://thegradient.pub/shape-symmetry-structure/>
2. Dev To Ai (2026-04-10). “Claude Code Cheatsheet: The Terminal AI Agent That Actually Codes.” <https://dev.to/mibii/claude-code-cheatsheet-the-terminal-ai-agent-that-actually-codes-36bk>
3. OpenAI (2026-04-10). “ChatGPT for research.” <https://openai.com/academy/research>
4. OpenAI (2026-04-10). “Research with ChatGPT.” <https://openai.com/academy/search-and-deep-research>
5. OpenAI (2026-04-10). “Analyzing data with ChatGPT.” <https://openai.com/academy/data-analysis>
6. Arxiv Cs Ai (2026-04-16). “GoodPoint: Learning Constructive Scientific Paper Feedback from Author Responses.” <https://arxiv.org/abs/2604.11924>
7. Simon Willison (2026-04-25). “Quoting Romain Huet.” <https://simonwillison.net/2026/Apr/25/ro-main-huet/#atom-everything>
8. Lesswrong (2026-04-29). “The fall of the theorem economy (David Bessis).” <https://www.lesswrong.com/posts/5CA4z7TumhrQu5Jr4/the-fall-of-the-theorem-economy-david-bessis>
9. Dev To Ai (2026-05-07). “From Reading to Reasoning: AI for Critical Synthesis.” https://dev.to/kenn_deng_ai/from-reading-to-reasoning-ai-for-critical-synthesis-25nb

Strategy Management

capability_id: strategy-management · created: 2026-09-30 · revised: 2026-09-30 · refs: 16

Near (2026–2031)

For institutions in China, state-funded research entities, public universities under fiscal pressure, and a subsegment of administrators adopting automated synthesis tools, strategy management is shifting toward a combination of external mandate alignment and tool-assisted synthesis. In China, national targets for AI deployment are integrated into strategic plan development and strategic plan management via state-led planning cycles [1]. For state-funded research bodies, as demonstrated by the Alan Turing Institute, funding agencies may mandate revisions to strategic plan development to ensure adherence to value-for-money metrics [2]. In public institutions facing budget constraints, such as Iowa State University and East Carolina University, corporate performance management is increasingly driven by mandated reviews that use viability data to execute program closures or mergers [6][10]. For institutions utilizing generative AI, the processes of business horizon scanning and strategic reporting may be streamlined through consumer insight tools [5], labor automation forecasting [7], or AGI scenario analysis [11][13]. Some administrators are transitioning strategic reporting from static dashboards to conversational intelligence to analyze institutional performance [17], while redefining corporate performance management via intent-driven budgeting [9] and KPIs that prioritize value over simple efficiency [8][16]. These changes represent delivery-mechanism changes in strategic reporting and horizon scanning, while shifting the educational substance of corporate performance management toward fiscal viability and redefined value metrics [6][8][10][16]. The fundamental legal fiduciary duties of governing boards and the statutory requirements of university charters are unlikely to change in this period.

Hypothesis: In the near term, strategy management in these settings may be increasingly defined by the transition of AI from a technical project to an executive mandate, utilizing automated synthesis to accelerate reporting and the redefinition of performance metrics.

Mid (2036–2041)

For institutions in China, state-funded research bodies, public universities under state oversight, and adopters of AI strategy tools, if current trajectories hold, corporate performance management may shift toward the quantitative measurement of external benchmarks and AI-generated KPIs. Extrapolating from state mandates [1] and mandated program reviews [6][10], strategic reporting could move from an internal reflection of institutional mission to a compliance verification tool, altering the educational substance of how institutional success is defined. For institutions using automated strategy platforms, the process of strategic plan development may shift from an act of primary synthesis to an act of auditing AI-generated outputs, such as automated SWOT matrices [14] or “vibe-based” consulting reports [3][4]. If the shift toward value-based KPIs persists [8][16], strategic vision development may diverge between institutions that prioritize AI-driven efficiency and those that prioritize human-centric value creation. The requirement for institutions to maintain a distinct institutional identity within their strategic vision development is likely to persist, though it will operate within boundaries set by external regulatory authorities and tool constraints.

Hypothesis: In the mid term, corporate performance management in these segments may transition into a process of algorithmic compliance and output auditing rather than primary strategic synthesis.

Far (2046–2051)

For the same subset of institutions identified in the near term, strategy management may bifurcate based on the stability of the institutional environment. If the trend toward agentic AI persists, strategic plan management could evolve into a real-time, continuous optimization process where agentic systems autonomously adjust resource allocation based on live performance data [15]. However, a rival outcome may emerge if the reliance on automated “vibe-style” reporting [3][4] leads to a loss of institutional legitimacy or a failure in long-term viability, triggering a return to manually curated, human-led strategic vision develop-

ment to restore trust. It remains unknown whether AI-driven KPIs can capture the qualitative essence of educational purpose or if they will inevitably drive counterproductive institutional behavior by optimizing for measurable metrics over educational substance. The necessity of final human sign-off for high-stakes strategic shifts is likely to resist change due to the requirement for legal and social accountability.

Hypothesis A: In the far term, strategy management for these institutions may shift toward a model of bounded autonomy, where AI agents manage the execution and reporting of strategy in real-time.

Hypothesis B (competing): In the far term, a failure of automated strategic synthesis may lead to a strategic regression, where human-led, slow-cycle planning is reinstated as a mechanism for institutional legitimacy.

References

1. Ai News (2026-04-02). "China's Five-Year Plan details the targets for AI deployment." <https://www.artificialintelligence-news.com/news/chinas-five-year-plan-details-the-targets-for-ai-deployment/>
2. Guardian Tech (2026-04-03). "UK's leading AI research institute told to make 'significant' changes." <https://www.theguardian.com/technology/2026/apr/03/uk-alan-turing-institute-ai-research-underperforming-ukri>
3. TechCrunch (2026-04-07). "AI startup Rocket offers vibe McKinsey-style reports at a fraction of the cost." <https://techcrunch.com/2026/04/06/indian-startup-rocket-wants-its-ai-to-do-mckinsey-style-consulting-at-a-fraction-of-the-cost/>
4. Dev To Ai (2026-04-07). "AI startup Rocket offers vibe McKinsey-style reports at a fraction of the cost." <https://dev.to/minimal-architect/ai-startup-rocket-offers-vibe-mckinsey-style-reports-at-a-fraction-of-the-cost-1edd>
5. Mit Sloan (2026-04-08). "Gain Consumer Insight With Generative AI." <https://sloanreview.mit.edu/article/gain-consumer-insight-with-generative-ai/>
6. Higher Ed Dive (2026-04-10). "Iowa State plans to cut or merge 23 programs after mandated review." <https://www.highereddive.com/news/iowa-state-university-program-cuts-review-regents/817262/>
7. Lesswrong (2026-04-21). "Introducing the Labor Automation Forecasting Hub." <https://www.lesswrong.com/posts/So7wr6djpQcoP4rz/introducing-the-labor-automation-forecasting-hub-1>
8. Fast Company (2026-04-23). "The Age of AI means we need to throw out our old KPIs and replace them with new ones." https://www.fastcompany.com/91521422/ai-kpis-new-metrics?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
9. Dzone Ai (2026-04-24). "88% of Executives Plan to Increase AI Budgets by 2026 - Are We Investing With Intent or Just Urgency?" <https://feeds.dzone.com/link/23558/17325011/increasing-ai-budgets-investing-intent-or-urgency>
10. Higher Ed Dive (2026-04-27). "East Carolina University plans to cut 44 academic programs." <https://www.highereddive.com/news/east-carolina-university-cut-around-44-academic-programs/818622/>
11. Arxiv Cs Cy (2026-04-28). "Artificial General Intelligence Forecasting and Scenario Analysis: State of the Field, Methodological Gaps, and Strategic Implications." <https://arxiv.org/abs/2604.22766>
12. Zdnet Ai (2026-05-01). "Moving from AI pilots to business-wide value requires a superhighway - how to ramp up." <https://www.zdnet.com/article/moving-from-ai-pilots-to-business-wide-value-accenture-research/>
13. Arxiv Cs Cy (2026-05-05). "Generative-AI and the transformation of workforce. A job postings-driven analysis." <https://arxiv.org/abs/2605.00843>
14. Dev To Ai (2026-05-05). "AI prompt for SWOT analysis: Build a Decision-Ready Strategy Matrix Faster." https://dev.to/asma_habib_1e94a3083c9049/ai-prompt-for-swot-analysis-build-a-decision-ready-strategy-matrix-faster-508o
15. Dev To Ai (2026-05-07). "The Future of Enterprise AI: Navigating the Agentic Shift Beyond RAG in 2027." <https://dev.to/workalizer/the-future-of-enterprise-ai-navigating-the-agentic-shift-beyond-rag-in-2027-41m8>
16. Dev To Ai (2026-05-07). "How AI and LLMs Are Changing the Business Metrics That Actually

- Matter.” <https://dev.to/xcelcorp/how-ai-and-llms-are-changing-the-business-metrics-that-actually-matter-57dk>
17. Dev To Ai (2026-05-10). “The End of the Static Dashboard: Why Generative AI is Eating BI.” <https://dev.to/dattasable/the-end-of-the-static-dashboard-why-generative-ai-is-eating-bi-4ce5>

Student Administration

capability_id: student-administration · created: 2026-06-17 · revised: 2026-06-17 · refs: 2

Near (2026–2031)

For a subset of institutions, including public universities subject to state-mandated program reviews [1] and those managing complex financial aid communications [2], administrative processes are facing increased volatility and regulatory pressure. For universities undergoing mandated closures of low-enrollment programs, as seen at Iowa State University [1], the process of programme transfer management may experience surges in volume, potentially driving the adoption of AI-based credit mapping to accelerate student reassignment. Simultaneously, legislative efforts to resolve jargon-heavy financial aid offers [2] may prompt student financial administration to implement AI-driven transparency tools to translate “total net expenses” into concrete out-of-pocket costs for families. These changes represent shifts in the delivery mechanism—specifically the speed of transfer processing and the clarity of financial communication—rather than changes to the educational substance. Enrolment status management will likely remain a manual, human-verified process to ensure legal compliance with state and federal funding mandates. It is unlikely that the core regulatory requirements for degree conferral will change, as these are anchored by external accreditation bodies.

Hypothesis: In the near term, for institutions facing programmatic volatility or legislative pressure on financial transparency, the delivery mechanisms of programme transfer management and student financial administration may shift toward AI-supported automation and translation.

Mid (2036–2041)

For institutions facing these specific regulatory and legislative pressures, if current trajectories of mandated program reviews and financial aid transparency hold, the load on programme transfer management could lead to a shift from AI-supported to AI-led automated matching. Extrapolating from current administrative constraints, such systems might automatically suggest alternative degree paths for displaced students based on completed credits and labor market demand. Within student financial administration, the requirement for standardized, jargon-free communication may lead to the adoption of algorithmic cost-projection tools that automate the calculation of individual family liability [2]. However, the final validation of enrolment status management will likely persist as a human-governed process to prevent fraud and ensure adherence to government funding audits.

Hypothesis: In the mid term, for institutions experiencing frequent program restructuring or strict financial reporting mandates, the delivery mechanism of student administration may move toward algorithmic decision-support for transfers and cost projections.

Far (2046–2051)

For institutions operating under these regimes of programmatic volatility and high financial transparency, one plausible outcome is that programme transfer management becomes a dynamic process that re-aligns student enrolments in real-time based on shifting state mandates and employment signals. A rival outcome is that the persistent instability of specific program offerings leads to a shift toward broader, degree-agnostic enrolment status management, which would reduce the frequency and necessity of specific programme transfer events. It remains unknown whether state regents will continue to maintain the authority to mandate such cuts or if institutional autonomy over program design will be restored. The legal requirement for permanent, immutable student records & details management will likely resist change even at this horizon due to the long-term necessity of degree verification.

Hypothesis A: In the far term, for public universities in volatile regulatory environments, programme transfer management may be integrated into dynamic, AI-driven labor-market alignment systems.

Hypothesis B (competing): In the far term, the administrative burden of frequent program cuts may drive a shift toward modular, degree-agnostic enrolment status management, reducing the operational reliance on

traditional programme transfer processes.

References

1. Higher Ed Dive (2026). “Iowa State plans to cut or merge 23 programs after mandated review.” <https://www.highereddive.com/news/iowa-state-university-program-cuts-review-regents/817262/>
2. The Hechinger Report (2026). “Confusing financial aid offers can leave families deeper in debt. Student groups say a new legislative fix doesn’t go far enough.” <https://hechingerreport.org/confusing-financial-aid-offers-can-leave-families-deeper-in-debt-a-fix-is-stalled-in-congress/>

Student Admission Management

capability_id: student-admission-management · created: 2026-09-01 · revised: 2026-09-01 · refs: 6

Near (2026–2031)

For a subsegment of institutions characterized by a reliance on digital identity verification, competitive portfolio review, and opaque financial aid communication, specific processes within this capability are under pressure. Because AI-generated imagery enables the creation of high-quality fraudulent identification [4] and real-time deepfake software is now available as a SaaS product [7], the identity verification process within student application management is undermined, increasing the risk of identity fraud during high-stakes live video interviews. Simultaneously, the emergence of AI-saturated markets creates a requirement for human-provenance verification [6], shifting the student application management process toward requiring evidence of human effort in portfolios and essays to distinguish human work from synthetic content. In the offer, acceptance & quota management process, prospective students may utilize large language models to normalize and compare disparate, confusing financial aid offer formats [1][2]. This represents a change in the delivery mechanism of the offer, as the interpretation of the cost of attendance moves from an institutionally controlled narrative to an AI-mediated analysis. For institutions implementing restrictive admission policies, the offer and quota management process is further constrained by regulatory and legal probes regarding eligibility criteria [5]. The requirement for a formally signed and dated offer of admission is unlikely to change in the near term due to the contractual nature of enrollment.

Hypothesis: In the near term, for a subsegment of institutions, AI may simultaneously undermine identity verification in application management and necessitate new provenance verification for portfolios, while increasing information symmetry for students decoding financial aid offers.

Mid (2036–2041)

For the same subsegment of institutions, if current trajectories hold, some may employ AI-driven personalization to optimize quota management by tailoring the presentation of aid to specific student psychographics. This would represent a change in the educational substance of the capability, as the logic of the offer shifts from transparent disclosure to yield optimization. Such a mechanism increases the risk of counterproductivity, where the pursuit of higher yield rates through calculated opacity undermines the social legitimacy of the admission process. To counter the proliferation of AI-generated fraudulent credentials and synthetic portfolios, institutions may shift identity and provenance verification from static document review to continuous or biometric authentication. Despite these shifts, the core regulatory requirements for federal financial aid reporting will likely remain a stabilizing constraint.

Hypothesis: In the mid term, some institutions may use AI to personalize financial aid offers to optimize yield, while shifting application management toward biometric and provenance-based identity verification to counter AI-generated fraud.

Far (2046–2051)

For this subsegment of institutions, two rival trajectories emerge regarding the delivery of admission offers. One outcome involves a regulatory response to AI-driven opacity, mandating a standardized, machine-readable format for all financial aid offers to ensure comparability and prevent predatory yield optimization. A competing outcome involves the obsolescence of the “offer letter” as a discrete artifact, replaced by real-time, AI-managed cost-of-attendance dashboards that update based on the student’s shifting financial profile. It remains unknown whether accreditation bodies will eventually categorize AI-optimized personalized pricing as a violation of equitable access. The fundamental necessity of a final selection event to manage physical campus capacity remains a resistant constraint even at this horizon.

Hypothesis A: In the far term, regulatory mandates may force the standardization of financial aid offers into machine-readable formats to eliminate AI-driven opacity.

Hypothesis B (competing): In the far term, the static admission offer may be replaced by dynamic, AI-

managed financial dashboards that update in real-time.

References

1. The Hechinger Report (2026). “How much will that college cost you? Good luck figuring it out.” <https://hechingerreport.org/how-much-will-that-college-cost-you-good-luck-figuring-it-out-2/>
2. The Hechinger Report (2026). “Confusing financial aid offers can leave families deeper in debt. Student groups say a new legislative fix doesn’t go far enough.” <https://hechingerreport.org/confusing-financial-aid-offers-can-leave-families-deeper-in-debt-a-fix-is-stalled-in-congress/>
3. Atlantic (2026). “Hey Chat, Make Me a Fake ID.” https://www.theatlantic.com/technology/2026/05/chatgpt-images-deepfakes-fraud/687023/?utm_source=feed
4. Higher Ed Dive (2026). “Smith College faces Title IX probe over policy of admitting trans students.” <https://www.highereddive.com/news/smith-college-faces-title-ix-probe-over-policy-of-admitting-trans-students/819358/>
5. Arxiv Cs Cy (2026). “Human-Provenance Verification should be Treated as Labor Infrastructure in AI-Saturated Markets.” <https://arxiv.org/abs/2605.03210>
6. Dev To Ai (2026). “Realtime deepfake software is a SaaS product now.” https://dev.to/thoams_ai/detection/realtime-deepfake-software-is-a-saas-product-now-13no

Student Assessment

capability_id: student-assessment · created: 2026-09-29 · revised: 2026-09-29 · refs: 22

Near (2026–2031)

For student assessment broadly, the delivery mechanism is shifting toward iterative writing cycles where students use AI as a partner [1], a transition that exposes structural flaws in the reliance on essay assessments within mass education systems [3]. This shift in delivery may induce “cognitive surrender,” where students offload logical thinking to synthetic agents [5], resulting in a homogenization of writing and thinking patterns that complicates assessment marking and feedback by eroding the distinct critical voice [9] and removing the cognitive struggle required to translate thought into prose [43]. In the process of assessment administration, the reliability of authorship verification is constrained by the persistence of false positives in AI detectors [13], the continuous decay of detection pipelines as models evolve [28], and the emergence of “humanizer” tools [15]. For institutions updating academic integrity policies, the definition of plagiarism is becoming fragmented as universities adopt divergent standards for AI usage [36]. For a subset of institutions utilizing remote proctoring, the integrity of assessment delivery is undermined by the availability of real-time deepfake software as a SaaS product [38]. For institutions adopting automated marking and feedback, the process is subject to algorithmic bias, with evidence that AI may provide disparate levels of praise and criticism based on student race [27]. High-stakes, invigilated examinations are unlikely to change in the near term due to professional accreditation requirements and the institutional need for authenticated baseline competency [19].

Hypothesis: In the near term, assessment delivery may shift from artifact-production to process-evaluation, while assessment administration is increasingly constrained by the technical failure of authorship detection and identity verification tools.

Mid (2036–2041)

For student assessment broadly, if current trajectories hold, the educational substance of assessment delivery may shift from evaluating a final product to evaluating a student’s capacity to audit, verify, and correct synthetic errors [7]. This represents a shift in substance from measuring output to measuring the ability to steer and validate synthetic agents [14], supported by the development of multi-agent approaches to validate LLM-generated problems [11] and reasoning-effort analysis in automated scoring [32]. To maintain the validity of assessment results management, institutions may implement deliverable-oriented governance frameworks to manage the inherent opacity of these AI systems [20]. For a subset of institutions focusing on multimodal pedagogical analysis, results management could incorporate the longitudinal mapping of cognitive effort via educational fMRIs [4]. The basic administrative recording of credits and formal result transmission will likely persist unchanged due to high regulatory inertia.

Hypothesis: In the mid term, the substance of student assessment may shift from measuring product output to measuring a student’s ability to audit and validate synthetic agents.

Far (2046–2051)

For student assessment broadly, the long-term trajectory depends on whether synthetic protocols can reliably replace human judgment in certifying higher-order cognitive capabilities [18]. It remains unknown if LLM-based protocols can meaningfully measure collaboration, creativity, and critical thinking without reducing these capabilities to observable patterns [17]. Because the “illusion of learning” may decouple task completion from actual capability [41], the biological nature of cognition remains a point of resistance that may force a return to strict provenance verification. The social legitimacy of a degree will depend on whether the market values the verification of biological cognition or the mastery of synthetic orchestration.

Hypothesis A: In the far term, human-provenance verification becomes treated as essential labor infrastructure to certify biological cognition in AI-saturated markets [35].

Hypothesis B (competing): In the far term, explainable AI certification protocols replace human judgment

as the primary mechanism for validating professional competency [18].

References

1. The Conversation (2026-03-26). “College students are writing with AI – but a pilot study finds they’re not simply letting it write for them.” <https://theconversation.com/college-students-are-writing-with-ai-but-a-pilot-study-finds-theyre-not-simply-letting-it-write-for-them-276856>
2. The Conversation Au (2026-04-01). “Does AI mean more uni students are plagiarising their work?” <https://theconversation.com/does-ai-mean-more-uni-students-are-plagiarising-their-work-279565>
3. Wonkhe (2026-04-01). “Generative AI exposes longstanding flaws in the use of essay assessments in a mass system.” <https://wonkhe.com/blogs/generative-ai-exposes-longstanding-flaws-in-the-use-of-essay-assessments-in-a-mass-system/>
4. Getting Smart (2026-04-01). “Educational fMRIs: Dynamic Pedagogy and Pedagogical Analysis in the Multimodal AI Era.” <https://www.gettingsmart.com/2026/04/01/educational-fmr-is-dynamic-pedagogy-and-pedagogical-analysis-in-the-multimodal-ai-era/>
5. Ars Technica (2026-04-03). ““Cognitive surrender” leads AI users to abandon logical thinking, research finds.” <https://arstechnica.com/ai/2026/04/research-finds-ai-users-scarily-willing-to-surrender-their-cognition-to-llms/>
6. Dev To Ai (2026-04-07). “Where Does Truth Live in AI-Generated Code?” <https://dev.to/wintrover/where-does-truth-live-in-ai-generated-code-5hmc>
7. BBC News (2026-04-07). ““Our teachers use AI to mark mock exams.”” https://www.bbc.com/news/articles/c6207xenjl1do?at_medium=RSS&at_campaign=rss
8. Hacker News (2026-04-07). “AI may be making us think and write more alike.” <https://dornsife.usc.edu/news/stories/ai-may-be-making-us-think-and-write-more-alike/>
9. Arxiv Cs Cy (2026-04-08). “A Multi-Agent Approach to Validate and Refine LLM-Generated Personalized Math Problems.” <https://arxiv.org/abs/2604.05160>
10. Dev To Ai (2026-04-09). “Why AI Detectors Produce False Positives: A Technical Analysis.” <https://dev.to/laakash/why-ai-detectors-produce-false-positives-a-technical-analysis-2gpc>
11. Atlantic (2026-04-10). “This Is What Fully Automated School Looks Like.” https://www.theatlantic.com/technology/2026/04/ai-agents-school-education/686754/?utm_source=feed
12. Dev To Ai (2026-04-11). “I Built a Free AI Detector + Humanizer with Sentence-Level Highlighting.” https://dev.to/fafa_ai/i-built-a-free-ai-detector-humanizer-with-sentence-level-highlighting-51jj
13. Marktechpost (2026-04-14). “Google AI Research Proposes Vantage: An LLM-Based Protocol for Measuring Collaboration, Creativity, and Critical Thinking.” <https://www.marktechpost.com/2026/04/13/google-ai-research-proposes-vantage-an-llm-based-protocol-for-measuring-collaboration-creativity-and-critical-thinking/>
14. Arxiv Cs Cy (2026-04-14). “Explainability and Certification of AI-Generated Educational Assessments.” <https://arxiv.org/abs/2604.09622>
15. University Affairs (2026-04-14). “Reconsidering the value of exams in the age of AI.” <https://universityaffairs.ca/career-advice/reconsidering-the-value-of-exams-in-the-age-of-ai/>
16. Arxiv Cs Ai (2026-04-23). “AI to Learn 2.0: A Deliverable-Oriented Governance Framework and Maturity Rubric for Opaque AI in Learning-Intensive Domains.” <https://arxiv.org/abs/2604.19751>
17. The Hechinger Report (2026-04-27). “AI gives more praise, less criticism to Black students.” <https://hechingerreport.org/proof-points-ai-bias-feedback/>
18. Dev To Ai (2026-04-27). “LLM Drift: Why Your AI Detection Pipeline is Quietly Decaying (Kimi K2 Benchmark).” https://dev.to/thoams_aidetection/llm-drift-why-your-ai-detection-pipeline-is-quietly-decaying-kimi-k2-benchmark-3gml
19. Arxiv Cs Cy (2026-05-01). “The Impact of LLM Self-Consistency and Reasoning Effort on Automated Scoring Accuracy and Cost.” <https://arxiv.org/abs/2604.26954>
20. Arxiv Cs Cy (2026-05-06). “Human-Provenance Verification should be Treated as Labor Infrastructure in AI-Saturated Markets.” <https://arxiv.org/abs/2605.03210>
21. Dev To Ai (2026-05-06). “How Different Universities Define Plagiarism in 2026.” <https://dev.to/m>

- [artiinn/how-different-universities-define-plagiarism-in-2026-9mi](#)
22. Dev To Ai (2026-05-07). “Realtime deepfake software is a SaaS product now.” https://dev.to/thoams_aidetection/realtime-deepfake-software-is-a-saas-product-now-13no
 23. Elearning Industry (2026-05-08). “Tunnel Vision, Consolidated Truth, And The Illusion Of Learning.” <https://elearningindustry.com/tunnel-vision-consolidated-truth-and-the-illusion-of-learning>
 24. The Guardian (2026-05-10). “I knew my writing students were using AI. Their confessions led to a powerful teaching moment | Micah Nathan.” <https://www.theguardian.com/us-news/ng-interactive/2026/may/10/fiction-writing-professor-ai>

Student Attraction & Recruitment

capability_id: student-attraction-recruitment · created: 2026-08-01 · revised: 2026-08-01 · refs: 11

Near (2026–2031)

For the capability of student attraction and recruitment broadly, the delivery mechanism for promotional content may shift as generative AI enables the production of comprehensive ad campaigns—including scripts and imagery—at negligible cost [10]. For a subset of institutions focusing on digital discovery, the process of prospective student engagement may shift from keyword-based search engine optimization toward generative engine optimization (GEO) to maintain visibility within AI-mediated search results [1][6][9]. Within the process of prospective student engagement at recruitment events, institutions may adopt AI to automate the conversion of raw lead notes into narratives [4], execute automated post-event follow-ups [7][11], and apply instant lead scoring to prioritize prospects [8]. For institutions utilizing multimedia outreach, the delivery of engagement videos may be streamlined through integrated AI video tools [2]. In specific regions where COVID-era remote learning has created a drain on college enrollment, the process of domestic student recruitment is constrained by a shrinking pool of traditional candidates [3]. Within scholarship and bursary management, prospective student engagement is hindered by a lack of transparency regarding the actual cost of attendance [5]. The core fiduciary criteria and legal frameworks used to determine scholarship eligibility are unlikely to change in this period due to rigid institutional governance and audit requirements.

Hypothesis: In the near term, institutions may shift recruitment delivery toward high-velocity AI content generation and automated lead-scoring pipelines, while regional enrollment declines and transparency failures in scholarship communication may limit the effectiveness of these tools.

Mid (2036–2041)

For the subset of institutions utilizing AI-driven engagement, if current trajectories hold, the process of prospective student engagement could shift toward the maintenance of structured data schemas optimized for machine readability rather than human-centric web narratives [1][6][9]. Extrapolating from current lead-intelligence deployments, the qualification of prospective students at recruitment fairs may move from an asynchronous post-event task to a real-time automated pipeline [4][7][8][11]. The proliferation of low-cost, AI-generated recruitment media may lead to content saturation, potentially shifting international and domestic student recruitment toward a higher reliance on human-led validation to establish institutional authenticity [2][10]. If regional enrollment drains persist, domestic student recruitment may pivot toward alternative lead-generation mechanisms that bypass traditional pipeline markers [3]. To address friction in scholarship and bursary management, institutions may deploy AI-mediated tools to provide personalized, transparent cost projections for prospective families [5]. The requirement for faculty-authored program descriptions will likely persist as a baseline for academic legitimacy.

Hypothesis: In the mid term, prospective student engagement for some institutions may shift toward the production of machine-optimized entity data, while AI may be used to reduce information asymmetry in scholarship and bursary management.

Far (2046–2051)

For institutions relying on AI-mediated prospective student engagement, it remains unknown whether AI discovery platforms will maintain neutral recommendation architectures or shift toward commercialized, pay-to-play visibility models. One possible outcome involves the bounded autonomy of recruitment pipelines, where AI manages the entirety of lead generation and scoring, but human intervention remains the sole mechanism for final conversion and social legitimacy. A rival outcome involves a shift in the educational substance of recruitment, where institutional brands become secondary to AI-curated “learning paths” that recommend institutions based on real-time outcome data. The legal requirement for accredited degree-granting status will likely resist change even at this horizon, as it remains the primary regulatory constraint on institutional legitimacy.

Hypothesis A: In the far term, recruitment may evolve into a highly automated pipeline of lead qualification where human interaction is reserved exclusively for high-value final conversion.

Hypothesis B (competing): In the far term, the process of student attraction may shift from institutional branding to AI-driven matching based on granular outcome data, reducing the autonomy of individual recruitment offices.

References

1. Dev To Ai (2026). “What is the difference between AI content and AI search optimised content?” <https://dev.to/rank4ai/what-is-the-difference-between-ai-content-and-ai-search-optimised-content-37bc>
2. Google AI Blog (2026). “Create, edit and share videos at no cost in Google Vids” <https://blog.google/products-and-platforms/products/workspace/google-vids-updates-lyria-veo/>
3. Higher Ed Dive (2026). “COVID remote learning put drain on college enrollment” <https://www.highereddive.com/news/covid-remote-learning-put-drain-on-college-enrollment/816899/>
4. Dev To Ai (2026). “From Notes to Narrative: How AI Automates Trade Show Lead Intelligence” https://dev.to/ken_deng_ai/from-notes-to-narrative-how-ai-automates-trade-show-lead-intelligence-1239
5. The Hechinger Report (2026). “How much will that college cost you? Good luck figuring it out” <https://hechingerreport.org/how-much-will-that-college-cost-you-good-luck-figuring-it-out-2/>
6. Dev To Ai (2026). “I Audited a Business’s AI Visibility Across Four Platforms. The Results Were Worse Than Expected.” <https://dev.to/edvisageglobal/i-audited-a-businesss-ai-visibility-across-four-platforms-the-results-were-worse-than-expected-4ml5>
7. Dev To Ai (2026). “Stop Chasing Ghosts: Automate Your Post-Trade Show Follow-Up” https://dev.to/ken_deng_ai/stop-chasing-ghosts-automate-your-post-trade-show-follow-up-5f65
8. Dev To Ai (2026). “Instant Lead Scoring: Teaching AI to Sort Trade Show Prospects” https://dev.to/ken_deng_ai/instant-lead-scoring-teaching-ai-to-sort-trade-show-prospects-4111
9. Dev To Ai (2026). “Write a 1,200-word blog post:”What is Generative Engine Optimization (GEO) and why SEO teams need i” <https://dev.to/sieok/write-a-1200-word-blog-post-what-is-generative-engine-optimization-geo-and-why-seo-teams-need-i-35a0>
10. Dev To Ai (2026). “How we shipped”One Brief Full Ad Campaign in 60s” for \$0.35 — a Dify + x402 + multi-engine AI build” <https://dev.to/aikitpros/how-we-shipped-one-brief-full-ad-campaign-in-60s-for-035-a-dify-x402-multi-engine-ai-ib>
11. Dev To Ai (2026). “From Booth to Bot: Automating Your Trade Show Follow-Up with AI” https://dev.to/ken_deng_ai/from-booth-to-bot-automating-your-trade-show-follow-up-with-ai-jmb

Student Completion & Graduation

capability_id: student-completion-graduation · created: 2026-06-29 · revised: 2026-06-29 · refs: 1

Near (2026–2031)

For the subsegment of institutions utilizing document-based identity verification within graduation record certificate management, the reliability of manual identity checks is declining. The deployment of generative AI tools capable of producing high-fidelity fraudulent identification [1] increases the risk of unauthorized individuals claiming certificates or degrees. This shift alters the delivery mechanism of identity verification by rendering visual inspection of IDs insufficient to prevent fraud. Because this specific pressure is external to the institution, graduation eligibility management—which relies on internal academic record audits—is unlikely to change in this period.

Hypothesis: In the near term, institutions managing certificate distribution may transition from simple document-based verification to multi-factor or digital identity checks to mitigate the risk of AI-generated identity fraud.

Mid (2036–2041)

For the same subsegment of institutions, if current trajectories in synthetic media hold, the use of traditional physical identification for the issuance of graduation records may become obsolete. Extrapolating from the increasing ease of creating deepfake credentials [1], the mechanism of trust may shift from visual verification to cryptographic verification of the recipient’s identity. This represents a change in the delivery mechanism of the capability, moving toward digital signatures and immutable identifiers. The social and ceremonial substance of graduation event management is likely to persist unchanged due to the high social legitimacy attached to the physical ritual.

Hypothesis: In the mid term, graduation record certificate management will likely move toward a model of cryptographic identity verification to maintain the integrity of the credential.

Far (2046–2051)

For the subsegment of institutions concerned with credential integrity, the ability to verify a graduate’s identity may diverge into two rival outcomes. One possibility involves the adoption of globally standardized, decentralized identity frameworks that eliminate the need for institution-specific identity verification. A competing outcome involves a return to high-friction, in-person biometric verification as a response to the total collapse of trust in all digital and physical documents. It remains unknown whether regulatory bodies will mandate a specific verification standard or allow institutional autonomy. Despite these shifts, the fundamental requirement for a verified terminal record to signal completion will likely resist change.

Hypothesis A: In the far term, identity verification for graduation records may be fully absorbed into a decentralized, third-party identity layer, removing the verification burden from the institution.

Hypothesis B (competing): In the far term, institutions may revert to strict, in-person biometric authentication for the issuance of degrees to counter persistent AI-driven identity fraud.

References

1. Atlantic (2026-05-02). “Hey Chat, Make Me a Fake ID.” https://www.theatlantic.com/technology/2026/05/chatgpt-images-deepfakes-fraud/687023/?utm_source=feed

Student Enrolment

capability_id: student-enrolment · created: 2026-04-10 · revised: 2026-04-10 · refs: 1

Near (2026–2031)

In the United States, programme enrolment volumes for first-year students have decreased [1]. This decline is linked to a lower rate of completion for financial aid applications (FAFSA) and standardized testing (ACT) [1]. These frictions at the entry point of the programme enrolment process reduce the pool of eligible candidates entering the pipeline [1]. Because this drain occurs at the initial sign-up phase, module selection management is unlikely to change in the near term as it is a downstream process. Student induction management remains largely unaffected by these volume shifts in the immediate term [1].

Hypothesis: In the near term, programme enrolment in the US region may experience continued volatility as institutions respond to changes in student application behaviors.

Mid (2036–2041)

In the United States, if current trajectories of declining first-year sign-ups hold, institutions may alter the delivery mechanism of programme enrolment to reduce friction [1]. This could involve the deployment of automated tools to assist students in navigating FAFSA and ACT requirements, shifting the process from a passive administrative task to an active, supported workflow. Extrapolating from the volume drain, institutions may also modify student induction management to prioritize retention of a smaller entering cohort [1]. The fundamental regulatory requirements for financial aid eligibility will likely persist unchanged due to government oversight.

Hypothesis: In the mid term, institutions in the affected region may adopt more automated delivery mechanisms for programme enrolment to mitigate volume loss.

Far (2046–2051)

In the United States, the long-term elasticity of student demand relative to remote learning formats remains unknown [1]. If the drain on traditional programme enrolment persists, the capability may diverge into two rival outcomes. One outcome involves the complete automation of the enrolment pipeline, where institutional barriers are minimized to maximize throughput. A competing outcome involves a return to high-touch, human-centric enrolment processes to restore social legitimacy and perceived value [1]. Legal requirements for institutional accreditation will resist change even at this horizon, as they provide the basis for enrolment legitimacy.

Hypothesis A: In the far term, programme enrolment may transition to a continuous, fully automated delivery mechanism to minimize student friction.

Hypothesis B (competing): In the far term, institutions may shift toward high-touch, human-led enrolment processes to counteract the perceived depersonalization of remote learning.

References

1. Higher Ed Dive (2026). “COVID remote learning put drain on college enrollment.” <https://www.highereddive.com/news/covid-remote-learning-put-drain-on-college-enrollment/816899/>

Student Support & Wellbeing Management

capability_id: student-support-wellbeing-management · created: 2026-09-03 · revised: 2026-09-03 · refs: 14

Near (2026–2031)

For a subset of institutions and specific student populations adopting AI-mediated workflows, shifts are emerging in both the delivery mechanism and the educational substance of support processes. In the process of personal learning management, the delivery mechanism is shifting toward AI-automated synthesis of study materials via tools such as Adobe Acrobat Spaces [4] and specific prompting strategies intended to enhance study efficiency for adult learners [12]. Within academic skill development, the educational substance is shifting toward the regulation of “cognitive partnership,” where institutions move students away from simple text generation toward iterative brainstorming [1] and the application of frameworks like TACO [8] or counterargument interventions [14] to prevent AI from substituting for logical thinking. This shift is a mechanistic response to the risk of “cognitive surrender,” where users abandon independent reasoning [2] and exhibit linguistic homogenization in their thought patterns [5]. In the process of academic advice management, the substance of guidance is vulnerable to sociodemographic biases inherent in LLM-based counselling, which may introduce inequities into student support [11]. Within disability support management, delivery is shifting toward the use of “inference engines” to assist students with learning differences, although these tools currently lack systematic safeguarding [9]. In the process of student health & wellbeing, delivery is incorporating AI despite workforce tension between clinical enthusiasm and fear of displacement [3], while the substance of crisis triage is adapting to address critical safety failures in LLMs handling suicide and self-harm [10] and risks of AI-induced delusions [6]. In the region of Florida, these processes are being formally constrained by state-directed partnerships to shield families from AI harm [7]. The processes of student financial advice and housing advice are unlikely to change in this horizon because they remain tied to external regulatory frameworks and material constraints, as evidenced by the persistence of non-technological interventions such as subsidized meal programs in France to combat food insecurity [13].

Hypothesis: In the near term, academic skill development for a subset of students may transition toward the formal regulation of cognitive partnership to mitigate cognitive surrender, while student health & wellbeing may prioritize the mitigation of critical AI failure modes in crisis triage.

Mid (2036–2041)

For the same subset of institutions and regions, if current trajectories hold, the academic skill development process may formalize the auditing of machine-generated logical structures as a core competency, extrapolating from early frameworks that distinguish between cognitive support and cognitive substitution [8][14]. Within disability support management, the implementation of “Responsible Inference Engines” may evolve from a policy recommendation into a standard operational requirement for safeguarding students with learning differences [9]. In student health & wellbeing, AI-driven documentation and triage may further integrate into the workforce, provided that institutional protocols resolve the tension between professional fear and operational efficiency [3] and establish redundant human-in-the-loop overrides for high-risk mental health interventions [10]. In regions implementing standardized harm reporting, such as Florida, these mechanisms may evolve into permanent regulatory requirements for student wellbeing management [7]. The human-led provision of personal tutors is likely to remain unchanged in its core substance because the social legitimacy of high-touch wellbeing support depends on human presence to prevent the psychological isolation associated with “cognitive surrender” [2].

Hypothesis: In the mid term, disability support management may shift toward standardized, safeguarded inference engines, while academic skill development may incorporate the auditing of AI-generated logic as a primary student capability.

Far (2046–2051)

For the same subset of institutions and regions, the scope of AI integration in wellbeing management remains uncertain due to the tension between operational efficiency and the risk of cognitive and linguistic homogenization [5]. One possible outcome involves the emergence of a highly regulated, AI-integrated wellbeing ecosystem where “Responsible Inference Engines” [9] and bias-corrected counselling models [11] provide the primary layer of student engagement. Conversely, a rival outcome may see a retrenchment toward “human-only” zones for mental health and academic advice, driven by systemic distrust following high-profile AI failures in crisis triage [10] or the observation of widespread cognitive decline [2]. The extent to which institutional governance can mitigate AI-induced delusions [6] remains an unresolvable uncertainty at this horizon. High-stakes clinical diagnosis within student health & wellbeing will likely resist full automation due to the professional liability and accreditation constraints governing healthcare practice.

Hypothesis A: In the far term, student support may evolve into a hybrid model where AI handles all low-stakes administrative and learning management, while human practitioners focus exclusively on high-stakes crisis and complex wellbeing interventions.

Hypothesis B (competing): In the far term, systemic failures in AI-driven mental health triage and the erosion of critical thinking may lead institutions to mandate human-exclusive support protocols for all wellbeing-related capabilities.

References

1. The Conversation (2026-03-26). “College students are writing with AI – but a pilot study finds they’re not simply letting it write for them.” <https://theconversation.com/college-students-are-writing-with-ai-but-a-pilot-study-finds-theyre-not-simply-letting-it-write-for-them-276856>
2. Ars Technica (2026-04-03). ““Cognitive surrender” leads AI users to abandon logical thinking, research finds.” <https://arstechnica.com/ai/2026/04/research-finds-ai-users-scarily-willing-to-surrender-their-cognition-to-llms/>
3. NPR (2026-04-07). “AI in the mental health care workforce is met with fear, pushback — and enthusiasm.” <https://www.npr.org/2026/04/07/nx-s1-5771707/mental-health-care-workforce-artificial-intelligence-ai>
4. TechCrunch (2026-04-07). “Adobe launches Acrobat Spaces, a free AI-powered study tool for students.” <https://techcrunch.com/2026/04/07/adobe-launches-acrobat-spaces-a-free-ai-powered-study-tool-for-students/>
5. Hacker News (2026-04-07). “AI may be making us think and write more alike.” <https://dornsife.usc.edu/news/stories/ai-may-be-making-us-think-and-write-more-alike/>
6. Daily Ai (2025-05-09). “ChatGPT Is Making People Think They’re Gods and Their Families Are Terrified.” <https://dailyai.com/2025/05/chatgpt-is-making-people-think-theyre-gods-and-their-families-are-terrified/>
7. Future Of Life (2026-03-09). “Governor DeSantis Directs Florida State Agencies to Partner with Future of Life Institute to Shield Families from AI Harm.” <https://futureoflife.org/press-release/de-santis-directs-florida-agencies-to-partner-with-fli/>
8. Arxiv Cs Cy (2026-04-22). “Students Know AI Should Not Replace Thinking, but How Do They Regulate It? The TACO Framework for Human-AI Cognitive Partnership.” <https://arxiv.org/abs/2604.18737>
9. Getting Smart (2026-04-22). “Responsible Inference Engines: Safeguarding Students with Learning Differences in the AI Era.” <https://www.gettingsmart.com/2026/04/22/responsible-inference-engines-safeguarding-students-with-learning-differences-in-the-ai-era/>
10. Fast Company (2026-04-24). “LLMs don’t get mental health right. We need a two-pronged approach to fix them.” https://www.fastcompany.com/91527462/llms-dont-get-mental-health-right-we-need-a-two-pronged-approach-to-fix-them-ai-suicide-self-harm?partner=rss&utm_source=rss&u

tm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss

11. Arxiv Cs Cy (2026-04-30). “Sociodemographic Biases in Educational Counselling by Large Language Models.” <https://arxiv.org/abs/2604.25932>
12. Elearning Industry (2026-05-04). “5 Ways Adult Learners Can Use AI To Study Smarter, Without Compromising Integrity.” <https://elearningindustry.com/5-ways-adult-learners-can-use-ai-to-study-smarter-without-compromising-integrity>
13. The Guardian (2026-05-05). “France serves up €1 meals to all university students in effort to cut hardship.” <https://www.theguardian.com/world/2026/may/05/france-serves-up-1-euro-meals-to-all-university-students-in-bid-to-cut-hardship>
14. Arxiv Cs Cl (2026-05-08). “Counterargument for Critical Thinking as Judged by AI and Humans.” <https://arxiv.org/abs/2605.05353>

Supporting Services

capability_id: supporting-services · created: 2026-07-25 · revised: 2026-07-25 · refs: 5

Near (2026–2031)

For a subset of supporting services, specifically food, retail, and venue management, the delivery mechanisms for logistics and spatial monitoring are shifting toward automated digital systems. In food and retail management, the adoption of machine learning algorithms for demand forecasting alters how institutions analyze purchasing patterns to manage inventory [4]. However, the implementation of automated approval systems for supply chain logistics creates a counterproductive risk where physical goods are wasted because digital validation systems fail to approve shipments [1]. In venue management, the integration of AI video analytics transforms traditional monitoring into systems capable of real-time object detection and tracking [3], while predictive intelligence allows a shift from reactive observation toward forecasting movement patterns and density [2]. In France, the implementation of €1 meals for all university students represents a change in the educational substance of food management by prioritizing student hardship reduction over cost-recovery models [5]. Processes such as religious support and child care management are unlikely to adopt these specific automated frameworks in the near term because their institutional value depends on human presence and social legitimacy.

Hypothesis: In the near term, food and retail management may experience increased operational fragility due to digital validation failures, while venue management may shift toward proactive risk detection and predictive spatial monitoring.

Mid (2036–2041)

For the food, retail, and venue management subsegments, if current trajectories of digital integration hold, the stability of these services will depend on the resilience of the underlying software. Extrapolating from current failures in automated approvals, the risk of systemic waste in food and retail management could become a structural vulnerability unless institutions implement manual override protocols for physical logistics [1]. In venue management, predictive intelligence and video analytics may evolve from monitoring density to automating real-time resource allocation and crowd control mechanisms [2][3]. However, the educational substance of complaint and compliment management is likely to remain unchanged, as the resolution of human grievances typically requires social empathy and discretionary authority.

Hypothesis: In the mid term, food logistics may face persistent risks of “digital waste,” while venue management may integrate predictive analytics into automated operational resource allocation.

Far (2046–2051)

For the food, retail, and venue management subsegments, the long-term outcome depends on whether the industry optimizes for centralized efficiency or distributed resilience. One possibility involves the development of high-resilience, autonomous validation systems that mitigate the single points of failure currently causing supply chain waste [1]. A rival outcome is a state of chronic operational instability where the gap between digital approval and physical reality creates permanent inefficiencies in food and retail management, potentially forcing a return to non-digital, localized procurement. It remains unknown if the cost of system-induced waste will outweigh the benefits of predictive demand forecasting [4]. Processes like religious support will likely resist this trajectory throughout this horizon as their primary function is the provision of human-centric spiritual care.

Hypothesis A: In the far term, these subsegments may achieve autonomous, high-resilience operational states via distributed intelligence and predictive validation.

Hypothesis B (competing): In the far term, supporting services may experience structural degradation and chronic waste due to an irreversible dependence on fragile automated validation systems.

References

1. Science Daily Comp (2026). “Truckloads of food are being wasted because computers won’t approve them.” <https://www.sciencedaily.com/releases/2026/04/260403224505.htm>
2. Dev To Ai (2026). “Architecting Predictive Intelligence for Smart Venues.” https://dev.to/tanisha_khiara_4b390b7946/architecting-predictive-intelligence-for-smart-venues-1lkj
3. Dev To Ai (2026). “Integrating AI Video Analytics into Venue Operations.” https://dev.to/tanisha_khiara_4b390b7946/integrating-ai-video-analytics-into-venue-operations-2m09
4. Dev To Ai (2026). “How Smart Technology Is Reshaping Retail.” https://dev.to/xit_vali_8353fdb3474555c/how-smart-technology-is-reshaping-retail-5308
5. The Guardian (2026). “France serves up €1 meals to all university students in effort to cut hardship.” <https://www.theguardian.com/world/2026/may/05/france-serves-up-1-euro-meals-to-all-university-students-in-bid-to-cut-hardship>

Teaching & Learning Delivery

capability_id: teaching-learning-delivery · created: 2026-09-29 · revised: 2026-09-29 · refs: 23

Near (2026–2031)

For teaching & learning delivery broadly, the educational substance is shifting toward skills-based learning frameworks to address workplace volatility [13] and a prioritization of learning processes over the production of academic artifacts [19]. Within the process of teaching & learning content management, the delivery mechanism is shifting from manual authoring to automated multimodal generation, utilizing tools for instructional videos [6], automated multimodal illustrations [7], and the conversion of URLs into presentations [16]. This delivery shift is supported by the treatment of instructional materials as raw material for multiple formats [30], the use of agentic co-authoring tools [20], the application of “vibe coding” to accelerate interactive digital materials [35], and AI avatars to convert static documents into interactive experiences [14]. In non presence based teaching, delivery mechanisms are being modified for specific subsegments through the deployment of intelligent tutoring systems for quantum education [26], Python [2], and low-resource languages [22], alongside 24/7 asynchronous doubt resolution [32], adaptive tutoring for specific examination levels [33], and temporal layers in retrieval-augmented generation (RAG) to prevent the delivery of outdated information [42]. Presence based teaching is being modified through the adoption of “AI-Active Lessons” [36] and “cognitive partnership” models [17], where the educational substance includes treating AI hallucinations as mechanisms to prompt higher-order critical thinking [41] and emphasizing the cognitive struggle of the writing process over the final product [43]. The professional socialization function of presence based teaching is unlikely to change yet because the social legitimacy of the human instructor remains a primary requirement for institutional accreditation and faculty labor contracts.

Hypothesis: In the near term, the capability shifts toward a process-oriented educational substance supported by a delivery mechanism of automated multimodal content and specialized, adaptive non presence based tutoring interfaces.

Mid (2036–2041)

For a subset of institutions adopting automated delivery models and STEM-focused subsegments, the delivery mechanism of non presence based teaching may shift toward agent-led orchestration of student pathways if current trajectories in specialized AI tutoring and autonomous schooling models hold [10]. Within teaching & learning content management, the process may evolve from a creation-centric activity into a verification-heavy auditing role, provided that governance frameworks for quality, accuracy, and transparency are standardized [5][11]. If the risk of “AI tunnel vision” [34] or “cognitive bypass” persists, presence based teaching could shift toward constrained delivery models that intentionally limit AI autonomy to protect the cognitive substance of the learning process [17]. The requirement for human-led mentorship in high-stakes professional socialization is likely to persist because it is deeply embedded in the regulatory regimes of professional licensing and accreditation.

Hypothesis: In the mid term, content management may transition into a verification and governance function, while non presence based teaching may shift toward agent-led pathway orchestration for a subset of institutions.

Far (2046–2051)

For institutions using automated delivery models, the delivery mechanism of non presence based teaching may bifurcate between fully autonomous skill-acquisition agents and human-led synthesis seminars. A primary unknown remains whether the cognitive substance of “learning how to learn” can be maintained when the struggle of information retrieval and synthesis is outsourced to agentic systems. One rival outcome suggests a total decoupling of skill acquisition from the campus, leaving presence based teaching exclusively for high-stakes social and professional validation. Conversely, a counter-movement may emerge that re-establishes “analog” presence based teaching as a premium, AI-restricted environment to ensure

cognitive depth. The institutional requirement for degrees as signals of social legitimacy is likely to resist change even at this horizon.

Hypothesis A: In the far term, the capability splits into a binary of autonomous agentic skill-acquisition and high-touch human synthesis.

Hypothesis B (competing): In the far term, a market for “AI-free” presence based teaching emerges as a premium educational substance to counteract cognitive bypass.

References

1. Elearning Industry (2026-04-05). “Leveraging AI And Automation To Streamline CTE Content Creation And Localization.” <https://elearningindustry.com/leveraging-ai-and-automation-to-streamline-cte-content-creation-and-localization>
2. The Hechinger Report (2026-04-06). “The quest to build a better AI tutor.” <https://hechingerreport.org/proof-points-ai-tutor-python/>
3. Elearning Industry (2026-04-07). “A 4-Step Framework For Using AI Transparently In Educational Content.” <https://elearningindustry.com/a-4-step-framework-for-using-ai-transparently-in-educational-content>
4. Google AI Blog (2026-04-02). “Create, edit and share videos at no cost in Google Vids.” <https://blog.google/products-and-platforms/products/workspace/google-vids-updates-lyria-veo/>
5. Arxiv Cs Cy (2026-04-08). “EduIllustrate: Towards Scalable Automated Generation Of Multimodal Educational Content.” <https://arxiv.org/abs/2604.05005>
6. Atlantic (2026-04-10). “This Is What Fully Automated School Looks Like.” https://www.theatlantic.com/technology/2026/04/ai-agents-school-education/686754/?utm_source=feed
7. Elearning Industry (2026-04-10). “AI-Generated Learning Content Governance: Ensuring Quality, Accuracy, And Trust.” <https://elearningindustry.com/ai-generated-learning-content-governance-ensuring-quality-accuracy-and-trust>
8. Elearning Industry (2026-04-15). “The Case For Skills-Based Learning In A Rapidly Changing Workplace.” <https://elearningindustry.com/the-case-for-skills-based-learning-in-a-rapidly-changing-workplace>
9. Elearning Industry (2026-04-21). “How AI Avatars Are Turning Static Documents Into Interactive Learning Experiences.” <https://elearningindustry.com/how-ai-avatars-are-turning-static-documents-into-interactive-learning-experiences>
10. Dev To Ai (2026-04-21). “Create Presentations from Any URL Automatically Using AI.” https://dev.to/ai_presentation_d9935ff95/create-presentations-from-any-url-automatically-using-ai-5hho
11. Arxiv Cs Cy (2026-04-22). “Students Know AI Should Not Replace Thinking, but How Do They Regulate It? The TACO Framework for Human-AI Cognitive Partnership.” <https://arxiv.org/abs/2604.18737>
12. EDUCAUSE Review (2026-04-23). “Teaching and Learning with AI: Three Priorities.” <https://er.educause.edu/podcasts/educause-exchange/teaching-and-learning-with-ai-three-priorities>
13. The Register (2026-04-23). “Microsoft gives your Word documents an AI co-author you didn’t ask for.” https://go.theregister.com/feed/www.theregister.com/2026/04/23/microsoft_gives_your_word_documents/
14. Arxiv Cs Cl (2026-04-24). “AFRILANGTUTOR: Advancing Language Tutoring and Culture Education in Low-Resource Languages with Large Language Models.” <https://arxiv.org/abs/2604.20996>
15. Arxiv Cs Cy (2026-04-29). “From Prototype to Classroom: An Intelligent Tutoring System for Quantum Education.” <https://arxiv.org/abs/2604.24807>
16. Fast Company (2026-05-01). “AI is turning every story into raw material.” https://www.fastcompany.com/91532289/ai-is-turning-every-story-into-raw-material?partner=rss&utm_source=rss&utm_medium=feed&utm_campaign=rss+fastcompany&utm_content=rss
17. Dev To Ai (2026-05-05). “Doubt Solver That Works at 2 AM — AI Never Sleeps When You Study.” https://dev.to/aditiya_samay_999/doubt-solver-that-works-at-2-am-ai-never-sleeps-when-you-

study-22c3

18. Dev To Ai (2026-05-05). “AI Tutor That Understands JEE Level vs CBSE Level — Smart Adaptation.” <https://dev.to/jaiswal007/ai-tutor-that-understands-jee-level-vs-cbse-level-smart-adaptation-19h1>
19. Elearning Industry (2026-05-05). “The Hidden Risk In AI-Driven Learning.” <https://elearningindustry.com/ai-tunnel-vision-the-hidden-risk-in-ai-driven-learning>
20. Elearning Industry (2026-05-05). “Vibe Coding In eLearning: How AI Is Transforming Interactive Course Creation.” <https://elearningindustry.com/vibe-coding-in-elearning-how-ai-is-transforming-interactive-course-creation>
21. Edtech Magazine (2026-05-05). “Bringing the AI-Active Lesson to Life in Higher Education.” <https://edtechmagazine.com/higher/article/2026/05/bringing-ai-active-lesson-life-higher-education>
22. Arxiv Cs Cy (2026-05-08). “The Pedagogy of AI Mistakes: Fostering Higher-Order Thinking.” <https://arxiv.org/abs/2605.05472>
23. Towards Data Science (2026-05-09). “RAG Is Blind to Time — I Built a Temporal Layer to Fix It in Production.” <https://towardsdatascience.com/rag-is-blind-to-time-i-built-a-temporal-layer-to-fix-it-in-production/>
24. The Guardian (2026-05-10). “I knew my writing students were using AI. Their confessions led to a powerful teaching moment | Micah Nathan.” <https://www.theguardian.com/us-news/ng-interactive/2026/may/10/fiction-writing-professor-ai>